Q :What is an on Demand instance?
On-Demand Instances let you pay for compute capacity by the hour or second (minimum of 60 seconds) with no long-term commitments. This frees you from the costs and complexities of planning, purchasing, and maintaining hardware and transforms what are commonly large fixed costs into much smaller variable costs.

Q. What is a reserved instance?
An Amazon Reserved Instance (RI) is a billing discount that allows you to save on your Amazon EC2 usage costs. When you purchase a Reserved Instance, you can set attributes such as instance type, platform, tenancy, Region, or Availability Zone (optional).

Q. What are the benefits of using reserved instances?
Image result for reserved instances
Reserved Instances provide you with significant savings on your Amazon EC2 costs compared to On-Demand Instance pricing. Reserved Instances are not physical instances, but rather a billing discount applied to the use of On-Demand Instances in your account.

Q. Does AWS have reserved instances?
Buying a Reserved Instance is easy. You can use either the AWS Management Console or API tools to complete your purchase. Reserved DB instances are available in three varieties - No Upfront, Partial Upfront, and All Upfront - that let you optimize your Amazon RDS costs based on your expected usage.

Q. What is the difference between On demand instances and Reserved instances?
In terms of compute options and configurations, Reserved Instances and On Demand instances are the same. The only difference between the two is that a Reserved Instance is one you rent (“reserve”) for a fixed duration, and in return you receive a discount on the base price of an On Demand instance.

Q. What is the difference between EC2 Dedicated instances and Reserved instances?
Dedicated instances may share hardware with other instances from the same AWS account that are not Dedicated instances. Pay for Dedicated Instances On-Demand, save up to 70% by purchasing Reserved Instances, or save up to 90% by purchasing Spot Instances.

Q. What are the three types of EC2 Instances?
They are On-Demand Instances, Reserved Instances, Spot Instances, and Savings Plans. Amazon also offers the additional option of Amazon EC2 Dedicated Hosts.

Q. IAM is a cloud service that controls the permissions and access for users and cloud resources. IAM policies are sets of permission policies that can be attached to either users or cloud resources to authorize what they access and what they can do with it.

Q. What is RDM used for?
Image result for RDM services in cloud computing
Raw device mapping (RDM) enables disk access in a virtual machine (VM) in the VMware server virtualization environment and allows a storage logical unit number (LUN) to be connected directly to a VM from the storage area network (SAN).

Q. What is difference between SQS and SNS?
SNS is a pub/sub system, while SQS is a queueing system. You'd typically use SNS to send the same message to multiple consumers via topics. In comparison, in most scenarios, each message in an SQS queue is processed by only one consumer.

Q.What is SNS and how it works?
How it works. Amazon Simple Notification Service (SNS) sends notifications two

ways, A2A and A2P. A2A provides high-throughput, push-based, many-to-many messaging between distributed systems, microservices, and event-driven serverless applications.

Q What is meant by Auto Scaling?
Image result for Auto scaling
Autoscaling provides users with an automated approach to increase or decrease the compute, memory or networking resources they have allocated, as traffic spikes and use patterns demand.

Q . The country currently is home to over 100 data centers, both operated locally and by large tech companies. Currently, India's data center industry is concentrated in four cities–Mumbai, Delhi, Bengaluru, and Chennai, accounting for most of the data center sites and the IT load capacities.

Q What is a datacenter region?
A Region is set of collaborating Zones (datacenters) grouped together based on their geographical proximity. The Zones within a Region are connected with high-speed networking to facilitate low-latency communication between compute instances and reliable replication of data