Custom View Settings

Topic 1 - Exam A

Question #1 *Topic 1*

A company collects data for temperature, humidity, and atmospheric pressure in cities across multiple continents. The average volume of data that the company collects from each site daily is 500 GB. Each site has a high-speed Internet connection.
The company wants to aggregate the data from all these global sites as quickly as possible in a single Amazon S3 bucket. The solution must minimize operational complexity.
Which solution meets these requirements?

A. Turn on S3 Transfer Acceleration on the destination S3 bucket. Use multipart uploads to directly upload site data to the destination S3 bucket.

B. Upload the data from each site to an S3 bucket in the closest Region. Use S3 Cross-Region Replication to copy objects to the destination S3 bucket. Then remove the data from the origin S3 bucket.

C. Schedule AWS Snowball Edge Storage Optimized device jobs daily to transfer data from each site to the closest Region. Use S3 Cross-Region Replication to copy objects to the destination S3 bucket.

D. Upload the data from each site to an Amazon EC2 instance in the closest Region. Store the data in an Amazon Elastic Block Store (Amazon EBS) volume. At regular intervals, take an EBS snapshot and copy it to the Region that contains the destination S3 bucket. Restore the EBS volume in that Region.

Correct Answer: *A*

Community vote distribution

A (100%)

D2w Highly Voted 1 month, 2 weeks ago

Selected Answer: A

S3 Transfer Acceleration is the best solution cz it's faster , good for high speed, Transfer Acceleration is designed to optimize transfer speeds from across the world into S3 buckets.

upvoted 12 times

BoboChow 1 month, 2 weeks ago

I thought S3 Transfer Acceleration is based on Cross Region Repilication, I made a mistake.

upvoted 1 times

jidexray Most Recent 5 days, 13 hours ago

A is the right answer

upvoted 1 times

Wpcorgan 1 week ago

A is correct

upvoted 1 times

JayanKuruwita 1 week, 1 day ago

Selected Answer: A

The only other answer give some compitition for A is C, but it said high speed internet and need to lower the operational overhead. So you be A.

upvoted 1 times

jidexray 1 week, 2 days ago

A is correct

upvoted 1 times

Buruguduystunstugudunstuy 2 weeks, 4 days ago

The correct Answer is A

http://lavnish.blogspot.com/2017/06/aws-s3-cross-region-replication.html

upvoted 1 times

Naneyerocky 3 weeks, 2 days ago

Amazon S3 Transfer Acceleration is a bucket-level feature that enables fast, easy, and secure transfers of files over long distances between your client and an S3 bucket. Transfer Acceleration is designed to optimize transfer speeds from across the world into S3 buckets.

upvoted 2 times

gulam11adam 3 weeks, 5 days ago

Selected Answer: A

A is correct

upvoted 1 times

17Master 1 month ago

Selected Answer: A

Ans is A

upvoted 1 times

cheese929 1 month ago

A is correct

upvoted 1 times

queen101 1 month, 1 week ago

AAAAAAAAAA

upvoted 1 times

bilel500 1 month, 2 weeks ago

Selected Answer: A

Answer is A.

upvoted 1 times

bilel500 1 month, 2 weeks ago

Answer is A.

upvoted 1 times

sba21 1 month, 2 weeks ago

Answer: A true

upvoted 1 times

Question #2 *Topic 1*

A company needs the ability to analyze the log files of its proprietary application. The logs are stored in JSON format in an Amazon S3 bucket. Queries will be simple and will run on-demand. A solutions architect needs to perform the analysis with minimal changes to the existing architecture.
What should the solutions architect do to meet these requirements with the LEAST amount of operational overhead?

A. Use Amazon Redshift to load all the content into one place and run the SQL queries as needed.

B. Use Amazon CloudWatch Logs to store the logs. Run SQL queries as needed from the Amazon CloudWatch console.

C. Use Amazon Athena directly with Amazon S3 to run the queries as needed.

D. Use AWS Glue to catalog the logs. Use a transient Apache Spark cluster on Amazon EMR to run the SQL queries as needed.

Correct Answer: *C*

Community vote distribution

C (100%)

airraid2010 Highly Voted 1 month, 2 weeks ago

Answer: C
https://docs.aws.amazon.com/athena/latest/ug/what-is.html
Amazon Athena is an interactive query service that makes it easy to analyze data directly in Amazon Simple Storage Service (Amazon S3) using standard SQL. With a few actions in the AWS Management Console, you can point Athena at your data stored in Amazon S3 and begin using standard SQL to run ad-hoc queries and get results in seconds.

upvoted 6 times

BoboChow 1 month, 2 weeks ago

I agree C is the answer

upvoted 1 times

tt79 1 month, 2 weeks ago

C is right.

upvoted 1 times

Drekorig Most Recent 6 days, 22 hours ago

Selected Answer: C

Athena allows query data stored in S3 with SQL

upvoted 1 times

Wpcorgan 1 week ago

C is correct

upvoted 1 times

pm2229 3 weeks, 1 day ago

Serverless query service to perform analytics on S3.

upvoted 1 times

Drchattss 4 weeks, 1 day ago

Can you pass test with ONLY this dump

upvoted 1 times

17Master 1 month ago

Selected Answer: C

Amazon Athena is an interactive query service that makes it easy to analyze data in Amazon S3 using standard SQL. Athena is serverless, so there is no infrastructure to manage, and you pay only for the queries that you run.
https://aws.amazon.com/athena/?nc1=h_ls&whats-new-cards.sort-by=item.additionalFields.postDateTime&whats-new-cards.sort-order=desc

upvoted 1 times

cheese929 1 month ago

Selected Answer: C

Amazon Athena is an interactive query service that makes it easy to analyze data directly in Amazon Simple Storage Service (Amazon S3) using standard SQL.

upvoted 1 times

ManoAni 1 month, 1 week ago

Are these the new set of questions as Amazon has updated their exam?

upvoted 1 times

dduque10 3 weeks ago

yes it is, please notice that the new version is SAA-C03, while the old version is SAA-C02

upvoted 1 times

queen101 1 month, 1 week ago

CCCCCCCCCC

upvoted 1 times

bilel500 1 month, 2 weeks ago

Selected Answer: C

Answer is C.

upvoted 1 times

sba21 1 month, 2 weeks ago

Selected Answer: C

Athena allows directly query data stored in S3

upvoted 2 times

sba21 1 month, 2 weeks ago

Answer: C

Athena is only service that allows directly query data stored in S3

upvoted 1 times

D2w 1 month, 2 weeks ago

Selected Answer: C

Amazon Athena is Serverless query service to perform analytics against S3 objects. And cz it wants minimal changes I'll definitely go with it.

upvoted 3 times

tt79 1 month, 2 weeks ago

Answer:C

upvoted 2 times

Question #3 *Topic 1*

A company uses AWS Organizations to manage multiple AWS accounts for different departments. The management account has an Amazon S3 bucket that contains project reports. The company wants to limit access to this S3 bucket to only users of accounts within the organization in AWS Organizations.
Which solution meets these requirements with the LEAST amount of operational overhead?

A. Add the aws PrincipalOrgID global condition key with a reference to the organization ID to the S3 bucket policy.

B. Create an organizational unit (OU) for each department. Add the aws:PrincipalOrgPaths global condition key to the S3 bucket policy.

C. Use AWS CloudTrail to monitor the CreateAccount, InviteAccountToOrganization, LeaveOrganization, and RemoveAccountFromOrganization events. Update the S3 bucket policy accordingly.

D. Tag each user that needs access to the S3 bucket. Add the aws:PrincipalTag global condition key to the S3 bucket policy.

Correct Answer: *A*

Community vote distribution

A (85%) B (15%)

ude Highly Voted 1 month, 2 weeks ago

Selected Answer: A

aws:PrincipalOrgID Validates if the principal accessing the resource belongs to an account in your organization.
https://aws.amazon.com/blogs/security/control-access-to-aws-resources-by-using-the-aws-organization-of-iam-principals/

upvoted 17 times

BoboChow 1 month, 2 weeks ago

the condition key aws:PrincipalOrgID can prevent the members who don't belong to your organization to access the resource

upvoted 1 times

Rock08 Highly Voted 1 month, 2 weeks ago

Selected Answer: B

Now you can arrange your AWS accounts into groups called organizational units (OUs) and apply policies to OUs or directly to accounts. For example, you can organize your accounts by application, environment, team, or any other grouping that makes sense for your business.

upvoted 5 times

thangvu1890 2 weeks, 6 days ago

It's ok. But It's not the LEAST amount of operational overhead.

upvoted 1 times

debillion 2 weeks, 6 days ago

Yeah! I thought as much.

upvoted 1 times

ArielSchivo 4 weeks ago

But the goal here is to limit access just for the users of the account, not groups or departments. That's why A is correct.

upvoted 3 times

Wpcorgan Most Recent 1 week ago

A is correct

upvoted 1 times

VTI_Training 2 weeks, 5 days ago

Selected Answer: A

A is correct

upvoted 1 times

Saiofy 2 weeks, 5 days ago

Selected Answer: A

.... and it's A

upvoted 1 times

pm2229 3 weeks, 1 day ago

It's A, IAM now makes it easier for you to control access to your AWS resources by using the AWS organization of IAM principals (users and roles). You can use the aws:PrincipalOrgID condition key in your resource-based policies to more easily restrict access to IAM principals from accounts in your AWS organization.

upvoted 1 times

Naneyerocky 3 weeks, 2 days ago

Selected Answer: A

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_permissions_overview.html
Condition keys: AWS provides condition keys that you can query to provide more granular control over certain actions.
The following condition keys are especially useful with AWS Organizations:

aws:PrincipalOrgID – Simplifies specifying the Principal element in a resource-based policy. This global key provides an alternative to listing all the account IDs for all AWS accounts in an organization. Instead of listing all of the accounts that are members of an organization, you can specify the organization ID in the Condition element.

aws:PrincipalOrgPaths – Use this condition key to match members of a specific organization root, an OU, or its children. The aws:PrincipalOrgPaths condition key returns true when the principal (root user, IAM user, or role) making the request is in the specified organization path. A path is a text representation of the structure of an AWS Organizations entity.

upvoted 3 times

17Master 1 month ago

Selected Answer: A

ans is A. The company wants to limit access to this S3 bucket to only users of accounts within the organization in AWS Organizations.
https://aws.amazon.com/blogs/security/control-access-to-aws-resources-by-using-the-aws-organization-of-iam-principals/

upvoted 1 times

JesseeS 1 month ago

https://aws.amazon.com/blogs/security/control-access-to-aws-resources-by-using-the-aws-organization-of-iam-principals/ Answer is A

upvoted 1 times

queen101 1 month, 1 week ago

AAAAAAAAAA

upvoted 1 times

123jhl0 1 month, 1 week ago

Selected Answer: A

A requires the LEAST effort

upvoted 2 times

bilel500 1 month, 2 weeks ago

Selected Answer: A

Answer is A.

upvoted 1 times

huiy 1 month, 2 weeks ago

Selected Answer: A

Agree A

upvoted 1 times

josh_fan 1 month, 2 weeks ago

Selected Answer: A

Now, you can use a new condition key, aws:PrincipalOrgID, in these policies to require all principals accessing the resource to be from an account (including the master account) in the organization.

upvoted 2 times

Question #4 *Topic 1*

An application runs on an Amazon EC2 instance in a VPC. The application processes logs that are stored in an Amazon S3 bucket. The EC2 instance needs to access the S3 bucket without connectivity to the internet.
Which solution will provide private network connectivity to Amazon S3?

A. Create a gateway VPC endpoint to the S3 bucket.

B. Stream the logs to Amazon CloudWatch Logs. Export the logs to the S3 bucket.

C. Create an instance profile on Amazon EC2 to allow S3 access.

D. Create an Amazon API Gateway API with a private link to access the S3 endpoint.

Correct Answer: *A*

Community vote distribution

A (100%)

D2w Highly Voted 1 month, 2 weeks ago

Selected Answer: A

VPC endpoint allows you to connect to AWS services using a private network instead of using the public Internet

upvoted 8 times

Drekorig Most Recent 5 days, 5 hours ago

Selected Answer: A

To provide connectivity the answer is "A". To authorize the connection we can use the instance profile.

upvoted 1 times

cheese929 5 days, 23 hours ago

Selected Answer: A

A is right.

upvoted 1 times

Wpcorgan 1 week ago

A is correct

upvoted 1 times

grzeev 1 week, 1 day ago

Selected Answer: A

https://docs.aws.amazon.com/vpc/latest/privatelink/gateway-endpoints.html

upvoted 1 times

pm2229 3 weeks, 1 day ago

It's A, Private endpoints within your VPC that allow AWS services to privately connect to resources within your VPC without traversing the public internet (cheaper)

upvoted 2 times

Naneyerocky 3 weeks, 2 days ago

Selected Answer: A

https://docs.aws.amazon.com/vpc/latest/privatelink/vpc-endpoints-s3.html
You can access Amazon S3 from your VPC using gateway VPC endpoints. After you create the gateway endpoint, you can add it as a target in your route table for traffic destined from your VPC to Amazon S3.
You can create a policy that restricts access to specific IP address ranges by using the aws:VpcSourceIp condition key.

upvoted 2 times

Adrianjavier 3 weeks, 6 days ago

Selected Answer: A

A for sure

upvoted 1 times

17Master 1 month ago

Selected Answer: A

ans is A

upvoted 1 times

bilel500 1 month, 1 week ago

Selected Answer: A

Answer is A.
upvoted 1 times

queen101 1 month, 1 week ago
AAAAAAAAAAA
upvoted 1 times

BoboChow 1 month, 2 weeks ago
Selected Answer: A
obviously it's A
upvoted 1 times

Mikedodo 1 month, 2 weeks ago
Selected Answer: A
Reduce Cost and Increase Security with Amazon VPC Endpoints
upvoted 2 times

Question #5 *Topic 1*

A company is hosting a web application on AWS using a single Amazon EC2 instance that stores user-uploaded documents in an Amazon EBS volume. For better scalability and availability, the company duplicated the architecture and created a second EC2 instance and EBS volume in another Availability Zone, placing both behind an Application Load Balancer. After completing this change, users reported that, each time they refreshed the website, they could see one subset of their documents or the other, but never all of the documents at the same time.
What should a solutions architect propose to ensure users see all of their documents at once?

A. Copy the data so both EBS volumes contain all the documents

B. Configure the Application Load Balancer to direct a user to the server with the documents

C. Copy the data from both EBS volumes to Amazon EFS. Modify the application to save new documents to Amazon EFS

D. Configure the Application Load Balancer to send the request to both servers. Return each document from the correct server

Correct Answer: *C*

Community vote distribution

C (100%)

D2w Highly Voted 1 month, 2 weeks ago

Selected Answer: C

Concurrent or at the same time key word for EFS

upvoted 11 times

cheese929 Most Recent 5 days, 23 hours ago

Selected Answer: C

C is the only solution that make sense.

upvoted 1 times

Wpcorgan 1 week ago

C is correct

upvoted 1 times

mikey2000 1 week, 5 days ago

Ebs doesnt support cross az only reside in one Az but Efs does, that why it's c

upvoted 4 times

Azeeza 2 weeks, 5 days ago

Amazon Elastic File System is a cloud storage service provided by Amazon Web Services designed to provide scalable, elastic and concurrency.
Answer is C

upvoted 3 times

pm2229 3 weeks, 1 day ago

It's C, EFS can be mounted to multiple EC2 instances across AZs. The Performance is higher latency & throughput.

upvoted 2 times

17Master 1 month ago

Selected Answer: C

Ans is C

upvoted 1 times

bilel500 1 month, 1 week ago

Selected Answer: C

Answer is C.
Adaptive throughput – EFS's performance can scale in-line with its storage, operating at a higher throughput for sudden, high-volume file dumps, reaching up to 500,000 IOPS or 10 GB per second
Totally elastic – once you've spun up an EFS instance, you can add add files without worrying about provisioning or disturbing your application's performance
Additional accessibility – EFS can be mounted from different EC2 instances, but it can also cross the AWS region boundary via the use of VPC peering

upvoted 3 times

queen101 1 month, 1 week ago

CCCCCCCCCCC

upvoted 1 times

BoboChow 1 month, 2 weeks ago

Selected Answer: C

I'm not sure if ALB could configure like session ID thing, so I choose C

upvoted 1 times

Question #6 *Topic 1*

A company uses NFS to store large video files in on-premises network attached storage. Each video file ranges in size from 1 MB to 500 GB. The total storage is 70 TB and is no longer growing. The company decides to migrate the video files to Amazon S3. The company must migrate the video files as soon as possible while using the least possible network bandwidth.
Which solution will meet these requirements?

A. Create an S3 bucket. Create an IAM role that has permissions to write to the S3 bucket. Use the AWS CLI to copy all files locally to the S3 bucket.

B. Create an AWS Snowball Edge job. Receive a Snowball Edge device on premises. Use the Snowball Edge client to transfer data to the device. Return the device so that AWS can import the data into Amazon S3.

C. Deploy an S3 File Gateway on premises. Create a public service endpoint to connect to the S3 File Gateway. Create an S3 bucket. Create a new NFS file share on the S3 File Gateway. Point the new file share to the S3 bucket. Transfer the data from the existing NFS file share to the S3 File Gateway.

D. Set up an AWS Direct Connect connection between the on-premises network and AWS. Deploy an S3 File Gateway on premises. Create a public virtual interface (VIF) to connect to the S3 File Gateway. Create an S3 bucket. Create a new NFS file share on the S3 File Gateway. Point the new file share to the S3 bucket. Transfer the data from the existing NFS file share to the S3 File Gateway.

Correct Answer: *C*

Community vote distribution

B (80%) C (17%)

tuloveu Highly Voted 1 month, 2 weeks ago

Selected Answer: B

As using the least possible network bandwidth.

upvoted 14 times

Gatt Highly Voted 1 month, 1 week ago

Selected Answer: B

Let's analyse this:

B. On a Snowball Edge device you can copy files with a speed of up to 100Gbps. 70TB will take around 5600 seconds, so very quickly, less than 2 hours. The downside is that it'll take between 4-6 working days to receive the device and then another 2-3 working days to send it back and for AWS to move the data onto S3 once it reaches them. Total time: 6-9 working days. Bandwidth used: 0.

C. File Gateway uses the Internet, so maximum speed will be at most 1Gbps, so it'll take a minimum of 6.5 days and you use 70TB of Internet bandwidth.

D. You can achieve speeds of up to 10Gbps with Direct Connect. Total time 15.5 hours and you will use 70TB of bandwidth. However, what's interesting is that the question does not specific what type of bandwidth? Direct Connect does not use your Internet bandwidth, as you will have a dedicate peer to peer connectivity between your on-prem and the AWS Cloud, so technically, you're not using your "public" bandwidth.

The requirements are a bit too vague but I think that B is the most appropriate answer, although D might also be correct if the bandwidth usage refers strictly to your public connectivity.

upvoted 9 times

JohnnyBG 6 days, 4 hours ago

So you decided the internet connectivity is at 1Gb at most ?!?!? I would select B also.

upvoted 1 times

abhishek_m89 1 week ago

and it says, "The total storage is 70 TB and is no longer growing". Thats why it should be B.

upvoted 1 times

Gatt 2 weeks, 5 days ago

I will add that the question does not specify if the company already has DA in place or not. If they don't have DA in place, it will take a long time (weeks) for DA connectivity to be setup. Another point for B here, as Snowball is much quicker from this perspective.

upvoted 2 times

th3cookie 1 month ago

The company must transfer the data asap. Direct connect takes a month to setup doesn't it?

upvoted 1 times

Gatt 2 weeks, 5 days ago

That's a good point, indeed, DA might take weeks to establish (depending on your local ISP). And the question does not state that DA has already been established for this company. If they are starting fresh, then certainly DA would be taking too long.

upvoted 1 times

Drekorig Most Recent 5 days, 5 hours ago

Selected Answer: C

This question is just terrible. I think the guys at AWS are talking about Amazon FSx File Gateway. So the reward is "C".
Amazon FSx File Gateway optimizes on-premises access to fully managed, highly reliable file shares in Amazon FSx for Windows File Server. Customers with unstructured or file data, whether from SMB-based group shares, or business applications, may require on-premises access to meet low-latency requirements. Amazon FSx File Gateway helps accelerate your file-based storage migration to the cloud to enable faster performance, improved data protection, and reduced cost.

upvoted 1 times

Wpcorgan 1 week ago

B is correct

upvoted 1 times

jackwang0616 1 week ago

Selected Answer: B

BBBBBBB

upvoted 1 times

manu427 2 weeks, 2 days ago

Direct Connect lead time more than a moth..so I will go with B

upvoted 1 times

Buruguduystunstugudunstuy 2 weeks, 4 days ago

The KEYWORD "The company must migrate the video files as soon as possible"

I would take the correct answer C.

upvoted 1 times

Azeeza 2 weeks, 5 days ago

The answer is C, because of time constraint identified in the question

upvoted 1 times

dduque10 3 weeks ago

Why does it say C if the most voted answer is B?

upvoted 1 times

KaityL 3 weeks, 1 day ago

Selected Answer: B

I chose B. A makes no sense. Direct connect takes very long. It said they wanted it done quickly. And its alot of data to be transfered.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: B

I go with B because out of all the options it uses the least bandwidth (none at all),
and 70 TB is pretty much the amount of data where you would want to use snowball

upvoted 2 times

Wajif 1 month, 1 week ago

Selected Answer: C

I am choosing C only because it says least network bandwidth didnt say "No" network bandwidth. It is extremely vague though. They could frame it a bit better.

upvoted 2 times

sensei705 1 month, 1 week ago

Option C is the first to be ruled out. File gateway (storage gateway) is not a migration tool. It simply bridges data from on-prem and AWS (e.g. S3) with local caching on-prem. They are asking about migrating data... file gateway does not migrate data, it is not a migration solution.

upvoted 4 times

17Master 1 week, 6 days ago

Its correct!.

upvoted 1 times

bilel500 1 month, 1 week ago

Selected Answer: B

Answer is B.

upvoted 1 times

ArielSchivo 1 month, 1 week ago

Selected Answer: D

A. No, it uses public connection bandwidth.
B. No, it takes at least 10 days to request, move files and return the device to AWS.
C. No, S3 File Gateway uses your public connection bandwidth.
D. CORRECT, with Direct Connect you are no using you public connection bandwidth.

upvoted 1 times

Foulstryke 1 month, 1 week ago

The problem with D is it takes over a month to get a new Direct Connection.
It's quite possibly the slowest option.

upvoted 5 times

ArielSchivo 3 weeks, 6 days ago

Yep, I think you are right, so I choose B.

upvoted 1 times

queen101 1 month, 1 week ago

BBBBBBBBBB

upvoted 1 times

123jhl0 1 month, 1 week ago

Selected Answer: B

We don't know the available bandwidth, BUT just suppose we can consume exclusively for migration 1Gbps. In this case the migration will take more than 6 days to move the 70TB. A week is also the time used by the snowball devices consuming no bandwidth at all.
B - If you choose any other option, you involve +1Gbps bandwidth to be used exclusively for migration, which goes against the LEAST BW requested. For instance, to migrate it in 1 day, you would need 10Gbps.

upvoted 2 times

Rachness 1 month, 2 weeks ago

Even with high-speed internet connections, it can take months to transfer large amounts of data. For example, it can take more than 100 days to transfer 100 TB of data over a dedicated 100-Mbps connection. You can accomplish the same transfer in less than one day, plus shipping time, using two AWS Snowball Edge devices. https://aws.amazon.com/blogs/storage/best-practices-for-accelerating-data-migrations-using-aws-snowball-edge/

upvoted 1 times

Question #7 *Topic 1*

A company has an application that ingests incoming messages. Dozens of other applications and microservices then quickly consume these messages. The number of messages varies drastically and sometimes increases suddenly to 100,000 each second. The company wants to decouple the solution and increase scalability.
Which solution meets these requirements?

A. Persist the messages to Amazon Kinesis Data Analytics. Configure the consumer applications to read and process the messages.

B. Deploy the ingestion application on Amazon EC2 instances in an Auto Scaling group to scale the number of EC2 instances based on CPU metrics.

C. Write the messages to Amazon Kinesis Data Streams with a single shard. Use an AWS Lambda function to preprocess messages and store them in Amazon DynamoDB. Configure the consumer applications to read from DynamoDB to process the messages.

D. Publish the messages to an Amazon Simple Notification Service (Amazon SNS) topic with multiple Amazon Simple Queue Service (Amazon SOS) subscriptions. Configure the consumer applications to process the messages from the queues.

Correct Answer: *A*

Community vote distribution

D (75%) A (22%)

rein_chau Highly Voted 1 month, 2 weeks ago

Selected Answer: D

D makes more sense to me.

upvoted 13 times

Bevemo Highly Voted 3 weeks ago

D. SNS Fan Out Pattern https://docs.aws.amazon.com/sns/latest/dg/sns-common-scenarios.html (A is wrong Kinesis Analysis does not 'persist' by itself.)

upvoted 6 times

yuchip Most Recent 5 days, 2 hours ago

A, Amazon Kinesis makes it easy to collect, process, and analyze real-time, streaming data so you can get timely insights and react quickly to new information.

upvoted 1 times

Wpcorgan 1 week ago

D is correct

upvoted 1 times

ABCMail 1 week, 2 days ago

Selected Answer: D

SQS is a better solution for decoupling

upvoted 1 times

renekton 1 week, 4 days ago

Selected Answer: D

decoupling an application using sqs and fanout using sns

upvoted 1 times

mikey2000 1 week, 5 days ago

Why is A the correct answer?

upvoted 1 times

azmne 3 weeks, 3 days ago

Selected Answer: A

D because decoupling = SQS in most cases,

upvoted 1 times

k0s3k 3 weeks, 6 days ago

Guys, I am not for option D. Look at those SNS limits: https://docs.aws.amazon.com/general/latest/gr/sns.html
Even so, why would you publish messages to SNS when you can do it directly to SQS

upvoted 3 times

bont 4 weeks ago

D fan out to SQS

upvoted 2 times

keezbadger 1 month ago

D: By default, FIFO queues support up to 3,000 messages per second with batching or up to 300 messages per second (300 send, receive, or delete operations per second) without batching. So with batching it solved the question and option D also mentions “multiple”. So D is the right choice.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: D

I go for D because decoupling = SQS in most cases,
and they did mention SNS into multiple SQS so with multiple SQS you can handle more than 100k requests

upvoted 2 times

cark0728 1 month, 1 week ago

Selected Answer: D

D가 정답이다

upvoted 2 times

Drekorig 5 days, 3 hours ago

D'accord

upvoted 1 times

Six_Fingered_Jose 1 month ago

yeah i agree

upvoted 3 times

17Master 1 week, 6 days ago

x2 -- XD jajajaja

upvoted 1 times

iCcma 1 month, 1 week ago

Selected Answer: A

https://docs.aws.amazon.com/general/latest/gr/sns.html
https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/quotas-messages.html
30,000 messages per second

upvoted 1 times

iCcma 1 month, 1 week ago

friends certainly D makes a lot of sense, so even with the answer I gave, I'm curious, here I leave a link with a debate with very good arguments.
https://www.examtopics.com/discussions/amazon/view/43777-exam-aws-certified-solutions-architect-associate-saa-c02/
I want to change my answer to D

upvoted 3 times

17Master 1 month ago

Ans is D

upvoted 1 times

PhilN135 1 month, 1 week ago

Selected Answer: D

D - it is

upvoted 1 times

queen101 1 month, 1 week ago

DDDDDDD

upvoted 1 times

brushek 1 month, 1 week ago

Selected Answer: A

Kinesis Data Analytics integrates with Amazon Kinesis Data Streams and Amazon Kinesis Data Firehose so you can readily ingest streaming data. Just point Kinesis Data Analytics at the input stream and it will automatically read the data, parse it, and make it available for processing. You can emit processed results to other AWS services including Amazon S3, Amazon Redshift, and Amazon OpenSearch Service through Kinesis Data Firehose. You can also send output data to Amazon Kinesis Data Streams to build advanced stream processing pipelines.

upvoted 1 times

Question #8 *Topic 1*

A company is migrating a distributed application to AWS. The application serves variable workloads. The legacy platform consists of a primary server that coordinates jobs across multiple compute nodes. The company wants to modernize the application with a solution that maximizes resiliency and scalability.
How should a solutions architect design the architecture to meet these requirements?

A. Configure an Amazon Simple Queue Service (Amazon SQS) queue as a destination for the jobs. Implement the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group. Configure EC2 Auto Scaling to use scheduled scaling.

B. Configure an Amazon Simple Queue Service (Amazon SQS) queue as a destination for the jobs. Implement the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group. Configure EC2 Auto Scaling based on the size of the queue.

C. Implement the primary server and the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group. Configure AWS CloudTrail as a destination for the jobs. Configure EC2 Auto Scaling based on the load on the primary server.

D. Implement the primary server and the compute nodes with Amazon EC2 instances that are managed in an Auto Scaling group. Configure Amazon EventBridge (Amazon CloudWatch Events) as a destination for the jobs. Configure EC2 Auto Scaling based on the load on the compute nodes.

Correct Answer: *C*

Community vote distribution

B (92%) 8%

rein_chau Highly Voted 1 month, 2 weeks ago

Selected Answer: B

A - incorrect: Schedule scaling policy doesn't make sense.
C, D - incorrect: Primary server should not be in same Auto Scaling group with compute nodes.
B is correct.

upvoted 20 times

Wilson_S 2 weeks, 4 days ago

https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-using-sqs-queue.html

upvoted 2 times

Sinaneos Highly Voted 1 month, 3 weeks ago

Selected Answer: B

The answer seems to be B for me:
A: doesn't make sense to schedule auto-scaling
C: Not sure how CloudTrail would be helpful in this case, at all.
D: EventBridge is not really used for this purpose, wouldn't be very reliable

upvoted 5 times

Wpcorgan Most Recent 1 week ago

B is correct

upvoted 1 times

ABCMail 1 week, 2 days ago

Selected Answer: B

The scalability can depend on the size of the queue

upvoted 1 times

mikey2000 1 week, 5 days ago

I'm so confuse, how is cloudtrail going to be a destination for Jobs.

upvoted 1 times

cinostroza 1 week, 6 days ago

should be B, is the C for real?

upvoted 1 times

dduque10 3 weeks ago

Selected Answer: A

Always go with the most voted, C is the worst one

upvoted 3 times

emyhe 3 weeks, 1 day ago

shoule be B

upvoted 2 times

Wajif 3 weeks, 2 days ago

Selected Answer: B

Why C? Surely thats a mistake.

upvoted 1 times

lampard999999999 3 weeks, 2 days ago

Selected Answer: B

B is corect

upvoted 1 times

17Master 1 month ago

Selected Answer: B

ans is B

upvoted 1 times

alexman94 1 month ago

C - the use of CloudTrail is for AWS events. It doesn't make sense in receiving jobs.
B is my answer.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: B

I go with B here
CloudTrail as a job destination doesnt make sense, thats a logging/auditing platform

upvoted 1 times

bilel500 1 month, 1 week ago

Selected Answer: B

Answer is B.

upvoted 1 times

BoboChow 1 month, 2 weeks ago

Selected Answer: B

B seems to work.
I.m not sure about D

upvoted 1 times

hanhdroid 1 month, 2 weeks ago

Selected Answer: B

Should be B

upvoted 1 times

Question #9 *Topic 1*

A company is running an SMB file server in its data center. The file server stores large files that are accessed frequently for the first few days after the files are created. After 7 days the files are rarely accessed.
The total data size is increasing and is close to the company's total storage capacity. A solutions architect must increase the company's available storage space without losing low-latency access to the most recently accessed files. The solutions architect must also provide file lifecycle management to avoid future storage issues.
Which solution will meet these requirements?

A. Use AWS DataSync to copy data that is older than 7 days from the SMB file server to AWS.

B. Create an Amazon S3 File Gateway to extend the company's storage space. Create an S3 Lifecycle policy to transition the data to S3 Glacier Deep Archive after 7 days.

C. Create an Amazon FSx for Windows File Server file system to extend the company's storage space.

D. Install a utility on each user's computer to access Amazon S3. Create an S3 Lifecycle policy to transition the data to S3 Glacier Flexible Retrieval after 7 days.

Correct Answer: *D*

Community vote distribution

B (85%) D (15%)

Sinaneos Highly Voted 1 month, 3 weeks ago

Answer directly points towards file gateway with lifecycles, https://docs.aws.amazon.com/filegateway/latest/files3/CreatingAnSMBFileShare.html

D is wrong because utility function is vague and there is no need for flexible storage.

upvoted 14 times

Wpcorgan Most Recent 1 week ago

B is correct for me

upvoted 1 times

ABCMail 1 week, 2 days ago

Selected Answer: B

Using aws S3 gateway seems to be a logical option

upvoted 1 times

digitalsee 1 week, 4 days ago

Selected Answer: D

After 7 days the files are rarely accessed, so they are still accessed and still need access within milliseconds. So B can't be correct.
Installing S3 clients on Clients is a weird solution but the others simply don't fit the case.

upvoted 1 times

mikey2000 1 week, 5 days ago

Selected Answer: D

I think the answer is D because the question said without losing low-latency access.With glacier flexible retrieval, data can but access within 1-5 mins unlike glacier deep archival, it take 12 hours to access your data, 12 hours would defeat trying to get a low- latency access

upvoted 1 times

Mihai7 1 week, 5 days ago

Selected Answer: D

Because it still requires "low latency access" and some files older than 7 days are infrequently, BUT still accessed only Flexible retrieval can retrieve the file in minutes. Deep archive can't. The other answers(A and C) are not any good.

upvoted 1 times

Ptopics 1 week, 1 day ago

Low latency access is for files less than 7 days old. File GW provides low latency access by caching frequently accessed files locally so answer is B.

upvoted 1 times

leonnnn 2 weeks ago

Selected Answer: B

A S3 gateway provides smb for end users and stores file in S3 so that it can use lifecycle management to transit old files.

upvoted 1 times

akellaaws 2 weeks, 1 day ago

Selected Answer: B

Utility is not necessary

upvoted 1 times

Azeeza 2 weeks, 5 days ago

Selected Answer: B

File Gateway will be the best bet to replace the SMB storage server

upvoted 2 times

VijayMeh 2 weeks, 5 days ago

Selected Answer: B

B is the right answer

upvoted 1 times

gokalpkocer3 3 weeks, 2 days ago

Selected Answer: B

lifecycle therefore B.

upvoted 1 times

keezbadger 1 month ago

D is tricky. Though the questions says that "A solutions architect must increase the company's available storage space without losing low-latency access to the most recently accessed files." Which falls under the category of S3 Glacier flexible retrieval. Amazon S3 File Gateway provides a seamless way to connect to the cloud in order to store application data files and backup images as durable objects in Amazon S3 cloud storage. Amazon S3 File Gateway offers SMB or NFS-based access to data in Amazon S3 with local caching. Therefore :
B is the right answer.

upvoted 1 times

17Master 1 month ago

Selected Answer: B

Ans is correct

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: B

I go with B with this one,
volume gateway should support SMB, and with cache volume you can cache most recently accessed files locally for minimum latency

upvoted 1 times

Six_Fingered_Jose 1 month ago

oh wait this is about filegateway, ignore my comment but answer is still B

upvoted 1 times

patni11 1 month, 1 week ago

eeeeeeeeeeeeeeeeee

upvoted 1 times

BoboChow 1 month, 2 weeks ago

Selected Answer: B

Amazon S3 File Gateway supports either SMB or NFS protocol

upvoted 4 times

pooppants 1 month, 2 weeks ago

Selected Answer: B

The only answer which supports lifecycles

upvoted 1 times

Question #10 *Topic 1*

A company is building an ecommerce web application on AWS. The application sends information about new orders to an Amazon API Gateway REST API to process. The company wants to ensure that orders are processed in the order that they are received.
Which solution will meet these requirements?

A. Use an API Gateway integration to publish a message to an Amazon Simple Notification Service (Amazon SNS) topic when the application receives an order. Subscribe an AWS Lambda function to the topic to perform processing.

B. Use an API Gateway integration to send a message to an Amazon Simple Queue Service (Amazon SQS) FIFO queue when the application receives an order. Configure the SQS FIFO queue to invoke an AWS Lambda function for processing.

C. Use an API Gateway authorizer to block any requests while the application processes an order.

D. Use an API Gateway integration to send a message to an Amazon Simple Queue Service (Amazon SQS) standard queue when the application receives an order. Configure the SQS standard queue to invoke an AWS Lambda function for processing.

Correct Answer: *A*

Community vote distribution

B (100%)

Sinaneos Highly Voted 1 month, 3 weeks ago

Selected Answer: B

B because FIFO is made for that specific purpose

upvoted 19 times

rein_chau Highly Voted 1 month, 2 weeks ago

Selected Answer: B

Should be B because SQS FIFO queue guarantees message order.

upvoted 14 times

ogwu2000 Most Recent 2 days, 6 hours ago

A. Can API Gateway integration sends a message or publishes a message?

upvoted 1 times

Drekorig 4 days, 3 hours ago

Selected Answer: B

Answers A, B and D are possible, but to guarantee the order of the messages the answer is B

upvoted 1 times

leonnnn 2 weeks ago

Selected Answer: B

SQS FIFO queue guraantees message order.

upvoted 1 times

akellaaws 2 weeks, 1 day ago

Selected Answer: B

SQS FIFO can be an event source for Lambda

upvoted 1 times

raqsa 2 weeks, 3 days ago

Selected Answer: B

yes B is correct

upvoted 1 times

Buruguduystunstugudunstuy 2 weeks, 4 days ago

KEYWORDS: "processed in the order that they are received"
FIFO, First·In·First·Out
The answer is B. Amazon SQS FIFO
https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/FIFO-queues.html

upvoted 1 times

emyhe 3 weeks, 1 day ago

should be B, because require message order, FIFO have First-in-First-Out Delivery feature.

upvoted 1 times

Paritoshjani 1 month ago

SQS is not a push service. Hence the consumer has to poll. SQS FIFO queue cannot invoke Lambda; it is reverse. While SNS can publish record to Lambda. Hence the more appropriate answer is A

upvoted 2 times

Hasu_Uchiha 2 weeks, 4 days ago

AWS Lambda now supports Amazon SQS First-In-First-Out (FIFO) as an event source.

upvoted 1 times

KADSM 3 weeks, 3 days ago

SQS FIFO can invoke lambda. https://aws.amazon.com/blogs/compute/new-for-aws-lambda-sqs-fifo-as-an-event-source/

upvoted 1 times

keezbadger 1 month ago

The right answer is B

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: B

answer is B here
> ensure that orders are processed in the order that they are received.
only SQS with FIFO here meets that condition

upvoted 1 times

PhilN135 1 month, 1 week ago

B - FIFO

upvoted 1 times

patni11 1 month, 1 week ago

galat answer hai bc

upvoted 1 times

alvarez100 1 month, 2 weeks ago

B FIFO !

upvoted 2 times

arthurfl 1 month, 2 weeks ago

Selected Answer: B

FIFO works here, son the answer is B

upvoted 2 times

BoboChow 1 month, 2 weeks ago

Selected Answer: B

SQS FIFO queue works for the secario

upvoted 2 times

Question #11 *Topic 1*

A company has an application that runs on Amazon EC2 instances and uses an Amazon Aurora database. The EC2 instances connect to the database by using user names and passwords that are stored locally in a file. The company wants to minimize the operational overhead of credential management.
What should a solutions architect do to accomplish this goal?

A. Use AWS Secrets Manager. Turn on automatic rotation.

B. Use AWS Systems Manager Parameter Store. Turn on automatic rotation.

C. Create an Amazon S3 bucket to store objects that are encrypted with an AWS Key Management Service (AWS KMS) encryption key. Migrate the credential file to the S3 bucket. Point the application to the S3 bucket.

D. Create an encrypted Amazon Elastic Block Store (Amazon EBS) volume for each EC2 instance. Attach the new EBS volume to each EC2 instance. Migrate the credential file to the new EBS volume. Point the application to the new EBS volume.

Correct Answer: *B*

Community vote distribution

A (100%)

Sinaneos Highly Voted 1 month, 3 weeks ago

Selected Answer: A

B is wrong because parameter store does not support auto rotation, unless the customer writes it themselves, A is the answer.

upvoted 13 times

17Master 1 month ago

READ!!! AWS Secrets Manager is a secrets management service that helps you protect access to your applications, services, and IT resources. This service enables you to rotate, manage, and retrieve database credentials, API keys, and other secrets throughout their lifecycle.
https://aws.amazon.com/cn/blogs/security/how-to-connect-to-aws-secrets-manager-service-within-a-virtual-private-cloud/ y
https://aws.amazon.com/secrets-manager/?nc1=h_ls

upvoted 4 times

iCcma 1 month, 1 week ago

ty bro, I was confused about that and you just mentioned the "key" phrase, B doesn't support autorotation

upvoted 1 times

Wpcorgan Most Recent 1 week ago

A is correct

upvoted 1 times

ABCMail 1 week, 2 days ago

Selected Answer: A

AWS secrets manager is most logical

upvoted 1 times

Kartikey140 1 week, 6 days ago

A option is correct.

upvoted 1 times

al20220711 2 weeks, 2 days ago

The advantage of Systems Manager over Secrets Manager is costing. But it doesn't mention in the question. SO A is better.

upvoted 1 times

17Master 1 month ago

Selected Answer: A

READ!!! AWS Secrets Manager is a secrets management service that helps you protect access to your applications, services, and IT resources. This service enables you to rotate, manage, and retrieve database credentials, API keys, and other secrets throughout their lifecycle.
https://aws.amazon.com/cn/blogs/security/how-to-connect-to-aws-secrets-manager-service-within-a-virtual-private-cloud/ y
https://aws.amazon.com/secrets-manager/?nc1=h_ls

upvoted 3 times

patni11 1 month, 1 week ago

dharma

upvoted 1 times

KVK16 1 month, 2 weeks ago

Selected Answer: A

AWS Secrets Manager, a service that makes it easier to rotate, manage, and retrieve database credentials, API keys, and other secrets throughout their lifecycle. You can configure Secrets Manager to rotate secrets automatically, which can help you meet your security and compliance needs. Secrets Manager offers built-in integrations for MySQL, PostgreSQL, and Amazon Aurora on Amazon RDS, and can rotate credentials for these databases natively. You can control access to your secrets by using fine-grained AWS Identity and Access Management (IAM) policies. To retrieve secrets, employees replace plaintext secrets with a call to Secrets Manager APIs, eliminating the need to hard-code secrets in source code or update configuration files and redeploy code when secrets are rotated.

upvoted 3 times

BoboChow 1 month, 2 weeks ago

Selected Answer: A

AWS Secrets Manager is newer to SSM Parameter Store, and I'm not sure the later one can turn on automatic rotation

upvoted 2 times

josh_fan 1 month, 2 weeks ago

Selected Answer: A

https://aws.amazon.com/cn/blogs/security/how-to-connect-to-aws-secrets-manager-service-within-a-virtual-private-cloud/

upvoted 1 times

ogerber 1 month, 2 weeks ago

Selected Answer: A

no rotation option on parameter store.
never store credentials on buckets, its a hazard.
its A!!

upvoted 4 times

D2w 1 month, 2 weeks ago

Selected Answer: A

A, is meant for storing secrets e.g password ect

upvoted 1 times

tuloveu 1 month, 2 weeks ago

Selected Answer: A

A is correct.

upvoted 3 times

rein_chau 1 month, 2 weeks ago

Selected Answer: A

A is correct.
https://aws.amazon.com/blogs/security/rotate-amazon-rds-database-credentials-automatically-with-aws-secrets-manager/

upvoted 3 times

Question #12 *Topic 1*

A global company hosts its web application on Amazon EC2 instances behind an Application Load Balancer (ALB). The web application has static data and dynamic data. The company stores its static data in an Amazon S3 bucket. The company wants to improve performance and reduce latency for the static data and dynamic data. The company is using its own domain name registered with Amazon Route 53.
What should a solutions architect do to meet these requirements?

A. Create an Amazon CloudFront distribution that has the S3 bucket and the ALB as origins. Configure Route 53 to route traffic to the CloudFront distribution.

B. Create an Amazon CloudFront distribution that has the ALB as an origin. Create an AWS Global Accelerator standard accelerator that has the S3 bucket as an endpoint Configure Route 53 to route traffic to the CloudFront distribution.

C. Create an Amazon CloudFront distribution that has the S3 bucket as an origin. Create an AWS Global Accelerator standard accelerator that has the ALB and the CloudFront distribution as endpoints. Create a custom domain name that points to the accelerator DNS name. Use the custom domain name as an endpoint for the web application.

D. Create an Amazon CloudFront distribution that has the ALB as an origin. Create an AWS Global Accelerator standard accelerator that has the S3 bucket as an endpoint. Create two domain names. Point one domain name to the CloudFront DNS name for dynamic content. Point the other domain name to the accelerator DNS name for static content. Use the domain names as endpoints for the web application.

Correct Answer: *C*

Community vote distribution

A (83%) C (17%)

Kartikey140 Highly Voted 1 week, 6 days ago

Answer is A
Explanation - AWS Global Accelerator vs CloudFront
• They both use the AWS global network and its edge locations around the world
• Both services integrate with AWS Shield for DDoS protection.
• CloudFront
• Improves performance for both cacheable content (such as images and videos)
• Dynamic content (such as API acceleration and dynamic site delivery)
• Content is served at the edge
• Global Accelerator
• Improves performance for a wide range of applications over TCP or UDP
• Proxying packets at the edge to applications running in one or more AWS Regions.
• Good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP
• Good for HTTP use cases that require static IP addresses
• Good for HTTP use cases that required deterministic, fast regional failover
upvoted 6 times

hakant Most Recent 19 hours, 3 minutes ago

Answer A
CF can have multiple origins just like ALB. Need to set the routes manually using Berhavior tab in the console
https://aws.amazon.com/premiumsupport/knowledge-center/cloudfront-distribution-serve-content/#:~:text=%20Yes%2C%20you%20can%20configure%20a%20single%20CloudFront,serve%20different%20types%20of%20requests%20from%20multiple%20origins.
upvoted 1 times

ogwu2000 2 days, 6 hours ago

C. Not A. Why should CloudFront distribution have two origins - S3 bucket and the ALB ?
upvoted 1 times

Wpcorgan 1 week ago

A is correct
upvoted 1 times

mikey2000 1 week, 5 days ago

Selected Answer: A
cf support s3, alb or lambda function url as origin
upvoted 2 times

kanweng 2 weeks ago

Selected Answer: A
Q: How is AWS Global Accelerator different from Amazon CloudFront?

A: AWS Global Accelerator and Amazon CloudFront are separate services that use the AWS global network and its edge locations around the world.

CloudFront improves performance for both cacheable content (such as images and videos) and dynamic content (such as API acceleration and dynamic site delivery). Global Accelerator improves performance for a wide range of applications over TCP or UDP by proxying packets at the edge to applications running in one or more AWS Regions. Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover. Both services integrate with AWS Shield for DDoS protection.

upvoted 2 times

yd_h 2 weeks ago

A
Endpoints for standard accelerators in AWS Global Accelerator can be Network Load Balancers, Application Load Balancers, Amazon EC2 instances, or Elastic IP addresses.
https://docs.aws.amazon.com/global-accelerator/latest/dg/about-endpoints.html

upvoted 2 times

goatbernard 2 weeks, 2 days ago

Selected Answer: A

answer is A

upvoted 1 times

yuantongxue 2 weeks, 2 days ago

When you set up your accelerator with Global Accelerator, you associate the static IP addresses to regional endpoints in one or more AWS Regions. For standard accelerators, the endpoints are Network Load Balancers, Application Load Balancers, Amazon EC2 instances, or Elastic IP addresses. For custom routing accelerators, endpoints are virtual private cloud (VPC) subnets with one or more EC2 instances.

upvoted 1 times

17Master 1 month ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/DownloadDistS3AndCustomOrigins.html#concept_elb_origin

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: A

This is a tough question, initially thought answer was C but after some research found A to be the most viable answer here.

CloudFront can be used to reduce latency for dynamic content (the ALB in this case) by using edge locations to connect via AWS network instead of the internet, thus reducing latency.
So it might not necessary cache the content but it still reduces latency and improve performance with the lowest amount of work.

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/DownloadDistS3AndCustomOrigins.html#concept_elb_origin

upvoted 2 times

kogs 1 month, 1 week ago

A makes more sense and its possible.
https://aws.amazon.com/premiumsupport/knowledge-center/cloudfront-distribution-serve-content/

upvoted 2 times

KVK16 1 month, 2 weeks ago

Selected Answer: C

Static content can be cached at Cloud front Edge locations from S3 and dynamic content
EC2 behind the ALB whose performance can be improved by Global Accelerator whose one endpoint is ALB and other Cloud front. So with regards to custom domain name endpoint is web application is R53 alias records for the custom domain point to web application
https://aws.amazon.com/blogs/networking-and-content-delivery/improving-availability-and-performance-for-application-load-balancers-using-one-click-integration-with-aws-global-accelerator/

S3 cannot be endpoint for Global Accelerator - (interesting to think how S3 Global transfer acceleration functions)

upvoted 2 times

BoboChow 1 month, 2 weeks ago

Selected Answer: A

B,D is out because I don't think S3 can be an endpoint to AWS Accelerator.
As for C, I'm wondring if it works that dynamic and static resource use different domain name.
So I am going with A

upvoted 2 times

BoboChow 1 month, 2 weeks ago

Neither CloudFront nor S3 can be as endpoint for AWS Accelerator, so B is out, too.

upvoted 1 times

sba21 1 month, 2 weeks ago

Selected Answer: A

https://www.examtopics.com/discussions/amazon/view/22044-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 3 times

Ekie 1 month, 2 weeks ago

can't have CloudFront nor S3 as endpoints for AWS Accelerator... that eliminate B,C,D..
also CloudFront can improve performance and latency for both static and dynamic contents.. so I am going with A

upvoted 3 times

ArielSchivo 1 month, 1 week ago

Option A includes an S3 bucket, how can that work with dynamic content? I'm confused here.

upvoted 1 times

17Master 1 month ago

For more information about using an Application Load Balancer as your origin for CloudFront, https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/DownloadDistS3AndCustomOrigins.html#concept_elb_origin

upvoted 1 times

17Master 1 month ago

CloudFront --> BackEnd(S3+ALB) its possible.

upvoted 1 times

ogerber 1 month, 2 weeks ago

Selected Answer: C

can an ALB be an origin to CF? dont think so, only data sources.

upvoted 1 times

rodriiviru 1 month, 2 weeks ago

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/DownloadDistS3AndCustomOrigins.html#concept_elb_origin

upvoted 1 times

Ekie 1 month, 2 weeks ago

yes you can

upvoted 1 times

Question #13 *Topic 1*

A company performs monthly maintenance on its AWS infrastructure. During these maintenance activities, the company needs to rotate the credentials for its Amazon RDS for MySQL databases across multiple AWS Regions.
Which solution will meet these requirements with the LEAST operational overhead?

A. Store the credentials as secrets in AWS Secrets Manager. Use multi-Region secret replication for the required Regions. Configure Secrets Manager to rotate the secrets on a schedule.

B. Store the credentials as secrets in AWS Systems Manager by creating a secure string parameter. Use multi-Region secret replication for the required Regions. Configure Systems Manager to rotate the secrets on a schedule.

C. Store the credentials in an Amazon S3 bucket that has server-side encryption (SSE) enabled. Use Amazon EventBridge (Amazon CloudWatch Events) to invoke an AWS Lambda function to rotate the credentials.

D. Encrypt the credentials as secrets by using AWS Key Management Service (AWS KMS) multi-Region customer managed keys. Store the secrets in an Amazon DynamoDB global table. Use an AWS Lambda function to retrieve the secrets from DynamoDB. Use the RDS API to rotate the secrets.

Correct Answer: *A*

Community vote distribution

A (100%)

rein_chau Highly Voted 1 month, 2 weeks ago

Selected Answer: A

A is correct.
https://aws.amazon.com/blogs/security/how-to-replicate-secrets-aws-secrets-manager-multiple-regions/

upvoted 6 times

Wpcorgan Most Recent 1 week ago

A is correct

upvoted 1 times

Megako 1 week ago

Selected Answer: A

Me Pick A

upvoted 1 times

ABCMail 1 week, 1 day ago

Selected Answer: A

AWS secrets manager

upvoted 1 times

17Master 1 month ago

Selected Answer: A

Ans is A

upvoted 2 times

GameDad09 1 month, 1 week ago

Selected Answer: A

A is correct

upvoted 1 times

KVK16 1 month, 2 weeks ago

AWS Secrets Manager is meant for storing secrets like credentials to RDS database etc.
Capability to force rotation of secrets every X days. Multi Region Secret Keys
https://youtu.be/GPab-mc-8nU

upvoted 4 times

BoboChow 1 month, 2 weeks ago

Selected Answer: A

AWS Secrets Manageris Newer service, meant for storing secrets
Capability to force rotation of scerets every X days
Intergraion with Amazon RDS
Secrets are encypted using KMS

upvoted 1 times

D2w 1 month, 2 weeks ago

Selected Answer: A

AWS Secrets Manager meant for storing secrets, Capability to force rotation of secrets every X days, Automate generation of secrets on rotation (uses Lambda), Integration with Amazon RDS (MySQL, PostgreSQL, Aurora).

upvoted 2 times

Question #14 *Topic 1*

A company runs an ecommerce application on Amazon EC2 instances behind an Application Load Balancer. The instances run in an Amazon EC2 Auto Scaling group across multiple Availability Zones. The Auto Scaling group scales based on CPU utilization metrics. The ecommerce application stores the transaction data in a MySQL 8.0 database that is hosted on a large EC2 instance.
The database's performance degrades quickly as application load increases. The application handles more read requests than write transactions. The company wants a solution that will automatically scale the database to meet the demand of unpredictable read workloads while maintaining high availability.
Which solution will meet these requirements?

A. Use Amazon Redshift with a single node for leader and compute functionality.

B. Use Amazon RDS with a Single-AZ deployment Configure Amazon RDS to add reader instances in a different Availability Zone.

C. Use Amazon Aurora with a Multi-AZ deployment. Configure Aurora Auto Scaling with Aurora Replicas.

D. Use Amazon ElastiCache for Memcached with EC2 Spot Instances.

Correct Answer: *C*

Community vote distribution

C (100%)

D2w Highly Voted 1 month, 2 weeks ago

Selected Answer: C

C, AURORA is 5x performance improvement over MySQL on RDS and handles more read requests than write,; maintaining high availability = Multi-AZ deployment

upvoted 13 times

Wpcorgan Most Recent 1 week ago

C is correct

upvoted 1 times

ABCMail 1 week, 1 day ago

Selected Answer: C

Aurora offers multi AZ for HA

upvoted 1 times

17Master 1 month ago

Selected Answer: C

Ans is Aurora

upvoted 2 times

keezbadger 1 month ago

C is the answer. Aurora is fast, and for this case will support unpredictable workloads through its read replicas. Simple!

upvoted 2 times

ukwafabian 1 month ago

Selected Answer: C

"Read workloads" "Maintaining high availability" = Read replica's

upvoted 1 times

GameDad09 1 month, 1 week ago

Selected Answer: C

C is correct.

upvoted 1 times

KVK16 1 month, 2 weeks ago

Selected Answer: C

1. Migration from My SQL, Postgres SQL to Aurora is 5x and 3x times improves performance . Also provision for Read replicas

upvoted 2 times

Sinaneos 1 month, 2 weeks ago

Selected Answer: C

High availability + SQL -> C

upvoted 1 times

BoboChow 1 month, 2 weeks ago

Selected Answer: C

C is better than B about availibility

upvoted 1 times

Question #15 *Topic 1*

A company recently migrated to AWS and wants to implement a solution to protect the traffic that flows in and out of the production VPC. The company had an inspection server in its on-premises data center. The inspection server performed specific operations such as traffic flow inspection and traffic filtering. The company wants to have the same functionalities in the AWS Cloud.
Which solution will meet these requirements?

A. Use Amazon GuardDuty for traffic inspection and traffic filtering in the production VPC.

B. Use Traffic Mirroring to mirror traffic from the production VPC for traffic inspection and filtering.

C. Use AWS Network Firewall to create the required rules for traffic inspection and traffic filtering for the production VPC.

D. Use AWS Firewall Manager to create the required rules for traffic inspection and traffic filtering for the production VPC.

Correct Answer: *C*

Community vote distribution

C (100%)

BoboChow Highly Voted 1 month, 2 weeks ago

Selected Answer: C

I agree with C.
AWS Network Firewall is a stateful, managed network firewall and intrusion detection and prevention service for your virtual private cloud (VPC) that you created in Amazon Virtual Private Cloud (Amazon VPC). With Network Firewall, you can filter traffic at the perimeter of your VPC. This includes filtering traffic going to and coming from an internet gateway, NAT gateway, or over VPN or AWS Direct Connect.

upvoted 8 times

BoboChow 1 month, 2 weeks ago

And I'm not sure Traffic Mirroring can be for filtering

upvoted 3 times

rein_chau Highly Voted 1 month, 2 weeks ago

Selected Answer: C

C is correct. AWS Network Firewall supports both inspection and filtering as required.
B is incorrect because Traffic Mirroring only for inspection.

upvoted 6 times

wisoxe8356 Most Recent 2 days, 7 hours ago

Selected Answer: C

Not A - Amazon GuardDuty is a threat detection service that continuously monitors your AWS accounts and workloads for malicious activity and delivers detailed security findings for visibility and remediation. like someone strange continuely download data from your s3

Not B - As it is monironing not filtering

C - good to do both

D - configure and manage firewall rules, not monitoring

upvoted 1 times

Wpcorgan 1 week ago

C is correct

upvoted 1 times

sandra42 2 weeks ago

Network Firewall

upvoted 1 times

17Master 1 month ago

Selected Answer: C

Ans is C (AWS Firewall Manager).
https://aws.amazon.com/network-firewall/?whats-new-cards.sort-by=item.additionalFields.postDateTime&whats-new-cards.sort-order=desc.
Traffic Mirroring = "copy". https://docs.aws.amazon.com/es_es/vpc/latest/mirroring/traffic-mirroring-how-it-works.html

upvoted 3 times

ArielSchivo 3 weeks, 6 days ago

Answer is Network Firewall, not Firewall Manager.

upvoted 3 times

17Master 1 week, 5 days ago

ups, its Network Firewall

upvoted 1 times

GameDad09 1 month, 1 week ago

Selected Answer: C

C is correct

upvoted 3 times

D2w 1 month, 2 weeks ago

Selected Answer: C

C, protections for all of your Amazon VPCs

upvoted 2 times

Question #16 *Topic 1*

A company hosts a data lake on AWS. The data lake consists of data in Amazon S3 and Amazon RDS for PostgreSQL. The company needs a reporting solution that provides data visualization and includes all the data sources within the data lake. Only the company's management team should have full access to all the visualizations. The rest of the company should have only limited access.
Which solution will meet these requirements?

A. Create an analysis in Amazon QuickSight. Connect all the data sources and create new datasets. Publish dashboards to visualize the data. Share the dashboards with the appropriate IAM roles.

B. Create an analysis in Amazon QuickSight. Connect all the data sources and create new datasets. Publish dashboards to visualize the data. Share the dashboards with the appropriate users and groups.

C. Create an AWS Glue table and crawler for the data in Amazon S3. Create an AWS Glue extract, transform, and load (ETL) job to produce reports. Publish the reports to Amazon S3. Use S3 bucket policies to limit access to the reports.

D. Create an AWS Glue table and crawler for the data in Amazon S3. Use Amazon Athena Federated Query to access data within Amazon RDS for PostgreSQL. Generate reports by using Amazon Athena. Publish the reports to Amazon S3. Use S3 bucket policies to limit access to the reports.

Correct Answer: *D*

Community vote distribution

B (72%) A (23%) 5%

rodriiviru Highly Voted 1 month, 2 weeks ago

Selected Answer: B

https://docs.aws.amazon.com/quicksight/latest/user/sharing-a-dashboard.html

upvoted 22 times

mattlai 1 month, 2 weeks ago

https://docs.aws.amazon.com/quicksight/latest/user/share-a-dashboard-grant-access-users.html
^ more percise link

upvoted 5 times

BoboChow 1 month, 2 weeks ago

Agree with you

upvoted 1 times

D2w Highly Voted 1 month, 2 weeks ago

Selected Answer: A

A, The rest of the company should have only limited access you have to create IAM role

upvoted 6 times

17Master 1 month ago

Answer is B. Permissions are handled directly.
https://docs.aws.amazon.com/quicksight/latest/user/share-a-dashboard-grant-access-users.html

upvoted 1 times

JayyRock 3 weeks ago

“Permissions are handled directly” is a broad response that doesn’t say anything or make a point. So you’re saying quicksight will automatically know which person is on the management team and which person isn’t. No it won’t without instructions! So you need to set up IAM groups and limit their access that way. IAM (identity and “ACCESS” management) That’s the other part of the question that needs to be addressed.

upvoted 1 times

17Master 1 week, 5 days ago

ajá...read...https://docs.aws.amazon.com/quicksight/latest/user/share-a-dashboard-grant-access-users.html -
In the Share dashboard page that opens, do the following:
For Invite users and groups to dashboard at left, enter a user email or group name in the search box.
Any users or groups that match your query appear in a list below the search box. Only active users and groups appear in the list.
For the user or group that you want to grant access to the dashboard, choose Add. Then choose the level of permissions that you want them to have. ********"it says NO here go to the IAM and assign the permissions." So you don't manage by IAM. Ok, correct answer is B

upvoted 4 times

ashish_t Most Recent 5 days, 17 hours ago

Selected Answer: B

https://docs.aws.amazon.com/quicksight/latest/user/sharing-a-dashboard.html
Option B is correct

upvoted 1 times

Wpcorgan 1 week ago

B is correct

upvoted 1 times

grzeev 2 weeks, 1 day ago

Selected Answer: B

https://stackshare.io/stackups/amazon-athena-vs-amazon-quicksight

upvoted 1 times

jidexray 2 weeks, 3 days ago

Please i am still preparing for my exams ..Please is the correct answer the exact answer provided or go with the majority.. its confusing sometimes because the answers are similar attimes

upvoted 2 times

PS_R 2 weeks, 6 days ago

Selected Answer: B

IAM role refers to direction communication between AWS services and IAM User and Gropus for human, so i believe the answer is B

upvoted 1 times

Parveen_88 3 weeks, 6 days ago

Question: In the exam, what would be considered as correct answer? The one which community voted for or what is considered as "Correct answer here?

upvoted 1 times

mma7999 3 weeks, 2 days ago

most voted

upvoted 3 times

17Master 1 month ago

Selected Answer: B

Granting individual Amazon QuickSight users and groups access to a dashboard in Amazon QuickSight- https://docs.aws.amazon.com/quicksight/latest/user/share-a-dashboard-grant-access-users.html

upvoted 1 times

keezbadger 1 month ago

B is the right answer.

Elimination process:
QuickSight gives decision-makers the opportunity to explore and interpret information in an interactive visual environment. QuickSight uses ML to generate accurate responses to natural language questions about data. Get started with data visualization on AWS

AWS Glue is a serverless data integration service that makes data preparation simpler, faster, and cheaper.

So AWS glue and crawler is out leaving the two option A and B,

B is the right answer because you after you publish a dashboard, you can share it with other users or groups in your QuickSight account.

upvoted 2 times

Six_Fingered_Jose 1 month ago

Selected Answer: B

agree with B
https://docs.aws.amazon.com/quicksight/latest/user/share-a-dashboard-grant-access-users.html
looked around the document and doesnt state anywhere that you can share access using IAM roles

upvoted 1 times

Wazhija 1 month, 1 week ago

Selected Answer: B

You can create users and groups and share the dashboards/analysis you create in Athena

upvoted 1 times

sba21 1 month, 2 weeks ago

Selected Answer: D

Answer: D

upvoted 2 times

17Master 1 month ago

Granting individual Amazon QuickSight users and groups access to a dashboard in Amazon QuickSight - https://docs.aws.amazon.com/es_es/quicksight/latest/user/share-a-dashboard-grant-access-users.html

upvoted 1 times

Six_Fingered_Jose 1 month ago

athena in itself does not provide visualization, which the question is looking for here

upvoted 2 times

rein_chau 1 month, 2 weeks ago

Selected Answer: A

A is correct.
D is incorrect. Although Amazon Athena Federated Query can read from relational database, option D doesn't provide a way to visualize the data as required.

upvoted 3 times

Question #17 *Topic 1*

A company is implementing a new business application. The application runs on two Amazon EC2 instances and uses an Amazon S3 bucket for document storage. A solutions architect needs to ensure that the EC2 instances can access the S3 bucket.
What should the solutions architect do to meet this requirement?

A. Create an IAM role that grants access to the S3 bucket. Attach the role to the EC2 instances.

B. Create an IAM policy that grants access to the S3 bucket. Attach the policy to the EC2 instances.

C. Create an IAM group that grants access to the S3 bucket. Attach the group to the EC2 instances.

D. Create an IAM user that grants access to the S3 bucket. Attach the user account to the EC2 instances.

Correct Answer: *A*

Community vote distribution

A (97%)

sba21 Highly Voted 1 month, 2 weeks ago

Selected Answer: A

Always remember that you should associate IAM roles to EC2 instances

upvoted 15 times

DerekMinstP Most Recent 1 day, 10 hours ago

A for sure

upvoted 1 times

sherbo 3 days, 17 hours ago

Selected Answer: A

A is correct

upvoted 1 times

ashish_t 5 days, 17 hours ago

Selected Answer: A

A is correct

upvoted 1 times

Wpcorgan 1 week ago

A is correct

upvoted 1 times

ABCMail 1 week ago

Selected Answer: A

IAM role is a better solution

upvoted 1 times

libz 2 weeks, 2 days ago

Selected Answer: A

IAM roles

upvoted 3 times

Six_Fingered_Jose 1 month ago

Selected Answer: A

agree with A

upvoted 3 times

Jahangeer_17 1 month, 1 week ago

Create one IAM role, give S3 access to it. Attach to EC2

upvoted 3 times

Chunsli 1 month, 2 weeks ago

Policy is conceptual, role or account is more physical.

upvoted 3 times

Sinaneos 1 month, 2 weeks ago

Selected Answer: A

IAM roles are the best and easiest way for this

upvoted 3 times

mattlai 1 month, 2 weeks ago

isnt the correct answer create role and attach policy?

upvoted 2 times

mattlai 1 month, 2 weeks ago

nah create role and attach role

upvoted 4 times

galbimandu 1 month, 2 weeks ago

Selected Answer: A

Create an AWS Identity and Access Management (IAM) profile role that grants access to Amazon S3. Attach the IAM instance profile to the instance.

upvoted 1 times

galbimandu 1 month, 2 weeks ago

IAM role, Bucket policy.

upvoted 2 times

Ekie 1 month, 2 weeks ago

A.. https://aws.amazon.com/premiumsupport/knowledge-center/ec2-instance-access-s3-bucket/

upvoted 3 times

ogerber 1 month, 2 weeks ago

Selected Answer: B

B is correct. you attach a policy to an instance not a role.

upvoted 1 times

ArielSchivo 1 month, 1 week ago

How's so? You can go to the instance, then "attach/replace IAM role".

upvoted 2 times

Question #18 *Topic 1*

An application development team is designing a microservice that will convert large images to smaller, compressed images. When a user uploads an image through the web interface, the microservice should store the image in an Amazon S3 bucket, process and compress the image with an AWS Lambda function, and store the image in its compressed form in a different S3 bucket.
A solutions architect needs to design a solution that uses durable, stateless components to process the images automatically.
Which combination of actions will meet these requirements? (Choose two.)

A. Create an Amazon Simple Queue Service (Amazon SQS) queue. Configure the S3 bucket to send a notification to the SQS queue when an image is uploaded to the S3 bucket.

B. Configure the Lambda function to use the Amazon Simple Queue Service (Amazon SQS) queue as the invocation source. When the SQS message is successfully processed, delete the message in the queue.

C. Configure the Lambda function to monitor the S3 bucket for new uploads. When an uploaded image is detected, write the file name to a text file in memory and use the text file to keep track of the images that were processed.

D. Launch an Amazon EC2 instance to monitor an Amazon Simple Queue Service (Amazon SQS) queue. When items are added to the queue, log the file name in a text file on the EC2 instance and invoke the Lambda function.

E. Configure an Amazon EventBridge (Amazon CloudWatch Events) event to monitor the S3 bucket. When an image is uploaded, send an alert to an Amazon ample Notification Service (Amazon SNS) topic with the application owner's email address for further processing.

Correct Answer: *AB*

Community vote distribution

AB (95%) 5%

sba21 Highly Voted 1 month, 2 weeks ago

Selected Answer: AB

It looks like A-B

upvoted 8 times

sherbo Most Recent 3 days, 17 hours ago

Selected Answer: AB

A and B are most correct

upvoted 1 times

Wpcorgan 1 week ago

A and B

upvoted 1 times

crystally77 2 weeks ago

How about "E". Amazon EventBridge can monitor S3 bucket and send an alert to an SNS.

upvoted 2 times

TuLe 2 days, 5 hours ago

it required the owner's app process image which is not realistic in usage. It's like automation all process and manual the last steps using human effort.

upvoted 1 times

iis 3 weeks, 4 days ago

Selected Answer: AB

AB is OK. It can be done more straightforwardly. Just connect the S3 event to Lambda, and it is done. I don't think we need SQS or anything.

upvoted 4 times

GameDad09 1 month, 1 week ago

Selected Answer: AB

A+B seems to be correct

upvoted 3 times

ogerber 1 month, 2 weeks ago

Selected Answer: AB

oops i ment A + B also

upvoted 3 times

ogerber 1 month, 2 weeks ago

Selected Answer: AC

why not?

upvoted 1 times

123jhl0 1 month, 1 week ago

If you go for A-C... what's the purpose of SQS in A?

upvoted 5 times

ocbn3wby 1 week, 3 days ago

Storing in-memory doesn't seem reliable for a possible huge workload. While SQS can handle bigger workloads.

upvoted 1 times

GameDad09 1 month, 1 week ago

Lambda to monitor S3 put event is incorrect.

upvoted 2 times

Question #19 *Topic 1*

A company has a three-tier web application that is deployed on AWS. The web servers are deployed in a public subnet in a VPC. The application servers and database servers are deployed in private subnets in the same VPC. The company has deployed a third-party virtual firewall appliance from AWS Marketplace in an inspection VPC. The appliance is configured with an IP interface that can accept IP packets.
A solutions architect needs to integrate the web application with the appliance to inspect all traffic to the application before the traffic reaches the web server.
Which solution will meet these requirements with the LEAST operational overhead?

A. Create a Network Load Balancer in the public subnet of the application's VPC to route the traffic to the appliance for packet inspection.

B. Create an Application Load Balancer in the public subnet of the application's VPC to route the traffic to the appliance for packet inspection.

C. Deploy a transit gateway in the inspection VPConfigure route tables to route the incoming packets through the transit gateway.

D. Deploy a Gateway Load Balancer in the inspection VPC. Create a Gateway Load Balancer endpoint to receive the incoming packets and forward the packets to the appliance.

Correct Answer: *B*

Community vote distribution

D (86%) 7%

CloudGuru99 Highly Voted 1 month, 2 weeks ago

Answer is D . Use Gateway Load balancer
REF: https://aws.amazon.com/blogs/networking-and-content-delivery/scaling-network-traffic-inspection-using-aws-gateway-load-balancer/

upvoted 10 times

rein_chau Highly Voted 1 month, 2 weeks ago

Selected Answer: D

D is correct because this is Gateway Load Balancer feature.

upvoted 5 times

sherbo Most Recent 3 days, 17 hours ago

Selected Answer: D

D is correct

upvoted 1 times

JohnnyBG 4 days, 4 hours ago

Selected Answer: A

Would Chose A since it can simply forward all packet at layer 4. Not D since GWLB is more costly and do not bring benefit compared to a simpler NLB.

upvoted 1 times

JohnnyBG 4 days, 3 hours ago

Sorry for my confusion, after more careful reading only GWLB will forward all IP packet where NLB is limited to TCP-UDP ...

upvoted 3 times

Wpcorgan 1 week ago

D is correct

upvoted 1 times

grzeev 2 weeks, 1 day ago

Selected Answer: D

https://piyushj02.medium.com/aws-gateway-load-balancer-1o1-26d7cf61b798

upvoted 1 times

Jtic 2 weeks, 6 days ago

Selected Answer: D

My Answer is D.
The use case falls under Gateway Load balancer capability.
https://aws.amazon.com/elasticloadbalancing/gateway-load-balancer/

upvoted 2 times

Newptone 3 weeks ago

Selected Answer: D

https://docs.aws.amazon.com/elasticloadbalancing/latest/gateway/getting-started.html
For option B, you need to consider more about how to handle routing traffic which is not LEAST effort.

upvoted 1 times

KADSM 3 weeks ago

Have doubt - whether D could be correct as it suggests created Gateway load balancer in Inspection VPC

upvoted 1 times

KADSM 3 weeks ago

Shouldn't be GWLB in Application VPC.

upvoted 1 times

pm2229 3 weeks, 1 day ago

It's D, Coz.. Gateway Load Balancer is a new type of load balancer that operates at layer 3 of the OSI model and is built on Hyperplane, which is capable of handling several thousands of connections per second. Gateway Load Balancer endpoints are configured in spoke VPCs originating or receiving traffic from the Internet. This architecture allows you to perform inline inspection of traffic from multiple spoke VPCs in a simplified and scalable fashion while still centralizing your virtual appliances.

upvoted 3 times

keezbadger 1 month ago

D: The question is quite tricky, but you have to be aware of the OSI Layers. The key emphasis here is "The appliance is configured with an IP interface that can accept IP packets." Hence, Option D with gateway mentions this. Thus Path Determination and IP (Network) is associated with IP Interface and IP packet

upvoted 1 times

ukwafabian 1 month ago

Selected Answer: B

"all traffic to the application before the traffic reaches the web server" Application is Layer 7 so it is B for me

upvoted 1 times

ManoAni 1 month ago

B is correct answer, they are saying with LEAST effort. If you deploy a Gateway load balancer infront of appliance, assuming only 1 instance running at the moment, you need to add secondary instance and its lot of work to re-configure the appliance to handle traffic. Since the web is already exposed to internet, deploying application load balancer at L7 is easier, from there you should be able to route the traffic for inspection before it reaches web server.

upvoted 2 times

Newptone 3 weeks ago

But if you choose B for forwarding the traffic, it then raises two questions:
1. How do you route the traffic from security appliance to your web servers?
2. How do you route the traffic from your web servers to your security appliance?
Actually it introduces much more efforts to do what GWLB has done by itself.

upvoted 1 times

GameDad09 1 month, 1 week ago

Selected Answer: D

D seems to be correct

upvoted 2 times

herculian_effort 1 month, 2 weeks ago

"The appliance is configured with an IP interface that can accept IP packets" which means Layer 3. Application Load Balancer is Layer 7. D is correct since it operates on Layer 3 as mentioned in the link provided by CloudGuru99

upvoted 3 times

tubtab 1 month, 2 weeks ago

D is correct

upvoted 2 times

Question #20 *Topic 1*

A company wants to improve its ability to clone large amounts of production data into a test environment in the same AWS Region. The data is stored in Amazon EC2 instances on Amazon Elastic Block Store (Amazon EBS) volumes. Modifications to the cloned data must not affect the production environment. The software that accesses this data requires consistently high I/O performance.
A solutions architect needs to minimize the time that is required to clone the production data into the test environment.
Which solution will meet these requirements?

A. Take EBS snapshots of the production EBS volumes. Restore the snapshots onto EC2 instance store volumes in the test environment.

B. Configure the production EBS volumes to use the EBS Multi-Attach feature. Take EBS snapshots of the production EBS volumes. Attach the production EBS volumes to the EC2 instances in the test environment.

C. Take EBS snapshots of the production EBS volumes. Create and initialize new EBS volumes. Attach the new EBS volumes to EC2 instances in the test environment before restoring the volumes from the production EBS snapshots.

D. Take EBS snapshots of the production EBS volumes. Turn on the EBS fast snapshot restore feature on the EBS snapshots. Restore the snapshots into new EBS volumes. Attach the new EBS volumes to EC2 instances in the test environment.

Correct Answer: *D*

Community vote distribution

D (100%)

BoboChow Highly Voted 1 month, 2 weeks ago

Selected Answer: D

https://aws.amazon.com/cn/about-aws/whats-new/2020/11/amazon-ebs-fast-snapshot-restore-now-available-us-govcloud-regions/

upvoted 5 times

Wpcorgan Most Recent 1 week ago

D is correct

upvoted 1 times

Wajif 2 weeks, 2 days ago

Selected Answer: D

Minimize the time is a key requirement. So D.

upvoted 2 times

UWSFish 1 month ago

Selected Answer: D

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-fast-snapshot-restore.html

Amazon EBS fast snapshot restore (FSR) enables you to create a volume from a snapshot that is fully initialized at creation. This eliminates the latency of I/O operations on a block when it is accessed for the first time. Volumes that are created using fast snapshot restore instantly deliver all of their provisioned performance.

upvoted 4 times

GameDad09 1 month, 1 week ago

Selected Answer: D

D seems to be the correct one.

upvoted 1 times

ninjawrz 1 month, 1 week ago

Selected Answer: D

D: Fast snapshots

Amazon EBS fast snapshot restore (FSR) enables you to create a volume from a snapshot that is fully initialized at creation. This eliminates the latency of I/O operations on a block when it is accessed for the first time. Volumes that are created using fast snapshot restore instantly deliver all of their provisioned performance.

upvoted 2 times

queen101 1 month, 1 week ago

DDDDDDDDDDDD

upvoted 1 times

123jhl0 1 month, 1 week ago

Selected Answer: D

The "fast snapshot restore" on EBS eliminates the initialisation time required for the EBS volumes, providing both requests: 1) reduces time to clone data from production to test and 2) provide consistently high performance (as volumes are already initialised)
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ebs-fast-snapshot-restore.html

upvoted 2 times

SAA003 1 month, 2 weeks ago

C is the answer
https://docs.aws.amazon.com/prescriptive-guidance/latest/backup-recovery/restore.html

upvoted 1 times

Question #21 *Topic 1*

An ecommerce company wants to launch a one-deal-a-day website on AWS. Each day will feature exactly one product on sale for a period of 24 hours. The company wants to be able to handle millions of requests each hour with millisecond latency during peak hours.
Which solution will meet these requirements with the LEAST operational overhead?

A. Use Amazon S3 to host the full website in different S3 buckets. Add Amazon CloudFront distributions. Set the S3 buckets as origins for the distributions. Store the order data in Amazon S3.

B. Deploy the full website on Amazon EC2 instances that run in Auto Scaling groups across multiple Availability Zones. Add an Application Load Balancer (ALB) to distribute the website traffic. Add another ALB for the backend APIs. Store the data in Amazon RDS for MySQL.

C. Migrate the full application to run in containers. Host the containers on Amazon Elastic Kubernetes Service (Amazon EKS). Use the Kubernetes Cluster Autoscaler to increase and decrease the number of pods to process bursts in traffic. Store the data in Amazon RDS for MySQL.

D. Use an Amazon S3 bucket to host the website's static content. Deploy an Amazon CloudFront distribution. Set the S3 bucket as the origin. Use Amazon API Gateway and AWS Lambda functions for the backend APIs. Store the data in Amazon DynamoDB.

Correct Answer: *D*

Community vote distribution

D (100%)

Sinaneos Highly Voted 1 month, 2 weeks ago

Selected Answer: D

D because all of the components are infinitely scalable
dynamoDB, API Gateway, Lambda, and of course s3+cloudfront

upvoted 10 times

Wpcorgan Most Recent 1 week ago

D is correct

upvoted 1 times

ABCMail 1 week ago

Selected Answer: D

Serverless technologies are better options

upvoted 1 times

Wajif 2 weeks, 2 days ago

Why not B? Application load balancer can accept millions of request/hr?

upvoted 1 times

keithkifo 2 weeks ago

For me, the keyword was millisecond latency. Option B suggests RDS as the database, but Option D is DynamoDB.

DynamoDB - Fast, flexible NoSQL database service for single-digit millisecond performance at any scale

upvoted 2 times

TuLe 2 days, 4 hours ago

Yes, and also LEAST operational overhead. Scaling the application on EC2 instance is hard work require the very good architect.

upvoted 1 times

sodyam 2 weeks, 3 days ago

D is the correct answer due to milliseconds latency which will involve cloud front.

upvoted 2 times

xeun88 2 weeks, 5 days ago

D is the correct answer due to milliseconds latency which will involve cloud front.

upvoted 1 times

17Master 4 weeks ago

Selected Answer: D

Ans is correct D

upvoted 1 times

GameDad09 1 month, 1 week ago

Selected Answer: D

D is the correct one.
upvoted 1 times

queen101 1 month, 1 week ago
DDDDDDDDDDDDDDD
upvoted 1 times

ninjawrz 1 month, 2 weeks ago
D: because of least operational overhead
upvoted 1 times

BoboChow 1 month, 2 weeks ago
Selected Answer: D
I feel like the scenario is not only static resource but also dynamic resources.
API Gateway + Lambda has a good scalibility
upvoted 4 times

Lilibell 1 month, 2 weeks ago
the answer is D
upvoted 2 times

Question #22 *Topic 1*

A solutions architect is using Amazon S3 to design the storage architecture of a new digital media application. The media files must be resilient to the loss of an Availability Zone. Some files are accessed frequently while other files are rarely accessed in an unpredictable pattern. The solutions architect must minimize the costs of storing and retrieving the media files.
Which storage option meets these requirements?

A. S3 Standard

B. S3 Intelligent-Tiering

C. S3 Standard-Infrequent Access (S3 Standard-IA)

D. S3 One Zone-Infrequent Access (S3 One Zone-IA)

Correct Answer: *B*

Community vote distribution

B (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: B

"unpredictable pattern" - always go for Intelligent Tiering of S3
It also meets the resiliency requirement: "S3 Standard, S3 Intelligent-Tiering, S3 Standard-IA, S3 Glacier Instant Retrieval, S3 Glacier Flexible Retrieval, and S3 Glacier Deep Archive redundantly store objects on multiple devices across a minimum of three Availability Zones in an AWS Region" https://docs.aws.amazon.com/AmazonS3/latest/userguide/DataDurability.html

upvoted 9 times

Wpcorgan Most Recent 1 week ago

B is correct

upvoted 1 times

AbhiJo 2 weeks, 3 days ago

B is correct, C is incorrect because of requirement for frequent access as well

upvoted 2 times

xeun88 2 weeks, 5 days ago

Since it said some data a access frequently and some are unpredictable, i will go for B.

upvoted 1 times

17Master 4 weeks ago

Selected Answer: B

ans is correct B

upvoted 2 times

GameDad09 1 month, 1 week ago

Selected Answer: B

B is the correct one.

upvoted 2 times

queen101 1 month, 1 week ago

BBBBBBBBBB

upvoted 1 times

ninjawrz 1 month, 2 weeks ago

B. S3 Intelligent-Tiering for unpredictable or vague usecase

upvoted 1 times

BoboChow 1 month, 2 weeks ago

Selected Answer: B

sure for B

upvoted 1 times

galbimandu 1 month, 2 weeks ago

Selected Answer: B

B
https://aws.amazon.com/getting-started/hands-on/getting-started-using-amazon-s3-intelligent-tiering/?nc1=h_ls

upvoted 3 times

tt79 1 month, 2 weeks ago

B is correct.

upvoted 1 times

Question #23 *Topic 1*

A company is storing backup files by using Amazon S3 Standard storage. The files are accessed frequently for 1 month. However, the files are not accessed after 1 month. The company must keep the files indefinitely.
Which storage solution will meet these requirements MOST cost-effectively?

A. Configure S3 Intelligent-Tiering to automatically migrate objects.

B. Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 Glacier Deep Archive after 1 month.

C. Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) after 1 month.

D. Create an S3 Lifecycle configuration to transition objects from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 1 month.

Correct Answer: *B*

Community vote distribution

B (100%)

Wpcorgan 1 week ago

B is correct

upvoted 1 times

Tsho 1 week ago

BBBBBBBBB

upvoted 1 times

renekton 1 week, 3 days ago

Selected Answer: B

B is the correct answer

upvoted 1 times

xeun88 2 weeks, 5 days ago

B is correct

upvoted 2 times

GameDad09 1 month, 1 week ago

Selected Answer: B

B is the correct one.

upvoted 2 times

queen101 1 month, 1 week ago

BBBBBBBBB

upvoted 2 times

ninjawrz 1 month, 2 weeks ago

B: Transition to Glacier deep archive for cost efficiency

upvoted 4 times

BoboChow 1 month, 2 weeks ago

Selected Answer: B

sure for B

upvoted 2 times

Ralston40 1 month, 2 weeks ago

The answer is B

upvoted 4 times

Question #24 *Topic 1*

A company observes an increase in Amazon EC2 costs in its most recent bill. The billing team notices unwanted vertical scaling of instance types for a couple of EC2 instances. A solutions architect needs to create a graph comparing the last 2 months of EC2 costs and perform an in-depth analysis to identify the root cause of the vertical scaling.
How should the solutions architect generate the information with the LEAST operational overhead?

A. Use AWS Budgets to create a budget report and compare EC2 costs based on instance types.

B. Use Cost Explorer's granular filtering feature to perform an in-depth analysis of EC2 costs based on instance types.

C. Use graphs from the AWS Billing and Cost Management dashboard to compare EC2 costs based on instance types for the last 2 months.

D. Use AWS Cost and Usage Reports to create a report and send it to an Amazon S3 bucket. Use Amazon QuickSight with Amazon S3 as a source to generate an interactive graph based on instance types.

Correct Answer: *C*

Community vote distribution

B (73%) C (18%) 9%

sba21 Highly Voted 1 month, 2 weeks ago

Selected Answer: B

https://www.examtopics.com/discussions/amazon/view/68306-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 10 times

Wpcorgan Most Recent 1 week ago

B is correct

upvoted 1 times

koreanmonkey 1 week, 4 days ago

Selected Answer: B

Because Cost Explorer also provide recommandation about after 12month, I think B is better than C

upvoted 1 times

pm2229 3 weeks, 1 day ago

B: AWS Cost Explorer is a tool that enables you to view and analyze your costs and usage. You can explore your usage and costs using the main graph, the Cost Explorer cost and usage reports, or the Cost Explorer RI reports. You can view data for up to the last 12 months.

Ref: https://docs.aws.amazon.com/cost-management/latest/userguide/ce-what-is.html

upvoted 2 times

keezbadger 4 weeks ago

C is the right answer. B talks about cost explorer granularity filtering and there is a cost assosicated with the usage of indepth analysis. To get indepth analysis you will have to set it to 1hour granularity and this has a cost, as against the questions that says "least operational overhead". For more information read through this: https://docs.aws.amazon.com/wellarchitected/latest/cost-optimization-pillar/cost_monitor_usage_detailed_source.html

upvoted 3 times

dduque10 3 weeks ago

I think you are confusing operational overhead with cost effectiveness

upvoted 4 times

Six_Fingered_Jose 1 month ago

This question is honestly too vague,
C does not state EXACTLY where from the dashboard they are pulling the graph from, and as for B, it is true that for granular filtering the function only allows for up to 14 days, but cost explorer itself should be able to easily view 2 months.

question states LEAST amount of overhead which means D should be out of the question since that requires a lot of work.

think it's either B or C but can't say for sure, why do AWS keeps coming up with these weird questions man

upvoted 4 times

123jhl0 1 month, 1 week ago

Selected Answer: C

The requested result is a graph, so...
A - can't be as the result is a report
B - can't be as it is limited to 14 days visibility and the graph has to cover 2 months

C - seems to provide graphs and the best option available, as...
D - could provide graphs, BUT involves operational overhead, which has been requested to be minimised.

upvoted 4 times

goku58 1 month, 1 week ago

12 months data visible on Cost Explorer.

upvoted 3 times

123jhl0 1 month, 1 week ago

B can't be the answer, as the functionality only works for the last 14 days and the task in the questions requires "create a graph comparing the last 2 months...". Check this restriction at https://aws.amazon.com/aws-cost-management/pricing/

upvoted 4 times

Valero_ 1 month, 2 weeks ago

Selected Answer: B

I think it's Cost Explorer too. https://aws.amazon.com/about-aws/whats-new/2019/11/aws-cost-explorer-supports-hourly-resource-level-granularity/

upvoted 3 times

ninjawrz 1 month, 2 weeks ago

B. AWS Cost Explorer is a tool that enables you to view and analyze your costs and usage. You can explore your usage and costs using the main graph, the Cost Explorer cost and usage reports, or the Cost Explorer RI reports. You can view data for up to the last 12 months, forecast how much you're likely to spend for the next 12 months, and get recommendations for what Reserved Instances to purchase. You can use Cost Explorer to identify areas that need further inquiry and see trends that you can use to understand your costs.
https://docs.aws.amazon.com/cost-management/latest/userguide/ce-what-is.html

upvoted 2 times

tubtab 1 month, 2 weeks ago

Selected Answer: B

i think b

upvoted 2 times

galbimandu 1 month, 2 weeks ago

Selected Answer: D

https://docs.aws.amazon.com/cur/latest/userguide/what-is-cur.html

upvoted 1 times

17Master 1 week, 5 days ago

it doesn't say anything about EC2.
Your report files consist of a .csv file or a collection of .csv files and a manifest file. You can choose to configure your report data for integration with Amazon Athena, Amazon Redshift, or Amazon QuickSight.

upvoted 1 times

ogerber 1 month, 2 weeks ago

Selected Answer: D

i think its D.

upvoted 1 times

Question #25 *Topic 1*

A company is designing an application. The application uses an AWS Lambda function to receive information through Amazon API Gateway and to store the information in an Amazon Aurora PostgreSQL database.
During the proof-of-concept stage, the company has to increase the Lambda quotas significantly to handle the high volumes of data that the company needs to load into the database. A solutions architect must recommend a new design to improve scalability and minimize the configuration effort.
Which solution will meet these requirements?

A. Refactor the Lambda function code to Apache Tomcat code that runs on Amazon EC2 instances. Connect the database by using native Java Database Connectivity (JDBC) drivers.

B. Change the platform from Aurora to Amazon DynamoDProvision a DynamoDB Accelerator (DAX) cluster. Use the DAX client SDK to point the existing DynamoDB API calls at the DAX cluster.

C. Set up two Lambda functions. Configure one function to receive the information. Configure the other function to load the information into the database. Integrate the Lambda functions by using Amazon Simple Notification Service (Amazon SNS).

D. Set up two Lambda functions. Configure one function to receive the information. Configure the other function to load the information into the database. Integrate the Lambda functions by using an Amazon Simple Queue Service (Amazon SQS) queue.

Correct Answer: *D*

Community vote distribution

D (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: D

A - refactoring can be a solution, BUT requires a LOT of effort - not the answer
B - DynamoDB is NoSQL and Aurora is SQL, so it requires a DB migration... again a LOT of effort, so no the answer
C and D are similar in structure, but...
C uses SNS, which would notify the 2nd Lambda function... provoking the same bottleneck... not the solution
D uses SQS, so the 2nd lambda function can go to the queue when responsive to keep with the DB load process.
Usually the app decoupling helps with the performance improvement by distributing load. In this case, the bottleneck is solved by uses queues... so D is the answer.

upvoted 16 times

Wpcorgan Most Recent 1 week ago

D is correct

upvoted 1 times

ABCMail 1 week ago

Selected Answer: D

Two single responsibility functions offer a better solution.

upvoted 1 times

akosigengen 1 week, 2 days ago

D. Keyword is to handle load which will be taking care of by SQS.

upvoted 2 times

Ajai23 1 month, 2 weeks ago

Selected Answer: D

Process of elimination, D

upvoted 2 times

BoboChow 1 month, 2 weeks ago

Selected Answer: D

Atually I'm really confused by those options.
A is not right obiously, but the remaining options don't make sense, either...

upvoted 1 times

Lilibell 1 month, 2 weeks ago

the answer is D

upvoted 2 times

Question #26 *Topic 1*

A company needs to review its AWS Cloud deployment to ensure that its Amazon S3 buckets do not have unauthorized configuration changes. What should a solutions architect do to accomplish this goal?

A. Turn on AWS Config with the appropriate rules.

B. Turn on AWS Trusted Advisor with the appropriate checks.

C. Turn on Amazon Inspector with the appropriate assessment template.

D. Turn on Amazon S3 server access logging. Configure Amazon EventBridge (Amazon Cloud Watch Events).

Correct Answer: *A*

Community vote distribution

A (87%) 13%

gokalpkocer3 Highly Voted 3 weeks, 2 days ago

Configuration changes= AWS Config

upvoted 5 times

Wpcorgan Most Recent 1 week ago

A is correct

upvoted 1 times

ABCMail 1 week ago

Selected Answer: A

AWS config allows scrutiny of past chnages

upvoted 1 times

grzeev 2 weeks, 1 day ago

Selected Answer: A

AWS Config is a fully managed service that provides you with an AWS resource inventory, configuration history, and configuration change notifications to enable security and governance

upvoted 2 times

pspinelli19 3 weeks ago

Selected Answer: A

With Config you can limit changes to your entire account/s.
https://www.examtopics.com/discussions/amazon/view/27941-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Solarch 3 weeks, 5 days ago

Answer is A. Trusted Advisor gives you a general check of your system and identifies ways to optimize your infrastructue and improve it. While AWS config is more about specific resource. Like stated (S3 bucket). Config lets you select particular resource you want to evaluate .

upvoted 1 times

keezbadger 4 weeks ago

A is the right answer. The key word in the question is "Review" Hence. AWS config use case here, "Evaluate resource configurations for potential vulnerabilities, and review your configuration history after potential incidents to examine your security posture."

Though Trusted advisor is similar but what it does is that, it provides important "recommendations" to optimize your cloud deployments, improve resilience, and address security gaps.

The keyword for Trusted advisor is Recommendation.

upvoted 1 times

123jhl0 1 month, 1 week ago

Selected Answer: A

A - according to the picture in the documentation... "AWS Config automatically evaluates the recorded configuration against the configuration that you specify."
https://d1.awsstatic.com/config-diagram-092122.974fe2a4cb6aae1fe564fdbbe30ab55841a9858e.png

upvoted 1 times

KVK16 1 month, 2 weeks ago

Selected Answer: A

Config - With AWS Config, you can dive deep into how your bucket was configured at any point in time. Additionally, Config rules enable you to check whether your S3 buckets have logging and versioning enabled
https://aws.amazon.com/about-aws/whats-new/2016/10/record-and-govern-s3-bucket-configurations-with-aws-config/

S3 only permissions check is done by Trust advisor - apart from other checks root MFA, SG open ports, RDS Public Snapshots , EBS Public Snapshots , IAM User - one min, Service limits

upvoted 1 times

Sinaneos 1 month, 2 weeks ago

Selected Answer: A

Config is better to PREVENT changes, Trusted advisor would review breaches that have already happened

upvoted 2 times

masetromain 1 month, 2 weeks ago

Selected Answer: A

Answere A : https://youtu.be/MJDuAvNEv64

upvoted 3 times

BoboChow 1 month, 2 weeks ago

Selected Answer: A

AWS Trusted Advisor is just to analyze your AWS accounts and providers reconmendation.

upvoted 2 times

BoboChow 1 month, 2 weeks ago

Inspector is to analyze the running OS against known vulnerabilities

upvoted 1 times

tuloveu 1 month, 2 weeks ago

Selected Answer: B

Trusted Advisor good option.

upvoted 2 times

17Master 4 weeks ago

why? - D is correct

upvoted 1 times

Question #27 *Topic 1*

A company is launching a new application and will display application metrics on an Amazon CloudWatch dashboard. The company's product manager needs to access this dashboard periodically. The product manager does not have an AWS account. A solutions architect must provide access to the product manager by following the principle of least privilege.
Which solution will meet these requirements?

A. Share the dashboard from the CloudWatch console. Enter the product manager's email address, and complete the sharing steps. Provide a shareable link for the dashboard to the product manager.

B. Create an IAM user specifically for the product manager. Attach the CloudWatchReadOnlyAccess AWS managed policy to the user. Share the new login credentials with the product manager. Share the browser URL of the correct dashboard with the product manager.

C. Create an IAM user for the company's employees. Attach the ViewOnlyAccess AWS managed policy to the IAM user. Share the new login credentials with the product manager. Ask the product manager to navigate to the CloudWatch console and locate the dashboard by name in the Dashboards section.

D. Deploy a bastion server in a public subnet. When the product manager requires access to the dashboard, start the server and share the RDP credentials. On the bastion server, ensure that the browser is configured to open the dashboard URL with cached AWS credentials that have appropriate permissions to view the dashboard.

Correct Answer: *B*

Community vote distribution

A (89%) 11%

masetromain Highly Voted 1 month, 2 weeks ago

Selected Answer: A

Answere A : https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/cloudwatch-dashboard-sharing.html

Share a single dashboard and designate specific email addresses of the people who can view the dashboard. Each of these users creates their own password that they must enter to view the dashboard.

upvoted 16 times

123jhl0 1 month, 1 week ago

Thanks for the link! No doubt A is the answer.

upvoted 2 times

Wpcorgan Most Recent 1 week ago

A is correct

upvoted 1 times

randall56 1 week ago

I think that the point in answer B is "Share the new login credentials with the product manager".That would imply that the PM, who didn't have a login before, would need a new one to access the dashboard. The PM receives an AWS account in this manner.

upvoted 1 times

htangga 1 week, 1 day ago

Selected Answer: A

the answer is A: product manager doesn't have AWS account

upvoted 1 times

pm2229 2 weeks, 3 days ago

A, Sharing a dashboard publicly makes it accessible to anyone who has the link, with no authentication. Do this only for dashboards that do not contain sensitive information.

upvoted 2 times

Pamban 2 weeks, 4 days ago

Selected Answer: A

"The product manager does not have an AWS account" so answer is A.. pretty straightforward

upvoted 2 times

Kartikey140 1 week, 5 days ago

Point.

upvoted 1 times

backbencher2022 3 weeks ago

Selected Answer: A

On second thoughts, A is the right answer. Please ignore my previous comment

upvoted 2 times

backbencher2022 3 weeks, 3 days ago

Selected Answer: B

B is the correct answer because the access must be given with least privilege and A gives more privileges than required as per this documentation.
https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/cloudwatch-dashboard-sharing.html

upvoted 1 times

17Master 1 week, 5 days ago

XD jajaa...Read "But the user doesn't have AWS account". Answer B may be questionable.

upvoted 1 times

backbencher2022 3 weeks, 3 days ago

B is the correct answer because the access must be given with least privilege and A gives more privileges than required as per this documentation.
https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/cloudwatch-dashboard-sharing.html

upvoted 1 times

KADSM 3 weeks ago

But the user doesn't have AWS account. Answer B may be questionable.

upvoted 1 times

keezbadger 4 weeks ago

This is a tricky one.
For the question the "product manager need access to the dashboard". However, if you share the dashboard. AWS documentation says:

Warning
All people who you share the dashboard with are granted these permissions for the account. If you share the dashboard publicly, then everyone who has the link to the dashboard has these permissions.

The cloudwatch:GetMetricData and ec2:DescribeTags permissions cannot be scoped down to specific metrics or EC2 instances, so the people with access to the dashboard can query all CloudWatch metrics and the names and tags of all EC2 instances in the account.

So the product manager can not only view, but query "ALL cloud watch metrics, Names, and tags of EC2 instances in the account" The question now is do we want the product manager to do that? Or just Read/view the dashboard? If it is just Read/view, then Option B, but if it is he needs to do the aformentioned, then Option A.

upvoted 3 times

tinyfoot 3 weeks ago

Initially I had the same idea as you but researching into the AWS documentations, actually, the CloudWatchReadOnlyAccess managed policy is much more permissive than that of the shared role that is created.

Based on (https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/iam-identity-based-access-control-cw.html#managed-policies-cloudwatch-CloudWatchReadOnlyAccess) you can see that the resource allowed is * and the various actions allowed are extremely wide including the cloudwatch:Get*. This essentially has the same rights as what you mentioned about having access to query all EC2 instances metrics in the account and more. If the option was to create a custom managed or inline policy that only grants access to that dashboard resource then it would be the least privilege method.

Whereas the CloudWatch Role is only granted 4 specific permissions on all resource.
https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/cloudwatch-dashboard-sharing.html

Hence I believe that A is a better answer.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: A

agree with A

upvoted 1 times

Sinaneos 1 month, 2 weeks ago

Selected Answer: A

it's definitely A, there's no point of creating IAM user if the client doesn't have an account. Dashboard share is least privileage

upvoted 1 times

BoboChow 1 month, 2 weeks ago

Selected Answer: B

It seems like B

upvoted 2 times

17Master 4 weeks ago

WHY? iam? - Read --> "The product manager does not have an AWS account" - https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/cloudwatch-dashboard-sharing.html -- ans is A

upvoted 2 times

sba21 1 month, 2 weeks ago

Selected Answer: A

CloudWatch > Dashboard > Select your board > Share Dashboard>Share your dashboard and require a username and password>Enter mail address

upvoted 1 times

Question #28 *Topic 1*

A company is migrating applications to AWS. The applications are deployed in different accounts. The company manages the accounts centrally by using AWS Organizations. The company's security team needs a single sign-on (SSO) solution across all the company's accounts. The company must continue managing the users and groups in its on-premises self-managed Microsoft Active Directory.
Which solution will meet these requirements?

A. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console. Create a one-way forest trust or a one-way domain trust to connect the company's self-managed Microsoft Active Directory with AWS SSO by using AWS Directory Service for Microsoft Active Directory.

B. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console. Create a two-way forest trust to connect the company's self-managed Microsoft Active Directory with AWS SSO by using AWS Directory Service for Microsoft Active Directory.

C. Use AWS Directory Service. Create a two-way trust relationship with the company's self-managed Microsoft Active Directory.

D. Deploy an identity provider (IdP) on premises. Enable AWS Single Sign-On (AWS SSO) from the AWS SSO console.

Correct Answer: *A*

Community vote distribution

B (93%) 7%

17Master Highly Voted 4 weeks ago

Selected Answer: B

Tricky question!!! forget one-way or two-way. In this scenario, AWS applications (Amazon Chime, Amazon Connect, Amazon QuickSight, AWS Single Sign-On, Amazon WorkDocs, Amazon WorkMail, Amazon WorkSpaces, AWS Client VPN, AWS Management Console, and AWS Transfer Family) need to be able to look up objects from the on-premises domain in order for them to function. This tells you that authentication needs to flow both ways. This scenario requires a two-way trust between the on-premises and AWS Managed Microsoft AD domains.
It is a requirement of the application
Scenario 2: https://aws.amazon.com/es/blogs/security/everything-you-wanted-to-know-about-trusts-with-aws-managed-microsoft-ad/

upvoted 7 times

Wpcorgan Most Recent 1 week ago

B is correct

upvoted 1 times

pepgua 1 week, 4 days ago

Selected Answer: B

https://docs.aws.amazon.com/singlesignon/latest/userguide/connectonpremad.html

upvoted 1 times

KADSM 3 weeks ago

Answer B as we have AWS SSO which requires two way trust. As per documentation - A two-way trust is required for AWS Enterprise Apps such as Amazon Chime, Amazon Connect, Amazon QuickSight, AWS IAM Identity Center (successor to AWS Single Sign-On), Amazon WorkDocs, Amazon WorkMail, Amazon WorkSpaces, and the AWS Management Console. AWS Managed Microsoft AD must be able to query the users and groups in your self-managed AD.

Amazon EC2, Amazon RDS, and Amazon FSx will work with either a one-way or two-way trust.

upvoted 3 times

dave9994 1 month, 1 week ago

B should be the correct answer. To access AWS Console, two-way trust is required.

upvoted 1 times

cark0728 1 month, 1 week ago

Selected Answer: B

b가 정답입니다

upvoted 2 times

KVK16 1 month, 2 weeks ago

Selected Answer: B

In this there is requirement of trust two ways - from On Prem domain to AWS domain. They both will trust each other. One-way trust : If Domain A is trusting Domain B , users from Domain B can access resources of Domain A but not vice-versa. So for A - its Outgoing trust and for B its incoming trust.

AWS applications (Amazon Chime, Amazon Connect, Amazon QuickSight, AWS Single Sign-On, Amazon WorkDocs, Amazon WorkMail, Amazon WorkSpaces, AWS Client VPN, AWS Management Console, and AWS Transfer Family) need to be able to look up objects from the on-premises domain in order for them to function. This tells you that authentication needs to flow both ways.

For Elaborate details refer
https://aws.amazon.com/blogs/security/everything-you-wanted-to-know-about-trusts-with-aws-managed-microsoft-ad/

upvoted 3 times

bluer1 1 month, 2 weeks ago

The answer should be B,

You can configure one and two-way external and forest trust relationships between your AWS Directory Service for Microsoft Active Directory and self-managed (on-premises) directories, as well as between multiple AWS Managed Microsoft AD directories in the AWS cloud. AWS Managed Microsoft AD supports all three trust relationship directions: Incoming, Outgoing and Two-way (Bi-directional).

upvoted 2 times

masetromain 1 month, 2 weeks ago

old question : https://aws.amazon.com/fr/about-aws/whats-new/2022/07/aws-single-sign-on-aws-sso-now-aws-iam-identity-center/

https://aws.amazon.com/fr/about-aws/whats-new/2022/04/aws-single-sign-on-configurable-synchronization-microsoft-active-directory/

upvoted 1 times

BoboChow 1 month, 2 weeks ago

Selected Answer: A

SSO is suitable for the senario. Because it's intergrated with AWS Organizations and on-premise Active Directory. You can centrally manage Single Sign-On to access multiple accounts .
But I don't know the one-way/two-way forest trust thing...

upvoted 1 times

Question #29 *Topic 1*

A company provides a Voice over Internet Protocol (VoIP) service that uses UDP connections. The service consists of Amazon EC2 instances that run in an Auto Scaling group. The company has deployments across multiple AWS Regions.
The company needs to route users to the Region with the lowest latency. The company also needs automated failover between Regions.
Which solution will meet these requirements?

A. Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Use the NLB as an AWS Global Accelerator endpoint in each Region.

B. Deploy an Application Load Balancer (ALB) and an associated target group. Associate the target group with the Auto Scaling group. Use the ALB as an AWS Global Accelerator endpoint in each Region.

C. Deploy a Network Load Balancer (NLB) and an associated target group. Associate the target group with the Auto Scaling group. Create an Amazon Route 53 latency record that points to aliases for each NLB. Create an Amazon CloudFront distribution that uses the latency record as an origin.

D. Deploy an Application Load Balancer (ALB) and an associated target group. Associate the target group with the Auto Scaling group. Create an Amazon Route 53 weighted record that points to aliases for each ALB. Deploy an Amazon CloudFront distribution that uses the weighted record as an origin.

Correct Answer: *C*

Community vote distribution

A (74%) C (24%)

Six_Fingered_Jose Highly Voted 1 month ago

Selected Answer: A

agree with A,
Global Accelerator has automatic failover and is perfect for this scenario with VoIP
https://aws.amazon.com/global-accelerator/faqs/

upvoted 7 times

BoboChow 1 month ago

Thank you for your link, it make me consolidate A.

upvoted 3 times

Wpcorgan Most Recent 1 week ago

A is correct

upvoted 1 times

grzeev 2 weeks ago

Selected Answer: A

https://youtu.be/GAxrPQ3ycsQ

upvoted 1 times

htangga 2 weeks ago

Selected Answer: C

C is correct answer: Because "The company needs to route users to the Region with the lowest latency", in this case, route53 is needed to do so

upvoted 2 times

vn_thanhtung 2 weeks, 3 days ago

Selected Answer: A

Global accelerator good fit for non-HTTP uses case, such as gaming UDP, Void over IP, IoT (MQTT)

upvoted 1 times

AbhiJo 2 weeks, 3 days ago

Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover

upvoted 2 times

mouhannadhaj 3 weeks, 2 days ago

Selected Answer: A

CloudFront uses Edge Locations to cache content while Global Accelerator uses Edge Locations to find an optimal pathway to the nearest regional endpoint. CloudFront is designed to handle HTTP protocol meanwhile Global Accelerator is best used for both HTTP and non-HTTP protocols such as TCP and UDP. so i think A is a better answer

upvoted 2 times

17Master 4 weeks ago

Selected Answer: A

Failover and fault tolerance - https://aws.amazon.com/es/global-accelerator/faqs/

upvoted 1 times

raffaello44 1 month ago

Selected Answer: A

not C because Route53 points to NLB but it should point to the Cloudfront.

upvoted 2 times

ManoAni 1 month ago

Selected Answer: C

Traffic must be routed with lower latency and deployments running in multiple regions. So its C

upvoted 1 times

iCcma 1 month ago

Selected Answer: A

"C" is very similar to the correct answer, it even sounds good to use Route 53, but it is VoIP not media that is going to be streamed, so cloudFront helps you rule out "C".
Correct answer A""

upvoted 3 times

sensei705 1 month ago

A is the correct answer as global accelerator performs automatic failover and due to only having 2 anycast IPs it automatically directs users to the closest region (edge location)

upvoted 1 times

GameDad09 1 month, 1 week ago

Selected Answer: A

A is the correct one.

upvoted 1 times

Valero_ 1 month, 2 weeks ago

Selected Answer: C

The question indicates that "The company needs to route users to the Region with the lowest latency. The company also needs automated failover between Regions.", so in my opinion, the only answer best fits here is C because of Route 53.
But I'm not sure about the CloudFront distribution setting.

upvoted 2 times

KVK16 1 month, 2 weeks ago

Selected Answer: A

HTTP /HTTPS - ALB ; TCP and UDP - NLB;
Lowest latency routing and more throughput. Also supports failover, uses Anycast Ip addressing - Global Accelerator
Caching at Egde Locations - Cloutfront

upvoted 1 times

Rachness 1 month, 2 weeks ago

WS Global Accelerator automatically checks the health of your applications and routes user traffic only to healthy application endpoints. If the health status changes or you make configuration updates, AWS Global Accelerator reacts instantaneously to route your users to the next available endpoint. https://aws.amazon.com/global-accelerator/faqs/

upvoted 1 times

ninjawrz 1 month, 2 weeks ago

Definitely A

upvoted 1 times

Question #30 *Topic 1*

A development team runs monthly resource-intensive tests on its general purpose Amazon RDS for MySQL DB instance with Performance Insights enabled. The testing lasts for 48 hours once a month and is the only process that uses the database. The team wants to reduce the cost of running the tests without reducing the compute and memory attributes of the DB instance.
Which solution meets these requirements MOST cost-effectively?

A. Stop the DB instance when tests are completed. Restart the DB instance when required.

B. Use an Auto Scaling policy with the DB instance to automatically scale when tests are completed.

C. Create a snapshot when tests are completed. Terminate the DB instance and restore the snapshot when required.

D. Modify the DB instance to a low-capacity instance when tests are completed. Modify the DB instance again when required.

Correct Answer: *C*

Community vote distribution

C (94%) 6%

hanhdroid Highly Voted 1 month, 2 weeks ago

Selected Answer: C

Answer C, you still pay for storage when an RDS database is stopped

upvoted 13 times

Wpcorgan Most Recent 1 week ago

C is correct

upvoted 1 times

17Master 4 weeks ago

Selected Answer: C

is correct

upvoted 1 times

KVK16 1 month, 2 weeks ago

Selected Answer: C

C - Create a manual Snapshot of DB and shift to S3- Standard and Restore form Manual Snapshot when required.

Not A - By stopping the DB although you are not paying for DB hours you are still paying for Provisioned IOPs , the storage for Stopped DB is more than Snapshot of underlying EBS vol. and Automated Back ups .
Not D - Is possible but not MOST cost effective, no need to run the RDS when not needed.

upvoted 2 times

BoboChow 1 month, 2 weeks ago

Selected Answer: A

If instance state is stopped, it's not billed.
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-lifecycle.html

upvoted 1 times

ArielSchivo 1 month, 1 week ago

It's a DB instance, not an EC2 instance. If the DB instance is stopped, you are still paying for the storage.

upvoted 6 times

Jerry84 2 weeks, 2 days ago

Thanks for your reply.

upvoted 1 times

BoboChow 1 month ago

Thank you for your explanation

upvoted 2 times

Rachness 1 month, 2 weeks ago

While your DB instance is stopped, you are charged for provisioned storage (including Provisioned IOPS)

upvoted 2 times

Question #31 *Topic 1*

A company that hosts its web application on AWS wants to ensure all Amazon EC2 instances. Amazon RDS DB instances. and Amazon Redshift clusters are configured with tags. The company wants to minimize the effort of configuring and operating this check.
What should a solutions architect do to accomplish this?

A. Use AWS Config rules to define and detect resources that are not properly tagged.

B. Use Cost Explorer to display resources that are not properly tagged. Tag those resources manually.

C. Write API calls to check all resources for proper tag allocation. Periodically run the code on an EC2 instance.

D. Write API calls to check all resources for proper tag allocation. Schedule an AWS Lambda function through Amazon CloudWatch to periodically run the code.

Correct Answer: *A*

Community vote distribution

A (89%) 11%

Wpcorgan 1 week ago

A is correct

upvoted 1 times

Buruguduystunstugudunstuy 1 week, 5 days ago

The correct answer is A.
https://docs.aws.amazon.com/config/latest/developerguide/tagging.html

upvoted 3 times

backbencher2022 3 weeks, 3 days ago

Selected Answer: A

Easiest option is A

upvoted 3 times

17Master 4 weeks ago

Selected Answer: R

Is correct

upvoted 1 times

123jhl0 1 month, 1 week ago

Selected Answer: A

A can do the task and is the one involving less effort.

upvoted 1 times

BoboChow 1 month, 2 weeks ago

Selected Answer: A

I think Config works

upvoted 4 times

Lilibell 1 month, 2 weeks ago

The answer is A

upvoted 2 times

Question #32 *Topic 1*

A development team needs to host a website that will be accessed by other teams. The website contents consist of HTML, CSS, client-side JavaScript, and images.
Which method is the MOST cost-effective for hosting the website?

A. Containerize the website and host it in AWS Fargate.

B. Create an Amazon S3 bucket and host the website there.

C. Deploy a web server on an Amazon EC2 instance to host the website.

D. Configure an Application Load Balancer with an AWS Lambda target that uses the Express.js framework.

Correct Answer: *B*

Community vote distribution

B (100%)

masetromain Highly Voted 1 month, 2 weeks ago

Selected Answer: B

Good answer is B: client-side JavaScript. the website is static, so it must be S3.

upvoted 6 times

Wpcorgan Most Recent 1 week ago

B is correct

upvoted 1 times

17Master 4 weeks ago

Selected Answer: B

Is correct

upvoted 1 times

ninjawrz 1 month, 2 weeks ago

Selected Answer: B

B: Host static website in S3

upvoted 2 times

BoboChow 1 month, 2 weeks ago

Selected Answer: B

HTML, CSS, client-side JavaScript, and images are all static resources.

upvoted 2 times

Lilibell 1 month, 2 weeks ago

The answer is B

upvoted 2 times

Question #33 *Topic 1*

A company runs an online marketplace web application on AWS. The application serves hundreds of thousands of users during peak hours. The company needs a scalable, near-real-time solution to share the details of millions of financial transactions with several other internal applications. Transactions also need to be processed to remove sensitive data before being stored in a document database for low-latency retrieval.
What should a solutions architect recommend to meet these requirements?

A. Store the transactions data into Amazon DynamoDB. Set up a rule in DynamoDB to remove sensitive data from every transaction upon write. Use DynamoDB Streams to share the transactions data with other applications.

B. Stream the transactions data into Amazon Kinesis Data Firehose to store data in Amazon DynamoDB and Amazon S3. Use AWS Lambda integration with Kinesis Data Firehose to remove sensitive data. Other applications can consume the data stored in Amazon S3.

C. Stream the transactions data into Amazon Kinesis Data Streams. Use AWS Lambda integration to remove sensitive data from every transaction and then store the transactions data in Amazon DynamoDB. Other applications can consume the transactions data off the Kinesis data stream.

D. Store the batched transactions data in Amazon S3 as files. Use AWS Lambda to process every file and remove sensitive data before updating the files in Amazon S3. The Lambda function then stores the data in Amazon DynamoDB. Other applications can consume transaction files stored in Amazon S3.

Correct Answer: *C*

Community vote distribution

C (86%) 14%

JesseeS Highly Voted 1 month, 1 week ago

The answer is C, because Firehose does not suppport DynamoDB and another key word is "data" Kinesis Data Streams is the correct choice. Pay attention to key words. AWS likes to trick you up to make sure you know the services.

upvoted 8 times

Wpcorgan Most Recent 1 week ago

C is correct

upvoted 1 times

Wajif 2 weeks, 1 day ago

Selected Answer: C

Will vote C due to lack of support of DynamoDB in Firehose. Solution A caught my eye for a second but then ignoring it becaue DynamoDB streams would only send data changes to an existing table.

upvoted 1 times

keezbadger 3 weeks, 2 days ago

This question is messed up. AWS question points out "Near real time" which points to Kinesis Firehose and then went ahead to mess it up with a solution support pointing to DynamoDB which is inaccurate and leave us with the C option with Kinesis Data Stream, which is a "real time" streams.

upvoted 4 times

ArielSchivo 1 month, 1 week ago

Selected Answer: C

I would go for C. The tricky phrase is "near-real-time solution", pointing to Firehouse, but it can't send data to DynamoDB, so it leaves us with C as best option.

Kinesis Data Firehose currently supports Amazon S3, Amazon Redshift, Amazon OpenSearch Service, Splunk, Datadog, NewRelic, Dynatrace, Sumologic, LogicMonitor, MongoDB, and HTTP End Point as destinations.

https://aws.amazon.com/kinesis/data-firehose/faqs/#:~:text=Kinesis%20Data%20Firehose%20currently%20supports,HTTP%20End%20Point%20as%20destinations.

upvoted 3 times

Gatt 1 month, 1 week ago

Selected Answer: C

A and D are out since they are bad answers.
B is impossible as Kinesis Firehose cannot directly store data in DynamoDB.
This only leaves C

upvoted 1 times

KVK16 1 month, 2 weeks ago

B vs C
B : Pro- Sensitive information removal by Lambda , near real-time
Con - Ingesting into Dynamo DB although has low latency is complex and other application using S3 for data

C: Pro - Other Apps using data off DB Streams
Con : Sensitive data removal lambda , near real-time
My Pick is B for being more managed service - a Rule of thumb
upvoted 1 times

BoboChow 1 month, 2 weeks ago

Selected Answer: C

https://docs.aws.amazon.com/zh_cn/streams/latest/dev/tutorial-stock-data-lambda.html
upvoted 1 times

340trunkhawk 1 month, 2 weeks ago

Please explain C
upvoted 1 times

Sinaneos 1 month, 2 weeks ago

Selected Answer: B

Near-real-time directly points to kinesis firehose, it's fully managed and scalable, along with dynamoDB. Transforming data with lambda function integration is one of its main use cases. Therefore the answer has got to be B
https://aws.amazon.com/kinesis/data-firehose/faqs/?nc=sn&loc=5
upvoted 1 times

herculian_effort 1 month, 2 weeks ago

It's C. There is no support for Firehose to DynamoDB.
upvoted 2 times

Sinaneos 1 month ago

that's correct, I missed that point, thanks :)
upvoted 2 times

Lilibell 1 month, 2 weeks ago

The answer is C
upvoted 4 times

Question #34 *Topic 1*

A company hosts its multi-tier applications on AWS. For compliance, governance, auditing, and security, the company must track configuration changes on its AWS resources and record a history of API calls made to these resources.
What should a solutions architect do to meet these requirements?

A. Use AWS CloudTrail to track configuration changes and AWS Config to record API calls.

B. Use AWS Config to track configuration changes and AWS CloudTrail to record API calls.

C. Use AWS Config to track configuration changes and Amazon CloudWatch to record API calls.

D. Use AWS CloudTrail to track configuration changes and Amazon CloudWatch to record API calls.

Correct Answer: *B*

Community vote distribution

B (100%)

airraid2010 Highly Voted 1 month ago

Selected Answer: B

CloudTrail - Track user activity and API call history.
Config - Assess, audits, and evaluates the configuration and relationships of tag resources.

Therefore, the answer is B

upvoted 11 times

Wpcorgan Most Recent 1 week ago

B is correct

upvoted 1 times

bansalhp 1 month, 1 week ago

Selected Answer: B

The answer is B

upvoted 2 times

Evangelia 1 month, 1 week ago

bbbbbbbb

upvoted 3 times

tubtab 1 month, 2 weeks ago

Selected Answer: B

bbbbbbbb

upvoted 4 times

Lilibell 1 month, 2 weeks ago

The answer is B

upvoted 2 times

Question #35 *Topic 1*

A company is preparing to launch a public-facing web application in the AWS Cloud. The architecture consists of Amazon EC2 instances within a VPC behind an Elastic Load Balancer (ELB). A third-party service is used for the DNS. The company's solutions architect must recommend a solution to detect and protect against large-scale DDoS attacks.
Which solution meets these requirements?

A. Enable Amazon GuardDuty on the account.

B. Enable Amazon Inspector on the EC2 instances.

C. Enable AWS Shield and assign Amazon Route 53 to it.

D. Enable AWS Shield Advanced and assign the ELB to it.

Correct Answer: *D*

Community vote distribution

D (100%)

BoboChow Highly Voted 1 month, 2 weeks ago

Selected Answer: D

AWS Shield Advanced provides expanded DDoS attack protection for your Amazon EC2 instances, Elastic Load Balancing load balancers, CloudFront distributions, Route 53 hosted zones, and AWS Global Accelerator standard accelerators.

upvoted 12 times

ninjawrz Highly Voted 1 month, 2 weeks ago

Selected Answer: D

Answer is D
C is incorrect because question says Third party DNS and route 53 is AWS proprietary

upvoted 7 times

Wpcorgan Most Recent 1 week ago

D is correct

upvoted 1 times

Buruguduystunstugudunstuy 1 week, 4 days ago

Large-scale DDoS attacks = AWS Shield Advanced
The correct answer is D
https://aws.amazon.com/shield/faqs/
https://docs.aws.amazon.com/whitepapers/latest/aws-best-practices-ddos-resiliency/elastic-load-balancing-bp6.html

upvoted 1 times

Wajif 2 weeks, 1 day ago

Selected Answer: D

Same reasoning as given by Ninjawarz

upvoted 1 times

Lilibell 1 month, 2 weeks ago

The answer is D

upvoted 3 times

Question #36 *Topic 1*

A company is building an application in the AWS Cloud. The application will store data in Amazon S3 buckets in two AWS Regions. The company must use an AWS Key Management Service (AWS KMS) customer managed key to encrypt all data that is stored in the S3 buckets. The data in both S3 buckets must be encrypted and decrypted with the same KMS key. The data and the key must be stored in each of the two Regions. Which solution will meet these requirements with the LEAST operational overhead?

A. Create an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Configure replication between the S3 buckets.

B. Create a customer managed multi-Region KMS key. Create an S3 bucket in each Region. Configure replication between the S3 buckets. Configure the application to use the KMS key with client-side encryption.

C. Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Configure replication between the S3 buckets.

D. Create a customer managed KMS key and an S3 bucket in each Region. Configure the S3 buckets to use server-side encryption with AWS KMS keys (SSE-KMS). Configure replication between the S3 buckets.

Correct Answer: *C*

Community vote distribution

B (67%) D (33%)

pooppants Highly Voted 1 month, 2 weeks ago

Selected Answer: B

KMS Multi-region keys are required https://docs.aws.amazon.com/kms/latest/developerguide/multi-region-keys-overview.html

upvoted 10 times

magazz 1 week ago

Amazon S3 cross-region replication decrypts and re-encrypts data under a KMS key in the destination Region, even when replicating objects protected by a multi-Region key. So stating that Amazon S3 cross-region replication decrypts and re-encrypts data under a KMS key in the destination Region, even when replicating objects protected by a multi-Region key is required is incorrect

upvoted 2 times

TuLe 1 day, 11 hours ago

@magazz: it's not true then. Based on the document from AWS https://docs.aws.amazon.com/AmazonS3/latest/userguide/replication-config-for-kms-objects.html , we will need to setup the replication rule with destination KMS. In order to have the key available in more than 2, then multi-region key should be required. But I'm still not favor option B - we can use server-side when why wasting effort to do client side encryption.

upvoted 1 times

TuLe 1 day, 11 hours ago

I would say it's true... Not sure the previous one say "not true" :D.

upvoted 1 times

KJa Highly Voted 1 month, 2 weeks ago

Selected Answer: D

Cannot be A - question says customer managed key
Cannot B - client side encryption is operational overhead
Cannot C -as it says SSE-S3 instead of customer managed
so the answer is D though it required one time setup of keys

upvoted 6 times

th3cookie 1 week, 4 days ago

How does client side encryption increase OPERATIONAL overhead? Do you think every connected client is sitting there with gpg cli, decrypting/encrypting every packet that comes in/out? No, it's done via SDK -> https://docs.aws.amazon.com/encryption-sdk/latest/developer-guide/introduction.html

The correct answer is B because that's the only way to actually get the same key across multiple regions with minimal operational overhead

upvoted 1 times

BoboChow 1 month, 2 weeks ago

The data in both S3 buckets must be encrypted and decrypted with the same KMS key.
AWS KMS supports multi-Region keys, which are AWS KMS keys in different AWS Regions that can be used interchangeably – as though you had the same key in multiple Regions.
"as though" means it's different.
So I agree with B

upvoted 3 times

BoboChow 1 month, 2 weeks ago

key change across regions unless you use multi-Region keys

upvoted 1 times

mattlai 1 month, 2 weeks ago

fun joke, if u dont do encryption on client side, where else could it be?

upvoted 1 times

Newptone 3 weeks ago

It could be server side. For client side, the application need to finish the encryption and decryption by itself. So S3 object encryption on the server side is less operational overhead. https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingClientSideEncryption.html

But for option B, the major issue is if you create KMS keys in 2 regions, they can not be the same.

upvoted 1 times

Newptone 3 weeks ago

Sorry for the typo, I mean option D.

upvoted 2 times

Mikkka Most Recent 1 week, 3 days ago

BBBBBB

upvoted 1 times

manu427 2 weeks, 2 days ago

Should be D (cant b B as B is client side encryption , but question says SSE-KMS)

upvoted 1 times

backbencher2022 3 weeks, 3 days ago

Selected Answer: B

B is the correct answer because of same key usage requirements at both regions and Multi-Region keys in AWS KMS fulfills this requirement. For more details check - https://docs.aws.amazon.com/kms/latest/developerguide/multi-region-keys-overview.html

upvoted 2 times

The_Body 1 month ago

B. “The company must use an AWS Key Management Service (AWS KMS) customer managed key to encrypt all data that is stored in the S3 buckets”. Customer managed keys can be multi-region. SSE-S3 doesn’t allow for customer managed key, AWS S3 manages the key— server side encryption.

upvoted 2 times

ManoAni 1 month ago

Selected Answer: B

Catch is multi-region keys and those must be same. so its B

upvoted 2 times

Six_Fingered_Jose 1 month ago

Selected Answer: D

answer should be D as question explicitly mentioned the use of AWS KMS key for the s3 encryption, while C is only mentioning the creation of the key but only encrypting with S3 encryption

upvoted 1 times

BoboChow 1 month ago

The data in both S3 buckets must be encrypted and decrypted with the same KMS key.
D creates KMS key in each region, which means to create two key in total

upvoted 1 times

Ekie 1 month, 2 weeks ago

B is right answer

upvoted 2 times

CloudGuru99 1 month, 2 weeks ago

B is the correct answer

upvoted 2 times

Question #37 *Topic 1*

A company recently launched a variety of new workloads on Amazon EC2 instances in its AWS account. The company needs to create a strategy to access and administer the instances remotely and securely. The company needs to implement a repeatable process that works with native AWS services and follows the AWS Well-Architected Framework.
Which solution will meet these requirements with the LEAST operational overhead?

A. Use the EC2 serial console to directly access the terminal interface of each instance for administration.

B. Attach the appropriate IAM role to each existing instance and new instance. Use AWS Systems Manager Session Manager to establish a remote SSH session.

C. Create an administrative SSH key pair. Load the public key into each EC2 instance. Deploy a bastion host in a public subnet to provide a tunnel for administration of each instance.

D. Establish an AWS Site-to-Site VPN connection. Instruct administrators to use their local on-premises machines to connect directly to the instances by using SSH keys across the VPN tunnel.

Correct Answer: *B*

Community vote distribution

B (75%) C (25%)

Wpcorgan 1 week ago
B is correct for me
upvoted 1 times

xeun88 1 week, 5 days ago
B is the right answer
upvoted 1 times

Keld 2 weeks, 3 days ago
Selected Answer: C
The answer is C, there is no indication of which type of EC2 Windows/Linux.
SSH only works for Windows
upvoted 1 times

Keld 2 weeks, 3 days ago
only works for *Linux*
upvoted 1 times

17Master 1 week, 5 days ago
Is correct C - https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html ------>
Cross-platform support for Windows, Linux, and macOS
Session Manager provides support for Windows, Linux, and macOS from a single tool. For example, you don't need to use an SSH client for Linux and macOS managed nodes or an RDP connection for Windows Server managed nodes.
upvoted 2 times

17Master 1 week, 5 days ago
B is correct
upvoted 1 times

BoboChow 1 month ago
Selected Answer: B
How can Session Manager benefit my organization?
Ans: No open inbound ports and no need to manage bastion hosts or SSH keys
https://docs.aws.amazon.com/systems-manager/latest/userguide/session-manager.html
upvoted 3 times

Nightducky 1 week, 6 days ago
Do you know what from the question is it Windows or Linux EC2. I think not so how you want to do SSH session for Windows?
Answer is C
upvoted 1 times

ManoAni 1 month ago
The answer is C, they mentioned that it must be native service, option B is not a service, it is one of the option to connect to instances.
upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: B

> works with native AWS services
think they want you to use AWS service so B seems to make the most sense.
and also least operational overhead

upvoted 1 times

Six_Fingered_Jose 1 month ago

in real world scenario C seems to be most common solution thought..

upvoted 1 times

ninjawrz 1 month, 2 weeks ago

Selected Answer: B

B for least operational overhead

but bastion is for best practice

upvoted 2 times

Sinaneos 1 month, 2 weeks ago

Selected Answer: B

It's either B or C, but B requires less operational overhead (just attaching the IAM role), so I'd pick B

upvoted 2 times

masetromain 1 month, 2 weeks ago

Selected Answer: C

good answer is C ! Bastion is all time the good answer for this case.

https://aws.amazon.com/fr/quickstart/architecture/linux-bastion/

upvoted 2 times

AbhiJo 2 weeks, 2 days ago

B is correct with less overhead in the question

upvoted 1 times

tubtab 1 month, 2 weeks ago

Selected Answer: B

I think b is the correct answer

upvoted 1 times

TaiTran1994 1 month, 2 weeks ago

I think C is the correct answer

upvoted 1 times

Lilibell 1 month, 2 weeks ago

they say the least operational overhead so B

upvoted 3 times

Question #38 *Topic 1*

A company is hosting a static website on Amazon S3 and is using Amazon Route 53 for DNS. The website is experiencing increased demand from around the world. The company must decrease latency for users who access the website.
Which solution meets these requirements MOST cost-effectively?

A. Replicate the S3 bucket that contains the website to all AWS Regions. Add Route 53 geolocation routing entries.

B. Provision accelerators in AWS Global Accelerator. Associate the supplied IP addresses with the S3 bucket. Edit the Route 53 entries to point to the IP addresses of the accelerators.

C. Add an Amazon CloudFront distribution in front of the S3 bucket. Edit the Route 53 entries to point to the CloudFront distribution.

D. Enable S3 Transfer Acceleration on the bucket. Edit the Route 53 entries to point to the new endpoint.

Correct Answer: *C*

Community vote distribution

C (100%)

Wpcorgan 1 week ago
C is correct
upvoted 1 times

17Master 1 week, 4 days ago
Selected Answer: C
ok CloudFront
upvoted 1 times

xeun88 1 week, 5 days ago
C is right
upvoted 1 times

Mordans 3 weeks ago
Selected Answer: C
ANSWER C
upvoted 1 times

ninjawrz 1 month, 2 weeks ago
Selected Answer: C
C: Cloudfront
upvoted 2 times

masetromain 1 month, 2 weeks ago
Selected Answer: C
of course cloudfront it's the answer
upvoted 3 times

Question #39 *Topic 1*

A company maintains a searchable repository of items on its website. The data is stored in an Amazon RDS for MySQL database table that contains more than 10 million rows. The database has 2 TB of General Purpose SSD storage. There are millions of updates against this data every day through the company's website.
The company has noticed that some insert operations are taking 10 seconds or longer. The company has determined that the database storage performance is the problem.
Which solution addresses this performance issue?

A. Change the storage type to Provisioned IOPS SSD.

B. Change the DB instance to a memory optimized instance class.

C. Change the DB instance to a burstable performance instance class.

D. Enable Multi-AZ RDS read replicas with MySQL native asynchronous replication.

Correct Answer: *B*

Community vote distribution

A (93%) 7%

Wpcorgan 1 week ago
A is correct
upvoted 1 times

17Master 1 week, 3 days ago
Selected Answer: A
SSD is correct
upvoted 1 times

TonyghostR05 1 week, 6 days ago
SSD is correct
upvoted 1 times

luvincanada 2 weeks, 5 days ago
A is correct.
upvoted 1 times

backbencher2022 3 weeks, 3 days ago
Selected Answer: A
A is the correct answer
upvoted 1 times

Six_Fingered_Jose 1 month ago
Selected Answer: A
> The company has determined that the database storage performance is the problem.
A seems to be the most feasible answer here
upvoted 1 times

Sarvar89 1 month, 1 week ago
i is indicating that its storage solution has a problem, so i think A should be correct.
upvoted 1 times

Anon9132656 1 month, 1 week ago
Selected Answer: A
A
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html
upvoted 1 times

Incognito013 1 month, 1 week ago
Answer: A
upvoted 1 times

ninjawrz 1 month, 2 weeks ago
Selected Answer: A
A: IOPS
Provisioned IOPS – Provisioned IOPS storage is designed to meet the needs of I/O-intensive workloads, particularly database workloads, that

require low I/O latency and consistent I/O throughput.

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html
upvoted 3 times

pooppants 1 month, 2 weeks ago

Selected Answer: A

A https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html
upvoted 2 times

pooppants 1 month, 2 weeks ago

Selected Answer: B

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_Storage.html
upvoted 1 times

pooppants 1 month, 2 weeks ago

ignore, answer is b. I misclicked
upvoted 1 times

pooppants 1 month, 2 weeks ago

ugh! A! the answer is A
upvoted 6 times

tuloveu 1 month, 2 weeks ago

Selected Answer: A

There are millions of updates against this data every day. The storage Provisioned IOPS SSD should be served.
upvoted 4 times

BoboChow 1 month, 2 weeks ago

You hit the point
upvoted 1 times

CloudGuru99 1 month, 2 weeks ago

A is the correct answer
upvoted 2 times

Question #40 *Topic 1*

A company has thousands of edge devices that collectively generate 1 TB of status alerts each day. Each alert is approximately 2 KB in size. A solutions architect needs to implement a solution to ingest and store the alerts for future analysis.
The company wants a highly available solution. However, the company needs to minimize costs and does not want to manage additional infrastructure. Additionally, the company wants to keep 14 days of data available for immediate analysis and archive any data older than 14 days.
What is the MOST operationally efficient solution that meets these requirements?

A. Create an Amazon Kinesis Data Firehose delivery stream to ingest the alerts. Configure the Kinesis Data Firehose stream to deliver the alerts to an Amazon S3 bucket. Set up an S3 Lifecycle configuration to transition data to Amazon S3 Glacier after 14 days.

B. Launch Amazon EC2 instances across two Availability Zones and place them behind an Elastic Load Balancer to ingest the alerts. Create a script on the EC2 instances that will store the alerts in an Amazon S3 bucket. Set up an S3 Lifecycle configuration to transition data to Amazon S3 Glacier after 14 days.

C. Create an Amazon Kinesis Data Firehose delivery stream to ingest the alerts. Configure the Kinesis Data Firehose stream to deliver the alerts to an Amazon OpenSearch Service (Amazon Elasticsearch Service) cluster. Set up the Amazon OpenSearch Service (Amazon Elasticsearch Service) cluster to take manual snapshots every day and delete data from the cluster that is older than 14 days.

D. Create an Amazon Simple Queue Service (Amazon SQS) standard queue to ingest the alerts, and set the message retention period to 14 days. Configure consumers to poll the SQS queue, check the age of the message, and analyze the message data as needed. If the message is 14 days old, the consumer should copy the message to an Amazon S3 bucket and delete the message from the SQS queue.

Correct Answer: *A*

Community vote distribution

A (85%) D (15%)

Sinaneos Highly Voted 1 month, 2 weeks ago

Selected Answer: A

Definitely A, it's the most operationally efficient compared to D, which requires a lot of code and infrastructure to maintain. A is mostly managed (firehose is fully managed and S3 lifecycles are also managed)

upvoted 14 times

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: A

Only A makes sense operationally.
If you think D, just consider what is needed to move the message from SQS to S3... you are polling daily 14 TB to take out 1 TB... that's no operationally efficient at all.

upvoted 7 times

Wpcorgan Most Recent 1 week ago

A is correct

upvoted 1 times

Incognito013 1 month, 1 week ago

A

Stroring the data in S3 and assign a policy to transfer the data to Glacier after 14 days

upvoted 1 times

KVK16 1 month, 1 week ago

Selected Answer: D

In most of the questions, first check the answers that are feasible and then check for the Well-Architected pillar emphasis in the question and hints pointing to it in solving Qs of SAA
D: SQS vs Kinesis
Both do support retention period 14 days , max record size [256 KB and 1MB] and
Total Data produced is 1TB/day

In Question there is "store the alerts for future analysis" "highly available solution. However, the company needs to minimize costs and does not want to manage additional infrastructure""MOST operationally efficient solution "
No requirement for real time and ordered processing. Also need for LEAST OPERATIONAL head. In Case of Kinesis one has to be watchful of shards capacity so no scope for Autoscaling like SQS and Cost Basis. No need for multi-consumers only one place to store S3. SQS- fully serverless
So I think its SQS . Incase there are even multi-consumers still consider SQS-SNS model.

upvoted 3 times

yd_h 1 month, 1 week ago

Nope, can't get immediate access to any data you want with SQS. Additionally, if you do somehow you have to stop calling the delete message API call for 14 hours, and then...

upvoted 1 times

ArielSchivo 1 month, 1 week ago

This is great, but Kinesis and S3 are managed services, so it should be option A as that one is the most operational.

upvoted 2 times

ninjawrz 1 month, 2 weeks ago

Selected Answer: A

This should be A

upvoted 2 times

tubtab 1 month, 2 weeks ago

Selected Answer: D

ddddddddd

upvoted 1 times

logicalbin 1 month, 2 weeks ago

A cannot be the right answer.
Reason - Fireshose cannot be used to store data. Only Kinesis Data Streams has that ability as it focuses on ingesting and storing data streams while Kinesis Data Firehose focuses on delivering data streams to select destinations.
SQS on the other hand can store data upto 14 days.
So option D is correct.

upvoted 1 times

ninjawrz 1 month, 2 weeks ago

you will be storing the data in S3 then with life cycle Glacier deep archive after 14 days

upvoted 3 times

mattlai 1 month, 2 weeks ago

option a does not have anything to do with storing data

upvoted 1 times

Lilibell 1 month, 2 weeks ago

The answer is A

upvoted 5 times

Question #41 *Topic 1*

A company's application integrates with multiple software-as-a-service (SaaS) sources for data collection. The company runs Amazon EC2 instances to receive the data and to upload the data to an Amazon S3 bucket for analysis. The same EC2 instance that receives and uploads the data also sends a notification to the user when an upload is complete. The company has noticed slow application performance and wants to improve the performance as much as possible.
Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Auto Scaling group so that EC2 instances can scale out. Configure an S3 event notification to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete.

B. Create an Amazon AppFlow flow to transfer data between each SaaS source and the S3 bucket. Configure an S3 event notification to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete.

C. Create an Amazon EventBridge (Amazon CloudWatch Events) rule for each SaaS source to send output data. Configure the S3 bucket as the rule's target. Create a second EventBridge (Cloud Watch Events) rule to send events when the upload to the S3 bucket is complete. Configure an Amazon Simple Notification Service (Amazon SNS) topic as the second rule's target.

D. Create a Docker container to use instead of an EC2 instance. Host the containerized application on Amazon Elastic Container Service (Amazon ECS). Configure Amazon CloudWatch Container Insights to send events to an Amazon Simple Notification Service (Amazon SNS) topic when the upload to the S3 bucket is complete.

Correct Answer: *B*

Community vote distribution

B (87%) 13%

Wpcorgan 1 week ago
B is Correct
upvoted 1 times

Wajif 1 week, 1 day ago
Selected Answer: B
Choosing B as it sounds simpler.
upvoted 1 times

peneloco 3 weeks, 4 days ago
Selected Answer: B
AppFlow is made for SaaS
upvoted 4 times

rob74 3 weeks, 5 days ago
Selected Answer: B
AppFlow , managed service SAAS-->Least effort
upvoted 3 times

Six_Fingered_Jose 1 month ago
Selected Answer: B
This question just screams AppFlow (SaaS integration)
https://aws.amazon.com/appflow/
upvoted 3 times

Six_Fingered_Jose 1 month ago
configuring Auto-Scaling also takes time when compared to AppFlow,
in AWS's words "in just a few clicks"
> Amazon AppFlow is a fully managed integration service that enables you to securely transfer data between Software-as-a-Service (SaaS) applications like Salesforce, SAP, Zendesk, Slack, and ServiceNow, and AWS services like Amazon S3 and Amazon Redshift, in just a few clicks
upvoted 3 times

Sinaneos 1 month ago
Appflow works very well with SaaS platforms, makes a lot more sense in this scenario. Using an ASG might improve the performance, but here it asks for THE BEST PERFORMANCE, hence ASG might not fix the underlying issue in an efficient manner.
upvoted 1 times

dave9994 1 month ago
A is the answer, as it is the LEAST ops. overhead as asked. Minimal changes on current system.
upvoted 1 times

yd_h 1 month, 1 week ago

Amazon AppFlow is a bi-directional data transfer service; however, not all source-destination combinations are currently supported. The question does not imply any SaaS providers. It could be any SaaS provider (https://docs.aws.amazon.com/appflow/latest/userguide/requirements.html)

upvoted 1 times

yd_h 1 month, 1 week ago

I will go with A. LEAST operational overhead to add an ASG to the existing ec2 instances let S3 handle the notification part.

upvoted 1 times

123jhl0 1 month, 1 week ago

Selected Answer: B

Amazon AppFlow is a fully managed integration service that enables you to securely transfer data between Software-as-a-Service (SaaS) applications like Salesforce, SAP, Zendesk, Slack, and ServiceNow, and AWS services like Amazon S3 and Amazon Redshift, in just a few clicks. https://aws.amazon.com/appflow/

upvoted 2 times

KVK16 1 month, 1 week ago

Amazon AppFlow is a fully managed integration service that enables you to securely transfer data between Software-as-a-Service (SaaS) applications like Salesforce, Marketo, Slack, and ServiceNow, and AWS services like Amazon S3 and Amazon Redshift, in just a few clicks. With AppFlow, you can run data flows at nearly any scale at the frequency you choose - on a schedule, in response to a business event, or on demand.

upvoted 2 times

BoboChow 1 month, 2 weeks ago

Selected Answer: A

How EventBridge loads data to S3?

upvoted 2 times

BoboChow 1 month, 2 weeks ago

I chang to B
https://aws.amazon.com/appflow/?nc1=h_ls

upvoted 4 times

alvarez100 1 month, 1 week ago

Yes I agree it is B.

upvoted 2 times

Question #42 *Topic 1*

A company runs a highly available image-processing application on Amazon EC2 instances in a single VPC. The EC2 instances run inside several subnets across multiple Availability Zones. The EC2 instances do not communicate with each other. However, the EC2 instances download images from Amazon S3 and upload images to Amazon S3 through a single NAT gateway. The company is concerned about data transfer charges.
What is the MOST cost-effective way for the company to avoid Regional data transfer charges?

A. Launch the NAT gateway in each Availability Zone.

B. Replace the NAT gateway with a NAT instance.

C. Deploy a gateway VPC endpoint for Amazon S3.

D. Provision an EC2 Dedicated Host to run the EC2 instances.

Correct Answer: *C*

Community vote distribution

C (90%) 10%

Shasha1 4 days, 23 hours ago

C is correct
https://docs.aws.amazon.com/vpc/latest/privatelink/gateway-endpoints.html

upvoted 1 times

Wpcorgan 1 week ago

C is Correct

upvoted 1 times

justsaysid 2 weeks, 3 days ago

Selected Answer: C

This link clearly states that "VPC gateway endpoints allow communication to Amazon S3 and Amazon DynamoDB without incurring data transfer charges within the same Region". On the other hand NAT gateway incurs additional data processing charges. Hence, C is the correct answer.
https://aws.amazon.com/blogs/architecture/overview-of-data-transfer-costs-for-common-architectures/

upvoted 3 times

dduque10 3 weeks ago

Selected Answer: A

Why not A?

upvoted 1 times

TuLe 1 day, 10 hours ago

using the NAT gateway you will be charge for data transfer out. When VPC gateway endpoint in place for S3, the service will use internal route inside AWS to send data to S3 -> no charge at all.

upvoted 1 times

airraid2010 1 month ago

Selected Answer: C

C is the answer

upvoted 4 times

Jahangeer_17 1 month, 1 week ago

If we deploy VPC Gateway Endpoint then data will be transfer through AWS network only.

upvoted 2 times

KADSM 2 weeks, 6 days ago

Though will it not incur regional data transfer cost ? Here the question is to avoid regional data transfer costs

upvoted 1 times

Wajif 1 week, 1 day ago

Here it also says "The company is concerned about data transfer charges". They just want to reduce costs hence it is C.

upvoted 2 times

Rachness 1 month, 2 weeks ago

Selected Answer: C

Gateway Endpoint

upvoted 2 times

Lilibell 1 month, 2 weeks ago

The answer is C
upvoted 3 times

Question #43 *Topic 1*

A company has an on-premises application that generates a large amount of time-sensitive data that is backed up to Amazon S3. The application has grown and there are user complaints about internet bandwidth limitations. A solutions architect needs to design a long-term solution that allows for both timely backups to Amazon S3 and with minimal impact on internet connectivity for internal users.
Which solution meets these requirements?

A. Establish AWS VPN connections and proxy all traffic through a VPC gateway endpoint.

B. Establish a new AWS Direct Connect connection and direct backup traffic through this new connection.

C. Order daily AWS Snowball devices. Load the data onto the Snowball devices and return the devices to AWS each day.

D. Submit a support ticket through the AWS Management Console. Request the removal of S3 service limits from the account.

Correct Answer: *B*

Community vote distribution

B (100%)

Sinaneos Highly Voted 1 month, 2 weeks ago

Selected Answer: B
A: VPN also goes through the internet and uses the bandwidth
C: daily Snowball transfer is not really a long-term solution when it comes to cost and efficiency
D: S3 limits don't change anything here

So the answer is B
upvoted 9 times

Wpcorgan Most Recent 1 week ago

B is Correct
upvoted 1 times

airraid2010 1 month ago

Selected Answer: B
B is the answer
upvoted 4 times

KVK16 1 month, 1 week ago

AWS Direct Connect and AWS Snowball Edge are primarily classified as "Cloud Dedicated Network Connection" and "Data Transfer" tools respectively.

Even if we say it takes 1/5th of cost for transfer of 250 TB data from on-prem to AWS in a week.
upvoted 1 times

KVK16 1 month, 1 week ago

Direct Connect vs Snowball
upvoted 1 times

oxfordcommaa 1 month, 2 weeks ago

B.
The keyword here is long term solution.
Direct connect is a dedicated connection between on-prem and AWS, this is the way to ensure stable network connectivity that will not wax and wane like internet connectivity.
upvoted 2 times

Lilibell 1 month, 2 weeks ago

The answer is B
upvoted 1 times

Question #44 *Topic 1*

A company has an Amazon S3 bucket that contains critical data. The company must protect the data from accidental deletion.
Which combination of steps should a solutions architect take to meet these requirements? (Choose two.)

A. Enable versioning on the S3 bucket.

B. Enable MFA Delete on the S3 bucket.

C. Create a bucket policy on the S3 bucket.

D. Enable default encryption on the S3 bucket.

E. Create a lifecycle policy for the objects in the S3 bucket.

Correct Answer: *BD*

Community vote distribution

AB (100%)

Uhrien Highly Voted 1 month, 2 weeks ago

Selected Answer: AB

The correct solution is AB, as you can see here:

https://aws.amazon.com/it/premiumsupport/knowledge-center/s3-audit-deleted-missing-objects/

It states the following:

To prevent or mitigate future accidental deletions, consider the following features:

Enable versioning to keep historical versions of an object.
Enable Cross-Region Replication of objects.
Enable MFA delete to require multi-factor authentication (MFA) when deleting an object version.

upvoted 16 times

Wpcorgan Most Recent 1 week ago

A and B

upvoted 1 times

Solarch 3 weeks, 4 days ago

AB, Versioning keeps a copy and can be retrieved. MFA ensures you have proper authorization to delete an item.

upvoted 2 times

ricenguyen208 3 weeks, 5 days ago

AB for sure =)))))))))))))

upvoted 1 times

raffaello44 1 month ago

Selected Answer: AB

there is no relation between cncryption and delete protection. so D is discarded.

upvoted 2 times

ukwafabian 1 month ago

Selected Answer: AB

There is no way it is not A and B.

upvoted 1 times

airraid2010 1 month ago

Selected Answer: AB

https://docs.aws.amazon.com/AmazonS3/latest/userguide/Versioning.html
You can use the S3 Versioning feature to preserve, retrieve, and restore every version of every object stored in your buckets

https://docs.aws.amazon.com/AmazonS3/latest/userguide/MultiFactorAuthenticationDelete.html
When working with S3 Versioning in Amazon S3 buckets, you can optionally add another layer of security by configuring a bucket to enable MFA (multi-factor authentication) delete.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: AB

obviously AB here

upvoted 1 times

Jahangeer_17 1 month, 1 week ago

Correct answer is A & B.

upvoted 1 times

Incognito013 1 month, 1 week ago

The answers are A and B

upvoted 1 times

Rachness 1 month, 2 weeks ago

Selected Answer: AB

AB-version and MFA to ensure.

upvoted 1 times

rajendradba 1 month, 2 weeks ago

Selected Answer: AB

https://aws.amazon.com/premiumsupport/knowledge-center/s3-audit-deleted-missing-objects/

upvoted 1 times

tt79 1 month, 2 weeks ago

AB is correct.

upvoted 3 times

tuloveu 1 month, 2 weeks ago

Selected Answer: AB

A,B is my option.

upvoted 2 times

Testtest123321 1 month, 2 weeks ago

I BELIVE A,C IS CORRECT!

upvoted 2 times

CloudGuru99 1 month, 2 weeks ago

AB is the correct answer

upvoted 3 times

Question #45 *Topic 1*

A company has a data ingestion workflow that consists of the following:
• An Amazon Simple Notification Service (Amazon SNS) topic for notifications about new data deliveries
• An AWS Lambda function to process the data and record metadata
The company observes that the ingestion workflow fails occasionally because of network connectivity issues. When such a failure occurs, the Lambda function does not ingest the corresponding data unless the company manually reruns the job.
Which combination of actions should a solutions architect take to ensure that the Lambda function ingests all data in the future? (Choose two.)

A. Deploy the Lambda function in multiple Availability Zones.

B. Create an Amazon Simple Queue Service (Amazon SQS) queue, and subscribe it to the SNS topic.

C. Increase the CPU and memory that are allocated to the Lambda function.

D. Increase provisioned throughput for the Lambda function.

E. Modify the Lambda function to read from an Amazon Simple Queue Service (Amazon SQS) queue.

Correct Answer: *BE*

Community vote distribution

BE (89%) 11%

Incognito013 Highly Voted 1 month, 1 week ago
A, C, D options are out, since Lambda is fully managed service which provides high availability and scalability by its own

Answers are B and E
upvoted 6 times

Sinaneos Highly Voted 1 month, 2 weeks ago
Selected Answer: BE
BE so that the lambda function reads the SQS queue and nothing gets lost
upvoted 5 times

Wpcorgan Most Recent 1 week ago
B and E
upvoted 1 times

Six_Fingered_Jose 1 month ago
Selected Answer: BE
B and E is the obvious answer here,
SQS ensures that message does not get lost
upvoted 3 times

D2w 1 month, 1 week ago
Selected Answer: AB
Why not AB
upvoted 1 times

Six_Fingered_Jose 1 month ago
lambda is serverless, it does not need to be multi-AZ..
upvoted 1 times

Question #46 *Topic 1*

A company has an application that provides marketing services to stores. The services are based on previous purchases by store customers. The stores upload transaction data to the company through SFTP, and the data is processed and analyzed to generate new marketing offers. Some of the files can exceed 200 GB in size.
Recently, the company discovered that some of the stores have uploaded files that contain personally identifiable information (PII) that should not have been included. The company wants administrators to be alerted if PII is shared again. The company also wants to automate remediation.
What should a solutions architect do to meet these requirements with the LEAST development effort?

A. Use an Amazon S3 bucket as a secure transfer point. Use Amazon Inspector to scan the objects in the bucket. If objects contain PII, trigger an S3 Lifecycle policy to remove the objects that contain PII.

B. Use an Amazon S3 bucket as a secure transfer point. Use Amazon Macie to scan the objects in the bucket. If objects contain PII, use Amazon Simple Notification Service (Amazon SNS) to trigger a notification to the administrators to remove the objects that contain PII.

C. Implement custom scanning algorithms in an AWS Lambda function. Trigger the function when objects are loaded into the bucket. If objects contain PII, use Amazon Simple Notification Service (Amazon SNS) to trigger a notification to the administrators to remove the objects that contain PII.

D. Implement custom scanning algorithms in an AWS Lambda function. Trigger the function when objects are loaded into the bucket. If objects contain PII, use Amazon Simple Email Service (Amazon SES) to trigger a notification to the administrators and trigger an S3 Lifecycle policy to remove the meats that contain PII.

Correct Answer: *B*

Community vote distribution

B (100%)

Gatt Highly Voted 2 weeks, 1 day ago

I have a problem with answer B. The question says: "automate remediation". B says that you inform the administrator and he removes the data manually, that's not automating remediation. Very weird, that would mean that D is correct - but it's so much harder to implement.

upvoted 8 times

Wpcorgan Most Recent 1 week ago

B is correct

upvoted 1 times

kanweng 1 week, 6 days ago

Selected Answer: B

https://aws.amazon.com/blogs/security/use-macie-to-discover-sensitive-data-as-part-of-automated-data-pipelines/

upvoted 1 times

grzeev 2 weeks ago

Selected Answer: B

Amazon Macie is a data security and data privacy service that uses machine learning (ML) and pattern matching to discover and protect your sensitive data

upvoted 2 times

grzeev 2 weeks ago

Macie automatically detects a large and growing list of sensitive data types, including personally identifiable information (PII) such as names, addresses, and credit card numbers. It also gives you constant visibility of the data security and data privacy of your data stored in Amazon S3

upvoted 2 times

17Master 3 weeks, 4 days ago

Selected Answer: B

https://aws.amazon.com/es/macie/faq/

upvoted 1 times

yd_h 1 month, 1 week ago

FYI: Macie doesn't directly integrate with SNS.
https://docs.aws.amazon.com/macie/latest/user/findings-monitor.html
"To support integration with other services and systems, Macie publishes findings to Amazon EventBridge as finding events."

upvoted 2 times

17Master 3 weeks, 4 days ago

Here - https://aws.amazon.com/es/macie/faq/

upvoted 1 times

Incognito013 1 month, 1 week ago

Macie is the solution to detect PII data

Answer B

upvoted 2 times

ninjawrz 1 month, 1 week ago

Selected Answer: B

B amazon macie
https://docs.aws.amazon.com/macie/latest/user/what-is-macie.html

upvoted 1 times

masetromain 1 month, 2 weeks ago

Selected Answer: B

Good answer is B : Amazon Macie is key

https://aws.amazon.com/fr/macie/

upvoted 2 times

Question #47 *Topic 1*

A company needs guaranteed Amazon EC2 capacity in three specific Availability Zones in a specific AWS Region for an upcoming event that will last 1 week.
What should the company do to guarantee the EC2 capacity?

A. Purchase Reserved Instances that specify the Region needed.

B. Create an On-Demand Capacity Reservation that specifies the Region needed.

C. Purchase Reserved Instances that specify the Region and three Availability Zones needed.

D. Create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones needed.

Correct Answer: *D*

Community vote distribution

D (100%)

Wajif 2 days, 9 hours ago

Selected Answer: D

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html#capacity-reservations-differences

upvoted 1 times

Wpcorgan 1 week ago

D is correct

upvoted 1 times

koreanmonkey 1 week, 3 days ago

Selected Answer: D

Absolutely D

upvoted 1 times

xeun88 1 week, 5 days ago

D is the correct answer

upvoted 1 times

MyNameIsJulien 2 weeks, 2 days ago

Selected Answer: D

Ans D for sure

upvoted 1 times

17Master 3 weeks, 4 days ago

Selected Answer: D

D. Create an On-Demand Capacity Reservation that specifies the Region and three Availability Zones needed

upvoted 1 times

Incognito013 1 month, 1 week ago

Reserved instances are for long term so on-demand will be the right choice - Answer D

upvoted 4 times

ninjawrz 1 month, 1 week ago

Selected Answer: D

Reserve instances: You will have to pay for the whole term (1 year or 3years) which is not cost effective
So answer is
D: on demand capacity region

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-capacity-reservations.html

upvoted 2 times

KVK16 1 month, 1 week ago

Selected Answer: D

on-demand Capacity reservation for a specific AZ for gamedays

upvoted 1 times

Question #48 *Topic 1*

A company's website uses an Amazon EC2 instance store for its catalog of items. The company wants to make sure that the catalog is highly available and that the catalog is stored in a durable location.
What should a solutions architect do to meet these requirements?

A. Move the catalog to Amazon ElastiCache for Redis.

B. Deploy a larger EC2 instance with a larger instance store.

C. Move the catalog from the instance store to Amazon S3 Glacier Deep Archive.

D. Move the catalog to an Amazon Elastic File System (Amazon EFS) file system.

Correct Answer: *A*

Community vote distribution

D (93%) 7%

Six_Fingered_Jose Highly Voted 1 month ago

Selected Answer: D

keyword is "durable" location
A and B is ephemeral storage
C takes forever so is not HA,
that leaves D

upvoted 6 times

rajendradba Highly Voted 1 month, 2 weeks ago

Selected Answer: D

Elasticache is in Memory, EFS is for durability

upvoted 6 times

Wpcorgan Most Recent 6 days, 23 hours ago

D is Correct

upvoted 1 times

renekton 1 week, 2 days ago

Selected Answer: A

Answer is A.
The company is using instance store which is ephemeral storage and they're looking for highly available storage.

ElastiCache for Redis is provides high availability and provides automatic failover by detecting primary node failures and promoting a replica to be primary with minimal impact.
https://aws.amazon.com/elasticache/redis/

upvoted 2 times

mabotega 3 weeks ago

Hello, answer A is correct with Redis AOF enabled (it is not by default), it will warmup the cache if one node would reboot.

https://docs.aws.amazon.com/AmazonElastiCache/latest/red-ug/RedisAOF.html

upvoted 2 times

17Master 3 weeks, 4 days ago

Selected Answer: D

But, only Linux - https://aws.amazon.com/getting-started/tutorials/create-network-file-system/?nc1=h_ls

upvoted 1 times

Incognito013 1 month, 1 week ago

instance store is not durable, if it goes down then instance store data is lost.
EFS is the only option here that will provide high availability and durability, plus it can be accessed by multiple instances at the same time

upvoted 4 times

Sinaneos 1 month, 2 weeks ago

Selected Answer: D

EFS for sure

upvoted 4 times

SAA003 1 month, 2 weeks ago

D for EFS

upvoted 3 times

MXB05 1 month, 2 weeks ago

Selected Answer: D

EFS is the only durable solution here.

upvoted 4 times

hanhdroid 1 month, 2 weeks ago

Selected Answer: D

I would say D for EFS

upvoted 5 times

Question #49 *Topic 1*

A company stores call transcript files on a monthly basis. Users access the files randomly within 1 year of the call, but users access the files infrequently after 1 year. The company wants to optimize its solution by giving users the ability to query and retrieve files that are less than 1-year-old as quickly as possible. A delay in retrieving older files is acceptable.
Which solution will meet these requirements MOST cost-effectively?

A. Store individual files with tags in Amazon S3 Glacier Instant Retrieval. Query the tags to retrieve the files from S3 Glacier Instant Retrieval.

B. Store individual files in Amazon S3 Intelligent-Tiering. Use S3 Lifecycle policies to move the files to S3 Glacier Flexible Retrieval after 1 year. Query and retrieve the files that are in Amazon S3 by using Amazon Athena. Query and retrieve the files that are in S3 Glacier by using S3 Glacier Select.

C. Store individual files with tags in Amazon S3 Standard storage. Store search metadata for each archive in Amazon S3 Standard storage. Use S3 Lifecycle policies to move the files to S3 Glacier Instant Retrieval after 1 year. Query and retrieve the files by searching for metadata from Amazon S3.

D. Store individual files in Amazon S3 Standard storage. Use S3 Lifecycle policies to move the files to S3 Glacier Deep Archive after 1 year. Store search metadata in Amazon RDS. Query the files from Amazon RDS. Retrieve the files from S3 Glacier Deep Archive.

Correct Answer: *C*

Community vote distribution

B (74%) 13% 13%

masetromain Highly Voted 1 month, 2 weeks ago

Selected Answer: B

I think the answer is B:
Users access the files randomly

S3 Intelligent-Tiering is the ideal storage class for data with unknown, changing, or unpredictable access patterns, independent of object size or retention period. You can use S3 Intelligent-Tiering as the default storage class for virtually any workload, especially data lakes, data analytics, new applications, and user-generated content.

https://aws.amazon.com/fr/s3/storage-classes/intelligent-tiering/

upvoted 12 times

Lilibell Highly Voted 1 month, 2 weeks ago

The answer is B

upvoted 6 times

Wajif Most Recent 2 days, 9 hours ago

Selected Answer: B

Between B and C I will choose B because it gives a way to query the file data using Athena. There is no way given in C to query the file but just the ability to retreive.

upvoted 1 times

ocbn3wby 6 days, 19 hours ago

Selected Answer: C

https://aws.amazon.com/s3/pricing/

A: S3 Glacier Instant Retrieval*** - For long-lived archive data accessed once a quarter with instant retrieval in milliseconds

While the users access the files randomly (not specified once/quarter) - Answer A does not fit.

B: When we are talking about S3 Intelligent Tiering - this is the most cost effective. However, Glacier Select defeats the purpose: it's focused on extracting data FROM a bigger file and not the file itself.
https://aws.amazon.com/blogs/aws/s3-glacier-select/
D: Involves RDS, increases the cost. Not much difference between Deep archive and instant retrieval (D vs C).

I'd go for C.

upvoted 1 times

Wpcorgan 6 days, 23 hours ago

B is correct

upvoted 1 times

pepgua 1 week, 1 day ago

B is correct because it's the most cost effective solution overall for random access (S3 Intelligent-tiering) for first year AND cheaper option for S3 Glacier storage classes after one year (Glacier-flexible is cheaper than Glacier Instant)

upvoted 1 times

hollie 1 week, 3 days ago

I don't think S3 Glacier can do metadata search. Thus B is more reasonable.

upvoted 1 times

grzeev 2 weeks ago

Selected Answer: B

S3 Intelligent-Tiering is the first cloud object storage class that delivers automatic cost savings by moving data between two access tiers — frequent access and infrequent access — when access patterns change, and is ideal for data with unknown or changing access patterns

upvoted 1 times

Adrianjavier 2 weeks, 1 day ago

Selected Answer: B

B for sure

upvoted 1 times

AbhiJo 2 weeks, 1 day ago

Flexible tiering wins

upvoted 2 times

mabotega 2 weeks, 2 days ago

Selected Answer: B

https://aws.amazon.com/pt/blogs/aws/s3-glacier-select/

upvoted 1 times

KADSM 2 weeks, 6 days ago

Answer should be C only.: Reasons:
File should be accessed as quickly as possible within 1 year. S3 Standard may be the choice. Overall solution should be cost efficient
A -> Uses Glacier Instant retreival but the put/get costs are more than S3 standard.
B -> Uses Intelligent tiering which may have overhead cost for intelligently doing the tiers. Glacier flexible retreival may have additional costs
D -> Ruled out - they are using RDS which may incur additional costs compared to other options.

upvoted 1 times

17Master 3 weeks, 4 days ago

Selected Answer: B

S3 Intelligent-Tiering + Glacier Flexible Retrieval is more cost-effective.

upvoted 1 times

SimonPark 1 month ago

The answer looks B

upvoted 1 times

BoboChow 1 month, 1 week ago

Selected Answer: C

A: you need to distinguish the storage within one year.
B:Select is not an archive retrieval option
C:seems like right
D:Querying the files from Amazon RDS is not a good choice

upvoted 2 times

BoboChow 1 month ago

I changed my mind.
Compared to Glacier Instant Retrieval, Glacier Flexible Retrieval is more cost-effective.

upvoted 1 times

Kikiokiki 1 month, 1 week ago

NOTE:"A delay in retrieving older files is acceptable."
"For archive data that needs immediate access, such as medical images, news media assets, or genomics data, choose the S3 Glacier Instant Retrieval storage class, an archive storage class that delivers the lowest cost storage with milliseconds retrieval. For archive data that does not require immediate access but needs the flexibility to retrieve large sets of data at no cost, such as backup or disaster recovery use cases, choose S3 Glacier Flexible Retrieval (formerly S3 Glacier), with retrieval in minutes or free bulk retrievals in 5-12 hours."

https://aws.amazon.com/about-aws/whats-new/2021/11/amazon-s3-glacier-instant-retrieval-storage-class/

upvoted 3 times

123jhl0 1 month, 1 week ago

Selected Answer: A

I think this is a bit tricky, as you have to pay attention to the sum of costs related at each answer:
A: S3 Glacier Instant Retrieval for all storage (>1yo & +1yo)
B: S3 Intelligent Tiering (>1yo storage) + S3 Glacier Flexible Retrieval (+1yo storage) + Amazon Athena + S3 Glacier Select !!!
C: S3 Standard (>1yo storage) + S3 Glacier Instant Retrieval (+1yo storage)
D: S3 Standard (>1yo storage) + S3 Glacier Deep Archive (+1yo storage) + Amazon RDS

Based on this, I'd choose A, but not sure I can retrieve files by tags... If it's not possible, then the cheapest after A is C.
upvoted 3 times

ocbn3wby 6 days, 19 hours ago

https://aws.amazon.com/s3/pricing/

A: S3 Glacier Instant Retrieval*** - For long-lived archive data accessed once a quarter with instant retrieval in milliseconds

While the users access the files randomly (not specified once/quarter) - Answer A does not fit.

B: When we are talking about S3 Intelligent Tiering - this is the most cost effective. However, Glacier Select defeats the purpose: it's focused on extracting data FROM a bigger file and not the file itself.
https://aws.amazon.com/blogs/aws/s3-glacier-select/
D: Involves RDS, increases the cost. Not much difference between Deep archive and instant retrieval (D vs C).

I'd go for C.
upvoted 1 times

Question #50 *Topic 1*

A company has a production workload that runs on 1,000 Amazon EC2 Linux instances. The workload is powered by third-party software. The company needs to patch the third-party software on all EC2 instances as quickly as possible to remediate a critical security vulnerability.
What should a solutions architect do to meet these requirements?

A. Create an AWS Lambda function to apply the patch to all EC2 instances.

B. Configure AWS Systems Manager Patch Manager to apply the patch to all EC2 instances.

C. Schedule an AWS Systems Manager maintenance window to apply the patch to all EC2 instances.

D. Use AWS Systems Manager Run Command to run a custom command that applies the patch to all EC2 instances.

Correct Answer: *D*

Community vote distribution

D (88%) 13%

brushek Highly Voted 1 month, 1 week ago

Selected Answer: D

It can't be B !
Look: https://docs.aws.amazon.com/systems-manager/latest/userguide/patch-manager-how-it-works-installation.html

and You need to patch third-party software ! It doesn't apply for yum,apt or dnf.

upvoted 5 times

Wpcorgan Most Recent 6 days, 23 hours ago

D is correct

upvoted 1 times

koreanmonkey 1 week, 3 days ago

Selected Answer: D

Because system use third-party software, it needs custom command. D is right.

upvoted 1 times

tinyfoot 2 weeks, 1 day ago

The primary focus of Patch Manager, a capability of AWS Systems Manager, is on installing operating systems security-related updates on managed nodes. By default, Patch Manager doesn't install all available patches, but rather a smaller set of patches focused on security. (Ref https://docs.aws.amazon.com/systems-manager/latest/userguide/patch-manager-how-it-works-selection.html)

Run Command allows you to automate common administrative tasks and perform one-time configuration changes at scale. (Ref https://docs.aws.amazon.com/systems-manager/latest/userguide/execute-remote-commands.html)

Seems like patch manager is meant for OS level patches and not 3rd party applications. And this falls under run command wheelhouse to carry out one-time configuration changes (update of 3rd part application) at scale.

upvoted 2 times

EKA_CloudGod 2 weeks, 5 days ago

I was torn between B and D ad after reviewing docs, I choose D, and here is why;
"For Linux-based operating system types that report a severity level for patches, Patch Manager uses the severity level reported by the software publisher for the update notice or individual patch. Patch Manager doesn't derive severity levels from third-party sources, such as the Common Vulnerability Scoring System (CVSS), or from metrics released by the National Vulnerability Database (NVD)."
https://docs.aws.amazon.com/systems-manager/latest/userguide/systems-manager-patch.html

upvoted 4 times

17Master 3 weeks, 4 days ago

Selected Answer: D

https://docs.aws.amazon.com/es_es/systems-manager/latest/userguide/execute-remote-commands.html

upvoted 3 times

Six_Fingered_Jose 1 month ago

Selected Answer: D

agree with D here

upvoted 2 times

dave9994 1 month ago

C should be the answer.

upvoted 1 times

Six_Fingered_Jose 1 month ago

question says as quickly as possible, C is not the best solution here

upvoted 1 times

dave9994 1 month ago

based on this : https://docs.aws.amazon.com/systems-manager/latest/userguide/systems-manager-maintenance.html

upvoted 1 times

17Master 3 weeks, 4 days ago

in your link it says: Commands in Run Command, a capability of Systems Manager - For more information about Run Command, see AWS Systems Manager Run Command. ANS is D.
:)

upvoted 2 times

ArielSchivo 1 month, 1 week ago

Selected Answer: D

It's not B for sure since you have to path 3rd party software and Systems Manager Patch Manager just works for OS owned software. I would go for D.

upvoted 3 times

KVK16 1 month, 1 week ago

Patch Manager: It automates the process of patching managed nodes with both security-related and other types of updates.

upvoted 1 times

FFORTUNATE 1 month, 1 week ago

its B.

https://docs.aws.amazon.com/systems-manager/latest/userguide/systems-manager-patch.html

upvoted 1 times

ninjawrz 1 month, 1 week ago

Looks B to me

upvoted 1 times

masetromain 1 month, 2 weeks ago

Selected Answer: B

I think the answer is B, but I'm not sure.
https://docs.aws.amazon.com/systems-manager/latest/userguide/systems-manager-patch.html

upvoted 2 times

Sinaneos 1 month ago

if you are not sure, pls don't vote, we need reliable answers in the website.
Reason B is wrong is that the systems manager patch operations are made for OS-related software, the question specifies "3rd party software", therefore we need a custom command, given in D

upvoted 7 times

pooppants 1 month, 2 weeks ago

This is a tough one. The only thing i can find is https://docs.aws.amazon.com/systems-manager/latest/userguide/about-windows-app-patching.html which supports answer D.

upvoted 4 times

Question #51 *Topic 1*

A company is developing an application that provides order shipping statistics for retrieval by a REST API. The company wants to extract the shipping statistics, organize the data into an easy-to-read HTML format, and send the report to several email addresses at the same time every morning.
Which combination of steps should a solutions architect take to meet these requirements? (Choose two.)

A. Configure the application to send the data to Amazon Kinesis Data Firehose.

B. Use Amazon Simple Email Service (Amazon SES) to format the data and to send the report by email.

C. Create an Amazon EventBridge (Amazon CloudWatch Events) scheduled event that invokes an AWS Glue job to query the application's API for the data.

D. Create an Amazon EventBridge (Amazon CloudWatch Events) scheduled event that invokes an AWS Lambda function to query the application's API for the data.

E. Store the application data in Amazon S3. Create an Amazon Simple Notification Service (Amazon SNS) topic as an S3 event destination to send the report by email.

Correct Answer: *DE*

Community vote distribution

BD (65%) DE (23%) 12%

whosawsome Highly Voted 1 month, 1 week ago

Selected Answer: BD

You can use SES to format the report in HTML.
https://docs.aws.amazon.com/ses/latest/dg/send-email-formatted.html

upvoted 12 times

Wpcorgan Most Recent 6 days, 23 hours ago

B and D

upvoted 1 times

pepgua 1 week, 1 day ago

We seem to agree option D is correct. The second choice is between B (SES) and E (SNS). SES is the best answer as it's specifically designed for Email service. SNS can also deliver notifications via email but it's not designed for that HTML format. BD is correct.
https://stackoverflow.com/questions/32241928/sending-html-content-in-aws-snssimple-notification-service-emails-notification

upvoted 1 times

baiy 3 weeks, 1 day ago

Selected Answer: BD

Not C because there is no direct connector available for Glue to connect to the internet world (REST API), you can set up a VPC, with a public and a private subnet.

upvoted 2 times

backbencher2022 3 weeks, 3 days ago

Selected Answer: BD

B&D are the only 2 correct options. If you are choosing option E then you missed the daily morning schedule requirement mentioned in the question which cant be achieved with S3 events for SNS. Event Bridge can used to configure scheduled events (every morning in this case). Option B fulfills the email in HTML format requirement (by SES) and D fulfills every morning schedule event requirement (by EventBridge)

upvoted 3 times

santosh2316 3 weeks, 3 days ago

Selected Answer: DE

You can use SES to format the report in HTML.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: BC

i think BC or CE
Glue should be the winner here with for ETL
i personally think SES is the winner here too because of this
https://docs.aws.amazon.com/ses/latest/dg/send-email-formatted.html

upvoted 3 times

BoboChow 1 month, 1 week ago

Selected Answer: DE

About sending the report to several email addresses, I would rather go for SNS than SES;
About extracting the shipping statistics, I would rather go for Lambda than Glue;

upvoted 2 times

KVK16 1 month, 1 week ago

Selected Answer: DE

Glue cannot output HTML Report, you'll need a Lambda. Also there is need to invoke the Lambda at a specific time - CRON Job - Thumbrule - Cloudwatch Events + Lambda. The reports are stored in S3 and SNS is invoke to email the reports to same selected emails everyday.

SNS vs SES -
SES Cannot extract data and it sends the emails without the consent of the receivers only take verification of the sender
Schedule Lambda to

upvoted 3 times

Newptone 2 weeks, 6 days ago

You didn't check SNS limitation: By contrast, Amazon Simple Notification Service (Amazon SNS) is for messaging-oriented applications, the body of an Amazon SNS notification is limited to 8192 characters of UTF-8 strings, and isn't intended to support multimedia content.
It does not support "easy-to-read HTML format" email.
https://aws.amazon.com/ses/faqs/#Amazon_SES_and_Other_AWS_Services

upvoted 2 times

FFORTUNATE 1 month, 1 week ago

C.E.
Glue job is good for ETL (extract transform load).

upvoted 2 times

Question #52 *Topic 1*

A company wants to migrate its on-premises application to AWS. The application produces output files that vary in size from tens of gigabytes to hundreds of terabytes. The application data must be stored in a standard file system structure. The company wants a solution that scales automatically. is highly available, and requires minimum operational overhead.
Which solution will meet these requirements?

A. Migrate the application to run as containers on Amazon Elastic Container Service (Amazon ECS). Use Amazon S3 for storage.

B. Migrate the application to run as containers on Amazon Elastic Kubernetes Service (Amazon EKS). Use Amazon Elastic Block Store (Amazon EBS) for storage.

C. Migrate the application to Amazon EC2 instances in a Multi-AZ Auto Scaling group. Use Amazon Elastic File System (Amazon EFS) for storage.

D. Migrate the application to Amazon EC2 instances in a Multi-AZ Auto Scaling group. Use Amazon Elastic Block Store (Amazon EBS) for storage.

Correct Answer: *C*

Community vote distribution

C (100%)

oxfordcommaa Highly Voted 1 month, 2 weeks ago

Selected Answer: C

chose this due to the key word "standard file system"

upvoted 5 times

masetromain Highly Voted 1 month, 2 weeks ago

I have absolutely no idea...

Output files that vary in size from tens of gigabytes to hundreds of terabytes

Simit size for a single object:
S3 5To TiB
https://aws.amazon.com/fr/blogs/aws/amazon-s3-object-size-limit/
EBS 64 Tib
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/volume_constraints.html
EFS 47.9 TiB
https://docs.aws.amazon.com/efs/latest/ug/limits.html

upvoted 5 times

Wpcorgan Most Recent 6 days, 23 hours ago

C is correct

upvoted 1 times

pspinelli19 2 weeks, 4 days ago

Selected Answer: C

https://www.examtopics.com/discussions/amazon/view/84147-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

BoboChow 1 month, 1 week ago

Selected Answer: C

standard file system => EFS rather than S3

upvoted 2 times

Kikiokiki 1 month, 1 week ago

EBS doesn't offer high availability, data is stored in one AZ.

upvoted 1 times

ArielSchivo 1 month, 1 week ago

Selected Answer: C

EFS is a standard file system, it scales automatically and is highly available.

upvoted 4 times

queen101 1 month, 1 week ago

cCCCCCCCCCC

upvoted 1 times

Question #53 *Topic 1*

A company needs to store its accounting records in Amazon S3. The records must be immediately accessible for 1 year and then must be archived for an additional 9 years. No one at the company, including administrative users and root users, can be able to delete the records during the entire 10-year period. The records must be stored with maximum resiliency.
Which solution will meet these requirements?

A. Store the records in S3 Glacier for the entire 10-year period. Use an access control policy to deny deletion of the records for a period of 10 years.

B. Store the records by using S3 Intelligent-Tiering. Use an IAM policy to deny deletion of the records. After 10 years, change the IAM policy to allow deletion.

C. Use an S3 Lifecycle policy to transition the records from S3 Standard to S3 Glacier Deep Archive after 1 year. Use S3 Object Lock in compliance mode for a period of 10 years.

D. Use an S3 Lifecycle policy to transition the records from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) after 1 year. Use S3 Object Lock in governance mode for a period of 10 years.

Correct Answer: *C*

Community vote distribution

C (100%)

Wpcorgan 6 days, 23 hours ago
C is correct
upvoted 1 times

airraid2010 3 weeks, 1 day ago
Selected Answer: C
A-Wrong as the records must be immediately accessble for the first year.
B-The question never mentioned about the records can be deleted or modified after 10-year period.
D-It does not fulfill the condition of securing resiliency; you need multi-AZ to guarantee it.

Therefore, the answer is C.
upvoted 1 times

17Master 3 weeks, 4 days ago
Selected Answer: C
ans is C
upvoted 1 times

BoboChow 1 month, 1 week ago
Selected Answer: C
sure for C
upvoted 1 times

queen101 1 month, 1 week ago
CCCCCCCCC
upvoted 1 times

ninjawrz 1 month, 1 week ago
Selected Answer: C
This is C
upvoted 1 times

Rachness 1 month, 1 week ago
Selected Answer: C
compliance lock cant be removed unlike governance
upvoted 1 times

Question #54 *Topic 1*

A company runs multiple Windows workloads on AWS. The company's employees use Windows file shares that are hosted on two Amazon EC2 instances. The file shares synchronize data between themselves and maintain duplicate copies. The company wants a highly available and durable storage solution that preserves how users currently access the files.
What should a solutions architect do to meet these requirements?

A. Migrate all the data to Amazon S3. Set up IAM authentication for users to access files.

B. Set up an Amazon S3 File Gateway. Mount the S3 File Gateway on the existing EC2 instances.

C. Extend the file share environment to Amazon FSx for Windows File Server with a Multi-AZ configuration. Migrate all the data to FSx for Windows File Server.

D. Extend the file share environment to Amazon Elastic File System (Amazon EFS) with a Multi-AZ configuration. Migrate all the data to Amazon EFS.

Correct Answer: *C*

Community vote distribution

C (100%)

Wpcorgan 6 days, 23 hours ago

C is correct

upvoted 1 times

Juhith 1 week, 2 days ago

Selected Answer: C

EFS is only for Linux.

upvoted 1 times

koreanmonkey 1 week, 3 days ago

Selected Answer: C

EFS is only for Linux.

upvoted 1 times

Buruguduystunstugudunstuy 1 week, 4 days ago

Selected Answer: C

Windows file shares = Amazon FSx for Windows File Server
Hence, the correct answer is C

upvoted 1 times

17Master 3 weeks, 3 days ago

Selected Answer: C

FSX---> SMB

upvoted 2 times

cark0728 1 month ago

Selected Answer: C

C가 올바릅니다

upvoted 3 times

KVK16 1 month, 1 week ago

Selected Answer: C

FSx- Windows File Share https://docs.aws.amazon.com/fsx/latest/WindowsGuide/managing-file-shares.html

upvoted 2 times

Question #55 *Topic 1*

A solutions architect is developing a VPC architecture that includes multiple subnets. The architecture will host applications that use Amazon EC2 instances and Amazon RDS DB instances. The architecture consists of six subnets in two Availability Zones. Each Availability Zone includes a public subnet, a private subnet, and a dedicated subnet for databases. Only EC2 instances that run in the private subnets can have access to the RDS databases.
Which solution will meet these requirements?

A. Create a new route table that excludes the route to the public subnets' CIDR blocks. Associate the route table with the database subnets.

B. Create a security group that denies inbound traffic from the security group that is assigned to instances in the public subnets. Attach the security group to the DB instances.

C. Create a security group that allows inbound traffic from the security group that is assigned to instances in the private subnets. Attach the security group to the DB instances.

D. Create a new peering connection between the public subnets and the private subnets. Create a different peering connection between the private subnets and the database subnets.

Correct Answer: *C*

Community vote distribution

C (100%)

Sinaneos Highly Voted 1 month, 2 weeks ago

Selected Answer: C

A: doesn't fully configure the traffic flow
B: security groups don't have deny rules
D: peering is mostly between VPCs, doesn't really help here

answer is C, most mainstream way

upvoted 10 times

ashish_t Most Recent 2 days, 20 hours ago

Selected Answer: C

The key is "Only EC2 instances that run in the private subnets can have access to the RDS databases"
The answer is C.

upvoted 1 times

Wpcorgan 6 days, 23 hours ago

C is correct

upvoted 1 times

17Master 3 weeks, 3 days ago

Selected Answer: C

Ans correct.

upvoted 2 times

KVK16 1 month, 1 week ago

Selected Answer: C

Inside a VPC, traffic locally between different subnets cannot be restricted by routing but incase they are in different VPCs then it would be possible. This is imp Gain in VPC
- So only method is Security Groups - like EC2 also RDS also has Security Groups to restrict traffic to database instances

upvoted 4 times

Question #56 *Topic 1*

A company has registered its domain name with Amazon Route 53. The company uses Amazon API Gateway in the ca-central-1 Region as a public interface for its backend microservice APIs. Third-party services consume the APIs securely. The company wants to design its API Gateway URL with the company's domain name and corresponding certificate so that the third-party services can use HTTPS.
Which solution will meet these requirements?

A. Create stage variables in API Gateway with Name="Endpoint-URL" and Value="Company Domain Name" to overwrite the default URL. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM).

B. Create Route 53 DNS records with the company's domain name. Point the alias record to the Regional API Gateway stage endpoint. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the us-east-1 Region.

C. Create a Regional API Gateway endpoint. Associate the API Gateway endpoint with the company's domain name. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the same Region. Attach the certificate to the API Gateway endpoint. Configure Route 53 to route traffic to the API Gateway endpoint.

D. Create a Regional API Gateway endpoint. Associate the API Gateway endpoint with the company's domain name. Import the public certificate associated with the company's domain name into AWS Certificate Manager (ACM) in the us-east-1 Region. Attach the certificate to the API Gateway APIs. Create Route 53 DNS records with the company's domain name. Point an A record to the company's domain name.

Correct Answer: *D*

Community vote distribution

C (94%) 6%

masetromain Highly Voted 1 month, 2 weeks ago

Selected Answer: C

I think the answer is C. we don't need to attach a certificate in us-east-1, if is not for cloudfront. In our case the target is ca-central-1.

upvoted 11 times

Valero_ 1 month, 1 week ago

I think that is C too, the target would be the same Region.
https://docs.aws.amazon.com/apigateway/latest/developerguide/apigateway-regional-api-custom-domain-create.html

upvoted 4 times

ashish_t Most Recent 2 days, 20 hours ago

Selected Answer: C

Cert should be in the same region.
Answer: C

upvoted 1 times

Vesperia 2 days, 23 hours ago

Selected Answer: D

I choose D since the company wants its own domain name - should not be a regional one. Even though the answer does not mention edge-optimized custom domain name, this setup has to use it.

upvoted 1 times

Vesperia 2 days, 23 hours ago

The only correct answer is D since the company wants to design its API Gateway URL with the company's domain name. Answer C supports only regional domain name.

upvoted 1 times

Wpcorgan 6 days, 23 hours ago

C is correct

upvoted 1 times

study_aws1 2 weeks, 2 days ago

Will change my earlier selection to C). Reason -
• If using Edge-Optimized endpoint, then the certificate must be in us-east-1
• If using Regional endpoint, the certificate must be in the API Gateway region

upvoted 1 times

Nirmal3331 2 weeks, 3 days ago

Answer is C:

Regional custom domain names must use an SSL/TLS certificate that's in the same AWS Region as your API.

Edge-optimized custom domain names must use a certificate that's in the following Region: US East (N. Virginia) (us-east-1)./

upvoted 1 times

luvincanada 2 weeks, 4 days ago

The question states..company uses Amazon API Gateway in the ca-central-1 Region. Answer D mentions region name as "us-east-1" Region. which does not match. Therefore C is the correct answer.

upvoted 1 times

17Master 3 weeks, 3 days ago

Selected Answer: C

same region

upvoted 1 times

study_aws1 4 weeks, 1 day ago

Nope option B) seems good. The below link will explain.
https://docs.amazonaws.cn/en_us/apigateway/latest/developerguide/how-to-custom-domains-prerequisites.html

Amazon API Gateway leverages Amazon CloudFront to support certificates for custom domain names. As such, the requirements and constraints of a custom domain name SSL/TLS certificate are dictated by CloudFront.

upvoted 1 times

Athena 3 weeks, 4 days ago

The certificate must be in the same region. as the API. Answer is C

upvoted 1 times

[Removed] 4 weeks, 1 day ago

Nope the below links will explain why it is option B)

https://docs.amazonaws.cn/en_us/apigateway/latest/developerguide/how-to-custom-domains-prerequisites.html

Amazon API Gateway leverages Amazon CloudFront to support certificates for custom domain names. As such, the requirements and constraints of a custom domain name SSL/TLS certificate are dictated by CloudFront.

upvoted 1 times

Nightducky 3 weeks ago

Read the question again. Cert has to be in the same region as API.

upvoted 1 times

BoboChow 1 month, 1 week ago

Selected Answer: C

The certificate must be in the same Region as the API.

upvoted 4 times

Question #57 *Topic 1*

A company is running a popular social media website. The website gives users the ability to upload images to share with other users. The company wants to make sure that the images do not contain inappropriate content. The company needs a solution that minimizes development effort.
What should a solutions architect do to meet these requirements?

A. Use Amazon Comprehend to detect inappropriate content. Use human review for low-confidence predictions.

B. Use Amazon Rekognition to detect inappropriate content. Use human review for low-confidence predictions.

C. Use Amazon SageMaker to detect inappropriate content. Use ground truth to label low-confidence predictions.

D. Use AWS Fargate to deploy a custom machine learning model to detect inappropriate content. Use ground truth to label low-confidence predictions.

Correct Answer: *B*

Community vote distribution

B (100%)

masetromain Highly Voted 1 month, 2 weeks ago

Selected Answer: B

Good Answer is B :
https://docs.aws.amazon.com/rekognition/latest/dg/moderation.html?pg=ln&sec=ft

upvoted 7 times

Wpcorgan Most Recent 6 days, 23 hours ago

B is correct

upvoted 1 times

ArielSchivo 1 month, 1 week ago

Selected Answer: B

Option B.

https://docs.aws.amazon.com/rekognition/latest/dg/a2i-rekognition.html

upvoted 1 times

Question #58 *Topic 1*

A company wants to run its critical applications in containers to meet requirements for scalability and availability. The company prefers to focus on maintenance of the critical applications. The company does not want to be responsible for provisioning and managing the underlying infrastructure that runs the containerized workload.
What should a solutions architect do to meet these requirements?

A. Use Amazon EC2 instances, and install Docker on the instances.

B. Use Amazon Elastic Container Service (Amazon ECS) on Amazon EC2 worker nodes.

C. Use Amazon Elastic Container Service (Amazon ECS) on AWS Fargate.

D. Use Amazon EC2 instances from an Amazon Elastic Container Service (Amazon ECS)-optimized Amazon Machine Image (AMI).

Correct Answer: *C*

Community vote distribution

C (100%)

masetromain Highly Voted 1 month, 2 weeks ago

Selected Answer: C

Good answer is C:
AWS Fargate is a serverless, pay-as-you-go compute engine that lets you focus on building applications without having to manage servers. AWS Fargate is compatible with Amazon Elastic Container Service (ECS) and Amazon Elastic Kubernetes Service (EKS).

https://aws.amazon.com/fr/fargate/

upvoted 6 times

ashish_t Most Recent 2 days, 20 hours ago

Selected Answer: C

The company does not want to be responsible for provisioning and managing the underlying infrastructure that runs the containerized workload. Fargate is serverless and no need to manage.
Answer: C

upvoted 1 times

Wpcorgan 6 days, 23 hours ago

C is correct

upvoted 1 times

PS_R 3 weeks ago

Selected Answer: C

Agree Serverless Containerization Think Fargate

upvoted 2 times

ArielSchivo 1 month, 1 week ago

Selected Answer: C

Option C. Fargate is serverless, no need to manage the underlying infrastructure.

upvoted 3 times

Question #59 *Topic 1*

A company hosts more than 300 global websites and applications. The company requires a platform to analyze more than 30 TB of clickstream data each day.
What should a solutions architect do to transmit and process the clickstream data?

A. Design an AWS Data Pipeline to archive the data to an Amazon S3 bucket and run an Amazon EMR cluster with the data to generate analytics.

B. Create an Auto Scaling group of Amazon EC2 instances to process the data and send it to an Amazon S3 data lake for Amazon Redshift to use for analysis.

C. Cache the data to Amazon CloudFront. Store the data in an Amazon S3 bucket. When an object is added to the S3 bucket. run an AWS Lambda function to process the data for analysis.

D. Collect the data from Amazon Kinesis Data Streams. Use Amazon Kinesis Data Firehose to transmit the data to an Amazon S3 data lake. Load the data in Amazon Redshift for analysis.

Correct Answer: *D*

Community vote distribution

D (100%)

ArielSchivo Highly Voted 1 month, 1 week ago

Selected Answer: D

Option D.

https://aws.amazon.com/es/blogs/big-data/real-time-analytics-with-amazon-redshift-streaming-ingestion/

upvoted 8 times

Wpcorgan Most Recent 6 days, 23 hours ago

D is correct

upvoted 1 times

PS_R 3 weeks ago

Click Stream & Analyse/ process- Think KDS,

upvoted 2 times

BoboChow 1 month, 1 week ago

Selected Answer: D

D seems to make sense

upvoted 4 times

JesseeS 1 month, 1 week ago

Option D is correct... See the resource. Thank you Ariel

upvoted 1 times

Question #60 *Topic 1*

A company has a website hosted on AWS. The website is behind an Application Load Balancer (ALB) that is configured to handle HTTP and HTTPS separately. The company wants to forward all requests to the website so that the requests will use HTTPS.
What should a solutions architect do to meet this requirement?

A. Update the ALB's network ACL to accept only HTTPS traffic.

B. Create a rule that replaces the HTTP in the URL with HTTPS.

C. Create a listener rule on the ALB to redirect HTTP traffic to HTTPS.

D. Replace the ALB with a Network Load Balancer configured to use Server Name Indication (SNI).

Correct Answer: *C*

Community vote distribution

C (100%)

masetromain Highly Voted 1 month, 2 weeks ago

Selected Answer: C

Answer C :
https://docs.aws.amazon.com/fr_fr/elasticloadbalancing/latest/application/create-https-listener.html
https://aws.amazon.com/fr/premiumsupport/knowledge-center/elb-redirect-http-to-https-using-alb/

upvoted 7 times

Wpcorgan Most Recent 6 days, 23 hours ago

C is correct

upvoted 1 times

hanhdroid 1 month, 2 weeks ago

Selected Answer: C

Answer C: https://aws.amazon.com/premiumsupport/knowledge-center/elb-redirect-http-to-https-using-alb/

upvoted 3 times

Question #61 *Topic 1*

A company is developing a two-tier web application on AWS. The company's developers have deployed the application on an Amazon EC2 instance that connects directly to a backend Amazon RDS database. The company must not hardcode database credentials in the application. The company must also implement a solution to automatically rotate the database credentials on a regular basis.
Which solution will meet these requirements with the LEAST operational overhead?

A. Store the database credentials in the instance metadata. Use Amazon EventBridge (Amazon CloudWatch Events) rules to run a scheduled AWS Lambda function that updates the RDS credentials and instance metadata at the same time.

B. Store the database credentials in a configuration file in an encrypted Amazon S3 bucket. Use Amazon EventBridge (Amazon CloudWatch Events) rules to run a scheduled AWS Lambda function that updates the RDS credentials and the credentials in the configuration file at the same time. Use S3 Versioning to ensure the ability to fall back to previous values.

C. Store the database credentials as a secret in AWS Secrets Manager. Turn on automatic rotation for the secret. Attach the required permission to the EC2 role to grant access to the secret.

D. Store the database credentials as encrypted parameters in AWS Systems Manager Parameter Store. Turn on automatic rotation for the encrypted parameters. Attach the required permission to the EC2 role to grant access to the encrypted parameters.

Correct Answer: *C*

Community vote distribution

C (100%)

Wpcorgan 6 days, 23 hours ago

C is correct

upvoted 2 times

BoboChow 1 month, 1 week ago

Selected Answer: C

AWS Secrets Manager is a newer service than SSM Parameter store

upvoted 1 times

ArielSchivo 1 month, 1 week ago

Selected Answer: C

Option C.

https://docs.aws.amazon.com/secretsmanager/latest/userguide/create_database_secret.html

upvoted 2 times

KVK16 1 month, 1 week ago

Selected Answer: C

Secrets manager supports Autorotation unlike Parameter store.

upvoted 4 times

JesseeS 1 month, 1 week ago

Parameter store does not support autorotation.

upvoted 4 times

Question #62 *Topic 1*

A company is deploying a new public web application to AWS. The application will run behind an Application Load Balancer (ALB). The application needs to be encrypted at the edge with an SSL/TLS certificate that is issued by an external certificate authority (CA). The certificate must be rotated each year before the certificate expires.
What should a solutions architect do to meet these requirements?

A. Use AWS Certificate Manager (ACM) to issue an SSL/TLS certificate. Apply the certificate to the ALB. Use the managed renewal feature to automatically rotate the certificate.

B. Use AWS Certificate Manager (ACM) to issue an SSL/TLS certificate. Import the key material from the certificate. Apply the certificate to the ALUse the managed renewal feature to automatically rotate the certificate.

C. Use AWS Certificate Manager (ACM) Private Certificate Authority to issue an SSL/TLS certificate from the root CA. Apply the certificate to the ALB. Use the managed renewal feature to automatically rotate the certificate.

D. Use AWS Certificate Manager (ACM) to import an SSL/TLS certificate. Apply the certificate to the ALB. Use Amazon EventBridge (Amazon CloudWatch Events) to send a notification when the certificate is nearing expiration. Rotate the certificate manually.

Correct Answer: *D*

Community vote distribution

D (93%) 7%

Sinaneos Highly Voted 1 month, 1 week ago

Selected Answer: D

It's a third-party certificate, hence AWS cannot manage renewal automatically. The closest thing you can do is to send a notification to renew the 3rd party certificate.

upvoted 6 times

mabotega Highly Voted 3 weeks ago

Selected Answer: D

It is D, because ACM does not manage the renewal process for imported certificates. You are responsible for monitoring the expiration date of your imported certificates and for renewing them before they expire.
Check this question on the link below:
Q: What types of certificates can I create and manage with ACM?
https://www.amazonaws.cn/en/certificate-manager/faqs/#Managed_renewal_and_deployment

upvoted 5 times

Wpcorgan Most Recent 6 days, 23 hours ago

D is correct

upvoted 2 times

17Master 1 week ago

Selected Answer: D

If issued by an external entity, the certificate must be imported.

upvoted 1 times

Ack3rman 1 week, 5 days ago

ACM certificates might be ineligible for renewal if:
The certificate isn't associated with another AWS service.
The certificate is expired.
The certificate is imported.
It's a private certificate issued with the IssueCertificate API call.

https://aws.amazon.com/tr/premiumsupport/knowledge-center/acm-certificate-ineligible/

upvoted 1 times

mabotega 3 weeks ago

It is D, because ACM does not manage the renewal process for imported certificates. You are responsible for monitoring the expiration date of your imported certificates and for renewing them before they expire.
Check this question on the link below:
Q: What types of certificates can I create and manage with ACM?
https://www.amazonaws.cn/en/certificate-manager/faqs/#Managed_renewal_and_deployment

upvoted 2 times

ManoAni 1 month ago

Selected Answer: D

When you have a cert issued by external CA, you can import and monitor for its expiration. AWS issued certificate contradicts the statement.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: D

> external certificate authority (CA)
answer is D here because question explicitly stated that they are using external CA

upvoted 1 times

dave9994 1 month ago

D is the Answer. https://docs.aws.amazon.com/acm/latest/userguide/import-certificate.html

upvoted 1 times

NIMIQ 1 month, 1 week ago

It is A: https://www.amazonaws.cn/en/certificate-manager/faqs/#Managed_renewal_and_deployment

upvoted 1 times

BoboChow 1 month, 1 week ago

Selected Answer: A

https://docs.aws.amazon.com/acm/latest/userguide/managed-renewal.html

upvoted 1 times

JesseeS 1 month, 1 week ago

It is option A
https://www.amazonaws.cn/en/certificate-manager/faqs/

upvoted 1 times

ManoAni 1 month ago

When you have a cert issued by external CA, you can import and monitor for its expiration. AWS issued certificate contradicts the statement.

upvoted 1 times

Question #63 *Topic 1*

A company runs its infrastructure on AWS and has a registered base of 700,000 users for its document management application. The company intends to create a product that converts large .pdf files to .jpg image files. The .pdf files average 5 MB in size. The company needs to store the original files and the converted files. A solutions architect must design a scalable solution to accommodate demand that will grow rapidly over time.
Which solution meets these requirements MOST cost-effectively?

A. Save the .pdf files to Amazon S3. Configure an S3 PUT event to invoke an AWS Lambda function to convert the files to .jpg format and store them back in Amazon S3.

B. Save the .pdf files to Amazon DynamoDUse the DynamoDB Streams feature to invoke an AWS Lambda function to convert the files to .jpg format and store them back in DynamoDB.

C. Upload the .pdf files to an AWS Elastic Beanstalk application that includes Amazon EC2 instances, Amazon Elastic Block Store (Amazon EBS) storage, and an Auto Scaling group. Use a program in the EC2 instances to convert the files to .jpg format. Save the .pdf files and the .jpg files in the EBS store.

D. Upload the .pdf files to an AWS Elastic Beanstalk application that includes Amazon EC2 instances, Amazon Elastic File System (Amazon EFS) storage, and an Auto Scaling group. Use a program in the EC2 instances to convert the file to .jpg format. Save the .pdf files and the .jpg files in the EBS store.

Correct Answer: *A*

Community vote distribution

A (94%) 6%

ArielSchivo Highly Voted 1 month, 1 week ago

Selected Answer: A

Option A. Elastic BeanStalk is expensive, and DocumentDB has a 400KB max to upload files. So Lambda and S3 should be the one.

upvoted 17 times

rob74 3 weeks, 5 days ago

In addition to this Lambda is paid only when used....

upvoted 3 times

raffaello44 1 month ago

is lambda scalable as an EC2 ?

upvoted 2 times

Wpcorgan Most Recent 6 days, 23 hours ago

A is correct

upvoted 1 times

TonyghostR05 1 week, 6 days ago

S3 is cost effective

upvoted 1 times

goku58 1 month ago

Selected Answer: B

For rapid scalability, B - DynamoDB looks to be a better solution.

upvoted 1 times

ludovikush 3 weeks, 5 days ago

It is not correct because the maximum item size in DynamoDB is 400 KB.

upvoted 5 times

Question #64 *Topic 1*

A company has more than 5 TB of file data on Windows file servers that run on premises. Users and applications interact with the data each day. The company is moving its Windows workloads to AWS. As the company continues this process, the company requires access to AWS and on-premises file storage with minimum latency. The company needs a solution that minimizes operational overhead and requires no significant changes to the existing file access patterns. The company uses an AWS Site-to-Site VPN connection for connectivity to AWS.
What should a solutions architect do to meet these requirements?

A. Deploy and configure Amazon FSx for Windows File Server on AWS. Move the on-premises file data to FSx for Windows File Server. Reconfigure the workloads to use FSx for Windows File Server on AWS.

B. Deploy and configure an Amazon S3 File Gateway on premises. Move the on-premises file data to the S3 File Gateway. Reconfigure the on-premises workloads and the cloud workloads to use the S3 File Gateway.

C. Deploy and configure an Amazon S3 File Gateway on premises. Move the on-premises file data to Amazon S3. Reconfigure the workloads to use either Amazon S3 directly or the S3 File Gateway. depending on each workload's location.

D. Deploy and configure Amazon FSx for Windows File Server on AWS. Deploy and configure an Amazon FSx File Gateway on premises. Move the on-premises file data to the FSx File Gateway. Configure the cloud workloads to use FSx for Windows File Server on AWS. Configure the on-premises workloads to use the FSx File Gateway.

Correct Answer: *A*

Community vote distribution

D (96%) 4%

sba21 Highly Voted 1 month, 2 weeks ago

Selected Answer: D

https://www.examtopics.com/discussions/amazon/view/83281-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 7 times

Vesperia Most Recent 2 days, 22 hours ago

Selected Answer: A

Answer A is correct. The company has a site to site VPN already.There is no need to install file gateway on-premise.
https://docs.aws.amazon.com/fsx/latest/LustreGuide/mounting-on-premises.html

upvoted 1 times

Wpcorgan 6 days, 23 hours ago

D is correct

upvoted 1 times

Kapello10 1 week, 6 days ago

Selected Answer: D

ddddddddd

upvoted 2 times

goku58 1 month ago

Selected Answer: D

Windows File server == FSx.
Since access from both on-prem and AWS is needed, A isn't sufficient. So D.

upvoted 2 times

17Master 2 weeks, 1 day ago

and VPN S2S?

upvoted 2 times

rewdboy 3 days, 12 hours ago

True, but the other requirement is no "significant changes to the existing file access patterns" which would mean mounting File Gateway shares in their on-premises location while they move their workloads to FSx during their migration. So D.

upvoted 1 times

BoboChow 1 month, 1 week ago

Selected Answer: D

https://docs.aws.amazon.com/filegateway/latest/filefsxw/what-is-file-fsxw.html

upvoted 3 times

tubtab 1 month, 2 weeks ago

Selected Answer: D

ddddddddd

upvoted 4 times

oldcardigan 1 month, 2 weeks ago

Selected Answer: D

i think its D

upvoted 4 times

Question #65 *Topic 1*

A hospital recently deployed a RESTful API with Amazon API Gateway and AWS Lambda. The hospital uses API Gateway and Lambda to upload reports that are in PDF format and JPEG format. The hospital needs to modify the Lambda code to identify protected health information (PHI) in the reports.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use existing Python libraries to extract the text from the reports and to identify the PHI from the extracted text.

B. Use Amazon Textract to extract the text from the reports. Use Amazon SageMaker to identify the PHI from the extracted text.

C. Use Amazon Textract to extract the text from the reports. Use Amazon Comprehend Medical to identify the PHI from the extracted text.

D. Use Amazon Rekognition to extract the text from the reports. Use Amazon Comprehend Medical to identify the PHI from the extracted text.

Correct Answer: *C*

Community vote distribution

C (100%)

Wpcorgan 6 days, 23 hours ago

C is correct

upvoted 1 times

bansalhp 1 month ago

Selected Answer: C

Textract -to extract textand Comprehend -to identify Medical info

upvoted 2 times

JesseeS 1 month, 1 week ago

Textract and Comprehend is HIPPA compliant
https://aws.amazon.com/blogs/machine-learning/amazon-textract-is-now-hipaa-eligible/

upvoted 1 times

KVK16 1 month, 1 week ago

Selected Answer: C

Textract - Comprehend Medical for PHI info

upvoted 3 times

Question #66 *Topic 1*

A company has an application that generates a large number of files, each approximately 5 MB in size. The files are stored in Amazon S3. Company policy requires the files to be stored for 4 years before they can be deleted. Immediate accessibility is always required as the files contain critical business data that is not easy to reproduce. The files are frequently accessed in the first 30 days of the object creation but are rarely accessed after the first 30 days.
Which storage solution is MOST cost-effective?

A. Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 Glacier 30 days from object creation. Delete the files 4 years after object creation.

B. Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 One Zone-Infrequent Access (S3 One Zone-IA) 30 days from object creation. Delete the files 4 years after object creation.

C. Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days from object creation. Delete the files 4 years after object creation.

D. Create an S3 bucket lifecycle policy to move files from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days from object creation. Move the files to S3 Glacier 4 years after object creation.

Correct Answer: *C*

Community vote distribution

C (71%) A (20%) 9%

Six_Fingered_Jose Highly Voted 1 month ago

Selected Answer: C

i think C should be the answer here,
> Immediate accessibility is always required as the files contain critical business data that is not easy to reproduce

If they do not explicitly mention that they are using Glacier Instant Retrieval, we should assume that Glacier -> takes more time to retrieve and may not meet the requirements

upvoted 13 times

ninjawrz Highly Voted 1 month, 1 week ago

Selected Answer: A

Most COST EFFECTIVE
A: S3 Glacier Instant Retrieval is a new storage class that delivers the fastest access to archive storage, with the same low latency and high-throughput performance as the S3 Standard and S3 Standard-IA storage classes. You can save up to 68 percent on storage costs as compared with using the S3 Standard-IA storage class when you use the S3 Glacier Instant Retrieval storage class and pay a low price to retrieve data.

upvoted 9 times

Pamban 2 weeks, 3 days ago

"Immediate accessibility is always required as the files contain critical business data that is not easy to reproduce" is the key sentence. answer is C.

upvoted 3 times

123jhl0 1 month, 1 week ago

I think A is the answer, too. As S3 Glacier has 3 different classes and one of them is able to retrieve objects in milliseconds (Instant Retrieval). As no class has been specified, we can assume the on that meets requirements is selected.

upvoted 5 times

GasGasJDM 1 month ago

I agree with this, since no specific information about Glacier modality, we can assume that Instant Retrieval is a possibility and as far as I know, this is cheaper than IA.

upvoted 1 times

jackwang0616 Most Recent 6 days, 9 hours ago

Selected Answer: C

"Immediate accessibility is always required as the files contain critical business data that is not easy to reproduce"
so can't move to the glacier

upvoted 1 times

Wpcorgan 6 days, 23 hours ago

C is correct

upvoted 1 times

mikey2000 1 week, 2 days ago

Selected Answer: C

I would choose A but that option didn't specify with glacier would be needed so C Is the answer because the file cannot be easily reproduced and s3 standard IA offers resiliency than s3 one zone IA(only uses one AZ)

upvoted 1 times

edk11z 1 week, 2 days ago

Selected Answer: B

B because most cost effective. and the data are not impossible to reproduce. they are just hard to reproduce.

upvoted 1 times

Ack3rman 1 week, 5 days ago

Since it is not specified which type of Glacier it is, the answer is C i believe.

upvoted 1 times

iCcma 1 week, 5 days ago

Selected Answer: C

A mentions "S3 Glacier" not "S3 Glacier Instant Retrieval"
Cheese and yogurt are made from milk but they are not the same.
Ans *C*

https://aws.amazon.com/es/s3/storage-classes/

upvoted 3 times

Kartikey140 1 week, 3 days ago

check out - https://aws.amazon.com/s3/pricing/
Cheese and yogurt 😂

upvoted 2 times

Pamban 2 weeks, 3 days ago

Selected Answer: C

"Immediate accessibility is always required as the files contain critical business data that is not easy to reproduce" is the key sentence. answer is C.

upvoted 2 times

backbencher2022 3 weeks ago

Selected Answer: C

C is the correct answer

upvoted 1 times

keezbadger 3 weeks ago

Let me help you settle this argument.
A is wrong as it only mentions S3 glacier and S3 glacier has 3 variations as shown below which conflicts the answer:
Amazon s3 Glacier Instant Retrieval: (Here it is millisecond retrieval and good for data accessed once a quarter and minimum duration is 90 days)
Amazon S3 Glacier Flexible Retrieval: (Here you have Expedited (1-5minutes), Standard (3-5hours), Bulk (5-12hours-free) retrieval and minimum duration is (90days)
Amazon S3 Glacier Deep Archive
B is wrong answer. S3 Zone IA is only in One Zone and data will be lost if AZ is destroyed
C is right as S3 Standard IA "data is less frequently accessed, but requires rapid access when needed"
I hope this helps :)

upvoted 2 times

PS_R 3 weeks ago

Selected Answer: C

Unless mentioned about Glacier Instant revival, i feel it is S3 Standard IA

upvoted 4 times

gokalpkocer3 3 weeks, 1 day ago

Selected Answer: B

Answers does not mention anything about Glacier Instant retrieval. Do not make assumptions if answer specifically does not say it. Answer also does not say anything about AZs. So you can store in infrequent one AZ for most cost effective solution. My answer is B.

upvoted 2 times

PS_R 3 weeks ago

Since it is Business Critical data, if S3 Onezone is not available for some reason, data will be available, so i would vote for C, S3 IA

upvoted 2 times

BoboChow 4 weeks, 1 day ago

Selected Answer: C

Before you transition objects from the S3 Standard or S3 Standard-IA storage classes to S3 Standard-IA or S3 One Zone-IA, you must store them at least 30 days in the S3 Standard storage class. For example, you cannot create a Lifecycle rule to transition objects to the S3 Standard-IA storage class one day after you create them. Amazon S3 doesn't transition objects within the first 30 days because newer objects are often accessed more frequently or deleted sooner than is suitable for S3 Standard-IA or S3 One Zone-IA storage.

At least C is not a wrong answer.

upvoted 1 times

SimonPark 1 month ago

Selected Answer: C

How S3 Glacier satisfies "Immediate accessibility is always required"

upvoted 4 times

ManoAni 1 month ago

Selected Answer: A

A is correct answer

upvoted 2 times

iCcma 1 week, 5 days ago

https://aws.amazon.com/es/s3/storage-classes/

upvoted 1 times

UWSFish 1 month ago

Its a horrible question. Are you supposed to assume S3 glacier instant retrieval or not. My gut tells me you shouldn't make that assumption but you can argue it either way.

upvoted 5 times

Question #67 *Topic 1*

A company hosts an application on multiple Amazon EC2 instances. The application processes messages from an Amazon SQS queue, writes to an Amazon RDS table, and deletes the message from the queue. Occasional duplicate records are found in the RDS table. The SQS queue does not contain any duplicate messages.
What should a solutions architect do to ensure messages are being processed once only?

A. Use the CreateQueue API call to create a new queue.

B. Use the AddPermission API call to add appropriate permissions.

C. Use the ReceiveMessage API call to set an appropriate wait time.

D. Use the ChangeMessageVisibility API call to increase the visibility timeout.

Correct Answer: *D*

Community vote distribution

D (100%)

KVK16 Highly Voted 1 month, 1 week ago

Selected Answer: D

In case of SQS - multi-consumers if one consumer has already picked the message and is processing, in meantime other consumer can pick it up and process the message there by two copies are added at the end. To avoid this the message is made invisible from the time its picked and deleted after processing. This visibility timeout is increased according to max time taken to process the message

upvoted 9 times

Wpcorgan Most Recent 6 days, 23 hours ago

D is correct

upvoted 1 times

mabotega 3 weeks ago

Selected Answer: D

D is the correct choise, increasing the visibility timeout according to max time taken to process the message on the RDS.

upvoted 1 times

Valero_ 1 month, 1 week ago

Selected Answer: D

True, it's D.
https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-visibility-timeout.html

upvoted 2 times

Question #68 *Topic 1*

A solutions architect is designing a new hybrid architecture to extend a company's on-premises infrastructure to AWS. The company requires a highly available connection with consistent low latency to an AWS Region. The company needs to minimize costs and is willing to accept slower traffic if the primary connection fails.
What should the solutions architect do to meet these requirements?

A. Provision an AWS Direct Connect connection to a Region. Provision a VPN connection as a backup if the primary Direct Connect connection fails.

B. Provision a VPN tunnel connection to a Region for private connectivity. Provision a second VPN tunnel for private connectivity and as a backup if the primary VPN connection fails.

C. Provision an AWS Direct Connect connection to a Region. Provision a second Direct Connect connection to the same Region as a backup if the primary Direct Connect connection fails.

D. Provision an AWS Direct Connect connection to a Region. Use the Direct Connect failover attribute from the AWS CLI to automatically create a backup connection if the primary Direct Connect connection fails.

Correct Answer: *A*

Community vote distribution

A (93%) 7%

KVK16 Highly Voted 1 month, 1 week ago

Selected Answer: A

Direct Connect + VPN best of both

upvoted 7 times

mabotega Highly Voted 3 weeks ago

Selected Answer: A

Direct Connect goes throught 1 Gbps, 10 Gbps or 100 Gbps and the VPN goes up to 1.25 Gbps.

https://docs.aws.amazon.com/whitepapers/latest/aws-vpc-connectivity-options/aws-direct-connect-vpn.html

upvoted 5 times

koreanmonkey Most Recent 1 day, 8 hours ago

Selected Answer: A

A is rigth I thought wrong

upvoted 1 times

koreanmonkey 1 day, 8 hours ago

Selected Answer: C

I think VPN is not right solution for "low latency"

So how about C?

upvoted 1 times

Wpcorgan 6 days, 23 hours ago

A is correct

upvoted 1 times

Question #69 *Topic 1*

A company is running a business-critical web application on Amazon EC2 instances behind an Application Load Balancer. The EC2 instances are in an Auto Scaling group. The application uses an Amazon Aurora PostgreSQL database that is deployed in a single Availability Zone. The company wants the application to be highly available with minimum downtime and minimum loss of data.
Which solution will meet these requirements with the LEAST operational effort?

A. Place the EC2 instances in different AWS Regions. Use Amazon Route 53 health checks to redirect traffic. Use Aurora PostgreSQL Cross-Region Replication.

B. Configure the Auto Scaling group to use multiple Availability Zones. Configure the database as Multi-AZ. Configure an Amazon RDS Proxy instance for the database.

C. Configure the Auto Scaling group to use one Availability Zone. Generate hourly snapshots of the database. Recover the database from the snapshots in the event of a failure.

D. Configure the Auto Scaling group to use multiple AWS Regions. Write the data from the application to Amazon S3. Use S3 Event Notifications to launch an AWS Lambda function to write the data to the database.

Correct Answer: *B*

Community vote distribution

B (88%) 13%

KVK16 Highly Voted 1 month, 1 week ago

Selected Answer: B

RDS Proxy for Aurora https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html

upvoted 5 times

Wajif Most Recent 2 days, 7 hours ago

Selected Answer: A

Why not A?

upvoted 1 times

koreanmonkey 1 day, 7 hours ago

maybe because of load balancer, diffrent region can't be answer.

upvoted 1 times

Wpcorgan 6 days, 22 hours ago

B is correct

upvoted 1 times

Ack3rman 1 week, 5 days ago

Important fact: EC2 Auto Scaling groups are regional constructs. They can span Availability Zones, but not AWS regions. So can't be D in case you are between B and D
https://aws.amazon.com/tr/ec2/autoscaling/faqs/

upvoted 1 times

Anji69659 2 weeks, 3 days ago

Selected Answer: B

MULTI-AZ FOR HIGH SCALABILITY .

upvoted 2 times

Question #70 *Topic 1*

A company's HTTP application is behind a Network Load Balancer (NLB). The NLB's target group is configured to use an Amazon EC2 Auto Scaling group with multiple EC2 instances that run the web service.
The company notices that the NLB is not detecting HTTP errors for the application. These errors require a manual restart of the EC2 instances that run the web service. The company needs to improve the application's availability without writing custom scripts or code.
What should a solutions architect do to meet these requirements?

A. Enable HTTP health checks on the NLB, supplying the URL of the company's application.

B. Add a cron job to the EC2 instances to check the local application's logs once each minute. If HTTP errors are detected. the application will restart.

C. Replace the NLB with an Application Load Balancer. Enable HTTP health checks by supplying the URL of the company's application. Configure an Auto Scaling action to replace unhealthy instances.

D. Create an Amazon Cloud Watch alarm that monitors the UnhealthyHostCount metric for the NLB. Configure an Auto Scaling action to replace unhealthy instances when the alarm is in the ALARM state.

Correct Answer: *C*

Community vote distribution

C (89%) 11%

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: C

I would choose A, as NLB supports HTTP and HTTPS Health Checks, BUT you can't put any URL (as proposed), only the node IP addresses. So, the solution is C.

upvoted 9 times

Ack3rman 1 week, 5 days ago

can you elaborate more pls

upvoted 1 times

LeGloupier Most Recent 3 days ago

Selected Answer: A

A is the best option.
NLB support http healthcheck, so why do we need to move to ALB ?
moreover the sentence "Configure an Auto Scaling action to replace unhealthy instances" in C seems to be wrong, as auto scaling remove any unhealthy instance by default, you do not need to configure it.

upvoted 1 times

Wpcorgan 6 days, 22 hours ago

C is correct

upvoted 1 times

mabotega 3 weeks ago

Selected Answer: C

C is the correct!
NLB does not handle HTTP (layer 7) listerns errors only TCP (layer 4) listeners.
https://docs.aws.amazon.com/elasticbeanstalk/latest/dg/environments-cfg-nlb.html

upvoted 3 times

Solarch 3 weeks, 3 days ago

Answer is A
NLB is ideal for TPC and UDP Traffic and checks operating in layer 4.
ALB- Supports HTTP and HTTPs traffics. Hence the ELB needs to be changed from NLB to ALB.

upvoted 1 times

Aman54 1 month, 1 week ago

Selected Answer: A

NLB supports HTTP health checks, they are part of the target group and the setting is the same for ALB and NLB HTTP/HTTPS health checks.

upvoted 1 times

Vesperia 2 days, 21 hours ago

A is incorrect. NLB cannot detect http errors. Adding health check only detects the healthiness of the instances, not http errors.

upvoted 2 times

oxfordcommaa 1 month ago

"The company needs to improve the application's availability"
Answer A does not address this. The auto scaling group in answer C does.

upvoted 1 times

Maharaja 3 weeks, 4 days ago

NLB is already configured with a target group supported by EC2 ASG "NLB's target group is configured to use an Amazon EC2 Auto Scaling group". NLB need to be configured to use http health check. Hence A

upvoted 2 times

ArielSchivo 1 month, 1 week ago

Selected Answer: C

Option C. NLB works at Layer 4 so it does not support HTTP/HTTPS. The replacement for the ALB is the best choice.

upvoted 4 times

Question #71 *Topic 1*

A company runs a shopping application that uses Amazon DynamoDB to store customer information. In case of data corruption, a solutions architect needs to design a solution that meets a recovery point objective (RPO) of 15 minutes and a recovery time objective (RTO) of 1 hour. What should the solutions architect recommend to meet these requirements?

A. Configure DynamoDB global tables. For RPO recovery, point the application to a different AWS Region.

B. Configure DynamoDB point-in-time recovery. For RPO recovery, restore to the desired point in time.

C. Export the DynamoDB data to Amazon S3 Glacier on a daily basis. For RPO recovery, import the data from S3 Glacier to DynamoDB.

D. Schedule Amazon Elastic Block Store (Amazon EBS) snapshots for the DynamoDB table every 15 minutes. For RPO recovery, restore the DynamoDB table by using the EBS snapshot.

Correct Answer: *B*

Community vote distribution

B (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: B

A - DynamoDB global tables provides multi-Region, and multi-active database, but it not valid "in case of data corruption". In this case, you need a backup. This solutions isn't valid.
B - Point in Time Recovery is designed as a continuous backup juts to recover it fast. It covers perfectly the RPO, and probably the RTO.
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/PointInTimeRecovery.html
C - A daily export will not cover the RPO of 15min.
D - DynamoDB is serverless... so what are these EBS snapshots taken from???

upvoted 17 times

Wpcorgan Most Recent 6 days, 22 hours ago

B is correct

upvoted 1 times

SimonPark 1 month ago

Selected Answer: B

B is the answer

upvoted 1 times

BoboChow 1 month, 1 week ago

Selected Answer: B

I think DynamoDB global tables also work here, but Point in Time Recovery is a better choice

upvoted 1 times

Kikiokiki 1 month, 1 week ago

I THINK B.
https://dynobase.dev/dynamodb-point-in-time-recovery/

upvoted 1 times

priya2224 1 month, 1 week ago

answer is D

upvoted 1 times

patni11 1 month, 1 week ago

bhk gandu chutiye glt ans btata hai

upvoted 1 times

Az900500 3 weeks, 2 days ago

Try communicate in English for audience

upvoted 3 times

123jhl0 1 month, 1 week ago

DynamoDB is serverless, so no storage snapshots available. https://aws.amazon.com/dynamodb/

upvoted 1 times

Question #72 *Topic 1*

A company runs a photo processing application that needs to frequently upload and download pictures from Amazon S3 buckets that are located in the same AWS Region. A solutions architect has noticed an increased cost in data transfer fees and needs to implement a solution to reduce these costs.
How can the solutions architect meet this requirement?

A. Deploy Amazon API Gateway into a public subnet and adjust the route table to route S3 calls through it.

B. Deploy a NAT gateway into a public subnet and attach an endpoint policy that allows access to the S3 buckets.

C. Deploy the application into a public subnet and allow it to route through an internet gateway to access the S3 buckets.

D. Deploy an S3 VPC gateway endpoint into the VPC and attach an endpoint policy that allows access to the S3 buckets.

Correct Answer: *D*

Community vote distribution

D (100%)

KVK16 Highly Voted 1 month, 1 week ago

Selected Answer: D

To reduce costs get rid of NAT Gateway , VPC endpoint to S3

upvoted 9 times

Wpcorgan Most Recent 6 days, 22 hours ago

D is correct

upvoted 1 times

yd_h 1 month ago

Selected Answer: D
FYI :
-There is no additional charge for using gateway endpoints.
-Interface endpoints are priced at ~ $0.01/per AZ/per hour. Cost depends on the Region
- S3 Interface Endpoints resolve to private VPC IP addresses and are routable from outside the VPC (e.g via VPN, Direct Connect, Transit Gateway, etc). S3 Gateway Endpoints use public IP ranges and are only routable from resources within the VPC.

upvoted 3 times

123jhl0 1 month, 1 week ago

Selected Answer: D

Close question to the Question #4, with same solution.

upvoted 3 times

Question #73 *Topic 1*

A company recently launched Linux-based application instances on Amazon EC2 in a private subnet and launched a Linux-based bastion host on an Amazon EC2 instance in a public subnet of a VPC. A solutions architect needs to connect from the on-premises network, through the company's internet connection, to the bastion host, and to the application servers. The solutions architect must make sure that the security groups of all the EC2 instances will allow that access.
Which combination of steps should the solutions architect take to meet these requirements? (Choose two.)

A. Replace the current security group of the bastion host with one that only allows inbound access from the application instances.

B. Replace the current security group of the bastion host with one that only allows inbound access from the internal IP range for the company.

C. Replace the current security group of the bastion host with one that only allows inbound access from the external IP range for the company.

D. Replace the current security group of the application instances with one that allows inbound SSH access from only the private IP address of the bastion host.

E. Replace the current security group of the application instances with one that allows inbound SSH access from only the public IP address of the bastion host.

Correct Answer: *CD*

Community vote distribution

CD (100%)

Six_Fingered_Jose Highly Voted 1 month ago

Selected Answer: CD

C because from on-prem network to bastion through internet (using on-prem resource's public IP),
D because bastion and ec2 is in same VPC, meaning bastion can communicate to EC2 via it's private IP address

upvoted 9 times

Wpcorgan Most Recent 6 days, 22 hours ago

C and D

upvoted 1 times

gcmrjbr 1 month ago

CD is Ok.

upvoted 1 times

Evangelia 1 month, 1 week ago

why C? External?

upvoted 2 times

ArielSchivo 1 month, 1 week ago

Selected Answer: CD

Option C (allow access just from the external IP) and D (allow inbound SSH from the private IP of the bastion host).

upvoted 2 times

ninjawrz 1 month, 1 week ago

Selected Answer: CD

CD is correct

upvoted 2 times

Question #74 *Topic 1*

A solutions architect is designing a two-tier web application. The application consists of a public-facing web tier hosted on Amazon EC2 in public subnets. The database tier consists of Microsoft SQL Server running on Amazon EC2 in a private subnet. Security is a high priority for the company.
How should security groups be configured in this situation? (Choose two.)

A. Configure the security group for the web tier to allow inbound traffic on port 443 from 0.0.0.0/0.

B. Configure the security group for the web tier to allow outbound traffic on port 443 from 0.0.0.0/0.

C. Configure the security group for the database tier to allow inbound traffic on port 1433 from the security group for the web tier.

D. Configure the security group for the database tier to allow outbound traffic on ports 443 and 1433 to the security group for the web tier.

E. Configure the security group for the database tier to allow inbound traffic on ports 443 and 1433 from the security group for the web tier.

Correct Answer: *AC*

Community vote distribution

AC (100%)

ArielSchivo Highly Voted 1 month, 1 week ago

Selected Answer: AC

EC2 web on public subnets + EC2 SQL on private subnet + security is high priority. So, Option A to allow HTTPS from everywhere. Plus option C to allow SQL connection from the web instance.

upvoted 7 times

Athena Highly Voted 3 weeks, 3 days ago

Selected Answer: AC

Web Server Rules: Inbound traffic from 443 (HTTPS) Source 0.0.0.0/0 - Allows inbound HTTPS access from any IPv4 address
Database Rules : 1433 (MS SQL)The default port to access a Microsoft SQL Server database, for example, on an Amazon RDS instance

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/security-group-rules-reference.html

upvoted 6 times

Wpcorgan Most Recent 6 days, 22 hours ago

A and C

upvoted 1 times

gcmrjbr 1 month ago

Agree with AC.

upvoted 2 times

srcshekhar 1 month, 2 weeks ago

Very good questions

upvoted 3 times

Question #75 *Topic 1*

A company wants to move a multi-tiered application from on premises to the AWS Cloud to improve the application's performance. The application consists of application tiers that communicate with each other by way of RESTful services. Transactions are dropped when one tier becomes overloaded. A solutions architect must design a solution that resolves these issues and modernizes the application.
Which solution meets these requirements and is the MOST operationally efficient?

A. Use Amazon API Gateway and direct transactions to the AWS Lambda functions as the application layer. Use Amazon Simple Queue Service (Amazon SQS) as the communication layer between application services.

B. Use Amazon CloudWatch metrics to analyze the application performance history to determine the servers' peak utilization during the performance failures. Increase the size of the application server's Amazon EC2 instances to meet the peak requirements.

C. Use Amazon Simple Notification Service (Amazon SNS) to handle the messaging between application servers running on Amazon EC2 in an Auto Scaling group. Use Amazon CloudWatch to monitor the SNS queue length and scale up and down as required.

D. Use Amazon Simple Queue Service (Amazon SQS) to handle the messaging between application servers running on Amazon EC2 in an Auto Scaling group. Use Amazon CloudWatch to monitor the SQS queue length and scale up when communication failures are detected.

Correct Answer: *A*

Community vote distribution

A (100%)

gcmrjbr Highly Voted 1 month ago
Agree with A>>> Lambda = serverless + autoscale (modernize), SQS= decouple (no more drops)
upvoted 7 times

Wpcorgan Most Recent 6 days, 22 hours ago
A is correct
upvoted 1 times

airraid2010 1 week, 5 days ago
Selected Answer: A
https://serverlessland.com/patterns/apigw-http-sqs-lambda-sls
upvoted 2 times

BoboChow 1 month ago
Selected Answer: A
Serverless + decouple
upvoted 3 times

cark0728 1 month, 1 week ago
Selected Answer: A
A가 올바른 정답이다
upvoted 3 times

Question #76 *Topic 1*

A company receives 10 TB of instrumentation data each day from several machines located at a single factory. The data consists of JSON files stored on a storage area network (SAN) in an on-premises data center located within the factory. The company wants to send this data to Amazon S3 where it can be accessed by several additional systems that provide critical near-real-time analytics. A secure transfer is important because the data is considered sensitive.
Which solution offers the MOST reliable data transfer?

A. AWS DataSync over public internet

B. AWS DataSync over AWS Direct Connect

C. AWS Database Migration Service (AWS DMS) over public internet

D. AWS Database Migration Service (AWS DMS) over AWS Direct Connect

Correct Answer: *B*

Community vote distribution

B (100%)

ArielSchivo Highly Voted 1 month, 1 week ago

Selected Answer: B

DMS is for databases and here refers to “JSON files”. Public internet is not reliable. So best option is B.

upvoted 9 times

Wpcorgan Most Recent 6 days, 22 hours ago

B is correct

upvoted 1 times

yd_h 1 month ago

B

- A SAN presents storage devices to a host such that the storage appears to be locally attached. (NFS is, or can be, a SAN - https://serverfault.com/questions/499185/is-san-storage-better-than-nfs)
- AWS Direct Connect does not encrypt your traffic that is in transit by default. But the connection is private (https://docs.aws.amazon.com/directconnect/latest/UserGuide/encryption-in-transit.html)

upvoted 3 times

Question #77 *Topic 1*

A company needs to configure a real-time data ingestion architecture for its application. The company needs an API, a process that transforms data as the data is streamed, and a storage solution for the data.
Which solution will meet these requirements with the LEAST operational overhead?

A. Deploy an Amazon EC2 instance to host an API that sends data to an Amazon Kinesis data stream. Create an Amazon Kinesis Data Firehose delivery stream that uses the Kinesis data stream as a data source. Use AWS Lambda functions to transform the data. Use the Kinesis Data Firehose delivery stream to send the data to Amazon S3.

B. Deploy an Amazon EC2 instance to host an API that sends data to AWS Glue. Stop source/destination checking on the EC2 instance. Use AWS Glue to transform the data and to send the data to Amazon S3.

C. Configure an Amazon API Gateway API to send data to an Amazon Kinesis data stream. Create an Amazon Kinesis Data Firehose delivery stream that uses the Kinesis data stream as a data source. Use AWS Lambda functions to transform the data. Use the Kinesis Data Firehose delivery stream to send the data to Amazon S3.

D. Configure an Amazon API Gateway API to send data to AWS Glue. Use AWS Lambda functions to transform the data. Use AWS Glue to send the data to Amazon S3.

Correct Answer: *C*

Community vote distribution

C (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: C

(A) - You don't need to deploy an EC2 instance to host an API - Operational overhead
(B) - Same as A
(**C**) - Is the answer
(D) - AWS Glue gets data from S3, not from API GW. AWS Glue could do ETL by itself, so don't need lambda. Non sense.
https://aws.amazon.com/glue/

upvoted 12 times

Wpcorgan Most Recent 6 days, 22 hours ago

C is correct

upvoted 1 times

Cristian93 1 month ago

Selected Answer: C

C is correct answer

upvoted 2 times

Question #78 *Topic 1*

A company needs to keep user transaction data in an Amazon DynamoDB table. The company must retain the data for 7 years.
What is the MOST operationally efficient solution that meets these requirements?

A. Use DynamoDB point-in-time recovery to back up the table continuously.

B. Use AWS Backup to create backup schedules and retention policies for the table.

C. Create an on-demand backup of the table by using the DynamoDB console. Store the backup in an Amazon S3 bucket. Set an S3 Lifecycle configuration for the S3 bucket.

D. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to invoke an AWS Lambda function. Configure the Lambda function to back up the table and to store the backup in an Amazon S3 bucket. Set an S3 Lifecycle configuration for the S3 bucket.

Correct Answer: *B*

Community vote distribution

B (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: B

Answer is B
"Amazon DynamoDB offers two types of backups: point-in-time recovery (PITR) and on-demand backups. (==> D is not the answer)
PITR is used to recover your table to any point in time in a rolling 35 day window, which is used to help customers mitigate accidental deletes or writes to their tables from bad code, malicious access, or user error. (==> A isn't the answer)
On demand backups are designed for long-term archiving and retention, which is typically used to help customers meet compliance and regulatory requirements.
This is the second of a series of two blog posts about using AWS Backup to set up scheduled on-demand backups for Amazon DynamoDB. Part 1 presents the steps to set up a scheduled backup for DynamoDB tables from the AWS Management Console." (==> Not the DynamoBD console and C isn't the answer either)
https://aws.amazon.com/blogs/database/part-2-set-up-scheduled-backups-for-amazon-dynamodb-using-aws-backup/

upvoted 17 times

Wpcorgan Most Recent 6 days, 22 hours ago

B is correct

upvoted 1 times

mabotega 1 week, 2 days ago

Selected Answer: B

We recommend you use AWS Backup to automatically delete the backups that you no longer need by configuring your lifecycle when you created your backup plan.
https://docs.aws.amazon.com/aws-backup/latest/devguide/deleting-backups.html

upvoted 1 times

SimonPark 4 weeks, 1 day ago

Selected Answer: B

B is clear

upvoted 2 times

Question #79 *Topic 1*

A company is planning to use an Amazon DynamoDB table for data storage. The company is concerned about cost optimization. The table will not be used on most mornings. In the evenings, the read and write traffic will often be unpredictable. When traffic spikes occur, they will happen very quickly.
What should a solutions architect recommend?

A. Create a DynamoDB table in on-demand capacity mode.

B. Create a DynamoDB table with a global secondary index.

C. Create a DynamoDB table with provisioned capacity and auto scaling.

D. Create a DynamoDB table in provisioned capacity mode, and configure it as a global table.

Correct Answer: *A*

Community vote distribution

A (100%)

SimonPark Highly Voted 4 weeks, 1 day ago

Selected Answer: A

On-demand mode is a good option if any of the following are true:
- You create new tables with unknown workloads.
- You have unpredictable application traffic.
- You prefer the ease of paying for only what you use.

upvoted 6 times

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: A

A - On demand is the answer -
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ReadWriteCapacityMode.html#HowItWorks.OnDemand
B - not related with the unpredictable traffic
C - provisioned capacity is recommended for known patterns. Not the case here.
D - same as C

upvoted 5 times

Wpcorgan Most Recent 6 days, 22 hours ago

A is correct

upvoted 1 times

Question #80 *Topic 1*

A company recently signed a contract with an AWS Managed Service Provider (MSP) Partner for help with an application migration initiative. A solutions architect needs ta share an Amazon Machine Image (AMI) from an existing AWS account with the MSP Partner's AWS account. The AMI is backed by Amazon Elastic Block Store (Amazon EBS) and uses an AWS Key Management Service (AWS KMS) customer managed key to encrypt EBS volume snapshots.
What is the MOST secure way for the solutions architect to share the AMI with the MSP Partner's AWS account?

A. Make the encrypted AMI and snapshots publicly available. Modify the key policy to allow the MSP Partner's AWS account to use the key.

B. Modify the launchPermission property of the AMI. Share the AMI with the MSP Partner's AWS account only. Modify the key policy to allow the MSP Partner's AWS account to use the key.

C. Modify the launchPermission property of the AMI. Share the AMI with the MSP Partner's AWS account only. Modify the key policy to trust a new KMS key that is owned by the MSP Partner for encryption.

D. Export the AMI from the source account to an Amazon S3 bucket in the MSP Partner's AWS account, Encrypt the S3 bucket with a new KMS key that is owned by the MSP Partner. Copy and launch the AMI in the MSP Partner's AWS account.

Correct Answer: *B*

Community vote distribution

B (91%) 9%

Sauran Highly Voted 1 month, 1 week ago

Selected Answer: B

Share the existing KMS key with the MSP external account because it has already been used to encrypt the AMI snapshot.

https://docs.aws.amazon.com/kms/latest/developerguide/key-policy-modifying-external-accounts.html

upvoted 6 times

Jtic Most Recent 2 weeks, 3 days ago

Selected Answer: B

Must use and share the existing KMS key to decrypt the same key

upvoted 2 times

flbcobra 2 weeks, 4 days ago

Selected Answer: B

https://aws.amazon.com/premiumsupport/knowledge-center/acm-certificate-expiration/

upvoted 1 times

ManoAni 1 month ago

Selected Answer: B

If EBS snapshots are encrypted, then we need to share the same KMS key to partners to be able to access it. Read the note section in the link

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/sharingamis-explicit.html

upvoted 1 times

tubtab 1 month, 1 week ago

Selected Answer: C

MOST secure way should be C

upvoted 1 times

Chunsli 1 month, 1 week ago

MOST secure way should be C, with a separate key, not the one already used.

upvoted 1 times

Jtic 2 weeks, 3 days ago

Must use and share the existing KMS key to decrypt the same key

upvoted 1 times

Sauran 1 month, 1 week ago

A seperate/new key is not possible because it won't be able to decrypt the AMI snapshot which was already encrypted with the existing/old key.

upvoted 6 times

UWSFish 1 month ago

This is truth

upvoted 2 times

Question #81 *Topic 1*

A solutions architect is designing the cloud architecture for a new application being deployed on AWS. The process should run in parallel while adding and removing application nodes as needed based on the number of jobs to be processed. The processor application is stateless. The solutions architect must ensure that the application is loosely coupled and the job items are durably stored.
Which design should the solutions architect use?

A. Create an Amazon SNS topic to send the jobs that need to be processed. Create an Amazon Machine Image (AMI) that consists of the processor application. Create a launch configuration that uses the AMI. Create an Auto Scaling group using the launch configuration. Set the scaling policy for the Auto Scaling group to add and remove nodes based on CPU usage.

B. Create an Amazon SQS queue to hold the jobs that need to be processed. Create an Amazon Machine Image (AMI) that consists of the processor application. Create a launch configuration that uses the AMI. Create an Auto Scaling group using the launch configuration. Set the scaling policy for the Auto Scaling group to add and remove nodes based on network usage.

C. Create an Amazon SQS queue to hold the jobs that need to be processed. Create an Amazon Machine Image (AMI) that consists of the processor application. Create a launch template that uses the AMI. Create an Auto Scaling group using the launch template. Set the scaling policy for the Auto Scaling group to add and remove nodes based on the number of items in the SQS queue.

D. Create an Amazon SNS topic to send the jobs that need to be processed. Create an Amazon Machine Image (AMI) that consists of the processor application. Create a launch template that uses the AMI. Create an Auto Scaling group using the launch template. Set the scaling policy for the Auto Scaling group to add and remove nodes based on the number of messages published to the SNS topic.

Correct Answer: *C*

Community vote distribution

C (100%)

Liliwood 3 days, 13 hours ago
I was not sure between Launch template and Launch configuration.
upvoted 1 times

Wpcorgan 6 days, 21 hours ago
C is correct
upvoted 1 times

devopspro 2 weeks, 6 days ago
Selected Answer: C
answer is c
upvoted 1 times

Wilson_S 3 weeks, 1 day ago
https://www.examtopics.com/discussions/amazon/view/22139-exam-aws-certified-solutions-architect-associate-saa-c02/
upvoted 2 times

wookchan 1 month ago
It looks like C
upvoted 1 times

dokaedu 1 month ago
Correct Answer: C
upvoted 1 times

Question #82 *Topic 1*

A company hosts its web applications in the AWS Cloud. The company configures Elastic Load Balancers to use certificates that are imported into AWS Certificate Manager (ACM). The company's security team must be notified 30 days before the expiration of each certificate.
What should a solutions architect recommend to meet this requirement?

A. Add a rule in ACM to publish a custom message to an Amazon Simple Notification Service (Amazon SNS) topic every day, beginning 30 days before any certificate will expire.

B. Create an AWS Config rule that checks for certificates that will expire within 30 days. Configure Amazon EventBridge (Amazon CloudWatch Events) to invoke a custom alert by way of Amazon Simple Notification Service (Amazon SNS) when AWS Config reports a noncompliant resource.

C. Use AWS Trusted Advisor to check for certificates that will expire within 30 days. Create an Amazon CloudWatch alarm that is based on Trusted Advisor metrics for check status changes. Configure the alarm to send a custom alert by way of Amazon Simple Notification Service (Amazon SNS).

D. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to detect any certificates that will expire within 30 days. Configure the rule to invoke an AWS Lambda function. Configure the Lambda function to send a custom alert by way of Amazon Simple Notification Service (Amazon SNS).

Correct Answer: *D*

Community vote distribution

B (100%)

LeGloupier Highly Voted 1 month, 1 week ago

B
AWS Config has a managed rule
named acm-certificate-expiration-check
to check for expiring certificates
(configurable number of days)

upvoted 8 times

LeGloupier 1 month, 1 week ago

https://aws.amazon.com/premiumsupport/knowledge-center/acm-certificate-expiration/

upvoted 4 times

ManoAni Highly Voted 1 month ago

Selected Answer: B

https://aws.amazon.com/premiumsupport/knowledge-center/acm-certificate-expiration/

upvoted 6 times

ashish_t Most Recent 2 days, 18 hours ago

Selected Answer: B

I think,
For SAA-C02: D
But for SAA-C03: B

upvoted 1 times

Wpcorgan 6 days, 21 hours ago

B is correct

upvoted 1 times

kvsomu 1 week, 1 day ago

With B option we won't get a notification, in the question we have option, must be notified.
So the answer is B

upvoted 1 times

htangga 1 week, 4 days ago

D is perfect : https://aws.amazon.com/blogs/security/how-to-monitor-expirations-of-imported-certificates-in-aws-certificate-manager-acm/

upvoted 2 times

Onimole 3 weeks, 4 days ago

weird but i could also configure this using eventbridge without needing config.config just makes it easier

upvoted 1 times

Vesperia 2 days, 20 hours ago

Eventbridge is just an event bus, requires event sources. For expiring certificate, Config or now ACM can be the source.
upvoted 1 times

jveluman 1 month ago

Selected Answer: B

B is Perfect
upvoted 1 times

Moonus 1 month, 1 week ago

Selected Answer: B

B is Perfect
upvoted 1 times

Question #83 *Topic 1*

A company's dynamic website is hosted using on-premises servers in the United States. The company is launching its product in Europe, and it wants to optimize site loading times for new European users. The site's backend must remain in the United States. The product is being launched in a few days, and an immediate solution is needed.
What should the solutions architect recommend?

A. Launch an Amazon EC2 instance in us-east-1 and migrate the site to it.

B. Move the website to Amazon S3. Use Cross-Region Replication between Regions.

C. Use Amazon CloudFront with a custom origin pointing to the on-premises servers.

D. Use an Amazon Route 53 geoproximity routing policy pointing to on-premises servers.

Correct Answer: *C*

Community vote distribution
C (100%)

koreanmonkey 8 hours, 45 minutes ago

Selected Answer: C

CloudFront = CDN Service
upvoted 1 times

Liliwood 3 days, 13 hours ago

C.
S3 Cross region Replication minimize latency but also copies objects across Amazon S3 buckets in different AWS Regions(data has to remain in origin thou) so B wrong.
Route 53 geo, does not help reducing the latency.
upvoted 1 times

Wpcorgan 6 days, 21 hours ago

C is correct
upvoted 1 times

Hunkie 3 weeks, 5 days ago

Same question with detailed explanation

https://www.examtopics.com/discussions/amazon/view/27898-exam-aws-certified-solutions-architect-associate-saa-c02/
upvoted 2 times

ArielSchivo 1 month, 1 week ago

Selected Answer: C

Option C, use CloudFront.
upvoted 3 times

Question #84 *Topic 1*

A company wants to reduce the cost of its existing three-tier web architecture. The web, application, and database servers are running on Amazon EC2 instances for the development, test, and production environments. The EC2 instances average 30% CPU utilization during peak hours and 10% CPU utilization during non-peak hours.
The production EC2 instances run 24 hours a day. The development and test EC2 instances run for at least 8 hours each day. The company plans to implement automation to stop the development and test EC2 instances when they are not in use.
Which EC2 instance purchasing solution will meet the company's requirements MOST cost-effectively?

A. Use Spot Instances for the production EC2 instances. Use Reserved Instances for the development and test EC2 instances.

B. Use Reserved Instances for the production EC2 instances. Use On-Demand Instances for the development and test EC2 instances.

C. Use Spot blocks for the production EC2 instances. Use Reserved Instances for the development and test EC2 instances.

D. Use On-Demand Instances for the production EC2 instances. Use Spot blocks for the development and test EC2 instances.

Correct Answer: *B*

Community vote distribution

B (100%)

ArielSchivo Highly Voted 1 month, 1 week ago

Selected Answer: B

Spot blocks are not longer available, and you can't use spot instances on Prod machines 24x7, so option B should be valid.

upvoted 6 times

Wpcorgan Most Recent 6 days, 21 hours ago

B is correct

upvoted 1 times

Jtic 2 weeks ago

Selected Answer: B

Reserved Instances and On-demand

Spot is out as the use case required continues instance running

upvoted 1 times

Nigma 3 weeks ago

B is the answer

https://www.examtopics.com/discussions/amazon/view/80956-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Question #85 *Topic 1*

A company has a production web application in which users upload documents through a web interface or a mobile app. According to a new regulatory requirement. new documents cannot be modified or deleted after they are stored.
What should a solutions architect do to meet this requirement?

A. Store the uploaded documents in an Amazon S3 bucket with S3 Versioning and S3 Object Lock enabled.

B. Store the uploaded documents in an Amazon S3 bucket. Configure an S3 Lifecycle policy to archive the documents periodically.

C. Store the uploaded documents in an Amazon S3 bucket with S3 Versioning enabled. Configure an ACL to restrict all access to read-only.

D. Store the uploaded documents on an Amazon Elastic File System (Amazon EFS) volume. Access the data by mounting the volume in read-only mode.

Correct Answer: *A*

Community vote distribution

A (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: A

You can use S3 Object Lock to store objects using a write-once-read-many (WORM) model. Object Lock can help prevent objects from being deleted or overwritten for a fixed amount of time or indefinitely. You can use S3 Object Lock to meet regulatory requirements that require WORM storage, or add an extra layer of protection against object changes and deletion.
Versioning is required and automatically activated as Object Lock is enabled.
https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock-overview.html

upvoted 9 times

Wpcorgan Most Recent 6 days, 21 hours ago

A is correct

upvoted 1 times

flbcobra 2 weeks, 4 days ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock-overview.html

upvoted 1 times

Evangelia 1 month, 1 week ago

Selected Answer: A

aaaaaaaaa

upvoted 1 times

Evangelia 1 month, 1 week ago

aaaaaaaaaa

upvoted 1 times

Question #86 *Topic 1*

A company has several web servers that need to frequently access a common Amazon RDS MySQL Multi-AZ DB instance. The company wants a secure method for the web servers to connect to the database while meeting a security requirement to rotate user credentials frequently.
Which solution meets these requirements?

A. Store the database user credentials in AWS Secrets Manager. Grant the necessary IAM permissions to allow the web servers to access AWS Secrets Manager.

B. Store the database user credentials in AWS Systems Manager OpsCenter. Grant the necessary IAM permissions to allow the web servers to access OpsCenter.

C. Store the database user credentials in a secure Amazon S3 bucket. Grant the necessary IAM permissions to allow the web servers to retrieve credentials and access the database.

D. Store the database user credentials in files encrypted with AWS Key Management Service (AWS KMS) on the web server file system. The web server should be able to decrypt the files and access the database.

Correct Answer: *A*

Community vote distribution

A (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: A

Secrets Manager enables you to replace hardcoded credentials in your code, including passwords, with an API call to Secrets Manager to retrieve the secret programmatically. This helps ensure the secret can't be compromised by someone examining your code, because the secret no longer exists in the code. Also, you can configure Secrets Manager to automatically rotate the secret for you according to a specified schedule. This enables you to replace long-term secrets with short-term ones, significantly reducing the risk of compromise.
https://docs.aws.amazon.com/secretsmanager/latest/userguide/intro.html

upvoted 8 times

Wpcorgan Most Recent 6 days, 21 hours ago

A is correct

upvoted 1 times

renekton 1 week, 3 days ago

Selected Answer: A

Answer is A

upvoted 2 times

Question #87 *Topic 1*

A company hosts an application on AWS Lambda functions that are invoked by an Amazon API Gateway API. The Lambda functions save customer data to an Amazon Aurora MySQL database. Whenever the company upgrades the database, the Lambda functions fail to establish database connections until the upgrade is complete. The result is that customer data is not recorded for some of the event.
A solutions architect needs to design a solution that stores customer data that is created during database upgrades.
Which solution will meet these requirements?

A. Provision an Amazon RDS proxy to sit between the Lambda functions and the database. Configure the Lambda functions to connect to the RDS proxy.

B. Increase the run time of the Lambda functions to the maximum. Create a retry mechanism in the code that stores the customer data in the database.

C. Persist the customer data to Lambda local storage. Configure new Lambda functions to scan the local storage to save the customer data to the database.

D. Store the customer data in an Amazon Simple Queue Service (Amazon SQS) FIFO queue. Create a new Lambda function that polls the queue and stores the customer data in the database.

Correct Answer: *A*

Community vote distribution

D (67%) A (33%)

brushek Highly Voted 1 month, 1 week ago

Selected Answer: A

https://aws.amazon.com/rds/proxy/

RDS Proxy minimizes application disruption from outages affecting the availability of your database by automatically connecting to a new database instance while preserving application connections. When failovers occur, RDS Proxy routes requests directly to the new database instance. This reduces failover times for Aurora and RDS databases by up to 66%.

upvoted 10 times

attila9778 5 days, 1 hour ago

Aurora supports RDS proxy!
https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html

upvoted 1 times

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: D

The answer is D.
RDS Proxy doesn't support Aurora DBs. See limitations at:
https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html

upvoted 8 times

gcmrjbr 4 days, 3 hours ago

You can use RDS Proxy with Aurora Serverless v2 clusters but not with Aurora Serverless v1 clusters.
https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html

upvoted 1 times

tinyfoot 1 week ago

Actually RDS Proxy supports Aurora DBs running on PostgreSQL and MySQL.
https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Concepts.Aurora_Fea_Regions_DB-eng.Feature.RDS_Proxy.html

With RDS proxy, you only expose a single endpoint for request to hit and any failure of the primary DB in a Multi-AZ configuration is will be managed automatically by RDS Proxy to point to the new primary DB. Hence RDS proxy is the most efficient way of solving the issue as additional code change is required.
https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.howitworks.html

upvoted 2 times

ocbn3wby Most Recent 2 days, 23 hours ago

Selected Answer: D

No specific information that supports answer A - Aurora has failover feature (promoting replica to master). However, in RDS Proxy documentation it's not specified how to benefit from the failover Aurora feature - nor the present question points to failing over master to replica.

Answer is D.

upvoted 1 times

hollie 1 week, 1 day ago

Selected Answer: A

"You can use RDS Proxy with Aurora Serverless v2 clusters but not with Aurora Serverless v1 clusters." according to https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html. Thus Aurora can be used with RDS proxy.

upvoted 2 times

pspinelli19 2 weeks ago

Selected Answer: D

dddd
https://www.examtopics.com/discussions/amazon/view/61069-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

study_aws1 2 weeks, 3 days ago

All choices point to D).
Only point is the specific reason for mention of Amazon SQS FIFO queue. Not sure how FIFO queue value adds in this case.

upvoted 1 times

flbcobra 2 weeks, 4 days ago

Selected Answer: D

Amazon RDS Proxy is a fully managed, highly available database proxy for Amazon Relational Database Service (RDS) that makes applications more scalable, more resilient to database failures, and more secure.

upvoted 1 times

backbencher2022 3 weeks, 2 days ago

Selected Answer: D

Solution requirement is to store customer data while upgrading. Proxy can't do that. We need SQS

upvoted 4 times

rob74 3 weeks, 5 days ago

Selected Answer: D

The question is " a solution that stores data"-->So for me answer is D-->Store data in SQS

upvoted 5 times

SimonPark 4 weeks, 1 day ago

Selected Answer: A

In my humble opinion, the answer is A.
RDS Proxy makes applications more resilient to database failures by automatically connecting to a standby DB instance while preserving application connections.

upvoted 2 times

ocbn3wby 2 days, 23 hours ago

Hi,

There is no strong documentation that points to this logic: RDS Proxy makes use of the Aurora Failover feature (which is failing over from primary write master to promoting read replicas).

It says it's resilient to failovers, but I could not find configuration on how to choose to what replica to make failover to. I'd still go with answer D.

upvoted 1 times

AbhiJo 1 week, 5 days ago

Using RDS Proxy with Aurora clusters that are part of an Aurora global database isn't currently supported.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: D

answer should be D,
SQS will keep the traffic in the queue while DB is in maintenance

upvoted 2 times

tubtab 1 month, 1 week ago

Selected Answer: D

ddddddddddddddd

upvoted 2 times

Evangelia 1 month, 1 week ago

Selected Answer: D

https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html

upvoted 3 times

palermo777 1 month, 1 week ago

The answer is D.
We don't have any information that customer has several RDS databases. Hence RDS Proxy might not help by redirecting the load to another DB

upvoted 2 times

Moonus 1 month, 1 week ago

Selected Answer: D

D is the answer

upvoted 2 times

Lilibell 1 month, 2 weeks ago

The answer is D

upvoted 2 times

Question #88 *Topic 1*

A survey company has gathered data for several years from areas in the United States. The company hosts the data in an Amazon S3 bucket that is 3 TB in size and growing. The company has started to share the data with a European marketing firm that has S3 buckets. The company wants to ensure that its data transfer costs remain as low as possible.
Which solution will meet these requirements?

A. Configure the Requester Pays feature on the company's S3 bucket.

B. Configure S3 Cross-Region Replication from the company's S3 bucket to one of the marketing firm's S3 buckets.

C. Configure cross-account access for the marketing firm so that the marketing firm has access to the company's S3 bucket.

D. Configure the company's S3 bucket to use S3 Intelligent-Tiering. Sync the S3 bucket to one of the marketing firm's S3 buckets.

Correct Answer: *B*

Community vote distribution

A (52%) B (48%)

Six_Fingered_Jose Highly Voted 1 month ago

Selected Answer: B

this question is too vague imho
if the question is looking for a way to incur charges to the European company instead of the US company, then requester pay makes sense.

if they are looking to reduce overall data transfer cost, then B makes sense because the data does not leave the AWS network, thus data transfer cost should be lower technically?

A. makes sense because the US company saves money, but the European company is paying for the charges so there is no overall saving in cost when you look at the big picture

I will go for B because they are not explicitly stating that they want the other company to pay for the charges

upvoted 9 times

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: A

"Typically, you configure buckets to be Requester Pays buckets when you want to share data but not incur charges associated with others accessing the data. For example, you might use Requester Pays buckets when making available large datasets, such as zip code directories, reference data, geospatial information, or web crawling data."
https://docs.aws.amazon.com/AmazonS3/latest/userguide/RequesterPaysBuckets.html

upvoted 8 times

Vesperia Most Recent 2 days, 20 hours ago

Selected Answer: B

B makes more sense. Using solution A, you have to code it for every single object in the bucket with the x-amz-request-payer. When you share data, it's not likely for per view retrieval only, European marketing firm would run analytics on the whole data set. I agree the question does not describe cleanly though.

upvoted 1 times

ocbn3wby 2 days, 22 hours ago

Selected Answer: B

The company has started to SHARE data - it's part of the company's business to send its data at minimal costs.

Let's think like this: If you own the European business which has own S3 buckets, and you are aware that there is CRR at low costs, but the provider of your data forces you to pay for all of the charges associated with "requester pays" - would you still do business with the US company?

Answer B is closer to my heart.

upvoted 1 times

Wpcorgan 6 days, 21 hours ago

A is correct

upvoted 1 times

edk11z 1 week, 1 day ago

Selected Answer: A

The company wants to reduce ITS data cost. Not the marketing firm's.

upvoted 1 times

renekton 1 week, 3 days ago

Selected Answer: A

Answer is A
"has started to share the data" so they already sharing and "The company wants to ensure that its data transfer costs remain as low as possible" this is A trust me.
upvoted 1 times

Ack3rman 1 week, 4 days ago

Selected Answer: B

If you are setting the replication configuration in a cross-account scenario, where the source and destination buckets are owned by different AWS accounts, the following additional requirement applies:

The owner of the destination buckets must grant the owner of the source bucket permissions to replicate objects with a bucket policy.

The destination buckets cannot be configured as Requester Pays buckets. (so can't be A)
https://docs.aws.amazon.com/AmazonS3/latest/userguide/replication.html
upvoted 1 times

pspinelli19 2 weeks ago

Selected Answer: A

aa
https://www.examtopics.com/discussions/amazon/view/82811-exam-aws-certified-solutions-architect-associate-saa-c02/
upvoted 2 times

vladislove 2 weeks ago

Selected Answer: A

The answer is A, because of the wording of the requirement: "The company wants to ensure that ITS data transfer costs remain as low as possible". If Requester Pays is enabled on the bucket, then the Marketing Firm will exclusively be paying the data transfer charges. The Company will only pay for data storage.
upvoted 3 times

EKA_CloudGod 2 weeks, 3 days ago

Answer is A.
https://www.examtopics.com/discussions/amazon/view/82811-exam-aws-certified-solutions-architect-associate-saa-c02/
upvoted 2 times

PS_R 2 weeks, 5 days ago

Selected Answer: B

Yes B Makes more sense, CRR - S3.
upvoted 2 times

vn_thanhtung 2 weeks, 3 days ago

When to use Cross-Region Replication
S3 Cross-Region Replication (CRR) is used to copy objects across Amazon S3 buckets in different AWS Regions. CRR can help you do the following:

Meet compliance requirements – Although Amazon S3 stores your data across multiple geographically distant Availability Zones by default, compliance requirements might dictate that you store data at even greater distances. To satisfy these requirements, use Cross-Region Replication to replicate data between distant AWS Regions.

Minimize latency – If your customers are in two geographic locations, you can minimize latency in accessing objects by maintaining object copies in AWS Regions that are geographically closer to your users.

Increase operational efficiency – If you have compute clusters in two different AWS Regions that analyze the same set of objects, you might choose to maintain object copies in those Regions.
upvoted 2 times

srao123 2 weeks, 6 days ago

B is the correct answer. As the data transfer costs (cross region replication)from S3 to another region is $0.02 which uses AWS backbone. Where option C uses internet for data transfer out from S3 which costs $0.09.
https://aws.amazon.com/s3/pricing/
upvoted 1 times

srao123 2 weeks, 6 days ago

A is incorrect because the requirement is to keep the data transfer costs as low as possible not who pays for it. We need to pick the option that costs less.
upvoted 1 times

Deplake 3 weeks ago

Selected Answer: A

B is not looking correct, because it says that you are going to replicate S3 from your account to another account, not sure it is possible
upvoted 2 times

raje7676 2 weeks, 2 days ago

https://docs.aws.amazon.com/AmazonS3/latest/userguide/replication-walkthrough-2.html

Setting up replication when source and destination buckets are owned by different AWS accounts is similar to setting replication when both

buckets are owned by the same account. The only difference is that the destination bucket owner must grant the source bucket owner permission to replicate objects by adding a bucket policy.

upvoted 1 times

SimonPark 4 weeks, 1 day ago

Selected Answer: B

B seems to make more sense to me

upvoted 2 times

Evangelia 1 month, 1 week ago

Why not C? Could someone explain?

upvoted 2 times

th3cookie 1 week, 3 days ago

If they give access for the bucket to the European company, they incur all of the data transfer charges (which for 3tb will be a lot). Using CRR to one of the European's owned buckets will be cheaper for the company hosting the data.

upvoted 1 times

RODROSKAR 3 weeks, 1 day ago

Agreed. B implies they are in the same AWS account. Which the question suggests, they are not. "With a European marketing firm"

upvoted 2 times

Awstech 1 month, 1 week ago

A. Requestor Pays

upvoted 3 times

Question #89 *Topic 1*

A company uses Amazon S3 to store its confidential audit documents. The S3 bucket uses bucket policies to restrict access to audit team IAM user credentials according to the principle of least privilege. Company managers are worried about accidental deletion of documents in the S3 bucket and want a more secure solution.
What should a solutions architect do to secure the audit documents?

A. Enable the versioning and MFA Delete features on the S3 bucket.

B. Enable multi-factor authentication (MFA) on the IAM user credentials for each audit team IAM user account.

C. Add an S3 Lifecycle policy to the audit team's IAM user accounts to deny the s3:DeleteObject action during audit dates.

D. Use AWS Key Management Service (AWS KMS) to encrypt the S3 bucket and restrict audit team IAM user accounts from accessing the KMS key.

Correct Answer: *A*

Community vote distribution

A (100%)

Wpcorgan 6 days, 21 hours ago

A is correct

upvoted 1 times

Jtic 2 weeks ago

Selected Answer: A

Enable the versioning and MFA Delete features on the S3 bucket.

upvoted 1 times

123jhl0 1 month, 1 week ago

Selected Answer: A

Same as Question #44

upvoted 4 times

Question #90 *Topic 1*

A company is using a SQL database to store movie data that is publicly accessible. The database runs on an Amazon RDS Single-AZ DB instance. A script runs queries at random intervals each day to record the number of new movies that have been added to the database. The script must report a final total during business hours.
The company's development team notices that the database performance is inadequate for development tasks when the script is running. A solutions architect must recommend a solution to resolve this issue.
Which solution will meet this requirement with the LEAST operational overhead?

A. Modify the DB instance to be a Multi-AZ deployment.

B. Create a read replica of the database. Configure the script to query only the read replica.

C. Instruct the development team to manually export the entries in the database at the end of each day.

D. Use Amazon ElastiCache to cache the common queries that the script runs against the database.

Correct Answer: *D*

Community vote distribution

B (92%) 8%

alvarez100 Highly Voted 1 month, 1 week ago

Selected Answer: B

Elasti Cache if for reading common results. The script is looking for new movies added. Read replica would be the best choice.

upvoted 10 times

Vesperia Most Recent 2 days, 19 hours ago

Caching works best for static contents. When you run a total, you need to go through all the records in a table. The question is not constructed properly. Best solution is to create an index on the added date, it won't take long, nor heavy io/cpu to get the total number of newly added total for the day. This approach takes minimal effort, does not incur any extra charge, better than both B and D.

upvoted 1 times

Vesperia 2 days, 19 hours ago

I would choose B as the answer. For the stated type of queries It's better than D .

upvoted 1 times

Wpcorgan 6 days, 21 hours ago

B is correct

upvoted 1 times

htangga 1 week, 4 days ago

Selected Answer: B

B is more make sense for me

upvoted 1 times

Bevemo 2 weeks, 6 days ago

Selected Answer: B

Not D as apps have to be re-written to take advantage of elasticache APIs - that is too much overhead.

upvoted 1 times

Gil80 3 weeks ago

Selected Answer: B

• You have a production DB that is taking on a normal load
• You want to run a reporting application to run some analytics
• You create a read replica to run the new workload there
• The prod application is unaffected
• Read replicas are used for SELECT (=read) only kind of statements
Therefore I believe B to be the better answer.

As for "D" - ElastiCache use cases are:
1. Your data is slow or expensive to get when compared to cache retrieval.
2. Users access your data often.
3. Your data stays relatively the same, or if it changes quickly staleness is not a large issue.

1 - Somewhat true.
2 - Not true for our case.
3 - Also not true. The data changes throughout the day.

For my understanding, caching has to do with millisecond results, high-performance reads. These are not the issues mentioned in the questions, therefore B.

upvoted 4 times

Cizzla7049 3 weeks, 1 day ago

Selected Answer: D

Even though B is correct, it says least operational overhead which is D. Like the other person said, AWS used similar use cases.

upvoted 1 times

Az900500 3 weeks, 1 day ago

When you want to free your database from occasional operation overhead, read replica is key ; the operations runs on the replica while the db run continue operation smoothly.. ElasticCahe would have been considered "If DB has issue before or after running the script"

B is the answer

upvoted 1 times

Hunkie 3 weeks, 4 days ago

Selected Answer: D

For me its D. It says least operation overhead. Its the use case of elastic cache.

https://docs.aws.amazon.com/AmazonElastiCache/latest/mem-ug/elasticache-use-cases.html#:~:text=Speed%20and%20expense,on%20other%0A%09%09%09%09factors

upvoted 1 times

Bevemo 2 weeks, 6 days ago

Apps have to be re-written to take advantage of elasticache APIs - that is too much overhead.

upvoted 1 times

ArielSchivo 3 weeks ago

Not sure about that since it says "Your data stays relatively the same, or if it changes quickly staleness is not a large issue", and this is a case related to new movies added to the database.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: B

agree with B, the script only reads the data to produce a report and the slowdown only occurs when the script is running, thus having a read replica solves the issue with the least overhead

upvoted 3 times

tubtab 1 month, 1 week ago

Selected Answer: B

bbbbbbbbbbbbb

upvoted 2 times

Chunsli 1 month, 1 week ago

B seems to make a better sense

upvoted 2 times

brushek 1 month, 2 weeks ago

Selected Answer: B

It should be read replica.

upvoted 2 times

Question #91 *Topic 1*

A company has applications that run on Amazon EC2 instances in a VPC. One of the applications needs to call the Amazon S3 API to store and read objects. According to the company's security regulations, no traffic from the applications is allowed to travel across the internet.
Which solution will meet these requirements?

A. Configure an S3 gateway endpoint.

B. Create an S3 bucket in a private subnet.

C. Create an S3 bucket in the same AWS Region as the EC2 instances.

D. Configure a NAT gateway in the same subnet as the EC2 instances.

Correct Answer: *A*

Community vote distribution

A (100%)

ArielSchivo Highly Voted 1 month, 1 week ago

Selected Answer: A

Gateway endpoints provide reliable connectivity to Amazon S3 and DynamoDB without requiring an internet gateway or a NAT device for your VPC. It should be option A.

https://docs.aws.amazon.com/vpc/latest/privatelink/gateway-endpoints.html

upvoted 7 times

Wpcorgan Most Recent 6 days, 20 hours ago

A is correct

upvoted 1 times

Question #92 *Topic 1*

A company is storing sensitive user information in an Amazon S3 bucket. The company wants to provide secure access to this bucket from the application tier running on Amazon EC2 instances inside a VPC.
Which combination of steps should a solutions architect take to accomplish this? (Choose two.)

A. Configure a VPC gateway endpoint for Amazon S3 within the VPC.

B. Create a bucket policy to make the objects in the S3 bucket public.

C. Create a bucket policy that limits access to only the application tier running in the VPC.

D. Create an IAM user with an S3 access policy and copy the IAM credentials to the EC2 instance.

E. Create a NAT instance and have the EC2 instances use the NAT instance to access the S3 bucket.

Correct Answer: *AC*

Community vote distribution

AC (100%)

DivaLight 1 day, 19 hours ago

Selected Answer: AC

Option AC

upvoted 1 times

Wpcorgan 6 days, 20 hours ago

A and C

upvoted 1 times

Jtic 2 weeks ago

Selected Answer: AC

AC is the correct answer in the use case

upvoted 1 times

rjam 2 weeks, 3 days ago

Options A and E

upvoted 1 times

rjam 2 weeks, 3 days ago

Typo it should be A and C

upvoted 1 times

ArielSchivo 1 month, 1 week ago

Selected Answer: AC

Options A and C.

https://aws.amazon.com/premiumsupport/knowledge-center/s3-private-connection-no-authentication/

upvoted 4 times

Question #93 *Topic 1*

A company runs an on-premises application that is powered by a MySQL database. The company is migrating the application to AWS to increase the application's elasticity and availability.
The current architecture shows heavy read activity on the database during times of normal operation. Every 4 hours, the company's development team pulls a full export of the production database to populate a database in the staging environment. During this period, users experience unacceptable application latency. The development team is unable to use the staging environment until the procedure completes.
A solutions architect must recommend replacement architecture that alleviates the application latency issue. The replacement architecture also must give the development team the ability to continue using the staging environment without delay.
Which solution meets these requirements?

A. Use Amazon Aurora MySQL with Multi-AZ Aurora Replicas for production. Populate the staging database by implementing a backup and restore process that uses the mysqldump utility.

B. Use Amazon Aurora MySQL with Multi-AZ Aurora Replicas for production. Use database cloning to create the staging database on-demand.

C. Use Amazon RDS for MySQL with a Multi-AZ deployment and read replicas for production. Use the standby instance for the staging database.

D. Use Amazon RDS for MySQL with a Multi-AZ deployment and read replicas for production. Populate the staging database by implementing a backup and restore process that uses the mysqldump utility.

Correct Answer: *B*

Community vote distribution

B (100%)

DivaLight 1 day, 19 hours ago

Selected Answer: B

Option B

upvoted 1 times

pspinelli19 2 weeks ago

Selected Answer: B

Amazon Aurora Fast Database Cloning is what is required here.
https://aws.amazon.com/blogs/aws/amazon-aurora-fast-database-cloning/

upvoted 1 times

KLLIM 4 weeks, 1 day ago

https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/Aurora.Managing.Clone.html

upvoted 1 times

LeGloupier 1 month, 1 week ago

Selected Answer: B

B
Database cloning

upvoted 4 times

Question #94 *Topic 1*

A company is designing an application where users upload small files into Amazon S3. After a user uploads a file, the file requires one-time simple processing to transform the data and save the data in JSON format for later analysis.
Each file must be processed as quickly as possible after it is uploaded. Demand will vary. On some days, users will upload a high number of files. On other days, users will upload a few files or no files.
Which solution meets these requirements with the LEAST operational overhead?

A. Configure Amazon EMR to read text files from Amazon S3. Run processing scripts to transform the data. Store the resulting JSON file in an Amazon Aurora DB cluster.

B. Configure Amazon S3 to send an event notification to an Amazon Simple Queue Service (Amazon SQS) queue. Use Amazon EC2 instances to read from the queue and process the data. Store the resulting JSON file in Amazon DynamoDB.

C. Configure Amazon S3 to send an event notification to an Amazon Simple Queue Service (Amazon SQS) queue. Use an AWS Lambda function to read from the queue and process the data. Store the resulting JSON file in Amazon DynamoDB.

D. Configure Amazon EventBridge (Amazon CloudWatch Events) to send an event to Amazon Kinesis Data Streams when a new file is uploaded. Use an AWS Lambda function to consume the event from the stream and process the data. Store the resulting JSON file in an Amazon Aurora DB cluster.

Correct Answer: *C*

Community vote distribution

C (100%)

Kapello10 1 day, 2 hours ago

Selected Answer: C

ccccccccccc

upvoted 1 times

DivaLight 1 day, 19 hours ago

Selected Answer: C

Option C

upvoted 1 times

Wpcorgan 6 days, 20 hours ago

C is correct

upvoted 1 times

Pamban 2 weeks ago

Selected Answer: C

https://www.examtopics.com/discussions/amazon/view/67958-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Jtic 2 weeks ago

Selected Answer: C

SQS + LAMDA + JSON to Dynamo DB

upvoted 1 times

Hunkie 3 weeks, 5 days ago

With explanations

https://www.examtopics.com/discussions/amazon/view/67958-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

rjam 4 weeks, 1 day ago

Option C
Dynamo DB is a NoSQL-JSON supported

upvoted 3 times

rjam 4 weeks, 1 day ago

also Use an AWS Lambda - serverless - less operational overhead

upvoted 3 times

Question #95 *Topic 1*

An application allows users at a company's headquarters to access product data. The product data is stored in an Amazon RDS MySQL DB instance. The operations team has isolated an application performance slowdown and wants to separate read traffic from write traffic. A solutions architect needs to optimize the application's performance quickly.
What should the solutions architect recommend?

A. Change the existing database to a Multi-AZ deployment. Serve the read requests from the primary Availability Zone.

B. Change the existing database to a Multi-AZ deployment. Serve the read requests from the secondary Availability Zone.

C. Create read replicas for the database. Configure the read replicas with half of the compute and storage resources as the source database.

D. Create read replicas for the database. Configure the read replicas with the same compute and storage resources as the source database.

Correct Answer: *D*

Community vote distribution

D (100%)

DivaLight 1 day, 19 hours ago

Selected Answer: D

Option D

upvoted 1 times

Wpcorgan 6 days, 20 hours ago

D is correct

upvoted 1 times

Nigma 3 weeks ago

D

https://www.examtopics.com/discussions/amazon/view/46461-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Hunkie 3 weeks, 5 days ago

Selected Answer: D

If you scale the source DB instance, also scale the read replicas.

upvoted 2 times

ArielSchivo 1 month, 1 week ago

Selected Answer: D

D is correct.

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_MySQL.Replication.ReadReplicas.html

upvoted 1 times

Question #96 *Topic 1*

An Amazon EC2 administrator created the following policy associated with an IAM group containing several users:

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "ec2:TerminateInstances",
            "Resource": "*",
            "Condition": {
                "IpAddress": {
                    "aws:SourceIp": "10.100.100.0/24"
                }
            }
        },
        {
            "Effect": "Deny",
            "Action": "ec2:*",
            "Resource": "*",
            "Condition": {
                "StringNotEquals": {
                    "ec2:Region": "us-east-1"
                }
            }
        }
    ]
}
```

What is the effect of this policy?

A. Users can terminate an EC2 instance in any AWS Region except us-east-1.

B. Users can terminate an EC2 instance with the IP address 10.100.100.1 in the us-east-1 Region.

C. Users can terminate an EC2 instance in the us-east-1 Region when the user's source IP is 10.100.100.254.

D. Users cannot terminate an EC2 instance in the us-east-1 Region when the user's source IP is 10.100.100.254.

Correct Answer: *C*

Community vote distribution

C (100%)

Wpcorgan 6 days, 20 hours ago

C is correct

upvoted 1 times

rjam 1 week, 5 days ago

Selected Answer: C

Option C is the correct answer.

if CIDR block 10.100.100.0/24, then reserved IP addresses are:
• 10.100.100.0 – Network Address
• 10.100.100.1 – reserved by AWS for the VPC router
• 10.100.100.2 – reserved by AWS for mapping to Amazon-provided DNS
• 10.100.100.3 – reserved by AWS for future use
• 10.100.100.255 – Network Broadcast Address. AWS does not support broadcast in a VPC,

The above numbers cannot be used. This rules out the option B

upvoted 2 times

Jtic 2 weeks ago

Selected Answer: C

should be Last IP 10.100.100.254

Not an option
10.100.100.1: Reserved by AWS for the VPC router

upvoted 2 times

masetromain 1 month ago

Selected Answer: C

https://www.examtopics.com/discussions/amazon/view/27814-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 3 times

Question #97 *Topic 1*

A company has a large Microsoft SharePoint deployment running on-premises that requires Microsoft Windows shared file storage. The company wants to migrate this workload to the AWS Cloud and is considering various storage options. The storage solution must be highly available and integrated with Active Directory for access control.
Which solution will satisfy these requirements?

A. Configure Amazon EFS storage and set the Active Directory domain for authentication.

B. Create an SMB file share on an AWS Storage Gateway file gateway in two Availability Zones.

C. Create an Amazon S3 bucket and configure Microsoft Windows Server to mount it as a volume.

D. Create an Amazon FSx for Windows File Server file system on AWS and set the Active Directory domain for authentication.

Correct Answer: *D*

Community vote distribution

D (100%)

Wpcorgan 6 days, 20 hours ago
D is correct
upvoted 1 times

TonyghostR05 1 week, 5 days ago
Window only available for using FSx
upvoted 2 times

Nigma 3 weeks ago
D. Windows is the keyword

https://www.examtopics.com/discussions/amazon/view/29780-exam-aws-certified-solutions-architect-associate-saa-c02/
upvoted 1 times

Nigma 3 weeks ago
EFS is for Linux
FSx is for Windows
upvoted 2 times

Hunkie 3 weeks, 5 days ago
Selected Answer: D
DDDDDDDD
upvoted 1 times

dokaedu 1 month ago
Correct Answer:D
https://docs.aws.amazon.com/fsx/latest/WindowsGuide/aws-ad-integration-fsxW.html
upvoted 2 times

Question #98 *Topic 1*

An image-processing company has a web application that users use to upload images. The application uploads the images into an Amazon S3 bucket. The company has set up S3 event notifications to publish the object creation events to an Amazon Simple Queue Service (Amazon SQS) standard queue. The SQS queue serves as the event source for an AWS Lambda function that processes the images and sends the results to users through email.
Users report that they are receiving multiple email messages for every uploaded image. A solutions architect determines that SQS messages are invoking the Lambda function more than once, resulting in multiple email messages.
What should the solutions architect do to resolve this issue with the LEAST operational overhead?

A. Set up long polling in the SQS queue by increasing the ReceiveMessage wait time to 30 seconds.

B. Change the SQS standard queue to an SQS FIFO queue. Use the message deduplication ID to discard duplicate messages.

C. Increase the visibility timeout in the SQS queue to a value that is greater than the total of the function timeout and the batch window timeout.

D. Modify the Lambda function to delete each message from the SQS queue immediately after the message is read before processing.

Correct Answer: *A*

Community vote distribution

C (100%)

Six_Fingered_Jose Highly Voted 1 month ago

Selected Answer: C

answer should be C,
users get duplicated messages because -> lambda polls the message, and starts processing the message.
However, before the first lambda can finish processing the message, the visibility timeout runs out on SQS, and SQS returns the message to the poll, causing another Lambda node to process that same message.
By increasing the visibility timeout, it should prevent SQS from returning a message back to the poll before Lambda can finish processing the message

upvoted 7 times

brushek Highly Voted 1 month, 2 weeks ago

Selected Answer: C

https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-visibility-timeout.html

this is important part:
Immediately after a message is received, it remains in the queue. To prevent other consumers from processing the message again, Amazon SQS sets a visibility timeout, a period of time during which Amazon SQS prevents other consumers from receiving and processing the message. The default visibility timeout for a message is 30 seconds. The minimum is 0 seconds. The maximum is 12 hours.

upvoted 5 times

Wpcorgan Most Recent 6 days, 20 hours ago

C is correct

upvoted 1 times

Hunkie 3 weeks, 5 days ago

Selected Answer: C

https://www.examtopics.com/discussions/amazon/view/83096-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

rob74 1 month, 1 week ago

I exlude Polling because-->"The maximum long polling wait time is 20 seconds"
https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-short-and-long-polling.html

upvoted 2 times

MXB05 1 month, 2 weeks ago

Selected Answer: C

A can not be correct, long polling will only ensure that all images are retrieved from all SQS servers in one query. If the same message triggers the lambda function twice it is likely because the visibility timeout isn't long enough and lambda didn't repsond in time with a deletion of the message
->> C is correct

upvoted 4 times

Question #99 *Topic 1*

A company is implementing a shared storage solution for a gaming application that is hosted in an on-premises data center. The company needs the ability to use Lustre clients to access data. The solution must be fully managed.
Which solution meets these requirements?

A. Create an AWS Storage Gateway file gateway. Create a file share that uses the required client protocol. Connect the application server to the file share.

B. Create an Amazon EC2 Windows instance. Install and configure a Windows file share role on the instance. Connect the application server to the file share.

C. Create an Amazon Elastic File System (Amazon EFS) file system, and configure it to support Lustre. Attach the file system to the origin server. Connect the application server to the file system.

D. Create an Amazon FSx for Lustre file system. Attach the file system to the origin server. Connect the application server to the file system.

Correct Answer: *D*

Community vote distribution

D (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: D

Answer is D.
Lustre in the question is only available as FSx
https://aws.amazon.com/fsx/lustre/

upvoted 8 times

Wpcorgan Most Recent 6 days, 20 hours ago

D is correct

upvoted 1 times

Question #100 *Topic 1*

A company's containerized application runs on an Amazon EC2 instance. The application needs to download security certificates before it can communicate with other business applications. The company wants a highly secure solution to encrypt and decrypt the certificates in near real time. The solution also needs to store data in highly available storage after the data is encrypted.
Which solution will meet these requirements with the LEAST operational overhead?

A. Create AWS Secrets Manager secrets for encrypted certificates. Manually update the certificates as needed. Control access to the data by using fine-grained IAM access.

B. Create an AWS Lambda function that uses the Python cryptography library to receive and perform encryption operations. Store the function in an Amazon S3 bucket.

C. Create an AWS Key Management Service (AWS KMS) customer managed key. Allow the EC2 role to use the KMS key for encryption operations. Store the encrypted data on Amazon S3.

D. Create an AWS Key Management Service (AWS KMS) customer managed key. Allow the EC2 role to use the KMS key for encryption operations. Store the encrypted data on Amazon Elastic Block Store (Amazon EBS) volumes.

Correct Answer: *D*

Community vote distribution

C (89%) 11%

Chunsli Highly Voted 1 month, 1 week ago
C makes a better sense. Between C (S3) and D (EBS), S3 is highly available with LEAST operational overhead.
upvoted 11 times

MXB05 Highly Voted 1 month, 2 weeks ago
Selected Answer: C
Correct Answer is C: EBS is not highly available
upvoted 10 times

Wpcorgan Most Recent 6 days, 20 hours ago
C is correct
upvoted 1 times

Cizzla7049 1 week, 5 days ago
Selected Answer: D
D is correct answer. Least overhead, just open EBS while launching EC2. It's not about which one is the most highly available, it's about which one fits the scenario best.
https://aws.amazon.com/blogs/compute/must-know-best-practices-for-amazon-ebs-encryption/
upvoted 1 times

SimonPark 4 weeks, 1 day ago
Selected Answer: C
S3 = highly available storage than EBS
https://www.justaftermidnight247.com/insights/ebs-efs-and-s3-when-to-use-awss-three-storage-solutions/
upvoted 2 times

Six_Fingered_Jose 1 month ago
Selected Answer: C
C should be the answer,
Highly available storage = S3
upvoted 4 times

nikerlas 1 month ago
Selected Answer: D
"store data in highly available storage after the data is encrypted." ==> EBS
upvoted 1 times

Question #101 *Topic 1*

A solutions architect is designing a VPC with public and private subnets. The VPC and subnets use IPv4 CIDR blocks. There is one public subnet and one private subnet in each of three Availability Zones (AZs) for high availability. An internet gateway is used to provide internet access for the public subnets. The private subnets require access to the internet to allow Amazon EC2 instances to download software updates.
What should the solutions architect do to enable Internet access for the private subnets?

A. Create three NAT gateways, one for each public subnet in each AZ. Create a private route table for each AZ that forwards non-VPC traffic to the NAT gateway in its AZ.

B. Create three NAT instances, one for each private subnet in each AZ. Create a private route table for each AZ that forwards non-VPC traffic to the NAT instance in its AZ.

C. Create a second internet gateway on one of the private subnets. Update the route table for the private subnets that forward non-VPC traffic to the private internet gateway.

D. Create an egress-only internet gateway on one of the public subnets. Update the route table for the private subnets that forward non-VPC traffic to the egress-only Internet gateway.

Correct Answer: *A*

Community vote distribution

A (83%) B (17%)

Gil80 2 weeks, 6 days ago

Selected Answer: A

NAT Instances - OUTDATED BUT CAN STILL APPEAR IN THE EXAM!
However, given that A provides the newer option of NAT Gateway, then A is the correct answer.

B would be correct if NAT Gateway wasn't an option.

upvoted 2 times

Deplake 3 weeks ago

Selected Answer: B

Should be B

upvoted 1 times

Nigma 3 weeks ago

https://www.examtopics.com/discussions/amazon/view/35679-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

dave9994 1 month ago

B should be the answer. https://docs.aws.amazon.com/vpc/latest/userguide/vpc-nat-gateway.html

upvoted 1 times

yalvar 5 days, 15 hours ago

Sir, you didn't even read the link you posted !! There it is clearly stated that when you need access to Internet from a private subnet you place the NAT gateway in a PUBLIC subnet.

upvoted 1 times

Gil80 2 weeks, 6 days ago

B is NAT Instances, which is outdated. The link you provided refers to NAT Gateways (the newer approach) - which means, A is the right answer.

upvoted 2 times

Evangelia 1 month, 1 week ago

Selected Answer: A

aaaaaaa

upvoted 3 times

Question #102 *Topic 1*

A company wants to migrate an on-premises data center to AWS. The data center hosts an SFTP server that stores its data on an NFS-based file system. The server holds 200 GB of data that needs to be transferred. The server must be hosted on an Amazon EC2 instance that uses an Amazon Elastic File System (Amazon EFS) file system.
Which combination of steps should a solutions architect take to automate this task? (Choose two.)

A. Launch the EC2 instance into the same Availability Zone as the EFS file system.

B. Install an AWS DataSync agent in the on-premises data center.

C. Create a secondary Amazon Elastic Block Store (Amazon EBS) volume on the EC2 instance for the data.

D. Manually use an operating system copy command to push the data to the EC2 instance.

E. Use AWS DataSync to create a suitable location configuration for the on-premises SFTP server.

Correct Answer: *AB*

Community vote distribution

AB (65%) BE (24%) 12%

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: AB

A. Launch the EC2 instance into the same Availability Zone as the EFS file system.
Makes sense to have the instance in the same AZ the EFS storage is.
B. Install an AWS DataSync agent in the on-premises data center.
The DataSync with move the data to the EFS, which already uses the EC2 instance (see the info provided). No more things are required...
C. Create a secondary Amazon Elastic Block Store (Amazon EBS) volume on the EC2 instance for the data.
This secondary EBS volume isn't required... the data should be move on to EFS...
D. Manually use an operating system copy command to push the data to the EC2 instance.
Potentially possible (instead of A), BUT the "automate this task" premise goes against any "manually" action. So, we should keep A.
E. Use AWS DataSync to create a suitable location configuration for the on-premises SFTP server.
I don't get the relationship between DataSync and the configuration for SFTP "on-prem"! Nonsense.
So, anwers are A&B

upvoted 11 times

attila9778 6 days, 6 hours ago

Can someone explain why A is correct?
EFS is spread across Availability Zones in a region, as per https://aws.amazon.com/blogs/gametech/gearbox-entertainment-goes-remote-with-aws-and-perforce/
My question then is whether it makes sense to launch EC2 instances in the *same Availability Zone as the EFS file system* ?

upvoted 2 times

Cizzla7049 1 week, 4 days ago

E is correct
https://aws.amazon.com/blogs/storage/migrating-storage-with-aws-datasync/

upvoted 1 times

Cizzla7049 Most Recent 8 hours, 21 minutes ago

Selected Answer: BE

https://aws.amazon.com/datasync/

upvoted 1 times

DivaLight 1 day, 17 hours ago

Selected Answer: BE

Option BE

upvoted 1 times

LeGloupier 3 days, 3 hours ago

Selected Answer: BE

B and E make more sens

upvoted 2 times

nVizzz 4 days, 21 hours ago

Selected Answer: AE

I chose A and E, because firstly we need to install DataSync agent to copy from SFTP and then create respective location.

upvoted 1 times

Wpcorgan 6 days, 19 hours ago

A and E
upvoted 1 times

kmliuy73 1 week, 3 days ago
BE are correct
upvoted 1 times

Cizzla7049 1 week, 4 days ago
Selected Answer: AE
AE are correct
https://aws.amazon.com/blogs/storage/migrating-storage-with-aws-datasync/
upvoted 1 times

Question #103 *Topic 1*

A company has an AWS Glue extract, transform, and load (ETL) job that runs every day at the same time. The job processes XML data that is in an Amazon S3 bucket. New data is added to the S3 bucket every day. A solutions architect notices that AWS Glue is processing all the data during each run.
What should the solutions architect do to prevent AWS Glue from reprocessing old data?

A. Edit the job to use job bookmarks.

B. Edit the job to delete data after the data is processed.

C. Edit the job by setting the NumberOfWorkers field to 1.

D. Use a FindMatches machine learning (ML) transform.

Correct Answer: *A*

Community vote distribution
A (100%)

123jhl0 Highly Voted 1 month, 1 week ago
Selected Answer: A
This is the purpose of bookmarks: "AWS Glue tracks data that has already been processed during a previous run of an ETL job by persisting state information from the job run. This persisted state information is called a job bookmark. Job bookmarks help AWS Glue maintain state information and prevent the reprocessing of old data."
https://docs.aws.amazon.com/glue/latest/dg/monitor-continuations.html
upvoted 13 times

Wpcorgan Most Recent 6 days, 19 hours ago
A is correct
upvoted 2 times

LeGloupier 1 month, 1 week ago
Selected Answer: A
A
https://docs.aws.amazon.com/glue/latest/dg/monitor-continuations.html
upvoted 3 times

Question #104 *Topic 1*

A solutions architect must design a highly available infrastructure for a website. The website is powered by Windows web servers that run on Amazon EC2 instances. The solutions architect must implement a solution that can mitigate a large-scale DDoS attack that originates from thousands of IP addresses. Downtime is not acceptable for the website.
Which actions should the solutions architect take to protect the website from such an attack? (Choose two.)

A. Use AWS Shield Advanced to stop the DDoS attack.

B. Configure Amazon GuardDuty to automatically block the attackers.

C. Configure the website to use Amazon CloudFront for both static and dynamic content.

D. Use an AWS Lambda function to automatically add attacker IP addresses to VPC network ACLs.

E. Use EC2 Spot Instances in an Auto Scaling group with a target tracking scaling policy that is set to 80% CPU utilization.

Correct Answer: *AC*

Community vote distribution

AC (80%) AB (20%)

alvarez100 Highly Voted 1 month, 1 week ago

Selected Answer: AC

I think it is AC, reason is they require a solution that is highly available. AWS Shield can handle the DDoS attacks. To make the solution HA you can use cloud front. AC seems to be the best answer imo.
AB seem like redundant answers. How do those answers make the solution HA?

upvoted 12 times

attila9778 6 days, 5 hours ago

A - AWS Shield Advanced
C - (protecting this option) IMO: AWS Shield Advanced has to be attached. But it can not be attached directly to EC2 instances.
According to the docs: https://aws.amazon.com/shield/
It requires to be attached to services such as CloudFront, Route 53, Global Accelerator, ELB or (in the most direct way using) Elastic IP (attached to the EC2 instance)

upvoted 2 times

Newptone Most Recent 2 days, 10 hours ago

Selected Answer: AC

I thought it was AB. But after I read the docs, I vote for AC.

Amazon GuardDuty is a threat detection service, it can NOT take action directly, it needs to work with Lambda.

upvoted 1 times

Wpcorgan 6 days, 19 hours ago

A and C

upvoted 1 times

rjam 1 week, 5 days ago

Selected Answer: AC

AWS Shield can handle the DDoS attacks
Amazon CloudFront supports DDoS protection, integration with Shield, AWS Web Application Firewall

upvoted 2 times

tubtab 2 weeks, 5 days ago

Selected Answer: AC

correct

upvoted 1 times

Az900500 3 weeks, 1 day ago

I believe it's A & E ; the questions speaks to two things.
1. That can mitigate large DDOS attack - (Ans A)
2. A solutions architect must design a highly available infrastructure for a website; Downtime is not acceptable (Ans E)
So Ans is AE

I guess we focus only on the DDOS attack aspect of the question

upvoted 1 times

ocbn3wby 2 days, 3 hours ago

So, spot instances mean HA for you?

upvoted 1 times

Cizzla7049 1 week, 4 days ago

spot instances aren't always going to be highly available enough for certain situations. its AC

upvoted 1 times

mm_ 3 weeks, 5 days ago

Selected Answer: AB

Amazon GuardDuty has Threat response and remediation automation.

upvoted 1 times

dokaedu 4 weeks, 1 day ago

A : handle DDoS
E: Use EC2 Spot Instances in an Auto Scaling group with a target tracking scaling policy that is set to 80% CPU utilization.

upvoted 1 times

Vic_d_gr8 4 days, 1 hour ago

spot instance are not reliable, they are for worlds which can tolerate downtime. So the Answer should be A & C

upvoted 2 times

Vic_d_gr8 4 days, 1 hour ago

*workloads

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: AB

CF doesn't help with preventing downtime with dynamic content, it improves latency yes but doesn't really help with this case imo

question is asking for ways to PROTECT the server and prevent downtime,
and if you read this, guardduty makes sense.

https://aws.amazon.com/guardduty/
> Gain insight of compromised credentials, unusual data access in Amazon S3, API calls from known malicious IP addresses, and more.

upvoted 1 times

dave9994 1 month ago

The question is about "Protect", not remediation. So, A and C are the possible answers.

upvoted 1 times

whosawsome 1 month, 1 week ago

Selected Answer: AB

GuardDuty allows you to take automated action to remedy an attack

upvoted 1 times

stjokerli 1 month, 2 weeks ago

Selected Answer: AB

Only a,b are related to security

upvoted 1 times

123jhl0 1 month, 1 week ago

Yes, BUT GuardDuty will help only with detection, not blocking or remediating. https://aws.amazon.com/guardduty/
On the other hand you miss the opportunity to introduce a massive HA which helps you to keep the website live (avoiding downtime) with CloudFront, valid for static and dynamic content (https://aws.amazon.com/cloudfront/)
I go for A and C

upvoted 6 times

Question #105 *Topic 1*

A company is preparing to deploy a new serverless workload. A solutions architect must use the principle of least privilege to configure permissions that will be used to run an AWS Lambda function. An Amazon EventBridge (Amazon CloudWatch Events) rule will invoke the function. Which solution meets these requirements?

A. Add an execution role to the function with lambda:InvokeFunction as the action and * as the principal.

B. Add an execution role to the function with lambda:InvokeFunction as the action and Service: lambda.amazonaws.com as the principal.

C. Add a resource-based policy to the function with lambda:* as the action and Service: events.amazonaws.com as the principal.

D. Add a resource-based policy to the function with lambda:InvokeFunction as the action and Service: events.amazonaws.com as the principal.

Correct Answer: *D*

Community vote distribution

D (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: D

Best way to check it... The question is taken from the example shown here in the documentation:
https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-use-resource-based.html#eb-lambda-permissions

upvoted 10 times

ocbn3wby Most Recent 2 days, 3 hours ago

Selected Answer: D

My answer was D, as this is the most specific answer.
And then there's this guy's answer (123jhl0) which provides more details.

upvoted 1 times

Question #106 *Topic 1*

A company is preparing to store confidential data in Amazon S3. For compliance reasons, the data must be encrypted at rest. Encryption key usage must be logged for auditing purposes. Keys must be rotated every year.
Which solution meets these requirements and is the MOST operationally efficient?

A. Server-side encryption with customer-provided keys (SSE-C)

B. Server-side encryption with Amazon S3 managed keys (SSE-S3)

C. Server-side encryption with AWS KMS keys (SSE-KMS) with manual rotation

D. Server-side encryption with AWS KMS keys (SSE-KMS) with automatic rotation

Correct Answer: *D*

Community vote distribution

D (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: D

The MOST operationally efficient one is D.
Automating the key rotation is the most efficient.
Just to confirm, the A and B options don't allow automate the rotation as explained here:
https://aws.amazon.com/kms/faqs/#:~:text=You%20can%20choose%20to%20have%20AWS%20KMS%20automatically%20rotate%20KMS,KMS%20custom%20key%20store%20feature

upvoted 5 times

ocbn3wby 2 days, 3 hours ago

Thank you for the explanation.

upvoted 1 times

PS_R Most Recent 2 weeks, 5 days ago

Selected Answer: D

Agree Also, SSE-S3 cannot be audited.

upvoted 1 times

Question #107 *Topic 1*

A bicycle sharing company is developing a multi-tier architecture to track the location of its bicycles during peak operating hours. The company wants to use these data points in its existing analytics platform. A solutions architect must determine the most viable multi-tier option to support this architecture. The data points must be accessible from the REST API.
Which action meets these requirements for storing and retrieving location data?

A. Use Amazon Athena with Amazon S3.

B. Use Amazon API Gateway with AWS Lambda.

C. Use Amazon QuickSight with Amazon Redshift.

D. Use Amazon API Gateway with Amazon Kinesis Data Analytics.

Correct Answer: *D*

Community vote distribution

B (58%) D (42%)

ArielSchivo Highly Voted 1 month, 1 week ago

Selected Answer: B

API Gateway is needed to get the data so option A and C are out.
"The company wants to use these data points in its existing analytics platform" so there is no need to add Kynesis. Option D is also out.
This leaves us with option B as the correct one.

upvoted 16 times

Six_Fingered_Jose Highly Voted 1 month ago

Selected Answer: D

I dont understand why you will vote B?
how are you going to store data with just lambda?
> Which action meets these requirements for storing and retrieving location data

In this use case there will obviously be a ton of data and you want to get real-time location data of the bicycles, and to analyze all these info kinesis is the one that makes most sense here.

upvoted 10 times

rob74 3 weeks, 4 days ago

I vote D because company HAS its analitcs Platform, Why pay?. Kinesis is for analys not for storing. Can you explain? Thanks

upvoted 2 times

Onimole 3 weeks, 4 days ago

Weird Q as they already have their own data analysis platform
Hopefully i dont see this question in the exam lol

upvoted 4 times

Onimole 3 weeks, 4 days ago

it can store according to the doc
There is no way to lambda to store data which is part of the requirements

upvoted 1 times

rob74 3 weeks, 4 days ago

B Lambda and API

upvoted 1 times

UWSFish 1 month ago

I don't think you need to worry about storing data. The question states their is an existing platform.

upvoted 2 times

Six_Fingered_Jose 1 month ago

https://aws.amazon.com/blogs/aws/real-time-hotspot-detection-in-amazon-kinesis-analytics/

upvoted 2 times

5up3rm4n Most Recent 3 days, 20 hours ago

Answer is B. Kinesis Data Streams does not persist data. It also only ingests data from Kinesis Data Stream and Firehose, and outputs to Kinesis Data Stream, Firehose and Lambda.

upvoted 1 times

mabotega 6 days, 2 hours ago

Selected Answer: D

This question is about location of company´s bicycles during peak operating hours", so, the data is coming from their application and must be checked in real-time during peak operating hours, so, it requires Kinesis Data Analytcs and API Gateway supports the Kinesyis.
https://docs.aws.amazon.com/apigateway/latest/developerguide/integrating-api-with-aws-services-kinesis.html
upvoted 1 times

rjam 1 week, 4 days ago

Selected Answer: D

B is wrong as the question specifies about storing and retrieving location data. APIGateway and lambda cannot store data.

Amazon API Gateway with Kinesis Data Analytics:
Amazon API Gateway is a fully managed service that allows you to publish, maintain, monitor, and secure APIs at any scale. Amazon API Gateway offers two options to create RESTful APIs, HTTP APIs and REST APIs, as well as an option to create WebSocket APIs.

You can use Kinesis Data Analytics to transform and analyze streaming data in real-time with Apache Flink. Kinesis Data Analytics enables you to quickly build end-to-end stream processing applications for log analytics, clickstream analytics, Internet of Things (IoT), ad tech, gaming, etc. The four most common use cases are streaming extract-transform-load (ETL), continuous metric generation, responsive real-time analytics, and interactive querying of data streams. Kinesis Data Analytics for Apache Flink applications provides your application 50 GB of running application storage per Kinesis Processing Unit (KPU).
upvoted 2 times

BD12 2 weeks, 5 days ago

B
If REST API means, there API gateway should be there.
upvoted 1 times

Gil80 2 weeks, 6 days ago

Selected Answer: B

Why B?
1. QuickSight doesn't support REST API.
2. They already have an analytics platform, A (Athena) and D (Kinesis Data Analytics) are out of the race even though S3 & APT Gateway Support REST API. B & C are an option. C will not support REST API. So answer should be B
upvoted 2 times

Nigma 3 weeks ago

D is the correct. The question is asking for storing and retrieving
upvoted 1 times

PS_R 3 weeks ago

Selected Answer: D

KDA also stores data for 365 days, so i feel D
upvoted 1 times

Onimole 3 weeks, 4 days ago

D might be the right choice
When Should I Use Amazon Kinesis Data Analytics?
Amazon Kinesis Data Analytics enables you to quickly author SQL code that continuously reads, processes, and stores data in near real time. Using standard SQL queries on the streaming data, you can construct applications that transform and provide insights into your data. Following are some of example scenarios for using Kinesis Data Analytics:
upvoted 3 times

Hunkie 3 weeks, 5 days ago

Selected Answer: B

They already have an existing analytics platform
upvoted 1 times

Dsouzaf 4 weeks, 1 day ago

to track the location of its bicycles during peak operating hours. This lines tells it is streaming data and hence Amazon Kinesis Data Analytics
upvoted 1 times

dave9994 1 month ago

The question is about storing and retrieving, API GW satisfies retrieving part, and Kinesys handles storing part. I would go with D.
upvoted 1 times

Lilibell 1 month, 2 weeks ago

The answer is B
upvoted 1 times

Stiffler1 1 month, 2 weeks ago

why d?
upvoted 1 times

Question #108 *Topic 1*

A company has an automobile sales website that stores its listings in a database on Amazon RDS. When an automobile is sold, the listing needs to be removed from the website and the data must be sent to multiple target systems.
Which design should a solutions architect recommend?

A. Create an AWS Lambda function triggered when the database on Amazon RDS is updated to send the information to an Amazon Simple Queue Service (Amazon SQS) queue for the targets to consume.

B. Create an AWS Lambda function triggered when the database on Amazon RDS is updated to send the information to an Amazon Simple Queue Service (Amazon SQS) FIFO queue for the targets to consume.

C. Subscribe to an RDS event notification and send an Amazon Simple Queue Service (Amazon SQS) queue fanned out to multiple Amazon Simple Notification Service (Amazon SNS) topics. Use AWS Lambda functions to update the targets.

D. Subscribe to an RDS event notification and send an Amazon Simple Notification Service (Amazon SNS) topic fanned out to multiple Amazon Simple Queue Service (Amazon SQS) queues. Use AWS Lambda functions to update the targets.

Correct Answer: *C*

Community vote distribution

D (75%) A (25%)

ocbn3wby 1 day, 5 hours ago

Selected Answer: A

There is RDS Fanout to SNS, but not specifically for DB level events (write, reads, etc).
It can fan out events at instance level (turn on, restart, update), cluster level (added to cluster, removed from cluster, etc). But not at DB level.

More detailed RDS event list here:
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_Events.Messages.html

Correct answer is A.

upvoted 1 times

anhdao1211 1 day, 6 hours ago

Selected Answer: A

A
RDS event notification ?

upvoted 1 times

Vic_d_gr8 4 days ago

Selected Answer: D

Amazon RDS uses the SNS to provide notification when an Amazon event occurs
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_Events.html

upvoted 2 times

Wpcorgan 6 days, 19 hours ago

D is correct

upvoted 1 times

romko 1 week ago

Selected Answer: A

Interesting point that Amazon RDS event notification doesn't support any notification when data inside DB is updated.
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_Events.overview.html
So subscription to RDS events doesn't give any value for Fanout = SNS => SQS

B is out because FIFO is not required here.

A is left as correct answer

upvoted 3 times

ocbn3wby 1 day, 5 hours ago

Romko, you are right pal. Nice research.

There is RDS Fanout to SNS, but not specifically for DB level events (write, reads, etc).
It can fan out events at instance level (turn on, restart, update), cluster level (added to cluster, removed from cluster, etc). But not at DB level.

More detailed event list here:
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_Events.Messages.html

Correct answer is A.
upvoted 1 times

Vic_d_gr8 4 days ago
Amazon RDS uses the SNS to provide notification when an Amazon event occurs
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_Events.html
upvoted 1 times

PS_R 2 weeks, 5 days ago
Selected Answer: D
Cannot be C as Fanout = SNS + SQS.
upvoted 2 times

Merchant 3 weeks, 5 days ago
Selected Answer: D
DDDDDD!
upvoted 2 times

Hunkie 3 weeks, 5 days ago
Selected Answer: D
The RDS event notifications sends the notification using SNS not SQS.

Similar to old question
upvoted 3 times

PS_R 2 weeks, 5 days ago
Agree, Also Fanout means SNS + SQS, which rules out option C, So the answer is D
upvoted 1 times

dokaedu 4 weeks, 1 day ago
answer should be D
https://betterprogramming.pub/how-to-fan-out-to-different-sqs-queues-using-sns-message-filtering-84cd23ed9d07
upvoted 1 times

Six_Fingered_Jose 1 month ago
Selected Answer: D
answer should be D
> the data must be sent to multiple target systems.

Fanning out -> SNS
upvoted 2 times

whosawsome 1 month, 1 week ago
Selected Answer: D
D if in doubt
upvoted 4 times

Moonus 1 month, 1 week ago
Answer is D
upvoted 3 times

Lilibell 1 month, 2 weeks ago
The answer is D
upvoted 4 times

Question #109 *Topic 1*

A company needs to store data in Amazon S3 and must prevent the data from being changed. The company wants new objects that are uploaded to Amazon S3 to remain unchangeable for a nonspecific amount of time until the company decides to modify the objects. Only specific users in the company's AWS account can have the ability 10 delete the objects.
What should a solutions architect do to meet these requirements?

A. Create an S3 Glacier vault. Apply a write-once, read-many (WORM) vault lock policy to the objects.

B. Create an S3 bucket with S3 Object Lock enabled. Enable versioning. Set a retention period of 100 years. Use governance mode as the S3 bucket's default retention mode for new objects.

C. Create an S3 bucket. Use AWS CloudTrail to track any S3 API events that modify the objects. Upon notification, restore the modified objects from any backup versions that the company has.

D. Create an S3 bucket with S3 Object Lock enabled. Enable versioning. Add a legal hold to the objects. Add the s3:PutObjectLegalHold permission to the IAM policies of users who need to delete the objects.

Correct Answer: *D*

Community vote distribution

D (83%) B (17%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: D

A - No as "specific users can delete"
B - No as "nonspecific amount of time"
C - No as "prevent the data from being change"
D - The answer: "The Object Lock legal hold operation enables you to place a legal hold on an object version. Like setting a retention period, a legal hold prevents an object version from being overwritten or deleted. However, a legal hold doesn't have an associated retention period and remains in effect until removed." https://docs.aws.amazon.com/AmazonS3/latest/userguide/batch-ops-legal-hold.html

upvoted 10 times

Chunsli Highly Voted 1 month, 1 week ago

typo -- 10 delete the objects => TO delete the objects

upvoted 5 times

nadjar007 Most Recent 1 week, 2 days ago

D
https://docs.aws.amazon.com/AmazonS3/latest/userguide/batch-ops-legal-hold.html

upvoted 1 times

Cizzla7049 1 week, 3 days ago

Selected Answer: B

https://aws.amazon.com/pt/blogs/storage/protecting-data-with-amazon-s3-object-lock/

upvoted 1 times

attila9778 6 days, 3 hours ago

Answer: D
Question says: "indefinite amount of time" -> "legal hold"
Your provided link describes this feature as well:
"A legal hold provides the same protection as a retention period, but it has no expiration date. Instead, a legal hold remains in place until you explicitly remove it."

upvoted 1 times

mabotega 2 weeks, 3 days ago

Selected Answer: B

Retention mode in governance mode is the correct way to some users have the ability to delete object in S3 ->
https://aws.amazon.com/pt/blogs/storage/protecting-data-with-amazon-s3-object-lock/

upvoted 1 times

attila9778 6 days, 3 hours ago

Question says: "indefinite amount of time" -> "legal hold"
Your provided link describes this feature as well:
"A legal hold provides the same protection as a retention period, but it has no expiration date. Instead, a legal hold remains in place until you explicitly remove it."

upvoted 1 times

attila9778 6 days, 3 hours ago

=> Answer: D

upvoted 1 times

Question #110 *Topic 1*

A social media company allows users to upload images to its website. The website runs on Amazon EC2 instances. During upload requests, the website resizes the images to a standard size and stores the resized images in Amazon S3. Users are experiencing slow upload requests to the website.
The company needs to reduce coupling within the application and improve website performance. A solutions architect must design the most operationally efficient process for image uploads.
Which combination of actions should the solutions architect take to meet these requirements? (Choose two.)

A. Configure the application to upload images to S3 Glacier.

B. Configure the web server to upload the original images to Amazon S3.

C. Configure the application to upload images directly from each user's browser to Amazon S3 through the use of a presigned URL

D. Configure S3 Event Notifications to invoke an AWS Lambda function when an image is uploaded. Use the function to resize the image.

E. Create an Amazon EventBridge (Amazon CloudWatch Events) rule that invokes an AWS Lambda function on a schedule to resize uploaded images.

Correct Answer: *BD*

Community vote distribution

BD (47%) CD (40%) 13%

leonnnn 1 day, 1 hour ago

Selected Answer: CD

It's not a common solution to upload from user browser directly to S3. But it says "reduce coupling". I think it means transferring between web server and S3 is unnecessary.

upvoted 1 times

ocbn3wby 1 day, 4 hours ago

Selected Answer: CD

B: UI -> Web Server -> S3 (this is not efficient and secure, you are transferring files twice).

C: based on the above, this is one transfer only (so more efficient!)

E: EventBridge + Lambda (this requires extra integration between S3 and EventBridge).

D: simply S3 event notifications to Lambda is more operationally efficient than E.

Answer: C+D

Supporting documentation:
https://docs.aws.amazon.com/lambda/latest/dg/with-s3.html
https://fullstackdojo.medium.com/s3-upload-with-presigned-url-react-and-nodejs-b77f348d54cc#:~:text=S3%20upload%20with%20presigned%20url%20%E2%80%94%20React%20and,upload%20the%20file%20with%20the%20returned%20presigned%20URL.

upvoted 1 times

Vic_d_gr8 4 days ago

Selected Answer: BD

Remember to have operationally efficiency in mind. It usually goes with configure rather than create. For Presigned URL's you have to assign it individually to the user. Example of Presiged URL's use case: For access to premium users on a premium content on a website .

upvoted 1 times

Wilson_S 6 days, 8 hours ago

Selected Answer: BD

"Opertationally efficient"

upvoted 1 times

Wpcorgan 6 days, 19 hours ago

B and D

upvoted 1 times

rjam 1 week, 5 days ago

Selected Answer: BE

Reason:
option B and E.
Configure the web server to upload the original images to Amazon S3.

Event Bridge has the capability to replay , reliable delivery, and archive so in case if lambda fails to process it can be retried in case of connection failures as the messages will be available in Event bridge.
A:: NO. cannot upload directly to Glacier. you need to do it first to S3
C: NO. Presigned URLS expire after some time
D: Event bridge is better . If you configure S3 Event Notifications directly to lambda, the event will be lost during connection failures are no way to reprocess in case of failure by lambda to process.

upvoted 2 times

Deplake 2 weeks, 6 days ago

Selected Answer: BD

C is not looking correct for me, you should handle somehow which image names are already uploaded in order to prevent conflicts and security could be the case as well

upvoted 2 times

baiy 3 weeks, 1 day ago

Selected Answer: BD

I go for BD, presigned URL does not improve performance, I don't see the meaning of using.

upvoted 3 times

gokalpkocer3 3 weeks, 1 day ago

presigned URL does not improve performance which is what question asking here.

upvoted 1 times

Stephhhh 4 weeks ago

Upload image directly from user browser? Will it incur extra support effort to the IT team? Does it means extra operational effort?

upvoted 2 times

Stephhhh 4 weeks ago

I will go for B and D, given the operation effort is a consideration

upvoted 3 times

Six_Fingered_Jose 1 month ago

Selected Answer: CD

This is a tough question but i will go with C and D
with B: client browser -> web server -> s3
with C: client browser -> s3 directly,

meaning C should be a better solution than B (setting up pre-signed URLs do feel like a chore though)

am not 100% sure on this.

upvoted 4 times

USalo 1 month ago

What about security? It can be potential issue if not properly covered on front-end, but handling that adding additional overhead.

upvoted 1 times

Question #111 *Topic 1*

A company recently migrated a message processing system to AWS. The system receives messages into an ActiveMQ queue running on an Amazon EC2 instance. Messages are processed by a consumer application running on Amazon EC2. The consumer application processes the messages and writes results to a MySQL database running on Amazon EC2. The company wants this application to be highly available with low operational complexity.
Which architecture offers the HIGHEST availability?

A. Add a second ActiveMQ server to another Availability Zone. Add an additional consumer EC2 instance in another Availability Zone. Replicate the MySQL database to another Availability Zone.

B. Use Amazon MQ with active/standby brokers configured across two Availability Zones. Add an additional consumer EC2 instance in another Availability Zone. Replicate the MySQL database to another Availability Zone.

C. Use Amazon MQ with active/standby brokers configured across two Availability Zones. Add an additional consumer EC2 instance in another Availability Zone. Use Amazon RDS for MySQL with Multi-AZ enabled.

D. Use Amazon MQ with active/standby brokers configured across two Availability Zones. Add an Auto Scaling group for the consumer EC2 instances across two Availability Zones. Use Amazon RDS for MySQL with Multi-AZ enabled.

Correct Answer: *D*

Community vote distribution

D (86%) 14%

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: D

Answer is D as the "HIGHEST available" and less "operational complex"
The "Amazon RDS for MySQL with Multi-AZ enabled" option excludes A and B
The "Auto Scaling group" is more available and reduces operational complexity in case of incidents (as remediation it is automated) than just adding one more instance. This excludes C.

C and D to choose from based on
D over C since is configured

upvoted 5 times

Wpcorgan Most Recent 6 days, 17 hours ago

D is correct

upvoted 1 times

UWSFish 1 month ago

Selected Answer: A

I don't know about D. Active/Standby adds to fault tolerance but does nothing for HA.

upvoted 1 times

nullvoiddeath 2 weeks, 4 days ago

Amazon RDS > MySQL, hence A and B are eliminated

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: D

agree with D

upvoted 1 times

Question #112 *Topic 1*

A company hosts a containerized web application on a fleet of on-premises servers that process incoming requests. The number of requests is growing quickly. The on-premises servers cannot handle the increased number of requests. The company wants to move the application to AWS with minimum code changes and minimum development effort.
Which solution will meet these requirements with the LEAST operational overhead?

A. Use AWS Fargate on Amazon Elastic Container Service (Amazon ECS) to run the containerized web application with Service Auto Scaling. Use an Application Load Balancer to distribute the incoming requests.

B. Use two Amazon EC2 instances to host the containerized web application. Use an Application Load Balancer to distribute the incoming requests.

C. Use AWS Lambda with a new code that uses one of the supported languages. Create multiple Lambda functions to support the load. Use Amazon API Gateway as an entry point to the Lambda functions.

D. Use a high performance computing (HPC) solution such as AWS ParallelCluster to establish an HPC cluster that can process the incoming requests at the appropriate scale.

Correct Answer: *A*

Community vote distribution

A (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: A

Less operational overhead means A: Fargate (no EC2), move the containers on ECS, autoscaling for growth and ALB to balance consumption.
B - requires configure EC2
C - requires add code (developpers)
D - seems like the most complex approach, like re-architecting the app to take advantage of an HPC platform.

upvoted 5 times

Wpcorgan Most Recent 6 days, 17 hours ago

A is correct

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: A

Agreed with A,
lambda will work too but requires more operational overhead (more chores)

with A, you are just moving from an on-prem container to AWS container

upvoted 2 times

Question #113 *Topic 1*

A company uses 50 TB of data for reporting. The company wants to move this data from on premises to AWS. A custom application in the company's data center runs a weekly data transformation job. The company plans to pause the application until the data transfer is complete and needs to begin the transfer process as soon as possible.
The data center does not have any available network bandwidth for additional workloads. A solutions architect must transfer the data and must configure the transformation job to continue to run in the AWS Cloud.
Which solution will meet these requirements with the LEAST operational overhead?

A. Use AWS DataSync to move the data. Create a custom transformation job by using AWS Glue.

B. Order an AWS Snowcone device to move the data. Deploy the transformation application to the device.

C. Order an AWS Snowball Edge Storage Optimized device. Copy the data to the device. Create a custom transformation job by using AWS Glue.

D. Order an AWS Snowball Edge Storage Optimized device that includes Amazon EC2 compute. Copy the data to the device. Create a new EC2 instance on AWS to run the transformation application.

Correct Answer: *C*

Community vote distribution

C (82%) D (18%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: C

A. Use AWS DataSync to move the data. Create a custom transformation job by using AWS Glue. - No BW available for DataSync, so "asap" will be weeks/months (?)
B. Order an AWS Snowcone device to move the data. Deploy the transformation application to the device. - Snowcone will just store 14TB (SSD configuration).
C. Order an AWS Snowball Edge Storage Optimized device. Copy the data to the device. Create a custom transformation job by using AWS Glue. - SnowBall can store 80TB (ok), takes around 1 week to move the device (faster than A), and AWS Glue allows to do ETL jobs. This is the answer.
D. Order an AWS Snowball Edge Storage Optimized device that includes Amazon EC2 compute. Copy the data to the device. Create a new EC2 instance on AWS to run the transformation application. - Same as C, but the ETL job requires the deployment/configuration/maintenance of an EC2 instance, while Glue is serverless. This means D has more operational overhead than C.

upvoted 7 times

ocbn3wby Most Recent 1 day, 4 hours ago

Selected Answer: D

I would stick to D answer.

In real life - this is what it would have happened. Maybe after the migration, the existing ETL application would be refactored to AWS services. But this takes development time vs "lift and shift" approach.

Edge Storage optimized offers EC2 compute functionality (with AMI directly integrated) https://docs.aws.amazon.com/snowball/latest/developer-guide/device-differences.html

upvoted 1 times

ocbn3wby 1 day, 4 hours ago

Also, the question clearly states they want to "pause the application" and not transform it to something more efficient.

upvoted 1 times

rewdboy 2 days, 16 hours ago

Selected Answer: D

D is correct
I'm a bit split on this, because C is also a good answer, especially if you consider that the operational overhead. On the other hand, it seems more operationally efficient to just run the custom ETL job that already exists on the on-prem server on an EC2 instance instead of using AWS Glue. The part of the question that says they already have a "custom application" tells me that it should stay 1:1 to make the migration easier.

upvoted 1 times

ocbn3wby 1 day, 4 hours ago

In real life - this is what it would have happened. Maybe after the migration, the existing ETL application would be refactored to AWS services (Glue). But this takes development time vs "lift and shift" approach.

I would stick to answer D.

upvoted 1 times

JayanKuruwita 2 days, 4 hours ago

As I know Snowball Edge dosn't come with EC2. So the answer should be C

upvoted 1 times

ocbn3wby 1 day, 4 hours ago

You can search this online.
https://docs.aws.amazon.com/snowball/latest/developer-guide/device-differences.html

upvoted 1 times

Wpcorgan 6 days, 17 hours ago

C is correct

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: C

agree with C
Glue is least effort

upvoted 2 times

Question #114 *Topic 1*

A company has created an image analysis application in which users can upload photos and add photo frames to their images. The users upload images and metadata to indicate which photo frames they want to add to their images. The application uses a single Amazon EC2 instance and Amazon DynamoDB to store the metadata.
The application is becoming more popular, and the number of users is increasing. The company expects the number of concurrent users to vary significantly depending on the time of day and day of week. The company must ensure that the application can scale to meet the needs of the growing user base.
Which solution meats these requirements?

A. Use AWS Lambda to process the photos. Store the photos and metadata in DynamoDB.

B. Use Amazon Kinesis Data Firehose to process the photos and to store the photos and metadata.

C. Use AWS Lambda to process the photos. Store the photos in Amazon S3. Retain DynamoDB to store the metadata.

D. Increase the number of EC2 instances to three. Use Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volumes to store the photos and metadata.

Correct Answer: *A*

Community vote distribution

C (100%)

MXB05 Highly Voted 1 month, 2 weeks ago

Selected Answer: C

Do not store images in databases ;)... correct answer should be C

upvoted 10 times

rewdboy Most Recent 2 days, 16 hours ago

Selected Answer: C

C for sure
I was originally leaning toward A because it seemed like a simpler setup to keep the images and metadata in the same service, but DynamoDB has a record limit of 64KB, so S3 would be better for image storage and then DynamoDB for metadata

upvoted 1 times

Wpcorgan 6 days, 17 hours ago

C is correct

upvoted 1 times

Pamban 2 weeks ago

Selected Answer: C

photo needs to be on S3

upvoted 1 times

mabotega 2 weeks, 3 days ago

Selected Answer: C

photos should be stored on S3

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: C

agree with C,
Storing image in DB wont be very scalable compared to S3
metadata does not take up much space and is more efficiently stored in DB

upvoted 2 times

tubtab 1 month ago

Selected Answer: C

ccccccccc

upvoted 1 times

Question #115 *Topic 1*

A medical records company is hosting an application on Amazon EC2 instances. The application processes customer data files that are stored on Amazon S3. The EC2 instances are hosted in public subnets. The EC2 instances access Amazon S3 over the internet, but they do not require any other network access.
A new requirement mandates that the network traffic for file transfers take a private route and not be sent over the internet.
Which change to the network architecture should a solutions architect recommend to meet this requirement?

A. Create a NAT gateway. Configure the route table for the public subnets to send traffic to Amazon S3 through the NAT gateway.

B. Configure the security group for the EC2 instances to restrict outbound traffic so that only traffic to the S3 prefix list is permitted.

C. Move the EC2 instances to private subnets. Create a VPC endpoint for Amazon S3, and link the endpoint to the route table for the private subnets.

D. Remove the internet gateway from the VPC. Set up an AWS Direct Connect connection, and route traffic to Amazon S3 over the Direct Connect connection.

Correct Answer: *C*

Community vote distribution

C (100%)

ocbn3wby 1 day, 3 hours ago

C is correct.
There is no requirement for public access from internet.

Application must be moved in Private subnet. This is a prerequisite in using VPC endpoints with S3
https://aws.amazon.com/blogs/storage/managing-amazon-s3-access-with-vpc-endpoints-and-s3-access-points/

upvoted 1 times

Wpcorgan 6 days, 17 hours ago

C is correct

upvoted 1 times

Jtic 2 weeks ago

Selected Answer: C

Use VPC endpoint

upvoted 1 times

Jtic 2 weeks ago

Selected Answer: C

User VPC endpoint and make the EC2 private

upvoted 1 times

Jtic 2 weeks ago

Use VPC endpoint

upvoted 1 times

backbencher2022 3 weeks, 1 day ago

Selected Answer: C

VPC endpoint is the best choice to route S3 traffic without traversing internet. Option A alone can't be used as NAT Gateway requires an Internet gateway for outbound internet traffic. Option B would still require traversing through internet and option D is also not a suitable solution

upvoted 1 times

Question #116 *Topic 1*

A company uses a popular content management system (CMS) for its corporate website. However, the required patching and maintenance are burdensome. The company is redesigning its website and wants anew solution. The website will be updated four times a year and does not need to have any dynamic content available. The solution must provide high scalability and enhanced security.
Which combination of changes will meet these requirements with the LEAST operational overhead? (Choose two.)

A. Configure Amazon CloudFront in front of the website to use HTTPS functionality.

B. Deploy an AWS WAF web ACL in front of the website to provide HTTPS functionality.

C. Create and deploy an AWS Lambda function to manage and serve the website content.

D. Create the new website and an Amazon S3 bucket. Deploy the website on the S3 bucket with static website hosting enabled.

E. Create the new website. Deploy the website by using an Auto Scaling group of Amazon EC2 instances behind an Application Load Balancer.

Correct Answer: *AD*

Community vote distribution

AD (70%) BD (20%) 10%

palermo777 Highly Voted 1 month, 1 week ago

A -> We can configure CloudFront to require HTTPS from clients (enhanced security)
https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/using-https-viewers-to-cloudfront.html
D -> storing static website on S3 provides scalability and less operational overhead, then configuration of Application LB and EC2 instances (hence E is out)

B is out since AWS WAF Web ACL does not to provide HTTPS functionality, but to protect HTTPS only.

upvoted 10 times

Wpcorgan Most Recent 6 days, 17 hours ago

A and D

upvoted 1 times

PS_R 2 weeks, 5 days ago

Selected Answer: AD

Cloudfront can do the WAF part so i chose A and D

upvoted 2 times

Bevemo 2 weeks, 5 days ago

Selected Answer: AD

Initially I thought B) WAF for HTTP to HTTPS redirect, but then I found CloudFront can do it so A) adds performance/scale and security.
https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/using-https.html

upvoted 2 times

ManoAni 1 month ago

Selected Answer: BD

For enhanced security B, and they mentioned patching is burdensome so if its E, then they must patch the EC2 instances. So hosting in S3 is ideal as it is static content.

upvoted 2 times

Six_Fingered_Jose 1 month ago

Selected Answer: AD

agree with A and D

static website -> obviously S3, and S3 is super scalable
CDN -> CloudFront obviously as well, and with HTTPS security is enhanced.

B does not make sense because you are not replacing the CDN with anything,
E works too but takes too much effort and compared to S3, S3 still wins in term of scalability. plus why use EC2 when you are only hosting static website

upvoted 3 times

rob74 1 month, 1 week ago

Selected Answer: BE

. The solution must provide high scalability and enhanced security
AWS WAF--> For enhanced security
high scalability -->behind an Application Load Balancer.

upvoted 1 times

ocbn3wby 1 day, 3 hours ago

Please provide informed answers. You are truly correct, but in this case, there is no specific need to host the website/cms on EC2 + ALB.

It only requires static website - which can be achieved with scalable S3.

upvoted 1 times

Question #117 *Topic 1*

A company stores its application logs in an Amazon CloudWatch Logs log group. A new policy requires the company to store all application logs in Amazon OpenSearch Service (Amazon Elasticsearch Service) in near-real time.
Which solution will meet this requirement with the LEAST operational overhead?

A. Configure a CloudWatch Logs subscription to stream the logs to Amazon OpenSearch Service (Amazon Elasticsearch Service).

B. Create an AWS Lambda function. Use the log group to invoke the function to write the logs to Amazon OpenSearch Service (Amazon Elasticsearch Service).

C. Create an Amazon Kinesis Data Firehose delivery stream. Configure the log group as the delivery streams sources. Configure Amazon OpenSearch Service (Amazon Elasticsearch Service) as the delivery stream's destination.

D. Install and configure Amazon Kinesis Agent on each application server to deliver the logs to Amazon Kinesis Data Streams. Configure Kinesis Data Streams to deliver the logs to Amazon OpenSearch Service (Amazon Elasticsearch Service).

Correct Answer: *C*

Community vote distribution

A (79%) 11% 11%

Six_Fingered_Jose Highly Voted 1 month ago

Selected Answer: A

answer is A
https://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/CWL_OpenSearch_Stream.html

> You can configure a CloudWatch Logs log group to stream data it receives to your Amazon OpenSearch Service cluster in NEAR REAL-TIME through a CloudWatch Logs subscription

least overhead compared to kinesis

upvoted 11 times

UWSFish 1 month ago

Great link. Convinced me

upvoted 4 times

study_aws1 Most Recent 2 weeks, 3 days ago

https://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/Subscriptions.html

You'll need to have destination arn (not mentioned under option A) - either Lambda or Kinesis Firehose.

The Amazon Resource Name (ARN) of the Kinesis stream, Kinesis Data Firehose stream, or Lambda function you want to use as the destination of the subscription feed.

Option B) does not mention the Subscription Filter. Looks more towards Option C)

upvoted 3 times

SimonPark 4 weeks, 1 day ago

Selected Answer: A

You can configure a CloudWatch Logs log group to stream data it receives to your Amazon OpenSearch Service cluster in near real-time through a CloudWatch Logs subscription.

upvoted 2 times

ManoAni 1 month ago

Selected Answer: C

They mentioned near real time

upvoted 2 times

capepenguin 1 month ago

Please tell me why not C?
https://docs.aws.amazon.com/opensearch-service/latest/developerguide/integrations.html#integrations-fh

upvoted 3 times

ericcloud20 1 month, 1 week ago

Answer is A

CloudWatch has a native feature to stream logs to OpenSearch, when you enable this setting it creates a Lambda Function automatically with pre-populated code which streams the logs to OpenSearch Cluster. The question here needs a solution with LEAST operational overhead, therefore the

answer should be A
REF: https://docs.aws.amazon.com/AmazonCloudWatch/latest/logs/CWL_OpenSearch_Stream.html
upvoted 4 times

palermo777 1 month, 1 week ago

B looks good:
https://d1.awsstatic.com/whitepapers/whitepaper-use-amazon-elasticsearch-to-log-and-monitor-almost-everything.pdf
Chapter: Pushing Amazon CloudWatch Logs into Amazon ES: "... integration makes it easy to send data to Elasticsearch if source data exists in CloudWatch Logs"

Approach with Amazon Kinesis Data Firehose requires installation of Amazon Kinesis agent on the EC2 instances, hence it's not considered as LEAST operational complex
upvoted 2 times

KJa 1 month, 1 week ago

Selected Answer: A
Should be A
upvoted 2 times

LeGloupier 1 month, 1 week ago

Selected Answer: B
B seems to be the right answer
https://computingforgeeks.com/stream-logs-in-aws-from-cloudwatch-to-elasticsearch/
upvoted 2 times

Question #118 *Topic 1*

A company is building a web-based application running on Amazon EC2 instances in multiple Availability Zones. The web application will provide access to a repository of text documents totaling about 900 TB in size. The company anticipates that the web application will experience periods of high demand. A solutions architect must ensure that the storage component for the text documents can scale to meet the demand of the application at all times. The company is concerned about the overall cost of the solution.
Which storage solution meets these requirements MOST cost-effectively?

A. Amazon Elastic Block Store (Amazon EBS)

B. Amazon Elastic File System (Amazon EFS)

C. Amazon OpenSearch Service (Amazon Elasticsearch Service)

D. Amazon S3

Correct Answer: *D*

Community vote distribution

D (100%)

Wpcorgan 6 days, 17 hours ago

D is correct
upvoted 1 times

PS_R 2 weeks, 5 days ago

Selected Answer: D
Only EFS and S3, Since EFS is make it much costly, S3 is the viable option
upvoted 2 times

UWSFish 1 month ago

Selected Answer: D
I originally thought C but the question is specific about wanting the storage to scale not the search capacity.
upvoted 2 times

Question #119 *Topic 1*

A global company is using Amazon API Gateway to design REST APIs for its loyalty club users in the us-east-1 Region and the ap-southeast-2 Region. A solutions architect must design a solution to protect these API Gateway managed REST APIs across multiple accounts from SQL injection and cross-site scripting attacks.
Which solution will meet these requirements with the LEAST amount of administrative effort?

A. Set up AWS WAF in both Regions. Associate Regional web ACLs with an API stage.

B. Set up AWS Firewall Manager in both Regions. Centrally configure AWS WAF rules.

C. Set up AWS Shield in bath Regions. Associate Regional web ACLs with an API stage.

D. Set up AWS Shield in one of the Regions. Associate Regional web ACLs with an API stage.

Correct Answer: *A*

Community vote distribution

B (85%) A (15%)

Gil80 Highly Voted 2 weeks, 5 days ago

Selected Answer: B

If you want to use AWS WAF across accounts, accelerate WAF configuration, automate the protection of new resources, use Firewall Manager with AWS WAF

upvoted 6 times

Wpcorgan Most Recent 6 days, 17 hours ago

B is correct

upvoted 1 times

PS_R 2 weeks, 5 days ago

Selected Answer: B

The Answer is AWS Firewall Manager, as it says multiple accounts - My bad.

upvoted 3 times

ArielSchivo 2 weeks, 5 days ago

Selected Answer: B

Option B.

https://aws.amazon.com/es/blogs/security/centrally-manage-aws-waf-api-v2-and-aws-managed-rules-at-scale-with-firewall-manager/

upvoted 2 times

PS_R 2 weeks, 5 days ago

Selected Answer: A

SQL Injections _ think WAF

upvoted 2 times

Nigma 2 weeks, 6 days ago

B

Using AWS WAF has several benefits. Additional protection against web attacks using criteria that you specify. You can define criteria using characteristics of web requests such as the following:
Presence of SQL code that is likely to be malicious (known as SQL injection).
Presence of a script that is likely to be malicious (known as cross-site scripting).

AWS Firewall Manager simplifies your administration and maintenance tasks across multiple accounts and resources for a variety of protections.

https://docs.aws.amazon.com/waf/latest/developerguide/what-is-aws-waf.html

upvoted 4 times

study_aws1 2 weeks, 6 days ago

The keyword is "protect API Gateway managed REST APIs across multiple accounts". Firewall Manager is used when it comes to managing multiple accounts. Option B)

upvoted 1 times

dokaedu 4 weeks, 1 day ago

Correct Answer: A
AWS Firewall Manager manage multiple AWS WAFs in many regions, Each Firewall Manager policy can only include resources available in that specified AWS Region, You can create a new policy for each region where you operate. The AWS WAF policies are reginal.

upvoted 1 times

envest 1 month ago

IMO: WAF & FW Mgr are regional.
https://aws.amazon.com/firewall-manager/faqs/#:~:text=No%2C%20AWS%20Firewall%20Manager%20security%20policies%20are%20region%20specific.

upvoted 1 times

Question #120 *Topic 1*

A company has implemented a self-managed DNS solution on three Amazon EC2 instances behind a Network Load Balancer (NLB) in the us-west-2 Region. Most of the company's users are located in the United States and Europe. The company wants to improve the performance and availability of the solution. The company launches and configures three EC2 instances in the eu-west-1 Region and adds the EC2 instances as targets for a new NLB.
Which solution can the company use to route traffic to all the EC2 instances?

A. Create an Amazon Route 53 geolocation routing policy to route requests to one of the two NLBs. Create an Amazon CloudFront distribution. Use the Route 53 record as the distribution's origin.

B. Create a standard accelerator in AWS Global Accelerator. Create endpoint groups in us-west-2 and eu-west-1. Add the two NLBs as endpoints for the endpoint groups.

C. Attach Elastic IP addresses to the six EC2 instances. Create an Amazon Route 53 geolocation routing policy to route requests to one of the six EC2 instances. Create an Amazon CloudFront distribution. Use the Route 53 record as the distribution's origin.

D. Replace the two NLBs with two Application Load Balancers (ALBs). Create an Amazon Route 53 latency routing policy to route requests to one of the two ALBs. Create an Amazon CloudFront distribution. Use the Route 53 record as the distribution's origin.

Correct Answer: *A*

Community vote distribution

B (69%) A (25%) 6%

LeGloupier Highly Voted 1 month, 1 week ago

Selected Answer: B

for me it is B

upvoted 8 times

Newptone Most Recent 2 days, 6 hours ago

Selected Answer: B

For option A: Create an Amazon CloudFront distribution.
DNS is tcp/udp only, you can not use CF to cache the traffic.

upvoted 1 times

jhxetc 6 days, 17 hours ago

Selected Answer: B

B makes the most sense since is not for a web app, but rather a DNS solution. We probably don't need to involve Route 53 (itself a DNS solution).

upvoted 1 times

grzeev 1 week ago

Selected Answer: D

for me it is D

upvoted 1 times

Jtic 2 weeks ago

Selected Answer: B

WS Global Accelerator is a networking service that helps you improve the availability and performance of the applications that you offer to your global users. AWS Global Accelerator is easy to set up, configure, and manage. It provides static IP addresses that provide a fixed entry point to your applications and eliminate the complexity of managing specific IP addresses for different AWS Regions and Availability Zones. AWS Global Accelerator always routes user traffic to the optimal endpoint based on performance, reacting instantly to changes in application health, your user's location, and policies that you configure

https://aws.amazon.com/global-accelerator/faqs/

upvoted 1 times

KADSM 2 weeks, 4 days ago

It should be A. As in B - The solution is about using Standard accelerator. It may need custom routing accelerator to route the traffic to specific EC2 instance.

Custom routing accelerator
Custom routing accelerators are a new type of accelerator in AWS Global Accelerator. This new accelerator lets you use your own application logic to route user traffic to a specific Amazon EC2 instance destination in a single or multiple AWS Regions. . A custom routing accelerator is an alternative to the standard accelerator, which automatically routes traffic to a healthy endpoint that is nearest to your users. Because standard accelerators are designed to load balance traffic, you can't use them to route users to a specific EC2 instance destination behind your accelerator.

upvoted 2 times

SimonPark 4 weeks, 1 day ago

Selected Answer: A

imho, it's A

upvoted 1 times

dokaedu 4 weeks, 1 day ago

B is the correct one for seld manage DNS
If need to use Route53, ALB (layar 7) needs to be used as end points for 2 reginal x 3 EC2s, if it the case answer would be the option 4

upvoted 4 times

Six_Fingered_Jose 1 month ago

Selected Answer: A

I think it is A,
if you carefully read the document below, with GA routing is managed manually in multi-region applications
> Your traffic routing is managed manually
https://www.examtopics.com/exams/amazon/aws-certified-solutions-architect-associate-saa-c03/view/30/

thus, Route53 geolocation seems to be the only solution to this problem (routing to different resources on different regions based on user's geolocation)

upvoted 3 times

Evangelia 1 month, 1 week ago

B for self managed DNS solution

upvoted 4 times

Question #121 *Topic 1*

A company is running an online transaction processing (OLTP) workload on AWS. This workload uses an unencrypted Amazon RDS DB instance in a Multi-AZ deployment. Daily database snapshots are taken from this instance.
What should a solutions architect do to ensure the database and snapshots are always encrypted moving forward?

A. Encrypt a copy of the latest DB snapshot. Replace existing DB instance by restoring the encrypted snapshot.

B. Create a new encrypted Amazon Elastic Block Store (Amazon EBS) volume and copy the snapshots to it. Enable encryption on the DB instance.

C. Copy the snapshots and enable encryption using AWS Key Management Service (AWS KMS) Restore encrypted snapshot to an existing DB instance.

D. Copy the snapshots to an Amazon S3 bucket that is encrypted using server-side encryption with AWS Key Management Service (AWS KMS) managed keys (SSE-KMS).

Correct Answer: *A*

Community vote distribution

A (71%) C (29%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: A

"You can enable encryption for an Amazon RDS DB instance when you create it, but not after it's created. However, you can add encryption to an unencrypted DB instance by creating a snapshot of your DB instance, and then creating an encrypted copy of that snapshot. You can then restore a DB instance from the encrypted snapshot to get an encrypted copy of your original DB instance."
https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/encrypt-an-existing-amazon-rds-for-postgresql-db-instance.html

upvoted 10 times

mricee9 Most Recent 6 days, 1 hour ago

Selected Answer: A

Cant be C - you cant restore it to an existing DB instance

upvoted 2 times

Wpcorgan 6 days, 16 hours ago

C is correct

upvoted 1 times

Jtic 2 weeks, 1 day ago

Selected Answer: C

It seems that D is the correct answer

Database and snapshots "are always encrypted moving forward?"
A. is only one time implementation

In the same document : https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/encrypt-an-existing-amazon-rds-for-postgresql-db-instance.html
It stated the steps moving forward under section "Encrypt the snapshot."

Select the Enable Encryption checkbox. For Master Key, specify the KMS key identifier to use to encrypt the DB snapshot copy

Prerequisites - Familiarity with AWS Key Management Service (AWS KMS) for encrypting databases

Architecture:
The destination RDS DB instance is created by restoring the DB snapshot copy of the source RDS DB instance.
An AWS KMS key is used for encryption while restoring the snapshot.

Tools: used to enable encryption
AWS KMS key for encryption

upvoted 1 times

Jtic 2 weeks, 1 day ago

sorry typo error: My vote is C

upvoted 1 times

tubtab 2 weeks, 5 days ago

Selected Answer: C

ccccccccc

upvoted 2 times

Bevemo 2 weeks, 5 days ago

Selected Answer: C

Restore to existing not replace instance. This link says copy and encrypt using KMS. https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/encrypt-an-existing-amazon-rds-for-postgresql-db-instance.html

upvoted 3 times

backbencher2022 3 weeks, 1 day ago

Selected Answer: A

A is the correct option based on this document - https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/encrypt-an-existing-amazon-rds-for-postgresql-db-instance.html

upvoted 3 times

Question #122 *Topic 1*

A company wants to build a scalable key management infrastructure to support developers who need to encrypt data in their applications.

What should a solutions architect do to reduce the operational burden?

A. Use multi-factor authentication (MFA) to protect the encryption keys.

B. Use AWS Key Management Service (AWS KMS) to protect the encryption keys.

C. Use AWS Certificate Manager (ACM) to create, store, and assign the encryption keys.

D. Use an IAM policy to limit the scope of users who have access permissions to protect the encryption keys.

Correct Answer: *B*

Community vote distribution

B (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: B

If you are a developer who needs to digitally sign or verify data using asymmetric keys, you should use the service to create and manage the private keys you'll need. If you're looking for a scalable key management infrastructure to support your developers and their growing number of applications, you should use it to reduce your licensing costs and operational burden...
https://aws.amazon.com/kms/faqs/#:~:text=If%20you%20are%20a%20developer%20who%20needs%20to%20digitally,a%20broad%20set%20of%20industry%20and%20regional%20compliance%20regimes.

upvoted 6 times

Wpcorgan Most Recent 6 days, 16 hours ago

B is correct

upvoted 1 times

Jtic 2 weeks ago

Selected Answer: B

If you are responsible for securing your data across AWS services, you should use it to centrally manage the encryption keys that control access to your data. If you are a developer who needs to encrypt data in your applications, you should use the AWS Encryption SDK with AWS KMS to easily generate, use and protect symmetric encryption keys in your code.

upvoted 1 times

Question #123 *Topic 1*

A company has a dynamic web application hosted on two Amazon EC2 instances. The company has its own SSL certificate, which is on each instance to perform SSL termination.
There has been an increase in traffic recently, and the operations team determined that SSL encryption and decryption is causing the compute capacity of the web servers to reach their maximum limit.
What should a solutions architect do to increase the application's performance?

A. Create a new SSL certificate using AWS Certificate Manager (ACM). Install the ACM certificate on each instance.

B. Create an Amazon S3 bucket Migrate the SSL certificate to the S3 bucket. Configure the EC2 instances to reference the bucket for SSL termination.

C. Create another EC2 instance as a proxy server. Migrate the SSL certificate to the new instance and configure it to direct connections to the existing EC2 instances.

D. Import the SSL certificate into AWS Certificate Manager (ACM). Create an Application Load Balancer with an HTTPS listener that uses the SSL certificate from ACM.

Correct Answer: *D*

Community vote distribution

D (100%)

Wpcorgan 6 days, 16 hours ago

D is correct

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: D

agree with D

upvoted 1 times

123jhl0 1 month, 1 week ago

Selected Answer: D

This issue is solved by SSL offloading, i.e. by moving the SSL termination task to the ALB.
https://aws.amazon.com/blogs/aws/elastic-load-balancer-support-for-ssl-termination/

upvoted 4 times

Question #124 *Topic 1*

A company has a highly dynamic batch processing job that uses many Amazon EC2 instances to complete it. The job is stateless in nature, can be started and stopped at any given time with no negative impact, and typically takes upwards of 60 minutes total to complete. The company has asked a solutions architect to design a scalable and cost-effective solution that meets the requirements of the job.
What should the solutions architect recommend?

A. Implement EC2 Spot Instances.

B. Purchase EC2 Reserved Instances.

C. Implement EC2 On-Demand Instances.

D. Implement the processing on AWS Lambda.

Correct Answer: *A*

Community vote distribution

A (100%)

Wpcorgan 6 days, 16 hours ago

A is correct

upvoted 1 times

SimonPark 4 weeks, 1 day ago

Selected Answer: A

A is the answer

upvoted 1 times

Evangelia 1 month, 1 week ago

spot instances

upvoted 3 times

Question #125 *Topic 1*

A company runs its two-tier ecommerce website on AWS. The web tier consists of a load balancer that sends traffic to Amazon EC2 instances. The database tier uses an Amazon RDS DB instance. The EC2 instances and the RDS DB instance should not be exposed to the public internet. The EC2 instances require internet access to complete payment processing of orders through a third-party web service. The application must be highly available.
Which combination of configuration options will meet these requirements? (Choose two.)

A. Use an Auto Scaling group to launch the EC2 instances in private subnets. Deploy an RDS Multi-AZ DB instance in private subnets.

B. Configure a VPC with two private subnets and two NAT gateways across two Availability Zones. Deploy an Application Load Balancer in the private subnets.

C. Use an Auto Scaling group to launch the EC2 instances in public subnets across two Availability Zones. Deploy an RDS Multi-AZ DB instance in private subnets.

D. Configure a VPC with one public subnet, one private subnet, and two NAT gateways across two Availability Zones. Deploy an Application Load Balancer in the public subnet.
D. Configure a VPC with two public subnets, two private subnets, and two NAT gateways across two Availability Zones. Deploy an Application Load Balancer in the public subnets.

Correct Answer: *CE*

Community vote distribution
A (64%) AD (21%) 14%

Wpcorgan 6 days, 16 hours ago
A and E
upvoted 2 times

AbhiJo 1 week, 2 days ago
Selected Answer: A
AE, E is not shown in voting,

E is correct, 2 subnets
upvoted 1 times

pspinelli19 1 week, 4 days ago
AE --> https://www.examtopics.com/discussions/amazon/view/60023-exam-aws-certified-solutions-architect-associate-saa-c02/
upvoted 1 times

Newptone 1 week, 5 days ago
AE
This link should be clear: https://aws.amazon.com/premiumsupport/knowledge-center/public-load-balancer-private-ec2/
upvoted 1 times

mabotega 1 week, 6 days ago
Selected Answer: AB
sorry guys, reading better, all instances should be in private subnets, so, answers should be A and B.
upvoted 1 times

mabotega 2 weeks, 2 days ago
Selected Answer: AD
Answer A for: The EC2 instances and the RDS DB instance should not be exposed to the public internet. Answer D for: The EC2 instances require internet access to complete payment processing of orders through a third-party web service. Answer A for: The application must be highly available.
upvoted 3 times

AbhiJo 1 week, 2 days ago
We will require 2 private subnets, D does mention 1 subnet
upvoted 1 times

Pamban 2 weeks, 2 days ago
Selected Answer: A
moderator, please correct the answers. there should be 2 answers and options are A B C D E. correct answer would be A &E. letter D has been duplicated

https://docs.aws.amazon.com/vpc/latest/userguide/vpc-nat-gateway.html

upvoted 3 times

mabotega 2 weeks, 3 days ago

Selected Answer: AB

A and B are the correct choices
https://www.examtopics.com/discussions/amazon/view/60023-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

PS_R 2 weeks, 5 days ago

Selected Answer: A

A E (Last option) looks to be more efficient

upvoted 1 times

Nigma 2 weeks, 6 days ago

"E" for request of network segmentation.
"A" for highly available
The two options together offer the indicated requirements

upvoted 1 times

backbencher2022 3 weeks, 1 day ago

First of all, due to typo error in the options, last option typed as D is actually E. For the sake of simplicity, we can say First (A) and Last option (E mistyped as D) are the correct options. Here is the explanation and reference to AWS document which helps us choose First (A) and Last (E) options as correct answer
https://aws.amazon.com/premiumsupport/knowledge-center/public-load-balancer-private-ec2/#

upvoted 2 times

TaiTran1994 3 weeks, 4 days ago

where option E ????

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: A

A and E is the answer

upvoted 2 times

ManoAni 1 month ago

Selected Answer: A

A & E is the answer

upvoted 1 times

dave9994 1 month ago

A & E should be the answer.

upvoted 1 times

palermo777 1 month, 1 week ago

A - since EC2 instances should not be exposed to the Internet
+E

upvoted 1 times

LukM 1 month, 1 week ago

E there is not. It would be:
E. Configure a VPC with two public subnets, two private subnets, and two NAT gateways across two Availability Zones. Deploy an Application Load Balancer in the public subnets.

upvoted 4 times

Question #126 *Topic 1*

A solutions architect needs to implement a solution to reduce a company's storage costs. All the company's data is in the Amazon S3 Standard storage class. The company must keep all data for at least 25 years. Data from the most recent 2 years must be highly available and immediately retrievable.
Which solution will meet these requirements?

A. Set up an S3 Lifecycle policy to transition objects to S3 Glacier Deep Archive immediately.

B. Set up an S3 Lifecycle policy to transition objects to S3 Glacier Deep Archive after 2 years.

C. Use S3 Intelligent-Tiering. Activate the archiving option to ensure that data is archived in S3 Glacier Deep Archive.

D. Set up an S3 Lifecycle policy to transition objects to S3 One Zone-Infrequent Access (S3 One Zone-IA) immediately and to S3 Glacier Deep Archive after 2 years.

Correct Answer: *B*

Community vote distribution

B (60%) C (40%)

Wpcorgan 6 days, 16 hours ago

B is correct

upvoted 1 times

TelaO 1 week, 3 days ago

Selected Answer: B

B is the only right answer. C does not indicate archiving after 2 years. If it did specify 2 years, then C would also be an option.

upvoted 2 times

rjam 1 week, 4 days ago

Selected Answer: B

Why Not C? Because Intelligent Tier the objects are automatically moved to different tiers.
The question says "the data from most recent 2 yrs should be highly available and immediately retrievable", which means in intelligent tier , if you activate archiving option(as Option C specifies) , the objects will be moved to Archive tiers(instant to access to deep archive access tiers) in 90 to 730 days. Remember these archive tiers performance will be similar to S3 glacier flexible and s3 deep archive which means files cannot be retrieved immediately within 2 yrs .

We have a hard requirement in question which says it should be retreivable immediately for the 2 yrs. which cannot be acheived in Intelligent tier. So B is the correct option imho.

Because of the above reason Its possible only in S3 standard and then configure lifecycle configuration to move to S3 Glacier Deep Archive after 2 yrs.

upvoted 2 times

Jtic 2 weeks ago

Selected Answer: C

C - S3 Intelligent-Tiering
Customers saving on storage with S3 Intelligent-Tiering

S3 Intelligent-Tiering automatically stores objects in three access tiers: one tier optimized for frequent access, a lower-cost tier optimized for infrequent access, and a very-low-cost tier optimized for rarely accessed data. For a small monthly object monitoring and automation charge, S3 Intelligent-Tiering moves objects that have not been accessed for 30 consecutive days to the Infrequent Access tier for savings of 40%; and after 90 days of no access, they're

There are no retrieval charges in S3 Intelligent-Tiering. S3 Intelligent-Tiering has no minimum eligible object size, but objects smaller than 128 KB are not eligible for auto tiering. These smaller objects may be stored, but they'll always be charged at the Frequent Access tier rates and don't incur the monitoring and automation charge

upvoted 1 times

Deplake 2 weeks ago

Selected Answer: B

Option C doesn't look correct for me because it is not clear when it will be moved to the Deep Archive. It could be earlier then 2 years, which is not correct

upvoted 2 times

Wilson_S 1 week, 5 days ago

https://docs.aws.amazon.com/AmazonS3/latest/userguide/intelligent-tiering-overview.html#:~:text=S3%20Intelligent%2DTiering%20provides%20you,minimum%20of%2090%20consecutive%20days. Option B / S3 Glacier Deep Archive seems correct to reduce a company's storage costs.

upvoted 1 times

MyNameIsJulien 2 weeks, 1 day ago

Selected Answer: C

The answer C seems correct

upvoted 3 times

ArielSchivo 2 weeks, 5 days ago

Glacier Deep Archive restores objects within 12 hours, so option A is out.
Option B could work but you will be paying S3 Standard for 2 years.
I would go with Option C then.
Option D is out since S3 One Zone IA is not highly available.

upvoted 1 times

rjam 3 weeks, 6 days ago

Option D as one-zone IA is cheaper than standard s3 . they never mentioned about multi zone. so we will go for one zone IA. The question mainly talks about reducing storage costs

upvoted 1 times

masetromain 3 weeks, 4 days ago

Data from the most recent 2 years must be highly available and immediately retrievable.

upvoted 4 times

Question #127 *Topic 1*

A media company is evaluating the possibility of moving its systems to the AWS Cloud. The company needs at least 10 TB of storage with the maximum possible I/O performance for video processing, 300 TB of very durable storage for storing media content, and 900 TB of storage to meet requirements for archival media that is not in use anymore.
Which set of services should a solutions architect recommend to meet these requirements?

A. Amazon EBS for maximum performance, Amazon S3 for durable data storage, and Amazon S3 Glacier for archival storage

B. Amazon EBS for maximum performance, Amazon EFS for durable data storage, and Amazon S3 Glacier for archival storage

C. Amazon EC2 instance store for maximum performance, Amazon EFS for durable data storage, and Amazon S3 for archival storage

D. Amazon EC2 instance store for maximum performance, Amazon S3 for durable data storage, and Amazon S3 Glacier for archival storage

Correct Answer: *A*

Community vote distribution

D (100%)

Sauran Highly Voted 1 month, 1 week ago

Selected Answer: D

Max instance store possible at this time is 30TB for NVMe which has the higher I/O compared to EBS.

is4gen.8xlarge 4 x 7,500 GB (30 TB) NVMe SSD

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/InstanceStorage.html#instance-store-volumes

upvoted 7 times

Wpcorgan Most Recent 6 days, 15 hours ago

D is corect

upvoted 1 times

mabotega 2 weeks, 2 days ago

Selected Answer: D

Answer: D

upvoted 1 times

dokaedu 4 weeks, 1 day ago

Selected Answer: D

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: D

agree with D, since it is only used for video processing instance store should be the fastest here (being ephemeral shouldnt be an issue because they are moving the data to S3 after processing)

upvoted 4 times

tubtab 1 month ago

Selected Answer: D

dddddd

upvoted 2 times

Lilibell 1 month, 2 weeks ago

The answer is D

upvoted 2 times

brushek 1 month, 1 week ago

wrong ! it is A, as there is 10TB needed:
taken from: https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/add-instance-store-volumes.html

A block device mapping always specifies the root volume for the instance. The root volume is either an Amazon EBS volume or an instance store volume. For more information, see Storage for the root device. The root volume is mounted automatically. For instances with an instance store volume for the root volume, the size of this volume varies by AMI, but the maximum size is 10 GB.

upvoted 1 times

Sauran 1 month, 1 week ago

It is possible to add an instance store volume up to 30TB at this time. The question doesn't say the instance store has to be the root volume. For instance stores, the root volume max size is 10GB but you can add instance stores up to 30TB.

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/InstanceStorage.html#instance-store-volumes
upvoted 4 times

Question #128 *Topic 1*

A company wants to run applications in containers in the AWS Cloud. These applications are stateless and can tolerate disruptions within the underlying infrastructure. The company needs a solution that minimizes cost and operational overhead.
What should a solutions architect do to meet these requirements?

A. Use Spot Instances in an Amazon EC2 Auto Scaling group to run the application containers.

B. Use Spot Instances in an Amazon Elastic Kubernetes Service (Amazon EKS) managed node group.

C. Use On-Demand Instances in an Amazon EC2 Auto Scaling group to run the application containers.

D. Use On-Demand Instances in an Amazon Elastic Kubernetes Service (Amazon EKS) managed node group.

Correct Answer: *A*

Community vote distribution

B (100%)

Wpcorgan 6 days, 15 hours ago
B is correct
upvoted 1 times

study_aws1 1 week, 4 days ago
Selected Answer: B
This should explain

https://docs.aws.amazon.com/eks/latest/userguide/managed-node-groups.html
upvoted 3 times

mabotega 2 weeks, 2 days ago
Selected Answer: B
Answer B
upvoted 1 times

Six_Fingered_Jose 1 month ago
Selected Answer: B
agreed with B cause container
upvoted 2 times

tubtab 1 month ago
Selected Answer: B
bbbbbb
upvoted 1 times

Lilibell 1 month, 2 weeks ago
The answer is B
upvoted 1 times

brushek 1 month, 2 weeks ago
Selected Answer: B
it should be B:

https://aws.amazon.com/about-aws/whats-new/2020/12/amazon-eks-support-ec2-spot-instances-managed-node-groups/
upvoted 3 times

Question #129 *Topic 1*

A company is running a multi-tier web application on premises. The web application is containerized and runs on a number of Linux hosts connected to a PostgreSQL database that contains user records. The operational overhead of maintaining the infrastructure and capacity planning is limiting the company's growth. A solutions architect must improve the application's infrastructure.
Which combination of actions should the solutions architect take to accomplish this? (Choose two.)

A. Migrate the PostgreSQL database to Amazon Aurora.

B. Migrate the web application to be hosted on Amazon EC2 instances.

C. Set up an Amazon CloudFront distribution for the web application content.

D. Set up Amazon ElastiCache between the web application and the PostgreSQL database.

E. Migrate the web application to be hosted on AWS Fargate with Amazon Elastic Container Service (Amazon ECS).

Correct Answer: *AE*

Community vote distribution

AE (100%)

Wpcorgan 6 days, 15 hours ago
A and E
upvoted 1 times

ArielSchivo 2 weeks, 5 days ago
Selected Answer: AE
I would say A and E since Aurora and Fargate are serverless (less operational overhead).
upvoted 2 times

Nigma 2 weeks, 6 days ago
https://www.examtopics.com/discussions/amazon/view/46457-exam-aws-certified-solutions-architect-associate-saa-c02/
upvoted 1 times

Nigma 2 weeks, 6 days ago
A and E

Aurora and serverless
upvoted 1 times

SimonPark 4 weeks, 1 day ago
Selected Answer: AE
B(X) E(O) not sure about A,C,D but A looks making sense
upvoted 1 times

Question #130 *Topic 1*

An application runs on Amazon EC2 instances across multiple Availability Zonas. The instances run in an Amazon EC2 Auto Scaling group behind an Application Load Balancer. The application performs best when the CPU utilization of the EC2 instances is at or near 40%.
What should a solutions architect do to maintain the desired performance across all instances in the group?

A. Use a simple scaling policy to dynamically scale the Auto Scaling group.

B. Use a target tracking policy to dynamically scale the Auto Scaling group.

C. Use an AWS Lambda function ta update the desired Auto Scaling group capacity.

D. Use scheduled scaling actions to scale up and scale down the Auto Scaling group.

Correct Answer: *B*

Community vote distribution

B (100%)

Wpcorgan 6 days, 15 hours ago

B is correct
upvoted 1 times

ArielSchivo 2 weeks, 5 days ago

Selected Answer: B

Option B. Target tracking policy.

https://docs.aws.amazon.com/autoscaling/ec2/userguide/as-scaling-target-tracking.html
upvoted 1 times

Nigma 2 weeks, 6 days ago

B

CPU utilization = target tracking
upvoted 1 times

SimonPark 4 weeks, 1 day ago

Selected Answer: B

B is the answer
upvoted 1 times

Question #131 *Topic 1*

A company is developing a file-sharing application that will use an Amazon S3 bucket for storage. The company wants to serve all the files through an Amazon CloudFront distribution. The company does not want the files to be accessible through direct navigation to the S3 URL. What should a solutions architect do to meet these requirements?

A. Write individual policies for each S3 bucket to grant read permission for only CloudFront access.

B. Create an IAM user. Grant the user read permission to objects in the S3 bucket. Assign the user to CloudFront.

C. Write an S3 bucket policy that assigns the CloudFront distribution ID as the Principal and assigns the target S3 bucket as the Amazon Resource Name (ARN).

D. Create an origin access identity (OAI). Assign the OAI to the CloudFront distribution. Configure the S3 bucket permissions so that only the OAI has read permission.

Correct Answer: *D*

Community vote distribution

D (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: D

I want to restrict access to my Amazon Simple Storage Service (Amazon S3) bucket so that objects can be accessed only through my Amazon CloudFront distribution. How can I do that?
Create a CloudFront origin access identity (OAI)
https://aws.amazon.com/premiumsupport/knowledge-center/cloudfront-access-to-amazon-s3/

upvoted 11 times

SimonPark 4 weeks, 1 day ago

Thanks it convinces me

upvoted 1 times

Wpcorgan Most Recent 6 days, 1 hour ago

D is correct

upvoted 1 times

Question #132 *Topic 1*

A company's website provides users with downloadable historical performance reports. The website needs a solution that will scale to meet the company's website demands globally. The solution should be cost-effective, limit the provisioning of infrastructure resources, and provide the fastest possible response time.
Which combination should a solutions architect recommend to meet these requirements?

A. Amazon CloudFront and Amazon S3

B. AWS Lambda and Amazon DynamoDB

C. Application Load Balancer with Amazon EC2 Auto Scaling

D. Amazon Route 53 with internal Application Load Balancers

Correct Answer: *A*

Community vote distribution

A (75%) 13% 13%

Wpcorgan 6 days, 1 hour ago

A is correct

upvoted 1 times

sdasdawa 2 weeks ago

Selected Answer: A

https://www.examtopics.com/discussions/amazon/view/27935-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 2 times

Nirmal3331 2 weeks ago

Selected Answer: A

https://www.examtopics.com/discussions/amazon/view/27935-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 2 times

samplunk 2 weeks, 1 day ago

Selected Answer: A

See this discussion:
https://www.examtopics.com/discussions/amazon/view/27935-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 2 times

manu427 2 weeks, 1 day ago

Selected Answer: C

load balancing + scalability + cost effective

upvoted 1 times

MyNameIsJulien 2 weeks, 1 day ago

Selected Answer: B

I think the answer is B

upvoted 1 times

dokaedu 4 weeks, 1 day ago

A is the correct answer
The solution should be cost-effective, limit the provisioning of infrastructure resources, and provide the fastest possible response time.

upvoted 4 times

Question #133 *Topic 1*

A company runs an Oracle database on premises. As part of the company's migration to AWS, the company wants to upgrade the database to the most recent available version. The company also wants to set up disaster recovery (DR) for the database. The company needs to minimize the operational overhead for normal operations and DR setup. The company also needs to maintain access to the database's underlying operating system.
Which solution will meet these requirements?

A. Migrate the Oracle database to an Amazon EC2 instance. Set up database replication to a different AWS Region.

B. Migrate the Oracle database to Amazon RDS for Oracle. Activate Cross-Region automated backups to replicate the snapshots to another AWS Region.

C. Migrate the Oracle database to Amazon RDS Custom for Oracle. Create a read replica for the database in another AWS Region.

D. Migrate the Oracle database to Amazon RDS for Oracle. Create a standby database in another Availability Zone.

Correct Answer: *D*

Community vote distribution

C (72%) D (17%) 10%

brushek Highly Voted 1 month, 2 weeks ago

Selected Answer: C

It should be C:

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/rds-custom.html
and
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/working-with-custom-oracle.html

upvoted 7 times

Gabs90 Most Recent 2 days, 3 hours ago

Maybe I missed the point, but why not A?
"The company also needs to maintain access to the database's underlying operating system."
If the company want to operate on an operating system level the only option here should be the EC2.

upvoted 2 times

leonnnn 23 hours, 26 minutes ago

Amazon RDS Custom for Oracle also provides ability to let user to access operating system.

upvoted 2 times

Wpcorgan 6 days, 1 hour ago

C is correct

upvoted 2 times

Gil80 2 weeks, 4 days ago

Selected Answer: C

Amazon RDS Custom for Oracle - • Managed Oracle and Microsoft SQL Server Database with OS and database customization.

If the company wants to reduce operational overhead, then RDS Custom is the way to go, since it's a managed service.

upvoted 1 times

Athena 2 weeks, 4 days ago

Selected Answer: D

Disaster Recovery = Standby Instance not read replica
The company also wants to set up disaster recovery (DR) for the database.

upvoted 1 times

Athena 2 weeks, 1 day ago

Answer is C for sure
https://aws.amazon.com/about-aws/whats-new/2021/10/amazon-rds-custom-oracle/ Z9 (Announced in 2021)

upvoted 1 times

Athena 2 weeks, 1 day ago

Still debating the correct answer on this one.....

upvoted 1 times

Maharaja 2 weeks ago

D also make sense to me. This is my understanding -
Muti-AZ gives you automated DR within a region. No need to change end point from your app end. But in case of Cross region read replica

you need to promote the same to main DB during DR and modify the endpoint of the DB in the application.
By doing Read replica this statement -"The company needs to minimize the operational overhead for normal operations and DR setup. " in the question is not meet which can be achieved by multi-AZ only. Moreover, there is no mention of global DR. DR can be global or even AZ failure in Single region as well. Hope that make sense.

upvoted 2 times

ArielSchivo 2 weeks, 5 days ago

Option C since RDS Custom has access to the underlying OS and it provides less operational overhead. Also, a read replica in another Region can be used for DR activities.

https://aws.amazon.com/blogs/database/implementing-a-disaster-recovery-strategy-with-amazon-rds/

upvoted 2 times

PS_R 3 weeks ago

Selected Answer: C

RDS Custom only allows access to underlying OS access, so it should be C

upvoted 1 times

backbencher2022 3 weeks, 1 day ago

Selected Answer: C

C is the correct answer indeed and here is the latest Oracle blog supporting this option - https://aws.amazon.com/blogs/database/build-high-availability-for-amazon-rds-custom-for-oracle-using-read-replicas/

B or D can't be the right option because RDS Oracle won't allow OS access. Only RDS Custom for Oracle allows OS access. Scenario is specifically talking about OS access.

Option A is ruled out for a different reason - more operational and admin overhead.

upvoted 2 times

backbencher2022 3 weeks, 1 day ago

Selected Answer: C

C is the correct answer indeed and here is the latest Oracle blog supporting this option - https://aws.amazon.com/blogs/database/build-high-availability-for-amazon-rds-custom-for-oracle-using-read-replicas/

B or D can't be the right option because RDS Oracle won't allow OS access. Only RDS Custom for Oracle allows OS access. Scenario is specifically talking about OS access

Option is ruled out for a different reason - more operational and admin overhead.

upvoted 2 times

Onimole 3 weeks, 2 days ago

Selected Answer: D

Maharaja appears to be right
Read replica wont help HA

upvoted 1 times

ArielSchivo 2 weeks, 5 days ago

No HA mentioned on the subject, this is about DR. Option C is the one.

upvoted 2 times

Maharaja 2 weeks, 1 day ago

Muti-AZ gives you automated DR within a region. No need to change end point from your app end. But in case of Cross region read replica you need to promote the same to main DB during DR and modify the endpoint of the DB in the application.
By doing Read replica this statement -"The company needs to minimize the operational overhead for normal operations and DR setup. " in the question is not meet which can be achieved by multi-AZ only. Moreover, there is no mention of global DR. DR can be global or even AZ failure in Single region as well. Hope that make sense.

upvoted 1 times

Maharaja 3 weeks, 4 days ago

D. Please have a look at this statement "To support replication between RDS Custom for Oracle instances, you can configure high availability (HA) with Oracle Data Guard. The primary DB instance automatically synchronizes data to the standby instances.

You can configure your high availability environment in the following ways:

Configure standby instances in different Availability Zones (AZs) to be resilient to AZ failures.

Place your standby databases in mounted or read-only mode.

Fail over or switch over from the primary database to a standby database with no data loss.

Migrate data by configuring high availability for your on-premises instance, and then failing over or switching over to the RDS Custom standby database." from https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/custom-managing.html

upvoted 2 times

ArielSchivo 2 weeks, 5 days ago

You are talking about HA but here we've been requested to create a DR environment, two different things. Also, with option D you can't have access to the underlying OS.

upvoted 1 times

SimonPark 4 weeks, 1 day ago

Selected Answer: D

Even though C seems to be right, I don't think that a 'read replica' can replace the primary DB instance in case of DR.

upvoted 3 times

ArielSchivo 2 weeks, 5 days ago

Yes, it can, check https://aws.amazon.com/blogs/database/implementing-a-disaster-recovery-strategy-with-amazon-rds/ .

upvoted 2 times

Reni_Varghese 4 weeks, 1 day ago

It is C for sure.
Please go thru the below documentation.
https://aws.amazon.com/blogs/aws/amazon-rds-custom-for-oracle-new-control-capabilities-in-database-environment/

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: C

agree with C
> The company also needs to maintain access to the database's underlying operating system

Amazon RDS Custom for Oracle provides that but not regular RDS
> https://aws.amazon.com/about-aws/whats-new/2021/10/amazon-rds-custom-oracle/

A has too much overhead so is not the optimal solution

upvoted 3 times

ManoAni 1 month ago

Selected Answer: C

https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/rds-custom.html

upvoted 3 times

dave9994 1 month ago

Selected Answer: C

C should be the answer. Not D, only creating a standby doesn't satisfy the condition. The database should also be in sync, so need to create a read replica, which is in sync with the master.

upvoted 2 times

herculian_effort 1 month ago

Selected Answer: A

A, because "The company also needs to maintain access to the database's underlying operating system".
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Welcome.html#Welcome.Concepts.RDS

upvoted 3 times

Question #134 *Topic 1*

A company wants to move its application to a serverless solution. The serverless solution needs to analyze existing and new data by using SL. The company stores the data in an Amazon S3 bucket. The data requires encryption and must be replicated to a different AWS Region.
Which solution will meet these requirements with the LEAST operational overhead?

A. Create a new S3 bucket. Load the data into the new S3 bucket. Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region. Use server-side encryption with AWS KMS multi-Region kays (SSE-KMS). Use Amazon Athena to query the data.

B. Create a new S3 bucket. Load the data into the new S3 bucket. Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region. Use server-side encryption with AWS KMS multi-Region keys (SSE-KMS). Use Amazon RDS to query the data.

C. Load the data into the existing S3 bucket. Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region. Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Use Amazon Athena to query the data.

D. Load the data into the existing S3 bucket. Use S3 Cross-Region Replication (CRR) to replicate encrypted objects to an S3 bucket in another Region. Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Use Amazon RDS to query the data.

Correct Answer: *A*

Community vote distribution

C (69%) A (31%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: C

SSE-KMS vs SSE-S3 - The last seems to have less overhead (as the keys are automatically generated by S3 and applied on data at upload, and don't require further actions. KMS provides more flexibility, but in turn involves a different service, which finally is more "complex" than just managing one (S3). So A and B are excluded. If you are in doubt, you are having 2 buckets in A and B, while just keeping one in C and D.
https://s3browser.com/server-side-encryption-types.aspx
Decide between C and D is deciding on Athena or RDS. RDS is a relational db, and we have documents on S3, which is the use case for Athena. Athena is also serverless, which eliminates the need of controlling the underlying infrastructure and capacity. So C is the answer.
https://aws.amazon.com/athena/

upvoted 8 times

dokaedu Highly Voted 4 weeks, 1 day ago

Answer is A:
Amazon S3 Bucket Keys reduce the cost of Amazon S3 server-side encryption using AWS Key Management Service (SSE-KMS). This new bucket-level key for SSE can reduce AWS KMS request costs by up to 99 percent by decreasing the request traffic from Amazon S3 to AWS KMS. With a few clicks in the AWS Management Console, and without any changes to your client applications, you can configure your bucket to use an S3 Bucket Key for AWS KMS-based encryption on new objects.
The Existing S3 bucket might have uncrypted data - encryption will apply new data received after the applying of encryption on the new bucket.

upvoted 6 times

RODROSKAR 2 weeks, 6 days ago

Reducing cost was never the target, it's LEAST operational. In that regard SSE-S3 AWS fully managed.

upvoted 1 times

JayanKuruwita Most Recent 1 day, 6 hours ago

Selected Answer: C

C would be hard if we have to replicate set of existing objects, because we need AWS support assistance for that. But this is about new objects that we are going to upload to S3. Thinking about setting up permissions for KMS key with SSE-KMS mode. I would go with C.

upvoted 1 times

justtry 2 days, 5 hours ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonS3/latest/userguide/default-bucket-encryption.html
There is no change to the encryption of the objects that existed in the bucket before default encryption was enabled.That's the reason why new bucket should be created.

upvoted 1 times

Wpcorgan 6 days, 1 hour ago

A is correct for me

upvoted 1 times

Onimole 1 week, 2 days ago

C appears to be right

upvoted 1 times

Ohnet 1 week, 6 days ago

It is A:
https://aws.amazon.com/blogs/storage/replicating-existing-objects-between-s3-buckets/
upvoted 1 times

study_aws1 2 weeks, 3 days ago
The issue with Option C) is the existing data before setting up CRR will not be replicated and that is why option A) looks good.
The question does not talk of employing S3 Batch Replication before CRR is setup.
upvoted 3 times

backbencher2022 3 weeks, 1 day ago
Selected Answer: A
A is the solution with least operational overhead
upvoted 1 times

masetromain 3 weeks, 5 days ago
Selected Answer: A
for me the answer is A:

The Existing S3 bucket might have uncrypted data - encryption will apply new data received after the applying of encryption on the new bucket.
upvoted 3 times

UWSFish 1 month ago
The question should read "SQL" not "SL"
upvoted 3 times

dave9994 1 month ago
Selected Answer: C
C has fewer ops. overhead. and the question says the customer has an existing S3 bucket. Then there is no need to create a new S3 bucket.
upvoted 2 times

Question #135 *Topic 1*

A company runs workloads on AWS. The company needs to connect to a service from an external provider. The service is hosted in the provider's VPC. According to the company's security team, the connectivity must be private and must be restricted to the target service. The connection must be initiated only from the company's VPC.
Which solution will mast these requirements?

A. Create a VPC peering connection between the company's VPC and the provider's VPC. Update the route table to connect to the target service.

B. Ask the provider to create a virtual private gateway in its VPC. Use AWS PrivateLink to connect to the target service.

C. Create a NAT gateway in a public subnet of the company's VPUpdate the route table to connect to the target service.

D. Ask the provider to create a VPC endpoint for the target service. Use AWS PrivateLink to connect to the target service.

Correct Answer: *D*

Community vote distribution
D (100%)

123jhl0 Highly Voted 1 month, 1 week ago
Selected Answer: D
AWS PrivateLink provides private connectivity between VPCs, AWS services, and your on-premises networks, without exposing your traffic to the public internet. AWS PrivateLink makes it easy to connect services across different accounts and VPCs to significantly simplify your network architecture.
Interface **VPC endpoints**, powered by AWS PrivateLink, connect you to services hosted by AWS Partners and supported solutions available in AWS Marketplace.
https://aws.amazon.com/privatelink/
upvoted 7 times

Question #136 *Topic 1*

A company is migrating its on-premises PostgreSQL database to Amazon Aurora PostgreSQL. The on-premises database must remain online and accessible during the migration. The Aurora database must remain synchronized with the on-premises database.
Which combination of actions must a solutions architect take to meet these requirements? (Choose two.)

A. Create an ongoing replication task.

B. Create a database backup of the on-premises database.

C. Create an AWS Database Migration Service (AWS DMS) replication server.

D. Convert the database schema by using the AWS Schema Conversion Tool (AWS SCT).

E. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to monitor the database synchronization.

Correct Answer: *CD*

Community vote distribution

AC (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: AC

AWS Database Migration Service (AWS DMS) helps you migrate databases to AWS quickly and securely. The source database remains fully operational during the migration, minimizing downtime to applications that rely on the database.
... With AWS Database Migration Service, you can also continuously replicate data with low latency from any supported source to any supported target.
https://aws.amazon.com/dms/

upvoted 6 times

JayanKuruwita Most Recent 1 day, 4 hours ago

Highly cofused why we ned SCT in here, because we don't have to use SCT for homogeneous migration. If someone knows please explain.

upvoted 1 times

Wpcorgan 6 days, 1 hour ago

Selected Answer: AC

A and C

upvoted 1 times

mabotega 2 weeks, 3 days ago

Selected Answer: AC

https://docs.aws.amazon.com/dms/latest/userguide/CHAP_Task.CDC.html

upvoted 2 times

SimonPark 4 weeks, 1 day ago

Selected Answer: AC

A and C

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: AC

They are migrating Postgres to Postgres,
thus D does not make sense and A and C seems to be the right answer to me

upvoted 2 times

dave9994 1 month ago

Selected Answer: AC

A & C are the correct combinations.

upvoted 2 times

trancex 1 month, 2 weeks ago

I think A and C

upvoted 1 times

huiy 1 month, 2 weeks ago

Selected Answer: AC

Same as this https://www.examtopics.com/discussions/amazon/view/81317-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 2 times

Question #137 *Topic 1*

A company uses AWS Organizations to create dedicated AWS accounts for each business unit to manage each business unit's account independently upon request. The root email recipient missed a notification that was sent to the root user email address of one account. The company wants to ensure that all future notifications are not missed. Future notifications must be limited to account administrators.
Which solution will meet these requirements?

A. Configure the company's email server to forward notification email messages that are sent to the AWS account root user email address to all users in the organization.

B. Configure all AWS account root user email addresses as distribution lists that go to a few administrators who can respond to alerts. Configure AWS account alternate contacts in the AWS Organizations console or programmatically.

C. Configure all AWS account root user email messages to be sent to one administrator who is responsible for monitoring alerts and forwarding those alerts to the appropriate groups.

D. Configure all existing AWS accounts and all newly created accounts to use the same root user email address. Configure AWS account alternate contacts in the AWS Organizations console or programmatically.

Correct Answer: *D*

Community vote distribution

B (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: B

Use a group email address for the management account's root user
https://docs.aws.amazon.com/organizations/latest/userguide/orgs_best-practices_mgmt-acct.html#best-practices_mgmt-acct_email-address

upvoted 11 times

PS_R Most Recent 3 weeks ago

Selected Answer: B

B makes more sense and is a best practise

upvoted 1 times

Chunsli 1 month, 1 week ago

Selected Answer: B

B makes better sense in the context

upvoted 3 times

Question #138 *Topic 1*

A company runs its ecommerce application on AWS. Every new order is published as a massage in a RabbitMQ queue that runs on an Amazon EC2 instance in a single Availability Zone. These messages are processed by a different application that runs on a separate EC2 instance. This application stores the details in a PostgreSQL database on another EC2 instance. All the EC2 instances are in the same Availability Zone.
The company needs to redesign its architecture to provide the highest availability with the least operational overhead.
What should a solutions architect do to meet these requirements?

A. Migrate the queue to a redundant pair (active/standby) of RabbitMQ instances on Amazon MQ. Create a Multi-AZ Auto Scaling group for EC2 instances that host the application. Create another Multi-AZ Auto Scaling group for EC2 instances that host the PostgreSQL database.

B. Migrate the queue to a redundant pair (active/standby) of RabbitMQ instances on Amazon MQ. Create a Multi-AZ Auto Scaling group for EC2 instances that host the application. Migrate the database to run on a Multi-AZ deployment of Amazon RDS for PostgreSQL.

C. Create a Multi-AZ Auto Scaling group for EC2 instances that host the RabbitMQ queue. Create another Multi-AZ Auto Scaling group for EC2 instances that host the application. Migrate the database to run on a Multi-AZ deployment of Amazon RDS for PostgreSQL.

D. Create a Multi-AZ Auto Scaling group for EC2 instances that host the RabbitMQ queue. Create another Multi-AZ Auto Scaling group for EC2 instances that host the application. Create a third Multi-AZ Auto Scaling group for EC2 instances that host the PostgreSQL database

Correct Answer: *B*

Community vote distribution

B (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: B

Migrating to Amazon MQ reduces the overhead on the queue management. C and D are dismissed.
Deciding between A and B means deciding to go for an AutoScaling group for EC2 or an RDS for Postgress (both multi- AZ). The RDS option has less operational impact, as provide as a service the tools and software required. Consider for instance, the effort to add an additional node like a read replica, to the DB.
https://docs.aws.amazon.com/amazon-mq/latest/developer-guide/active-standby-broker-deployment.html
https://aws.amazon.com/rds/postgresql/

upvoted 5 times

UWSFish 1 month ago

Yes but active/standby is fault tolerance, not HA. I would concede after thinking about it that B is probably the answer that will be marked correct but its not a great question.

upvoted 1 times

Wpcorgan Most Recent 6 days, 1 hour ago

B for me

upvoted 1 times

Question #139 *Topic 1*

A reporting team receives files each day in an Amazon S3 bucket. The reporting team manually reviews and copies the files from this initial S3 bucket to an analysis S3 bucket each day at the same time to use with Amazon QuickSight. Additional teams are starting to send more files in larger sizes to the initial S3 bucket.
The reporting team wants to move the files automatically analysis S3 bucket as the files enter the initial S3 bucket. The reporting team also wants to use AWS Lambda functions to run pattern-matching code on the copied data. In addition, the reporting team wants to send the data files to a pipeline in Amazon SageMaker Pipelines.
What should a solutions architect do to meet these requirements with the LEAST operational overhead?

A. Create a Lambda function to copy the files to the analysis S3 bucket. Create an S3 event notification for the analysis S3 bucket. Configure Lambda and SageMaker Pipelines as destinations of the event notification. Configure s3:ObjectCreated:Put as the event type.

B. Create a Lambda function to copy the files to the analysis S3 bucket. Configure the analysis S3 bucket to send event notifications to Amazon EventBridge (Amazon CloudWatch Events). Configure an ObjectCreated rule in EventBridge (CloudWatch Events). Configure Lambda and SageMaker Pipelines as targets for the rule.

C. Configure S3 replication between the S3 buckets. Create an S3 event notification for the analysis S3 bucket. Configure Lambda and SageMaker Pipelines as destinations of the event notification. Configure s3:ObjectCreated:Put as the event type.

D. Configure S3 replication between the S3 buckets. Configure the analysis S3 bucket to send event notifications to Amazon EventBridge (Amazon CloudWatch Events). Configure an ObjectCreated rule in EventBridge (CloudWatch Events). Configure Lambda and SageMaker Pipelines as targets for the rule.

Correct Answer: *A*

Community vote distribution

D (64%) B (36%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: B

C and D aren't answers as replicating the S3 bucket isn't efficient, as other teams are starting to use it to store larger docs not related to the reporting, making replication not useful.
As Amazon SageMaker Pipelines, ..., is now supported as a target for routing events in Amazon EventBridge, means the answer is B
https://aws.amazon.com/about-aws/whats-new/2021/04/new-options-trigger-amazon-sagemaker-pipeline-executions/

upvoted 8 times

KADSM 2 weeks, 3 days ago

Not sure how far lambda will cope up with larger files with the timelimit in place.

upvoted 1 times

LeGloupier Highly Voted 1 month, 1 week ago

Selected Answer: D

D make more sense

upvoted 6 times

123jhl0 1 month, 1 week ago

Replicating docs not related to the reporting team, you are consuming more resources (duplicating storage of docs not needed), and very probably introducing discarding overhead, noise or breaking the reporting result, as these additional docs are introduced in the analytic process.

upvoted 1 times

Onimole 3 weeks, 2 days ago

arent they already duplicating it and need it to be in another bucket?

upvoted 1 times

Onimole 3 weeks, 2 days ago

someone creating a function is additional overhead. CRR replication is an aws feature

upvoted 3 times

nhlegend Most Recent 2 weeks, 2 days ago

Should be D
They manually copy data, uses S3 replication to copy the files (eliminate A and B)
C is incorrect since S3 event notification can only send to SNS, SQS and Lambda
https://docs.aws.amazon.com/AmazonS3/latest/userguide/NotificationHowTo.html
So D is correct

upvoted 2 times

KADSM 2 weeks, 3 days ago

Answer D - The files are getting large, less operational overhead - so will choose S3 replication. Event bridge is far more advanced than S3 event notification and they support multiple targets. S3 Event notification may not support Sagemaker. Also filtering and pattern matching available in Event bridge. So answer D

upvoted 3 times

backbencher2022 3 weeks, 1 day ago

Selected Answer: D

Option D has the least overhead and should be the correct answer in my opinion.

upvoted 3 times

Six_Fingered_Jose 1 month ago

Selected Answer: D

i go for D here
A and B says you are copying the file to another bucket using lambda,
C an D just uses S3 replication to copy the files,

They are doing exactly the same thing while C and D do not require setting up of lambda, which should be more efficient

The question says the team is manually copying the files, automatically replicating the files should be the most efficient method vs manually copying or copying with lambda.

upvoted 5 times

Question #140 *Topic 1*

A solutions architect needs to help a company optimize the cost of running an application on AWS. The application will use Amazon EC2 instances, AWS Fargate, and AWS Lambda for compute within the architecture.
The EC2 instances will run the data ingestion layer of the application. EC2 usage will be sporadic and unpredictable. Workloads that run on EC2 instances can be interrupted at any time. The application front end will run on Fargate, and Lambda will serve the API layer. The front-end utilization and API layer utilization will be predictable over the course of the next year.
Which combination of purchasing options will provide the MOST cost-effective solution for hosting this application? (Choose two.)

A. Use Spot Instances for the data ingestion layer

B. Use On-Demand Instances for the data ingestion layer

C. Purchase a 1-year Compute Savings Plan for the front end and API layer.

D. Purchase 1-year All Upfront Reserved instances for the data ingestion layer.

E. Purchase a 1-year EC2 instance Savings Plan for the front end and API layer.

Correct Answer: *AC*

Community vote distribution

AC (100%)

rjam 1 week, 6 days ago

its A and C . https://www.densify.com/finops/aws-savings-plan

upvoted 1 times

SimonPark 4 weeks, 1 day ago

Selected Answer: AC

EC2 instance Savings Plan saves 72% while Compute Savings Plans saves 66%. But according to link, it says "Compute Savings Plans provide the most flexibility and help to reduce your costs by up to 66%. These plans automatically apply to EC2 instance usage regardless of instance family, size, AZ, region, OS or tenancy, and also apply to Fargate and Lambda usage." EC2 instance Savings Plans are not applied to Fargate or Lambda

upvoted 3 times

bunnychip 1 month ago

Selected Answer: AC

api is not EC2.need to use compute savings plan

upvoted 3 times

Chunsli 1 month, 1 week ago

E makes more sense than C. See https://aws.amazon.com/savingsplans/faq/, EC2 instance Savings Plan (up to 72% saving) costs less than Compute Savings Plan (up to 66% saving)

upvoted 3 times

capepenguin 1 month ago

Isn't the EC2 Instance Savings Plan not applicable to Fargate and Lambda?
https://aws.amazon.com/savingsplans/compute-pricing/

upvoted 4 times

Question #141 *Topic 1*

A company runs a web-based portal that provides users with global breaking news, local alerts, and weather updates. The portal delivers each user a personalized view by using mixture of static and dynamic content. Content is served over HTTPS through an API server running on an Amazon EC2 instance behind an Application Load Balancer (ALB). The company wants the portal to provide this content to its users across the world as quickly as possible.
How should a solutions architect design the application to ensure the LEAST amount of latency for all users?

A. Deploy the application stack in a single AWS Region. Use Amazon CloudFront to serve all static and dynamic content by specifying the ALB as an origin.

B. Deploy the application stack in two AWS Regions. Use an Amazon Route 53 latency routing policy to serve all content from the ALB in the closest Region.

C. Deploy the application stack in a single AWS Region. Use Amazon CloudFront to serve the static content. Serve the dynamic content directly from the ALB.

D. Deploy the application stack in two AWS Regions. Use an Amazon Route 53 geolocation routing policy to serve all content from the ALB in the closest Region.

Correct Answer: *B*

Community vote distribution

A (72%) B (28%)

huiy Highly Voted 1 month, 2 weeks ago

Selected Answer: A

Answer is A.
Amazon CloudFront is a web service that speeds up distribution of your static and dynamic web content
https://www.examtopics.com/discussions/amazon/view/81081-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 9 times

Pamban Most Recent 2 weeks, 1 day ago

Selected Answer: A

Answer A. refer Question #: 612 in SAA-C002
https://www.examtopics.com/discussions/amazon/view/81081-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Bevemo 2 weeks, 5 days ago

Selected Answer: B

B = as quickly as possible = lowest latency dynamic + static content.
CloudFront works for static and dynamic content, but the dynamic content delivery will not be 'as fast as possible' if it's coming from another region.

upvoted 1 times

tubtab 1 month ago

Selected Answer: A

aaaaaaaaaaaa

upvoted 3 times

Six_Fingered_Jose 1 month ago

Selected Answer: B

Answer should be B,

CloudFront reduces latency if its only static content, which is not the case here.
For Dynamic content, CF cant cache the content so it sends the traffic through the AWS Network which does reduces latency, but it still has to travel through another region.

For the case with 2 region and Route 53 latency routing, Route 53 detects the nearest resouce (with lowest latency) and routes the traffic there. Because the traffic does not have to travel to resources far away, it should have the least latency in this case here.

upvoted 4 times

Onimole 3 weeks, 2 days ago

Cf works for both static and dynamic content

upvoted 2 times

trancex 1 month, 2 weeks ago

not sure but A https://aws.amazon.com/blogs/networking-and-content-delivery/deliver-your-apps-dynamic-content-using-amazon-cloudfront-getting-started-template/

upvoted 1 times

Question #142 *Topic 1*

A gaming company is designing a highly available architecture. The application runs on a modified Linux kernel and supports only UDP-based traffic. The company needs the front-end tier to provide the best possible user experience. That tier must have low latency, route traffic to the nearest edge location, and provide static IP addresses for entry into the application endpoints.
What should a solutions architect do to meet these requirements?

A. Configure Amazon Route 53 to forward requests to an Application Load Balancer. Use AWS Lambda for the application in AWS Application Auto Scaling.

B. Configure Amazon CloudFront to forward requests to a Network Load Balancer. Use AWS Lambda for the application in an AWS Application Auto Scaling group.

C. Configure AWS Global Accelerator to forward requests to a Network Load Balancer. Use Amazon EC2 instances for the application in an EC2 Auto Scaling group.

D. Configure Amazon API Gateway to forward requests to an Application Load Balancer. Use Amazon EC2 instances for the application in an EC2 Auto Scaling group.

Correct Answer: *C*

Community vote distribution

C (100%)

dokaedu Highly Voted 4 weeks, 1 day ago

Correct Answer: C
AWS Global Accelerator and Amazon CloudFront are separate services that use the AWS global network and its edge locations around the world. CloudFront improves performance for both cacheable content (such as images and videos) and dynamic content (such as API acceleration and dynamic site delivery). Global Accelerator improves performance for a wide range of applications over TCP or UDP by proxying packets at the edge to applications running in one or more AWS Regions. Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses or deterministic, fast regional failover. Both services integrate with AWS Shield for DDoS protection.

upvoted 8 times

iCcma 5 days, 8 hours ago

Thank you, your explanation helped me to better understand even the answer of question 29

upvoted 1 times

Wpcorgan Most Recent 3 days, 19 hours ago

C is correct

upvoted 1 times

PS_R 3 weeks ago

Selected Answer: C

Cloud Fronts supports both Static and Dynamic and Global Accelerator means low latency over UDP

upvoted 1 times

Question #143 *Topic 1*

A company wants to migrate its existing on-premises monolithic application to AWS. The company wants to keep as much of the front-end code and the backend code as possible. However, the company wants to break the application into smaller applications. A different team will manage each application. The company needs a highly scalable solution that minimizes operational overhead.
Which solution will meet these requirements?

A. Host the application on AWS Lambda. Integrate the application with Amazon API Gateway.

B. Host the application with AWS Amplify. Connect the application to an Amazon API Gateway API that is integrated with AWS Lambda.

C. Host the application on Amazon EC2 instances. Set up an Application Load Balancer with EC2 instances in an Auto Scaling group as targets.

D. Host the application on Amazon Elastic Container Service (Amazon ECS). Set up an Application Load Balancer with Amazon ECS as the target.

Correct Answer: *D*

Community vote distribution

D (75%) A (25%)

Ken701 Highly Voted 4 weeks, 1 day ago

I think the answer here is "D" because usually when you see terms like "monolithic" the answer will likely refer to microservices.

upvoted 9 times

Wpcorgan Most Recent 3 days, 19 hours ago

D is correct

upvoted 1 times

Bevemo 2 weeks, 5 days ago

Selected Answer: D

D is organic pattern, lift and shift, decompose to containers, first making most use of existing code, whilst new features can be added over time with lambda+api gw later.
A is leapfrog pattern. requiring refactoring all code up front.

upvoted 3 times

backbencher2022 3 weeks, 1 day ago

Selected Answer: D

I think D is the right choice as they want application to be managed by different people which could be enabled by breaking it into different containers

upvoted 1 times

SimonPark 4 weeks, 1 day ago

Selected Answer: D

imho, it's D because "break the application into smaller applications" doesn't mean it has to be 'serverless'. Rather it can be divided into smaller application running on containers.

upvoted 2 times

Six_Fingered_Jose 1 month ago

Selected Answer: A

I think A is the answer here, breaking into smaller pieces so lambda makes the most sense.
I don't see any restrictions in the question that forbids the usage of lambda

upvoted 2 times

Newptone 4 days, 13 hours ago

The reason for not choosing A: "The company wants to keep as much of the front-end code and the backend code as possible"

upvoted 1 times

Question #144 *Topic 1*

A company recently started using Amazon Aurora as the data store for its global ecommerce application. When large reports are run, developers report that the ecommerce application is performing poorly. After reviewing metrics in Amazon CloudWatch, a solutions architect finds that the ReadIOPS and CPUUtilizalion metrics are spiking when monthly reports run.
What is the MOST cost-effective solution?

A. Migrate the monthly reporting to Amazon Redshift.

B. Migrate the monthly reporting to an Aurora Replica.

C. Migrate the Aurora database to a larger instance class.

D. Increase the Provisioned IOPS on the Aurora instance.

Correct Answer: *B*

Community vote distribution

B (100%)

Wpcorgan 5 days, 22 hours ago

B is correct

upvoted 1 times

backbencher2022 3 weeks, 1 day ago

Selected Answer: B

ReadIOPS issue inclining towards Read Replica as the most cost effective solution here

upvoted 3 times

rjam 3 weeks, 6 days ago

Answer B

upvoted 1 times

Question #145 *Topic 1*

A company hosts a website analytics application on a single Amazon EC2 On-Demand Instance. The analytics software is written in PHP and uses a MySQL database. The analytics software, the web server that provides PHP, and the database server are all hosted on the EC2 instance. The application is showing signs of performance degradation during busy times and is presenting 5xx errors. The company needs to make the application scale seamlessly.
Which solution will meet these requirements MOST cost-effectively?

A. Migrate the database to an Amazon RDS for MySQL DB instance. Create an AMI of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use an Application Load Balancer to distribute the load to each EC2 instance.

B. Migrate the database to an Amazon RDS for MySQL DB instance. Create an AMI of the web application. Use the AMI to launch a second EC2 On-Demand Instance. Use Amazon Route 53 weighted routing to distribute the load across the two EC2 instances.

C. Migrate the database to an Amazon Aurora MySQL DB instance. Create an AWS Lambda function to stop the EC2 instance and change the instance type. Create an Amazon CloudWatch alarm to invoke the Lambda function when CPU utilization surpasses 75%.

D. Migrate the database to an Amazon Aurora MySQL DB instance. Create an AMI of the web application. Apply the AMI to a launch template. Create an Auto Scaling group with the launch template Configure the launch template to use a Spot Fleet. Attach an Application Load Balancer to the Auto Scaling group.

Correct Answer: *D*

Community vote distribution

D (60%) A (40%)

ds0321 4 days, 1 hour ago
D is correct
upvoted 1 times

Wpcorgan 5 days, 22 hours ago
D is correct
upvoted 1 times

dokaedu 4 weeks, 1 day ago
Answer: D
upvoted 1 times

SimonPark 4 weeks, 1 day ago
Selected Answer: D
I think D is the answer
upvoted 1 times

USalo 4 weeks, 1 day ago
Selected Answer: D
Agreed with D as Spot Fleet can leverage both spot+on-demand instances, it should be the most cost-effective.
https://www.youtube.com/watch?v=rlYLbs33Ofs&ab_channel=AmazonWebServices
upvoted 2 times

Six_Fingered_Jose 1 month ago
Selected Answer: A
Answer should be A,
question is asking for most cost effective way to scale the application layer, which means AS on EC2 is enough.

Moreover, Option D is saying to use Spot Fleet, which has potential to cause downtime in the app layer
upvoted 2 times

rewdboy 2 days, 12 hours ago
I see what you're saying, but the question also requires that "The company needs to make the application scale seamlessly." This tells me that it needs Auto Scaling. Auto Scaling with Spot fleet will allow for automatic scaling up and scaling down, which will save money over having two 24x7 EC2 instances.
upvoted 1 times

Six_Fingered_Jose 1 month ago
Actually on second thought, adding optional on-demand instances might solve the spot fleet problem, so maybe leaning back to D again. Wish the questions would be more specific
upvoted 7 times

Question #146 *Topic 1*

A company runs a stateless web application in production on a group of Amazon EC2 On-Demand Instances behind an Application Load Balancer. The application experiences heavy usage during an 8-hour period each business day. Application usage is moderate and steady overnight. Application usage is low during weekends.
The company wants to minimize its EC2 costs without affecting the availability of the application.
Which solution will meet these requirements?

A. Use Spot Instances for the entire workload.

B. Use Reserved Instances for the baseline level of usage. Use Spot instances for any additional capacity that the application needs.

C. Use On-Demand Instances for the baseline level of usage. Use Spot Instances for any additional capacity that the application needs.

D. Use Dedicated Instances for the baseline level of usage. Use On-Demand Instances for any additional capacity that the application needs.

Correct Answer: *B*

Community vote distribution

B (70%) C (30%)

Wpcorgan 5 days, 21 hours ago

B IS CORRECT

upvoted 1 times

ArielSchivo 2 weeks, 4 days ago

Selected Answer: B

They are currently using On Demand instances, so option C is out.
A uses Spot instances which is not recommended for PROD and D uses Dedicated instances which are expensive.
So option B should be the one.

upvoted 3 times

rob74 3 weeks, 5 days ago

Selected Answer: B

In the Question is mentioned that it has o Demand instances...so I think is more cheapest reserved and spot

upvoted 4 times

Dsouzaf 3 weeks, 6 days ago

If we select B, Spot instance are reliable though it saves cost.
In D: base line & additional capacity is also On-Demand. Expensive than Reserve Instance but will not bring down Production

upvoted 1 times

TaiTran1994 3 weeks, 6 days ago

Selected Answer: C

I think C should be corrected.

upvoted 3 times

Question #147 *Topic 1*

A company needs to retain application log files for a critical application for 10 years. The application team regularly accesses logs from the past month for troubleshooting, but logs older than 1 month are rarely accessed. The application generates more than 10 TB of logs per month.
Which storage option meets these requirements MOST cost-effectively?

A. Store the logs in Amazon S3. Use AWS Backup to move logs more than 1 month old to S3 Glacier Deep Archive.

B. Store the logs in Amazon S3. Use S3 Lifecycle policies to move logs more than 1 month old to S3 Glacier Deep Archive.

C. Store the logs in Amazon CloudWatch Logs. Use AWS Backup to move logs more than 1 month old to S3 Glacier Deep Archive.

D. Store the logs in Amazon CloudWatch Logs. Use Amazon S3 Lifecycle policies to move logs more than 1 month old to S3 Glacier Deep Archive.

Correct Answer: *B*

Community vote distribution

B (100%)

Wpcorgan 5 days, 21 hours ago

B is correct

upvoted 1 times

rjam 1 week, 4 days ago

Selected Answer: B

Why not AwsBackup? No Glacier Deep is supported by AWS Backup

https://docs.aws.amazon.com/aws-backup/latest/devguide/s3-backups.html
AWS Backup allows you to backup your S3 data stored in the following S3 Storage Classes:
• S3 Standard
• S3 Standard - Infrequently Access (IA)
• S3 One Zone-IA
• S3 Glacier Instant Retrieval
• S3 Intelligent-Tiering (S3 INT)

upvoted 1 times

tdkcumberland 1 day, 17 hours ago

AWS BackUp costs something, setting up S3 LCP doesn't.

upvoted 1 times

ArielSchivo 2 weeks, 4 days ago

Selected Answer: B

S3 + Glacier is the most cost effective.

upvoted 2 times

Bevemo 2 weeks, 5 days ago

Selected Answer: B

D works, archive cloudwatch logs to S3 .... but is an additional service to pay for over B.

upvoted 1 times

masetromain 3 weeks, 4 days ago

Selected Answer: B

https://www.examtopics.com/discussions/amazon/view/80772-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Question #148 *Topic 1*

A company has a data ingestion workflow that includes the following components:
An Amazon Simple Notification Service (Amazon SNS) topic that receives notifications about new data deliveries
An AWS Lambda function that processes and stores the data
The ingestion workflow occasionally fails because of network connectivity issues. When failure occurs, the corresponding data is not ingested unless the company manually reruns the job.
What should a solutions architect do to ensure that all notifications are eventually processed?

A. Configure the Lambda function for deployment across multiple Availability Zones.

B. Modify the Lambda function's configuration to increase the CPU and memory allocations for the function.

C. Configure the SNS topic's retry strategy to increase both the number of retries and the wait time between retries.

D. Configure an Amazon Simple Queue Service (Amazon SQS) queue as the on-failure destination. Modify the Lambda function to process messages in the queue.

Correct Answer: *D*

Community vote distribution

D (75%) C (25%)

bunnychip Highly Voted 1 month ago

Selected Answer: D

ensure that all notifications are eventually processed

upvoted 6 times

LeGloupier Most Recent 5 days, 21 hours ago

Selected Answer: C

"on-failure destination" doesn't exist, only dead letter queue exist.
that's why I am leaning for C

upvoted 1 times

Wpcorgan 5 days, 21 hours ago

D is correct

upvoted 1 times

ds0321 1 week, 2 days ago

Selected Answer: D

D is the answer

upvoted 1 times

ArielSchivo 2 weeks, 4 days ago

Selected Answer: D

Option C could work but the max retries attempts is 23 days. After that messages are deleted. And you do not want that to happen! So, Option D.

upvoted 1 times

SimonPark 4 weeks, 1 day ago

Selected Answer: D

imho, D is the answer

upvoted 1 times

brushek 1 month, 2 weeks ago

Selected Answer: C

should be C:

https://docs.aws.amazon.com/sns/latest/dg/sns-message-delivery-retries.html

upvoted 2 times

Question #149 *Topic 1*

A company has a service that produces event data. The company wants to use AWS to process the event data as it is received. The data is written in a specific order that must be maintained throughout processing. The company wants to implement a solution that minimizes operational overhead.
How should a solutions architect accomplish this?

A. Create an Amazon Simple Queue Service (Amazon SQS) FIFO queue to hold messages. Set up an AWS Lambda function to process messages from the queue.

B. Create an Amazon Simple Notification Service (Amazon SNS) topic to deliver notifications containing payloads to process. Configure an AWS Lambda function as a subscriber.

C. Create an Amazon Simple Queue Service (Amazon SQS) standard queue to hold messages. Set up an AWS Lambda function to process messages from the queue independently.

D. Create an Amazon Simple Notification Service (Amazon SNS) topic to deliver notifications containing payloads to process. Configure an Amazon Simple Queue Service (Amazon SQS) queue as a subscriber.

Correct Answer: *A*

Community vote distribution

A (100%)

Wpcorgan 5 days, 21 hours ago
A is correct
upvoted 1 times

ArielSchivo 2 weeks, 4 days ago
Selected Answer: A
FIFO means order, so Option A.
upvoted 3 times

rjam 3 weeks, 6 days ago
Order --- means FIFO option A
upvoted 2 times

Question #150 *Topic 1*

A company is migrating an application from on-premises servers to Amazon EC2 instances. As part of the migration design requirements, a solutions architect must implement infrastructure metric alarms. The company does not need to take action if CPU utilization increases to more than 50% for a short burst of time. However, if the CPU utilization increases to more than 50% and read IOPS on the disk are high at the same time, the company needs to act as soon as possible. The solutions architect also must reduce false alarms.
What should the solutions architect do to meet these requirements?

A. Create Amazon CloudWatch composite alarms where possible.

B. Create Amazon CloudWatch dashboards to visualize the metrics and react to issues quickly.

C. Create Amazon CloudWatch Synthetics canaries to monitor the application and raise an alarm.

D. Create single Amazon CloudWatch metric alarms with multiple metric thresholds where possible.

Correct Answer: *A*

Community vote distribution

A (100%)

123jhl0 Highly Voted 1 month, 1 week ago

Selected Answer: A

Composite alarms determine their states by monitoring the states of other alarms. You can **use composite alarms to reduce alarm noise**. For example, you can create a composite alarm where the underlying metric alarms go into ALARM when they meet specific conditions. You then can set up your composite alarm to go into ALARM and send you notifications when the underlying metric alarms go into ALARM by configuring the underlying metric alarms never to take actions. Currently, composite alarms can take the following actions:
https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/Create_Composite_Alarm.html

upvoted 8 times

Wpcorgan Most Recent 5 days, 21 hours ago

A is correct

upvoted 1 times

Question #151 *Topic 1*

A company wants to migrate its on-premises data center to AWS. According to the company's compliance requirements, the company can use only the ap-northeast-3 Region. Company administrators are not permitted to connect VPCs to the internet.
Which solutions will meet these requirements? (Choose two.)

A. Use AWS Control Tower to implement data residency guardrails to deny internet access and deny access to all AWS Regions except ap-northeast-3.

B. Use rules in AWS WAF to prevent internet access. Deny access to all AWS Regions except ap-northeast-3 in the AWS account settings.

C. Use AWS Organizations to configure service control policies (SCPS) that prevent VPCs from gaining internet access. Deny access to all AWS Regions except ap-northeast-3.

D. Create an outbound rule for the network ACL in each VPC to deny all traffic from 0.0.0.0/0. Create an IAM policy for each user to prevent the use of any AWS Region other than ap-northeast-3.

E. Use AWS Config to activate managed rules to detect and alert for internet gateways and to detect and alert for new resources deployed outside of ap-northeast-3.

Correct Answer: *AC*

Community vote distribution

AC (100%)

Wpcorgan 5 days, 21 hours ago

A and C

upvoted 1 times

rjam 1 week, 6 days ago

https://aws.amazon.com/blogs/aws/new-for-aws-control-tower-region-deny-and-guardrails-to-help-you-meet-data-residency-requirements/
*Disallow internet access for an Amazon VPC instance managed by a customer

upvoted 3 times

rjam 1 week, 6 days ago

Option A and C

upvoted 1 times

rjam 1 week, 6 days ago

*You can use data-residency guardrails to control resources in any AWS Region.

upvoted 1 times

Six_Fingered_Jose 1 month ago

Selected Answer: AC

agree with A and C
https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scps_examples_vpc.html#example_vpc_2

upvoted 3 times

Question #152 *Topic 1*

A company uses a three-tier web application to provide training to new employees. The application is accessed for only 12 hours every day. The company is using an Amazon RDS for MySQL DB instance to store information and wants to minimize costs.
What should a solutions architect do to meet these requirements?

A. Configure an IAM policy for AWS Systems Manager Session Manager. Create an IAM role for the policy. Update the trust relationship of the role. Set up automatic start and stop for the DB instance.

B. Create an Amazon ElastiCache for Redis cache cluster that gives users the ability to access the data from the cache when the DB instance is stopped. Invalidate the cache after the DB instance is started.

C. Launch an Amazon EC2 instance. Create an IAM role that grants access to Amazon RDS. Attach the role to the EC2 instance. Configure a cron job to start and stop the EC2 instance on the desired schedule.

D. Create AWS Lambda functions to start and stop the DB instance. Create Amazon EventBridge (Amazon CloudWatch Events) scheduled rules to invoke the Lambda functions. Configure the Lambda functions as event targets for the rules.

Correct Answer: *D*

Community vote distribution

D (69%) A (31%)

study_aws1 Highly Voted 4 weeks ago

https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-lambda/

It is option D. Option A could have been applicable had it been AWS Systems Manager State Manager & not AWS Systems Manager Session Manager

upvoted 7 times

ArielSchivo Most Recent 2 weeks, 4 days ago

Selected Answer: D

Option D is the one. Option A could be as well if it referred to State Manager instead of Session Manager.

upvoted 3 times

rob74 3 weeks, 5 days ago

Selected Answer: D

I think A or D but D is cheaper (mimimize costs) because you pay Lambda only if you use it.

upvoted 1 times

rob74 3 weeks, 5 days ago

I think A or D but D is cheaper (mimimize costs) because you pay Lambda only if you use it.

upvoted 2 times

SimonPark 4 weeks, 1 day ago

Selected Answer: D

voted d

upvoted 2 times

Kien048 1 month ago

Selected Answer: D

Vote D

upvoted 3 times

Six_Fingered_Jose 1 month ago

Selected Answer: A

agreed with A

upvoted 1 times

123jhl0 1 month, 1 week ago

Selected Answer: A

A is true for sure. "Schedule Amazon RDS stop and start using AWS Systems Manager" Steps in the documentation:
1. Configure an AWS Identity and Access Management (IAM) policy for State Manager.
2. Create an IAM role for the new policy.
3. Update the trust relationship of the role so Systems Manager can use it.
4. Set up the automatic stop with State Manager.
5. Set up the automatic start with State Manager.
https://aws.amazon.com/blogs/database/schedule-amazon-rds-stop-and-start-using-aws-systems-manager/

upvoted 3 times

ArielSchivo 2 weeks, 4 days ago

Option A refers to Session Manager, not State Manager as you pointed, so it is wrong. Option D is valid.

upvoted 1 times

Bevemo 2 weeks, 5 days ago

Agree A, free to use state manager within limits, and don't need to code or manage lambda.

upvoted 1 times

Kien048 1 month ago

Look like State manager and Session manager use for difference purpose even both in same dashboard console.

upvoted 1 times

Kien048 1 month ago

And ofcause, D is working, so if A also right, the question is wrong.

upvoted 2 times

Question #153 *Topic 1*

A company sells ringtones created from clips of popular songs. The files containing the ringtones are stored in Amazon S3 Standard and are at least 128 KB in size. The company has millions of files, but downloads are infrequent for ringtones older than 90 days. The company needs to save money on storage while keeping the most accessed files readily available for its users.
Which action should the company take to meet these requirements MOST cost-effectively?

A. Configure S3 Standard-Infrequent Access (S3 Standard-IA) storage for the initial storage tier of the objects.

B. Move the files to S3 Intelligent-Tiering and configure it to move objects to a less expensive storage tier after 90 days.

C. Configure S3 inventory to manage objects and move them to S3 Standard-Infrequent Access (S3 Standard-1A) after 90 days.

D. Implement an S3 Lifecycle policy that moves the objects from S3 Standard to S3 Standard-Infrequent Access (S3 Standard-1A) after 90 days.

Correct Answer: *D*

Community vote distribution

B (55%) D (45%)

taichun 3 days, 20 hours ago

Selected Answer: B

"128KB" and "90 Days" keywords should point to Answer B

upvoted 1 times

zeronine75 5 days, 9 hours ago

Selected Answer: B

B/D seems possible answer. But, I'll go with "B".
In the following table, S3 Intelligent-Tiering seems not so expansive than S3 Standard.
https://aws.amazon.com/s3/pricing/?nc1=h_ls
And, in the question "128KB" size is talking about S3 Intelligent-Tiering stuff.

upvoted 3 times

Wilson_S 4 days, 8 hours ago

This link also has me going with “B.” Specifying 128 KB in size is not a coincidence. https://aws.amazon.com/s3/storage-classes/intelligent-tiering/

upvoted 1 times

Wpcorgan 5 days, 21 hours ago

D is correct

upvoted 1 times

rjam 1 week, 5 days ago

Selected Answer: D

Answer D
Why Optoin D ?
The Question talks about downloads are infrequent older than 90 days which means files less than 90 days are accessed frequently. Standard-Infrequent Access (S3 Standard-IA) needs a minimum 30 days if accessed before, it costs more.
So to access the files frequently you need a S3 Standard . After 90 days you can move it to Standard-Infrequent Access (S3 Standard-IA) as its going to be less frequently accessed

upvoted 3 times

TaiTran1994 2 weeks, 4 days ago

Selected Answer: D

I think D is answer

upvoted 2 times

PS_R 2 weeks, 5 days ago

Selected Answer: B

Another thought on B, Because it is asking for cost effective, it is possible to move the data to Glacier instant retrival after 90 days

upvoted 2 times

PS_R 2 weeks, 4 days ago

but first 90 days we dont have infrequent data and hence intelligent tiering might not be required.

upvoted 3 times

Cynthia19 3 weeks, 3 days ago

answer : D

upvoted 3 times

Question #154 *Topic 1*

A company needs to save the results from a medical trial to an Amazon S3 repository. The repository must allow a few scientists to add new files and must restrict all other users to read-only access. No users can have the ability to modify or delete any files in the repository. The company must keep every file in the repository for a minimum of 1 year after its creation date.
Which solution will meet these requirements?

A. Use S3 Object Lock in governance mode with a legal hold of 1 year.

B. Use S3 Object Lock in compliance mode with a retention period of 365 days.

C. Use an IAM role to restrict all users from deleting or changing objects in the S3 bucket. Use an S3 bucket policy to only allow the IAM role.

D. Configure the S3 bucket to invoke an AWS Lambda function every time an object is added. Configure the function to track the hash of the saved object so that modified objects can be marked accordingly.

Correct Answer: *B*

Community vote distribution

B (69%) A (31%)

Cizzla7049 3 days, 21 hours ago

Selected Answer: A

ANSWER IS DEFINITELY A

upvoted 1 times

Wpcorgan 5 days, 21 hours ago

B i think. im not sure..thougts?

upvoted 1 times

Qjb8m9h 1 week, 6 days ago

Answer : B
Reason: Compliance Mode. The key difference between Compliance Mode and Governance Mode is that there are NO users that can override the retention periods set or delete an object, and that also includes your AWS root account which has the highest privileges.

upvoted 3 times

mabotega 2 weeks, 3 days ago

Selected Answer: A

https://cloudacademy.com/course/using-amazon-s3-bucket-properties-management-features-maintain-data/object-lock/#:~:text=be%20deleted%20again.-,Compliance%20Mode.,which%20has%20the%20highest%20privileges.

upvoted 1 times

ArielSchivo 2 weeks, 4 days ago

Selected Answer: B

"No users can have the ability to modify or delete any files in the repository" = Compliance mode.

upvoted 3 times

USalo 3 weeks, 3 days ago

Selected Answer: B

B. Due to compliance

upvoted 2 times

nikerlas 3 weeks, 4 days ago

A is Correct

"In governance mode, users can't overwrite or delete an object version or alter its lock settings unless they have special permissions. With governance mode, you protect objects against being deleted by most users, but you can still grant some users permission to alter the retention settings or delete the object if necessary."

https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock-overview.html

upvoted 1 times

Bobbybash 6 days, 7 hours ago

if you have very specific permissions, including s3:BypassGovernanceMode, s3:GetObjectLockConfiguration, s3:GetObjectRetention, then a user will still be able to delete an object version within the retention period or change any retention settings set on the bucket.

upvoted 1 times

bunnychip 1 month ago

Selected Answer: B

'No users" can have the ability to modify or delete any files in the repository
upvoted 4 times

Six_Fingered_Jose 1 month ago

Selected Answer: A

Answer should be A because a few scientist must be able to edit the file
> In governance mode, users can't overwrite or delete an object version or alter its lock settings unless they have special permissions.

It cant be B because in compliance mode, absolutely nobody can touch the file during its period
> In compliance mode, a protected object version can't be overwritten or deleted by any user, including the root user in your AWS account

https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock-overview.html#object-lock-retention-modes
upvoted 2 times

Six_Fingered_Jose 1 month ago

actually i read the question again
> No users can have the ability to modify or delete any files in the repository.

answer should be B ignore my comment
upvoted 9 times

dave9994 1 month ago

Compliance mode is more restrictive : https://docs.aws.amazon.com/AmazonS3/latest/userguide/object-lock-overview.html
upvoted 3 times

Question #155 *Topic 1*

A large media company hosts a web application on AWS. The company wants to start caching confidential media files so that users around the world will have reliable access to the files. The content is stored in Amazon S3 buckets. The company must deliver the content quickly, regardless of where the requests originate geographically.
Which solution will meet these requirements?

A. Use AWS DataSync to connect the S3 buckets to the web application.

B. Deploy AWS Global Accelerator to connect the S3 buckets to the web application.

C. Deploy Amazon CloudFront to connect the S3 buckets to CloudFront edge servers.

D. Use Amazon Simple Queue Service (Amazon SQS) to connect the S3 buckets to the web application.

Correct Answer: *C*

Community vote distribution
C (100%)

rjam Highly Voted 3 weeks, 5 days ago

key :caching
Option C
upvoted 5 times

Wpcorgan Most Recent 5 days, 21 hours ago

C is correct
upvoted 1 times

MyNameIsJulien 2 weeks, 1 day ago

Selected Answer: C

Answer is C
upvoted 1 times

Question #156 *Topic 1*

A company produces batch data that comes from different databases. The company also produces live stream data from network sensors and application APIs. The company needs to consolidate all the data into one place for business analytics. The company needs to process the incoming data and then stage the data in different Amazon S3 buckets. Teams will later run one-time queries and import the data into a business intelligence tool to show key performance indicators (KPIs).
Which combination of steps will meet these requirements with the LEAST operational overhead? (Choose two.)

A. Use Amazon Athena for one-time queries. Use Amazon QuickSight to create dashboards for KPIs.

B. Use Amazon Kinesis Data Analytics for one-time queries. Use Amazon QuickSight to create dashboards for KPIs.

C. Create custom AWS Lambda functions to move the individual records from the databases to an Amazon Redshift cluster.

D. Use an AWS Glue extract, transform, and load (ETL) job to convert the data into JSON format. Load the data into multiple Amazon OpenSearch Service (Amazon Elasticsearch Service) clusters.

E. Use blueprints in AWS Lake Formation to identify the data that can be ingested into a data lake. Use AWS Glue to crawl the source, extract the data, and load the data into Amazon S3 in Apache Parquet format.

Correct Answer: *AC*

Community vote distribution

AE (93%) 7%

Wazhija Highly Voted 1 month, 1 week ago

Selected Answer: AE

I believe AE makes the most sense

upvoted 6 times

Six_Fingered_Jose Highly Voted 1 month ago

Selected Answer: AE

yeah AE makes sense, only E is working with S3 here and questions wants them to be in S3

upvoted 5 times

DivaLight Most Recent 21 hours, 57 minutes ago

Selected Answer: AE

Option AE

upvoted 1 times

Cizzla7049 3 days, 21 hours ago

Selected Answer: AC

A and C are correct

upvoted 1 times

backbencher2022 3 weeks ago

Selected Answer: AE

A&E is the correct answer

upvoted 1 times

Dsouzaf 3 weeks, 6 days ago

AC is correct. Ans E is also correct But in ans E: since Apache Parquer format is used, this is not correct answer as per AWS exam answer Six_Fingered_Jose

upvoted 3 times

Question #157 *Topic 1*

A company stores data in an Amazon Aurora PostgreSQL DB cluster. The company must store all the data for 5 years and must delete all the data after 5 years. The company also must indefinitely keep audit logs of actions that are performed within the database. Currently, the company has automated backups configured for Aurora.

Which combination of steps should a solutions architect take to meet these requirements? (Choose two.)

A. Take a manual snapshot of the DB cluster.

B. Create a lifecycle policy for the automated backups.

C. Configure automated backup retention for 5 years.

D. Configure an Amazon CloudWatch Logs export for the DB cluster.

E. Use AWS Backup to take the backups and to keep the backups for 5 years.

Correct Answer: *BE*

Community vote distribution

DE (58%) AD (42%)

DivaLight 21 hours, 54 minutes ago

Selected Answer: DE

DE Option

upvoted 1 times

Phinx 1 day, 6 hours ago

Selected Answer: DE

D and E is the most sensible options here.

upvoted 1 times

justtry 2 days, 3 hours ago

Selected Answer: DE

https://aws.amazon.com/about-aws/whats-new/2020/06/amazon-aurora-snapshots-can-be-managed-via-aws-backup/?nc1=h_ls
AWS Backup adds Amazon Aurora database cluster snapshots as its latest protected resource

upvoted 1 times

Nightducky 3 days, 5 hours ago

Selected Answer: DE

There is no sense with A if you can use AWS backup and keep snapshot for 5 years.

upvoted 1 times

TECHNOWARRIOR 4 days, 6 hours ago

DE, AWS Backup adds Amazon Aurora database cluster snapshots as its latest protected resource. Starting today, you can use AWS Backup to manage Amazon Aurora database cluster snapshots. AWS Backup can centrally configure backup policies, monitor backup activity, copy a snapshot within and across AWS regions, except for China regions, where snapshots can only be copied from one China region to another.

upvoted 1 times

ds0321 4 days, 19 hours ago

Selected Answer: AD

35 days is the Maximum time for Backups aurora

upvoted 1 times

Nigma 5 days, 7 hours ago

Selected Answer: AD

A and D

upvoted 1 times

Wpcorgan 5 days, 21 hours ago

D and E

upvoted 1 times

peneloco 1 week, 3 days ago

Selected Answer: DE

I'm going for DE. Picked E because AWS backup does work to create and manage snapshots of Aurora DB instances.

upvoted 1 times

LeGloupier 1 week, 3 days ago

Selected Answer: DE

D for audit log
E for backup

upvoted 2 times

rjam 1 week, 5 days ago

Answer A and D
Manual DB Snapshots. Retention of backup for as long as you want
Audit Logs can be enabled and sent to CloudWatch Logs for longer retention

upvoted 3 times

babaxoxo 1 week, 5 days ago

Selected Answer: AD

Answer A&D
If you want to retain Aurora Backup beyond the maximum retention day (35) -> do manual snapshot

upvoted 3 times

Question #158 *Topic 1*

A solutions architect is optimizing a website for an upcoming musical event. Videos of the performances will be streamed in real time and then will be available on demand. The event is expected to attract a global online audience.

Which service will improve the performance of both the real-time and on-demand streaming?

A. Amazon CloudFront

B. AWS Global Accelerator

C. Amazon Route 53

D. Amazon S3 Transfer Acceleration

Correct Answer: *A*

Community vote distribution

A (100%)

Nigma Highly Voted 1 week, 5 days ago

A is right

You can use CloudFront to deliver video on demand (VOD) or live streaming video using any HTTP origin

Global Accelerator is a good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP, as well as for HTTP use cases that specifically require static IP addresses

upvoted 5 times

Wpcorgan Most Recent 5 days, 21 hours ago

A is correct

upvoted 1 times

babaxoxo 1 week, 5 days ago

Selected Answer: A

CloudFront for sure

upvoted 1 times

Question #159 *Topic 1*

A company is running a publicly accessible serverless application that uses Amazon API Gateway and AWS Lambda. The application's traffic recently spiked due to fraudulent requests from botnets.

Which steps should a solutions architect take to block requests from unauthorized users? (Choose two.)

A. Create a usage plan with an API key that is shared with genuine users only.

B. Integrate logic within the Lambda function to ignore the requests from fraudulent IP addresses.

C. Implement an AWS WAF rule to target malicious requests and trigger actions to filter them out.

D. Convert the existing public API to a private API. Update the DNS records to redirect users to the new API endpoint.

E. Create an IAM role for each user attempting to access the API. A user will assume the role when making the API call.

Correct Answer: *CD*

Community vote distribution

AC (100%)

Phinx 1 day, 6 hours ago

Selected Answer: AC

A and C are the correct choices.

upvoted 1 times

justtry 2 days, 3 hours ago

Selected Answer: AC

https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-usage-plans.html

upvoted 1 times

5up3rm4n 3 days, 5 hours ago

Only answer C is an obviouis choice. B and D are clearly not right and A is the only remotely viable other answer but even then the documentation on API Keys and Usage quotas states not to rely on it to block API requests;

Usage plan throttling and quotas are not hard limits, and are applied on a best-effort basis. In some cases, clients can exceed the quotas that you set. Don't rely on usage plan quotas or throttling to control costs or block access to an API. Consider using AWS Budgets to monitor costs and AWS WAF to manage API requests.

https://docs.aws.amazon.com/apigateway/latest/developerguide/api-gateway-api-usage-plans.html

upvoted 2 times

ds0321 4 days, 1 hour ago

Selected Answer: AC

A and C

upvoted 1 times

babaxoxo 1 week, 5 days ago

Selected Answer: AC

use usage plan API key

upvoted 2 times

Nigma 1 week, 5 days ago

A and C

upvoted 3 times

Question #160 *Topic 1*

An ecommerce company hosts its analytics application in the AWS Cloud. The application generates about 300 MB of data each month. The data is stored in JSON format. The company is evaluating a disaster recovery solution to back up the data. The data must be accessible in milliseconds if it is needed, and the data must be kept for 30 days.

Which solution meets these requirements MOST cost-effectively?

A. Amazon OpenSearch Service (Amazon Elasticsearch Service)

B. Amazon S3 Glacier

C. Amazon S3 Standard

D. Amazon RDS for PostgreSQL

Correct Answer: *C*

Community vote distribution

C (100%)

Wpcorgan 5 days, 21 hours ago

C is correct

upvoted 1 times

sdasdawa 1 week, 5 days ago

Selected Answer: C

IMHO
Normally ElasticSearch would be ideal here, however as question states "Most cost-effective"
S3 is the best choice in this case

upvoted 2 times

babaxoxo 1 week, 5 days ago

Selected Answer: C

Ans C:
Cost-effective solution with milliseconds of retrieval -> it should be s3 standard

upvoted 1 times

Question #161 *Topic 1*

A company has a small Python application that processes JSON documents and outputs the results to an on-premises SQL database. The application runs thousands of times each day. The company wants to move the application to the AWS Cloud. The company needs a highly available solution that maximizes scalability and minimizes operational overhead.

Which solution will meet these requirements?

A. Place the JSON documents in an Amazon S3 bucket. Run the Python code on multiple Amazon EC2 instances to process the documents. Store the results in an Amazon Aurora DB cluster.

B. Place the JSON documents in an Amazon S3 bucket. Create an AWS Lambda function that runs the Python code to process the documents as they arrive in the S3 bucket. Store the results in an Amazon Aurora DB cluster.

C. Place the JSON documents in an Amazon Elastic Block Store (Amazon EBS) volume. Use the EBS Multi-Attach feature to attach the volume to multiple Amazon EC2 instances. Run the Python code on the EC2 instances to process the documents. Store the results on an Amazon RDS DB instance.

D. Place the JSON documents in an Amazon Simple Queue Service (Amazon SQS) queue as messages. Deploy the Python code as a container on an Amazon Elastic Container Service (Amazon ECS) cluster that is configured with the Amazon EC2 launch type. Use the container to process the SQS messages. Store the results on an Amazon RDS DB instance.

Correct Answer: *D*

Community vote distribution

B (100%)

babaxoxo Highly Voted 1 week, 5 days ago

Selected Answer: B

solution should remove operation overhead -> s3 -> lambda -> aurora

upvoted 6 times

Phinx Most Recent 1 day, 6 hours ago

Selected Answer: B

D is incorrect because using ECS entails a lot of admin overhead. so B is the correct one.

upvoted 1 times

Wpcorgan 5 days, 21 hours ago

B is correct

upvoted 1 times

EKA_CloudGod 1 week ago

Selected Answer: B

B is the answer
https://aws.amazon.com/rds/aurora/

upvoted 1 times

BENICE 1 week, 4 days ago

D is correct option

upvoted 1 times

Nightducky 1 week, 4 days ago

ehhhhhh

upvoted 3 times

Question #162 *Topic 1*

A company wants to use high performance computing (HPC) infrastructure on AWS for financial risk modeling. The company's HPC workloads run on Linux. Each HPC workflow runs on hundreds of Amazon EC2 Spot Instances, is short-lived, and generates thousands of output files that are ultimately stored in persistent storage for analytics and long-term future use.

The company seeks a cloud storage solution that permits the copying of on-premises data to long-term persistent storage to make data available for processing by all EC2 instances. The solution should also be a high performance file system that is integrated with persistent storage to read and write datasets and output files.

Which combination of AWS services meets these requirements?

A. Amazon FSx for Lustre integrated with Amazon S3

B. Amazon FSx for Windows File Server integrated with Amazon S3

C. Amazon S3 Glacier integrated with Amazon Elastic Block Store (Amazon EBS)

D. Amazon S3 bucket with a VPC endpoint integrated with an Amazon Elastic Block Store (Amazon EBS) General Purpose SSD (gp2) volume

Correct Answer: *A*

Community vote distribution

A (100%)

Wpcorgan 5 days, 21 hours ago

A is correct

upvoted 1 times

BENICE 1 week, 4 days ago

A - for HPC "Amazon FSx for Lustre" and long-term persistence "S3"

upvoted 1 times

rjam 1 week, 5 days ago

Amazon FSx for Lustre:
• HPC optimized distributed file system, millions of IOPS
• Backed by S3

upvoted 2 times

rjam 1 week, 5 days ago

Answer A

upvoted 1 times

babaxoxo 1 week, 5 days ago

Selected Answer: A

FxS Lustre integrated with S3

upvoted 1 times

Question #163 *Topic 1*

A company is building a containerized application on premises and decides to move the application to AWS. The application will have thousands of users soon after it is deployed. The company is unsure how to manage the deployment of containers at scale. The company needs to deploy the containerized application in a highly available architecture that minimizes operational overhead.

Which solution will meet these requirements?

A. Store container images in an Amazon Elastic Container Registry (Amazon ECR) repository. Use an Amazon Elastic Container Service (Amazon ECS) cluster with the AWS Fargate launch type to run the containers. Use target tracking to scale automatically based on demand.

B. Store container images in an Amazon Elastic Container Registry (Amazon ECR) repository. Use an Amazon Elastic Container Service (Amazon ECS) cluster with the Amazon EC2 launch type to run the containers. Use target tracking to scale automatically based on demand.

C. Store container images in a repository that runs on an Amazon EC2 instance. Run the containers on EC2 instances that are spread across multiple Availability Zones. Monitor the average CPU utilization in Amazon CloudWatch. Launch new EC2 instances as needed.

D. Create an Amazon EC2 Amazon Machine Image (AMI) that contains the container image. Launch EC2 instances in an Auto Scaling group across multiple Availability Zones. Use an Amazon CloudWatch alarm to scale out EC2 instances when the average CPU utilization threshold is breached.

Correct Answer: *C*

Community vote distribution

A (100%)

goatbernard Highly Voted 1 week, 6 days ago

Selected Answer: A

AWS Fargate

upvoted 6 times

Phinx Most Recent 1 day, 6 hours ago

Selected Answer: A

Fargate is the only serverless option.

upvoted 1 times

Wpcorgan 5 days, 21 hours ago

A is correct

upvoted 1 times

ds0321 1 week, 3 days ago

Selected Answer: A

AWS Fargate

upvoted 1 times

BENICE 1 week, 4 days ago

I think A is the correct option. AWS Farget

upvoted 1 times

mricee9 1 week, 5 days ago

Selected Answer: A

A seems right

upvoted 2 times

Question #164 *Topic 1*

A company has two applications: a sender application that sends messages with payloads to be processed and a processing application intended to receive the messages with payloads. The company wants to implement an AWS service to handle messages between the two applications. The sender application can send about 1,000 messages each hour. The messages may take up to 2 days to be processed: If the messages fail to process, they must be retained so that they do not impact the processing of any remaining messages.

Which solution meets these requirements and is the MOST operationally efficient?

A. Set up an Amazon EC2 instance running a Redis database. Configure both applications to use the instance. Store, process, and delete the messages, respectively.

B. Use an Amazon Kinesis data stream to receive the messages from the sender application. Integrate the processing application with the Kinesis Client Library (KCL).

C. Integrate the sender and processor applications with an Amazon Simple Queue Service (Amazon SQS) queue. Configure a dead-letter queue to collect the messages that failed to process.

D. Subscribe the processing application to an Amazon Simple Notification Service (Amazon SNS) topic to receive notifications to process. Integrate the sender application to write to the SNS topic.

Correct Answer: *C*

Community vote distribution

B (50%) C (50%)

TelaO 1 day, 20 hours ago

Selected Answer: B

Please explain by "B" is incorrect? How does SQS process data?

"KCL helps you consume and process data from a Kinesis data stream by taking care of many of the complex tasks associated with distributed computing."

https://docs.aws.amazon.com/streams/latest/dev/shared-throughput-kcl-consumers.html

upvoted 1 times

Wpcorgan 5 days, 21 hours ago

C is correct

upvoted 1 times

mabotega 1 week, 4 days ago

Selected Answer: C

https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-dead-letter-queues.html

upvoted 1 times

BENICE 1 week, 4 days ago

Option: C
"Amazon FSx for Lustre" ---> Dead Letter Queue

upvoted 1 times

Nigma 1 week, 5 days ago

Ans: C
https://aws.amazon.com/blogs/compute/building-loosely-coupled-scalable-c-applications-with-amazon-sqs-and-amazon-sns/
https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-dead-letter-queues.html

upvoted 3 times

Question #165 *Topic 1*

A solutions architect must design a solution that uses Amazon CloudFront with an Amazon S3 origin to store a static website. The company's security policy requires that all website traffic be inspected by AWS WAF.

How should the solutions architect comply with these requirements?

A. Configure an S3 bucket policy to accept requests coming from the AWS WAF Amazon Resource Name (ARN) only.

B. Configure Amazon CloudFront to forward all incoming requests to AWS WAF before requesting content from the S3 origin.

C. Configure a security group that allows Amazon CloudFront IP addresses to access Amazon S3 only. Associate AWS WAF to CloudFront.

D. Configure Amazon CloudFront and Amazon S3 to use an origin access identity (OAI) to restrict access to the S3 bucket. Enable AWS WAF on the distribution.

Correct Answer: *D*

Community vote distribution

D (75%) B (25%)

Nigma Highly Voted 1 week, 5 days ago

Answer D. Use an OAI to lockdown CloudFront to S3 origin & enable WAF on CF distribution

upvoted 5 times

Wpcorgan Most Recent 5 days, 21 hours ago

D is correct

upvoted 1 times

TonyghostR05 6 days, 11 hours ago

D using OAI

upvoted 2 times

handyplazt 1 week ago

Selected Answer: D

Answer D

not B because it is not supported to forward S3 requests to WAF

upvoted 2 times

EKA_CloudGod 1 week ago

Selected Answer: B

https://aws.amazon.com/premiumsupport/knowledge-center/cloudfront-serve-static-website/

upvoted 1 times

sdasdawa 1 week, 5 days ago

Selected Answer: D

https://www.examtopics.com/discussions/amazon/view/35639-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Question #166 *Topic 1*

Organizers for a global event want to put daily reports online as static HTML pages. The pages are expected to generate millions of views from users around the world. The files are stored in an Amazon S3 bucket. A solutions architect has been asked to design an efficient and effective solution.

Which action should the solutions architect take to accomplish this?

A. Generate presigned URLs for the files.

B. Use cross-Region replication to all Regions.

C. Use the geoproximity feature of Amazon Route 53.

D. Use Amazon CloudFront with the S3 bucket as its origin.

Correct Answer: *D*

Community vote distribution

D (100%)

Wpcorgan 5 days, 21 hours ago

D is correct

upvoted 1 times

Nigma 1 week, 5 days ago

D

Static content on S3 and hence Cloudfront is the best way

upvoted 2 times

Pamban 1 week, 5 days ago

Selected Answer: D

D is the correct answer

upvoted 2 times

Question #167 *Topic 1*

A company runs a production application on a fleet of Amazon EC2 instances. The application reads the data from an Amazon SQS queue and processes the messages in parallel. The message volume is unpredictable and often has intermittent traffic. This application should continually process messages without any downtime.

Which solution meets these requirements MOST cost-effectively?

A. Use Spot Instances exclusively to handle the maximum capacity required.

B. Use Reserved Instances exclusively to handle the maximum capacity required.

C. Use Reserved Instances for the baseline capacity and use Spot Instances to handle additional capacity.

D. Use Reserved Instances for the baseline capacity and use On-Demand Instances to handle additional capacity.

Correct Answer: *C*

Community vote distribution

D (75%) C (25%)

romko 5 days, 5 hours ago

Selected Answer: C

A - it's out because it's not ok to use full spot coverage.
B - it's hard to predict how much resources are needed to buy ahead, so it's suitable to no have any down time but not the best from cost perspective
C - possible to be correct answer such as use cover baseline with RI and rest with spot that is cheaper. Regarding don't time. there is no downtime because vaseline covered with RI and all communication is via SQS (distributed model)
D - possible but less cost effective then C

upvoted 1 times

Wpcorgan 5 days, 21 hours ago

D is correct

upvoted 1 times

ds0321 1 week, 3 days ago

Selected Answer: D

D is the correct answer

upvoted 1 times

taer 1 week, 4 days ago

Selected Answer: D

D is the correct answer

upvoted 4 times

Danny23132412141_2312 1 week, 4 days ago

D is the correct answer

upvoted 2 times

mabotega 1 week, 4 days ago

Selected Answer: D

Answer D: unpredictable and often has intermittent traffic

upvoted 3 times

babaxoxo 1 week, 5 days ago

Selected Answer: D

No downtime required so the answer is D

upvoted 4 times

mricee9 1 week, 5 days ago

Selected Answer: C

C
https://www.examtopics.com/discussions/amazon/view/35772-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Nigma 1 week, 5 days ago

Answer is D because you CAN NOT use spot instances for unpredictability

upvoted 3 times

goatbernard 1 week, 5 days ago

Selected Answer: C

reserved instance and spot instance

upvoted 2 times

Question #168 *Topic 1*

A security team wants to limit access to specific services or actions in all of the team's AWS accounts. All accounts belong to a large organization in AWS Organizations. The solution must be scalable and there must be a single point where permissions can be maintained.

What should a solutions architect do to accomplish this?

A. Create an ACL to provide access to the services or actions.

B. Create a security group to allow accounts and attach it to user groups.

C. Create cross-account roles in each account to deny access to the services or actions.

D. Create a service control policy in the root organizational unit to deny access to the services or actions.

Correct Answer: *D*

Community vote distribution

D (100%)

Wpcorgan 5 days, 21 hours ago

D iscorrect

upvoted 1 times

babaxoxo 1 week, 5 days ago

an organization and requires single point place to manage permissions

upvoted 2 times

Nigma 1 week, 5 days ago

D. Service control policies (SCPs) are one type of policy that you can use to manage your organization. SCPs offer central control over the maximum available permissions for all accounts in your organization, allowing you to ensure your accounts stay within your organization's access control guidelines. See https://docs.aws.amazon.com/organizations/latest/userguide/orgs_manage_policies_scp.html.

upvoted 3 times

goatbernard 1 week, 5 days ago

Selected Answer: D

SCP for organization

upvoted 2 times

Question #169 *Topic 1*

A company is concerned about the security of its public web application due to recent web attacks. The application uses an Application Load Balancer (ALB). A solutions architect must reduce the risk of DDoS attacks against the application.

What should the solutions architect do to meet this requirement?

A. Add an Amazon Inspector agent to the ALB.

B. Configure Amazon Macie to prevent attacks.

C. Enable AWS Shield Advanced to prevent attacks.

D. Configure Amazon GuardDuty to monitor the ALB.

Correct Answer: *C*

Community vote distribution

C (100%)

Wpcorgan 5 days, 21 hours ago

C is correct

upvoted 1 times

goatbernard 1 week, 5 days ago

Selected Answer: C

AWS Shield Advanced

upvoted 2 times

Nigma 1 week, 5 days ago

DDOS = AWS Shield

upvoted 2 times

Question #170 *Topic 1*

A company's web application is running on Amazon EC2 instances behind an Application Load Balancer. The company recently changed its policy, which now requires the application to be accessed from one specific country only.

Which configuration will meet this requirement?

A. Configure the security group for the EC2 instances.

B. Configure the security group on the Application Load Balancer.

C. Configure AWS WAF on the Application Load Balancer in a VPC.

D. Configure the network ACL for the subnet that contains the EC2 instances.

Correct Answer: *C*

Community vote distribution

C (100%)

Wpcorgan 5 days, 20 hours ago

C is correct

upvoted 1 times

handyplazt 1 week ago

Selected Answer: C

Geographic (Geo) Match Conditions in AWS WAF. This new condition type allows you to use AWS WAF to restrict application access based on the geographic location of your viewers. With geo match conditions you can choose the countries from which AWS WAF should allow access.
https://aws.amazon.com/about-aws/whats-new/2017/10/aws-waf-now-supports-geographic-match/

upvoted 1 times

mricee9 1 week, 5 days ago

Selected Answer: C

C
https://aws.amazon.com/about-aws/whats-new/2017/10/aws-waf-now-supports-geographic-match/

upvoted 2 times

Nigma 1 week, 5 days ago

C. WAF with ALB is the right option

upvoted 1 times

Question #171 *Topic 1*

A company provides an API to its users that automates inquiries for tax computations based on item prices. The company experiences a larger number of inquiries during the holiday season only that cause slower response times. A solutions architect needs to design a solution that is scalable and elastic.

What should the solutions architect do to accomplish this?

A. Provide an API hosted on an Amazon EC2 instance. The EC2 instance performs the required computations when the API request is made.

B. Design a REST API using Amazon API Gateway that accepts the item names. API Gateway passes item names to AWS Lambda for tax computations.

C. Create an Application Load Balancer that has two Amazon EC2 instances behind it. The EC2 instances will compute the tax on the received item names.

D. Design a REST API using Amazon API Gateway that connects with an API hosted on an Amazon EC2 instance. API Gateway accepts and passes the item names to the EC2 instance for tax computations.

Correct Answer: *D*

Community vote distribution

B (100%)

VISHNUKANDH 2 days, 23 hours ago

B is the right answer

upvoted 1 times

Wpcorgan 5 days, 20 hours ago

B is correct

upvoted 1 times

BENICE 1 week, 4 days ago

Seems like B is the correct option

upvoted 3 times

goatbernard 1 week, 5 days ago

Selected Answer: B

Lambda

upvoted 1 times

sdasdawa 1 week, 5 days ago

Selected Answer: B

https://www.examtopics.com/discussions/amazon/view/35849-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Ohnet 1 week, 5 days ago

It should be B,Lambda server-less is scalable and elastic than EC2 api gateway solution

upvoted 3 times

Nigma 1 week, 5 days ago

B. Lambda serverless is scalable and elastic than EC2 api gateway solution

upvoted 3 times

Question #172 *Topic 1*

A solutions architect is creating a new Amazon CloudFront distribution for an application. Some of the information submitted by users is sensitive. The application uses HTTPS but needs another layer of security. The sensitive information should.be protected throughout the entire application stack, and access to the information should be restricted to certain applications.

Which action should the solutions architect take?

A. Configure a CloudFront signed URL.

B. Configure a CloudFront signed cookie.

C. Configure a CloudFront field-level encryption profile.

D. Configure CloudFront and set the Origin Protocol Policy setting to HTTPS Only for the Viewer Protocol Policy.

Correct Answer: *A*

Community vote distribution

C (100%)

Wpcorgan 5 days, 20 hours ago

C is correct

upvoted 1 times

Bobbybash 6 days, 6 hours ago

CCCCCCCCC

Field-level encryption allows you to enable your users to securely upload sensitive information to your web servers. The sensitive information provided by your users is encrypted at the edge, close to the user, and remains encrypted throughout your entire application stack. This encryption ensures that only applications that need the data—and have the credentials to decrypt it—are able to do so.

upvoted 1 times

sdasdawa 1 week, 5 days ago

Selected Answer: C

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/field-level-encryption.html

upvoted 1 times

Ohnet 1 week, 5 days ago

It should be C.https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/field-level-encryption.html

upvoted 2 times

Nigma 1 week, 5 days ago

C

https://docs.aws.amazon.com/AmazonCloudFront/latest/DeveloperGuide/field-level-encryption.html

upvoted 2 times

goatbernard 1 week, 5 days ago

Selected Answer: C

CloudFront signed URLs provide a mechanism to control access to the content served through a distribution. Unlike the Origin Access Identity, it restricts access to which users can see the content. When you create a distribution, by default, it is open to everybody who knows the URL. But sometimes you want to limit that.

upvoted 1 times

Question #173 *Topic 1*

A gaming company hosts a browser-based application on AWS. The users of the application consume a large number of videos and images that are stored in Amazon S3. This content is the same for all users.

The application has increased in popularity, and millions of users worldwide accessing these media files. The company wants to provide the files to the users while reducing the load on the origin.

Which solution meets these requirements MOST cost-effectively?

A. Deploy an AWS Global Accelerator accelerator in front of the web servers.

B. Deploy an Amazon CloudFront web distribution in front of the S3 bucket.

C. Deploy an Amazon ElastiCache for Redis instance in front of the web servers.

D. Deploy an Amazon ElastiCache for Memcached instance in front of the web servers.

Correct Answer: *B*

Community vote distribution

B (100%)

rewdboy 2 days, 9 hours ago

Selected Answer: B

B is the correct answer

upvoted 1 times

Wpcorgan 5 days, 20 hours ago

B is correct

upvoted 1 times

Nigma 1 week, 5 days ago

B. Cloud front is best for content delivery. Global Accelerator is best for non-HTTP (TCP/UDP) cases and supports HTTP cases as well but with static IP (elastic IP) or anycast IP address only.

upvoted 4 times

Question #174 *Topic 1*

A company has a multi-tier application that runs six front-end web servers in an Amazon EC2 Auto Scaling group in a single Availability Zone behind an Application Load Balancer (ALB). A solutions architect needs to modify the infrastructure to be highly available without modifying the application.

Which architecture should the solutions architect choose that provides high availability?

A. Create an Auto Scaling group that uses three instances across each of two Regions.

B. Modify the Auto Scaling group to use three instances across each of two Availability Zones.

C. Create an Auto Scaling template that can be used to quickly create more instances in another Region.

D. Change the ALB in front of the Amazon EC2 instances in a round-robin configuration to balance traffic to the web tier.

Correct Answer: *B*

Community vote distribution

B (100%)

mricee9 2 days, 20 hours ago

Selected Answer: B

B is rightt

upvoted 1 times

Wpcorgan 5 days, 20 hours ago

B is correct

upvoted 1 times

xua81376 1 week, 2 days ago

B auto scaling i multiple AZ

upvoted 1 times

Nigma 1 week, 5 days ago

B. auto scaling groups can not span multi region

upvoted 4 times

Question #175 *Topic 1*

An ecommerce company has an order-processing application that uses Amazon API Gateway and an AWS Lambda function. The application stores data in an Amazon Aurora PostgreSQL database. During a recent sales event, a sudden surge in customer orders occurred. Some customers experienced timeouts, and the application did not process the orders of those customers.

A solutions architect determined that the CPU utilization and memory utilization were high on the database because of a large number of open connections. The solutions architect needs to prevent the timeout errors while making the least possible changes to the application.

Which solution will meet these requirements?

A. Configure provisioned concurrency for the Lambda function. Modify the database to be a global database in multiple AWS Regions.

B. Use Amazon RDS Proxy to create a proxy for the database. Modify the Lambda function to use the RDS Proxy endpoint instead of the database endpoint.

C. Create a read replica for the database in a different AWS Region. Use query string parameters in API Gateway to route traffic to the read replica.

D. Migrate the data from Aurora PostgreSQL to Amazon DynamoDB by using AWS Database Migration Service (AWS DMS). Modify the Lambda function to use the DynamoDB table.

Correct Answer: *B*

Community vote distribution

B (100%)

babaxoxo Highly Voted 1 week, 5 days ago

Selected Answer: B

Issue related to opening many connections and the solution requires least code changes so B satisfies the conditions

upvoted 5 times

Wpcorgan Most Recent 5 days, 20 hours ago

B is correct

upvoted 1 times

handyplazt 1 week ago

Selected Answer: B

Many applications, including those built on modern serverless architectures, can have a large number of open connections to the database server and may open and close database connections at a high rate, exhausting database memory and compute resources. Amazon RDS Proxy allows applications to pool and share connections established with the database, improving database efficiency and application scalability.
https://aws.amazon.com/id/rds/proxy/

upvoted 3 times

xua81376 1 week, 2 days ago

B - Proxy to manage connections

upvoted 2 times

Nigma 1 week, 5 days ago

Correct B

upvoted 1 times

Question #176 *Topic 1*

An application runs on Amazon EC2 instances in private subnets. The application needs to access an Amazon DynamoDB table.

What is the MOST secure way to access the table while ensuring that the traffic does not leave the AWS network?

A. Use a VPC endpoint for DynamoDB.

B. Use a NAT gateway in a public subnet.

C. Use a NAT instance in a private subnet.

D. Use the internet gateway attached to the VPC.

Correct Answer: *D*

Community vote distribution

A (100%)

Wpcorgan 5 days, 20 hours ago

A is correct

upvoted 1 times

xua81376 1 week, 2 days ago

Sure A

upvoted 1 times

ds0321 1 week, 4 days ago

Selected Answer: A

A - VPC endpoint

upvoted 2 times

goatbernard 1 week, 5 days ago

Selected Answer: A

A - VPC endpoint

upvoted 3 times

mabotega 1 week, 5 days ago

Selected Answer: A

VPC endpoints for service in private subnets
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html

upvoted 2 times

sdasdawa 1 week, 5 days ago

Selected Answer: A

https://www.examtopics.com/discussions/amazon/view/27700-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Nigma 1 week, 5 days ago

A for sure. https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/vpc-endpoints-dynamodb.html

upvoted 3 times

Ohnet 1 week, 5 days ago

Its A.

upvoted 1 times

Question #177 *Topic 1*

An entertainment company is using Amazon DynamoDB to store media metadata. The application is read intensive and experiencing delays. The company does not have staff to handle additional operational overhead and needs to improve the performance efficiency of DynamoDB without reconfiguring the application.

What should a solutions architect recommend to meet this requirement?

A. Use Amazon ElastiCache for Redis.

B. Use Amazon DynamoDB Accelerator (DAX).

C. Replicate data by using DynamoDB global tables.

D. Use Amazon ElastiCache for Memcached with Auto Discovery enabled.

Correct Answer: *B*

Community vote distribution

B (100%)

Wpcorgan 5 days, 20 hours ago

B is correct

upvoted 1 times

goatbernard 1 week, 5 days ago

Selected Answer: B

DAX is the cache for this

upvoted 1 times

nhlegend 1 week, 5 days ago

B is correct, DAX provides caching + no changes

upvoted 2 times

Question #178 *Topic 1*

A company's infrastructure consists of Amazon EC2 instances and an Amazon RDS DB instance in a single AWS Region. The company wants to back up its data in a separate Region.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use AWS Backup to copy EC2 backups and RDS backups to the separate Region.

B. Use Amazon Data Lifecycle Manager (Amazon DLM) to copy EC2 backups and RDS backups to the separate Region.

C. Create Amazon Machine Images (AMIs) of the EC2 instances. Copy the AMIs to the separate Region. Create a read replica for the RDS DB instance in the separate Region.

D. Create Amazon Elastic Block Store (Amazon EBS) snapshots. Copy the EBS snapshots to the separate Region. Create RDS snapshots. Export the RDS snapshots to Amazon S3. Configure S3 Cross-Region Replication (CRR) to the separate Region.

Correct Answer: *A*

Community vote distribution

A (100%)

babaxoxo 1 week, 5 days ago

Selected Answer: A

Ans A with least operational overhead

upvoted 1 times

rjam 1 week, 5 days ago

AWS Backup supports Supports cross-region backups

upvoted 2 times

rjam 1 week, 5 days ago

Selected Answer: A

Option A

Aws back up supports , EC2, RDS

upvoted 2 times

rjam 1 week, 5 days ago

AWS Backup suports Supports cross-region backups

upvoted 1 times

Question #179 *Topic 1*

A solutions architect needs to securely store a database user name and password that an application uses to access an Amazon RDS DB instance. The application that accesses the database runs on an Amazon EC2 instance. The solutions architect wants to create a secure parameter in AWS Systems Manager Parameter Store.

What should the solutions architect do to meet this requirement?

A. Create an IAM role that has read access to the Parameter Store parameter. Allow Decrypt access to an AWS Key Management Service (AWS KMS) key that is used to encrypt the parameter. Assign this IAM role to the EC2 instance.

B. Create an IAM policy that allows read access to the Parameter Store parameter. Allow Decrypt access to an AWS Key Management Service (AWS KMS) key that is used to encrypt the parameter. Assign this IAM policy to the EC2 instance.

C. Create an IAM trust relationship between the Parameter Store parameter and the EC2 instance. Specify Amazon RDS as a principal in the trust policy.

D. Create an IAM trust relationship between the DB instance and the EC2 instance. Specify Systems Manager as a principal in the trust policy.

Correct Answer: *A*

Community vote distribution

A (86%) 14%

EKA_CloudGod 6 days, 10 hours ago

Selected Answer: A

A. Attach IAM role to EC2 Instance
https://aws.amazon.com/blogs/security/digital-signing-asymmetric-keys-aws-kms/

upvoted 1 times

sdasdawa 1 week, 4 days ago

Selected Answer: A

Agree with A, IAM role is for services (EC2 for example)
IAM policy is more for users and groups

upvoted 2 times

babaxoxo 1 week, 5 days ago

Selected Answer: A

Attach IAM role to EC2 Instance profile

upvoted 3 times

goatbernard 1 week, 5 days ago

Selected Answer: B

IAM policy

upvoted 1 times

Question #180 *Topic 1*

A company is designing a cloud communications platform that is driven by APIs. The application is hosted on Amazon EC2 instances behind a Network Load Balancer (NLB). The company uses Amazon API Gateway to provide external users with access to the application through APIs. The company wants to protect the platform against web exploits like SQL injection and also wants to detect and mitigate large, sophisticated DDoS attacks.

Which combination of solutions provides the MOST protection? (Choose two.)

A. Use AWS WAF to protect the NLB.

B. Use AWS Shield Advanced with the NLB.

C. Use AWS WAF to protect Amazon API Gateway.

D. Use Amazon GuardDuty with AWS Shield Standard

E. Use AWS Shield Standard with Amazon API Gateway.

Correct Answer: *BC*

Community vote distribution

BC (100%)

babaxoxo Highly Voted 1 week, 5 days ago

Selected Answer: BC

Shield - Load Balancer, CF, Route53
AWF - CF, ALB, API Gateway

upvoted 6 times

Wpcorgan Most Recent 5 days, 20 hours ago

B and C

upvoted 1 times

BENICE 1 week, 3 days ago

B and C
"AWS Shield Advanced" for "sophisticated DDoS attacks"
"AWS WAF" for "NLB

upvoted 2 times

Nigma 1 week, 5 days ago

B and C

upvoted 1 times

rjam 1 week, 5 days ago

Selected Answer: BC

AWS Shield Advanced - DDos attacks
AWS WAF to protect Amazon API Gateway, because WAF sits before the API Gateway and then comes NLB.

upvoted 2 times

Question #181 *Topic 1*

A company has a legacy data processing application that runs on Amazon EC2 instances. Data is processed sequentially, but the order of results does not matter. The application uses a monolithic architecture. The only way that the company can scale the application to meet increased demand is to increase the size of the instances.

The company's developers have decided to rewrite the application to use a microservices architecture on Amazon Elastic Container Service (Amazon ECS).

What should a solutions architect recommend for communication between the microservices?

A. Create an Amazon Simple Queue Service (Amazon SQS) queue. Add code to the data producers, and send data to the queue. Add code to the data consumers to process data from the queue.

B. Create an Amazon Simple Notification Service (Amazon SNS) topic. Add code to the data producers, and publish notifications to the topic. Add code to the data consumers to subscribe to the topic.

C. Create an AWS Lambda function to pass messages. Add code to the data producers to call the Lambda function with a data object. Add code to the data consumers to receive a data object that is passed from the Lambda function.

D. Create an Amazon DynamoDB table. Enable DynamoDB Streams. Add code to the data producers to insert data into the table. Add code to the data consumers to use the DynamoDB Streams API to detect new table entries and retrieve the data.

Correct Answer: *A*

Community vote distribution

A (100%)

xua81376 1 week, 2 days ago

We need decoupling so ok to use SQS

upvoted 1 times

BENICE 1 week, 3 days ago

Can someone explain it bit more? Not able to understand it.

upvoted 1 times

EKA_CloudGod 6 days, 10 hours ago

As monolithic systems become too large to deal with, many enterprises are drawn to breaking them down into the microservices architectural style by means of decoupling. Amazon Simple Queue Service (Amazon SQS) is a fully managed message queuing service that makes it easy to decouple and scale microservices, distributed systems, and serverless applications

upvoted 2 times

taer 1 week, 4 days ago

Selected Answer: A

Answer is A

upvoted 1 times

Nigma 1 week, 5 days ago

SQS to decouple.

upvoted 1 times

Question #182 *Topic 1*

A company wants to migrate its MySQL database from on premises to AWS. The company recently experienced a database outage that significantly impacted the business. To ensure this does not happen again, the company wants a reliable database solution on AWS that minimizes data loss and stores every transaction on at least two nodes.

Which solution meets these requirements?

A. Create an Amazon RDS DB instance with synchronous replication to three nodes in three Availability Zones.

B. Create an Amazon RDS MySQL DB instance with Multi-AZ functionality enabled to synchronously replicate the data.

C. Create an Amazon RDS MySQL DB instance and then create a read replica in a separate AWS Region that synchronously replicates the data.

D. Create an Amazon EC2 instance with a MySQL engine installed that triggers an AWS Lambda function to synchronously replicate the data to an Amazon RDS MySQL DB instance.

Correct Answer: *B*

Community vote distribution

B (100%)

Wpcorgan 5 days, 20 hours ago

B is correct

upvoted 1 times

EKA_CloudGod 6 days, 10 hours ago

Selected Answer: B

Option B is the correct answer:
https://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZSingleStandby.html

upvoted 1 times

Nigma 1 week, 5 days ago

B is the answer

upvoted 2 times

rjam 1 week, 5 days ago

Selected Answer: B

Amazon RDS MySQL DB instance with Multi-AZ functionality enabled to synchronously replicate the data
Standby DB in Multi-AZ- synchronous replication

Read Replica always asynchronous. so option C is ignored.

upvoted 4 times

Question #183 *Topic 1*

A company is building a new dynamic ordering website. The company wants to minimize server maintenance and patching. The website must be highly available and must scale read and write capacity as quickly as possible to meet changes in user demand.

Which solution will meet these requirements?

A. Host static content in Amazon S3. Host dynamic content by using Amazon API Gateway and AWS Lambda. Use Amazon DynamoDB with on-demand capacity for the database. Configure Amazon CloudFront to deliver the website content.

B. Host static content in Amazon S3. Host dynamic content by using Amazon API Gateway and AWS Lambda. Use Amazon Aurora with Aurora Auto Scaling for the database. Configure Amazon CloudFront to deliver the website content.

C. Host all the website content on Amazon EC2 instances. Create an Auto Scaling group to scale the EC2 instances. Use an Application Load Balancer to distribute traffic. Use Amazon DynamoDB with provisioned write capacity for the database.

D. Host all the website content on Amazon EC2 instances. Create an Auto Scaling group to scale the EC2 instances. Use an Application Load Balancer to distribute traffic. Use Amazon Aurora with Aurora Auto Scaling for the database.

Correct Answer: *A*

Community vote distribution

A (88%) 13%

romko 5 days, 1 hour ago

Selected Answer: A

A - is correct, because Dynamodb on-demand scales write and read capacity
B - Aurora auto scaling scales only read replicas

upvoted 2 times

mabotega 5 days, 3 hours ago

Selected Answer: A

On-demand mode is a good option if any of the following are true:

You create new tables with unknown workloads.

You have unpredictable application traffic.

You prefer the ease of paying for only what you use.
https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.ReadWriteCapacityMode.html

upvoted 1 times

Wpcorgan 5 days, 20 hours ago

A is correct

upvoted 1 times

Manlikeleke 1 week, 1 day ago

please is this dump enough to pass the exam?

upvoted 4 times

Bobbybash 6 days, 5 hours ago

I HOPE SO

upvoted 1 times

Az900500 1 week, 3 days ago

Selected Answer A

"Read write capacity = DynamoDb" Read Replica mostly Aurora .. @nhlegend yes DynampDB has 400KB maximum but in the answer neither Dynamo or Aurora was used as primary storage

upvoted 4 times

sdasdawa 1 week, 4 days ago

Selected Answer: A

Agree with A, DynamoDB is perfect for storing ordering data (key-values)

upvoted 4 times

Nigma 1 week, 5 days ago

A is the answer

upvoted 2 times

rjam 1 week, 5 days ago

Selected Answer: B

option B . Aurora is better than DynamoDB

upvoted 1 times

rjam 1 week, 5 days ago

amazon aurora - highly available, self-healing, auto-scaling

upvoted 1 times

nhlegend 1 week, 5 days ago

B is correct, DynampDB has 400KB maximum

upvoted 1 times

nhlegend 1 week, 5 days ago

typo, I mean A is correct

upvoted 3 times

Question #184 *Topic 1*

A company has an AWS account used for software engineering. The AWS account has access to the company's on-premises data center through a pair of AWS Direct Connect connections. All non-VPC traffic routes to the virtual private gateway.

A development team recently created an AWS Lambda function through the console. The development team needs to allow the function to access a database that runs in a private subnet in the company's data center.

Which solution will meet these requirements?

A. Configure the Lambda function to run in the VPC with the appropriate security group.

B. Set up a VPN connection from AWS to the data center. Route the traffic from the Lambda function through the VPN.

C. Update the route tables in the VPC to allow the Lambda function to access the on-premises data center through Direct Connect.

D. Create an Elastic IP address. Configure the Lambda function to send traffic through the Elastic IP address without an elastic network interface.

Correct Answer: *C*

Community vote distribution

A (60%) C (40%)

Newptone 3 days, 15 hours ago

Selected Answer: A

When you connect a function to a VPC, Lambda assigns your function to a Hyperplane ENI (elastic network interface) for each subnet in your function's VPC configuration. Lambda creates a Hyperplane ENI the first time a unique subnet and security group combination is defined for a VPC-enabled function in an account.

upvoted 1 times

romko 5 days ago

Selected Answer: A

lambda by default runs out of vpc, so without A lambda is out of vpc.

C is incorrect, because don't matter how you change route tables in VPC it doesn't make sense while lambda is out of vpc.

So the correct answer is A

upvoted 1 times

Wpcorgan 5 days, 20 hours ago

C is correct

upvoted 1 times

taer 1 week, 4 days ago

Selected Answer: C

Answer is C

upvoted 1 times

mricee9 1 week, 4 days ago

Selected Answer: C

C
https://www.examtopics.com/discussions/amazon/view/68069-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Ohnet 1 week, 5 days ago

Its A.Deploy the Lambda Function in the VPC with a security group.
https://docs.aws.amazon.com/lambda/latest/dg/configuration-vpc.html#vpc-managing-eni

upvoted 3 times

sdasdawa 1 week, 5 days ago

Selected Answer: A

1st section in
https://docs.aws.amazon.com/lambda/latest/dg/configuration-vpc.html#vpc-managing-eni

upvoted 1 times

Question #185 *Topic 1*

A company runs an application using Amazon ECS. The application creates resized versions of an original image and then makes Amazon S3 API calls to store the resized images in Amazon S3.

How can a solutions architect ensure that the application has permission to access Amazon S3?

A. Update the S3 role in AWS IAM to allow read/write access from Amazon ECS, and then relaunch the container.

B. Create an IAM role with S3 permissions, and then specify that role as the taskRoleArn in the task definition.

C. Create a security group that allows access from Amazon ECS to Amazon S3, and update the launch configuration used by the ECS cluster.

D. Create an IAM user with S3 permissions, and then relaunch the Amazon EC2 instances for the ECS cluster while logged in as this account.

Correct Answer: *B*

Community vote distribution

B (100%)

Wpcorgan 5 days, 20 hours ago

B is correct

upvoted 1 times

taer 1 week, 4 days ago

Selected Answer: B

The answer is B.

upvoted 1 times

Nigma 1 week, 5 days ago

B is the answer

upvoted 2 times

Question #186 *Topic 1*

A company has a Windows-based application that must be migrated to AWS. The application requires the use of a shared Windows file system attached to multiple Amazon EC2 Windows instances that are deployed across multiple Availability Zone:

What should a solutions architect do to meet this requirement?

A. Configure AWS Storage Gateway in volume gateway mode. Mount the volume to each Windows instance.

B. Configure Amazon FSx for Windows File Server. Mount the Amazon FSx file system to each Windows instance.

C. Configure a file system by using Amazon Elastic File System (Amazon EFS). Mount the EFS file system to each Windows instance.

D. Configure an Amazon Elastic Block Store (Amazon EBS) volume with the required size. Attach each EC2 instance to the volume. Mount the file system within the volume to each Windows instance.

Correct Answer: *B*

Community vote distribution

B (100%)

Wpcorgan 5 days, 20 hours ago

B is correct

upvoted 1 times

xua81376 1 week, 1 day ago

B FSx for windows

upvoted 1 times

BENICE 1 week, 3 days ago

B is correct option

upvoted 1 times

rjam 1 week, 4 days ago

Selected Answer: B

Amazon FSx for Windows File Server

upvoted 2 times

Nigma 1 week, 5 days ago

Correct is B
FSx --> shared Windows file system (SMB)
EFS --> Linux NFS

upvoted 2 times

Question #187 *Topic 1*

A company is developing an ecommerce application that will consist of a load-balanced front end, a container-based application, and a relational database. A solutions architect needs to create a highly available solution that operates with as little manual intervention as possible.

Which solutions meet these requirements? (Choose two.)

A. Create an Amazon RDS DB instance in Multi-AZ mode.

B. Create an Amazon RDS DB instance and one or more replicas in another Availability Zone.

C. Create an Amazon EC2 instance-based Docker cluster to handle the dynamic application load.

D. Create an Amazon Elastic Container Service (Amazon ECS) cluster with a Fargate launch type to handle the dynamic application load.

E. Create an Amazon Elastic Container Service (Amazon ECS) cluster with an Amazon EC2 launch type to handle the dynamic application load.

Correct Answer: *AD*

Community vote distribution

AD (100%)

Gabs90 1 day, 20 hours ago

Selected Answer: AD

A and D

upvoted 1 times

Wpcorgan 3 days, 16 hours ago

A and D

upvoted 1 times

BENICE 1 week, 3 days ago

A and D are the options

upvoted 1 times

Danny23132412141_2312 1 week, 4 days ago

AD for sure

Link: https://www.examtopics.com/discussions/amazon/view/43729-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 2 times

Question #188 *Topic 1*

A company uses Amazon S3 as its data lake. The company has a new partner that must use SFTP to upload data files. A solutions architect needs to implement a highly available SFTP solution that minimizes operational overhead.

Which solution will meet these requirements?

A. Use AWS Transfer Family to configure an SFTP-enabled server with a publicly accessible endpoint. Choose the S3 data lake as the destination.

B. Use Amazon S3 File Gateway as an SFTP server. Expose the S3 File Gateway endpoint URL to the new partner. Share the S3 File Gateway endpoint with the new partner.

C. Launch an Amazon EC2 instance in a private subnet in a VPInstruct the new partner to upload files to the EC2 instance by using a VPN. Run a cron job script, on the EC2 instance to upload files to the S3 data lake.

D. Launch Amazon EC2 instances in a private subnet in a VPC. Place a Network Load Balancer (NLB) in front of the EC2 instances. Create an SFTP listener port for the NLB. Share the NLB hostname with the new partner. Run a cron job script on the EC2 instances to upload files to the S3 data lake.

Correct Answer: *D*

Community vote distribution

A (100%)

Bobbybash 6 days, 4 hours ago

Selected Answer: A

AAAAAAAA
AWS Transfer for SFTP, a fully-managed, highly-available SFTP service. You simply create a server, set up user accounts, and associate the server with one or more Amazon Simple Storage Service (Amazon S3) buckets

upvoted 1 times

Bobbybash 6 days, 4 hours ago

AAAAAAAA
AWS Transfer for SFTP, a fully-managed, highly-available SFTP service. You simply create a server, set up user accounts, and associate the server with one or more Amazon Simple Storage Service (Amazon S3) buckets.

upvoted 1 times

mabotega 1 week, 4 days ago

Selected Answer: A

A is the answer - https://docs.aws.amazon.com/transfer/latest/userguide/create-server-sftp.html

upvoted 1 times

Nigma 1 week, 5 days ago

A is the answer

upvoted 1 times

LeGloupier 1 week, 5 days ago

Selected Answer: A

answer is A

upvoted 1 times

mabotega 1 week, 5 days ago

Selected Answer: A

https://www.examtopics.com/discussions/amazon/view/83197-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

Question #189 *Topic 1*

A company needs to store contract documents. A contract lasts for 5 years. During the 5-year period, the company must ensure that the documents cannot be overwritten or deleted. The company needs to encrypt the documents at rest and rotate the encryption keys automatically every year.

Which combination of steps should a solutions architect take to meet these requirements with the LEAST operational overhead? (Choose two.)

A. Store the documents in Amazon S3. Use S3 Object Lock in governance mode.

B. Store the documents in Amazon S3. Use S3 Object Lock in compliance mode.

C. Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Configure key rotation.

D. Use server-side encryption with AWS Key Management Service (AWS KMS) customer managed keys. Configure key rotation.

E. Use server-side encryption with AWS Key Management Service (AWS KMS) customer provided (imported) keys. Configure key rotation.

Correct Answer: *CE*

Community vote distribution

BC (53%) BD (47%)

Cizzla7049 1 day, 21 hours ago

Selected Answer: BC

I thought it was B and C but after some research, i believed B and D but after more research and recent changes by AWS (may 2022), it is B and C. It is D because the question says key must be rotated yearly. So you have to use CMK and configure key rotation. It automatically rotates yearly, while S3 managed keys rotate every 3 years.

https://stackoverflow.com/questions/63478626/which-aws-s3-encryption-technique-provides-rotation-policy-for-encryption-keys

After further research it is B and . AWS Managed keys changed from 3yrs to 1yr rotation.
https://docs.aws.amazon.com/kms/latest/developerguide/concepts.html#kms_keys

upvoted 1 times

peneloco 2 days, 16 hours ago

Selected Answer: BC

It's B and C. C because the question does say LEAST operational overhead, customer managed anything equals more overhead.

upvoted 2 times

Gil80 3 days, 21 hours ago

why not B? SSE-S3:
• Encryption using keys handled & managed by Amazon S3
• Object is encrypted server side
• AES-256 encryption type
• Must set header: "x-amz-server-side-encryption":"AES256"
Can someone please explain?

upvoted 1 times

BENICE 1 week, 3 days ago

B and D

upvoted 3 times

ds0321 1 week, 4 days ago

Selected Answer: BD

B and D

upvoted 3 times

Nigma 1 week, 5 days ago

B and D

upvoted 1 times

LeGloupier 1 week, 5 days ago

Selected Answer: BD

should be BD
C could have been fine, but key rotation is activate per default on SSE-S3, and no way to deactivate it if I am not wrong

upvoted 3 times

taer 1 week, 5 days ago

Selected Answer: BD

BD
https://www.examtopics.com/discussions/amazon/view/81524-exam-aws-certified-solutions-architect-associate-saa-c02/
upvoted 2 times

yuantongxue 1 week, 5 days ago

Selected Answer: BC

Lock + Key
upvoted 3 times

sdasdawa 1 week, 5 days ago

Selected Answer: BC

B due to compliance mode no user (including root) can delete files
C due to SSE-S3 automatic key rotation
upvoted 3 times

Question #190 *Topic 1*

A company has a web application that is based on Java and PHP. The company plans to move the application from on premises to AWS. The company needs the ability to test new site features frequently. The company also needs a highly available and managed solution that requires minimum operational overhead.

Which solution will meet these requirements?

A. Create an Amazon S3 bucket. Enable static web hosting on the S3 bucket. Upload the static content to the S3 bucket. Use AWS Lambda to process all dynamic content.

B. Deploy the web application to an AWS Elastic Beanstalk environment. Use URL swapping to switch between multiple Elastic Beanstalk environments for feature testing.

C. Deploy the web application to Amazon EC2 instances that are configured with Java and PHP. Use Auto Scaling groups and an Application Load Balancer to manage the website's availability.

D. Containerize the web application. Deploy the web application to Amazon EC2 instances. Use the AWS Load Balancer Controller to dynamically route traffic between containers that contain the new site features for testing.

Correct Answer: *D*

Community vote distribution

B (75%) D (25%)

romko 5 days ago

Selected Answer: D

wow, so many votes for B.

B will be correct if application requires one of runtime java or php, elastic Beanstallk allows to specify only one runtime. In requirement is "web application that is based on Java and PHP"
so B is out.

D allows to setup own container and there you may install as many as system needs

upvoted 1 times

Cizzla7049 1 day, 21 hours ago

D can also be done by Elastic Beanstalk. Answer is B, as it using beanstalk removes the overhead

AWS Elastic Beanstalk is the fastest way to get web applications up and running on AWS. You can simply upload your application code, and the service automatically handles details such as resource provisioning, load balancing, auto scaling, and monitoring. Elastic Beanstalk is ideal if you have a PHP, Java, Python, Ruby, Node.js, .NET, Go, or Docker web application. Elastic Beanstalk uses core AWS services such as Amazon Elastic Compute Cloud (EC2), Amazon Elastic Container Service (ECS), AWS Auto Scaling, and Elastic Load Balancing (ELB) to easily support applications that need to scale to serve millions of users.

upvoted 1 times

Wpcorgan 5 days, 20 hours ago

B is correct

upvoted 1 times

rjam 1 week, 4 days ago

Selected Answer: B

Swapping URL : ElasticBeanStalk
https://docs.aws.amazon.com/elasticbeanstalk/latest/dg/using-features.CNAMESwap.html

upvoted 1 times

Nigma 1 week, 4 days ago

B is the answer

upvoted 1 times

LeGloupier 1 week, 5 days ago

isn't it B ?

upvoted 1 times

taer 1 week, 5 days ago

Selected Answer: B

B
https://www.examtopics.com/discussions/amazon/view/81534-exam-aws-certified-solutions-architect-associate-saa-c02/

upvoted 1 times

sdasdawa 1 week, 5 days ago

Selected Answer: B

Elastic Beanstalk seems to be perfect for this

upvoted 1 times

Question #191 *Topic 1*

A company has an ordering application that stores customer information in Amazon RDS for MySQL. During regular business hours, employees run one-time queries for reporting purposes. Timeouts are occurring during order processing because the reporting queries are taking a long time to run. The company needs to eliminate the timeouts without preventing employees from performing queries.

What should a solutions architect do to meet these requirements?

A. Create a read replica. Move reporting queries to the read replica.

B. Create a read replica. Distribute the ordering application to the primary DB instance and the read replica.

C. Migrate the ordering application to Amazon DynamoDB with on-demand capacity.

D. Schedule the reporting queries for non-peak hours.

Correct Answer: *B*

Question #192 *Topic 1*

A hospital wants to create digital copies for its large collection of historical written records. The hospital will continue to add hundreds of new documents each day. The hospital's data team will scan the documents and will upload the documents to the AWS Cloud.

A solutions architect must implement a solution to analyze the documents, extract the medical information, and store the documents so that an application can run SQL queries on the data. The solution must maximize scalability and operational efficiency.

Which combination of steps should the solutions architect take to meet these requirements? (Choose two.)

A. Write the document information to an Amazon EC2 instance that runs a MySQL database.

B. Write the document information to an Amazon S3 bucket. Use Amazon Athena to query the data.

C. Create an Auto Scaling group of Amazon EC2 instances to run a custom application that processes the scanned files and extracts the medical information.

D. Create an AWS Lambda function that runs when new documents are uploaded. Use Amazon Rekognition to convert the documents to raw text. Use Amazon Transcribe Medical to detect and extract relevant medical information from the text.

E. Create an AWS Lambda function that runs when new documents are uploaded. Use Amazon Textract to convert the documents to raw text. Use Amazon Comprehend Medical to detect and extract relevant medical information from the text.

Correct Answer: *CD*

Question #193 *Topic 1*

A company is running a batch application on Amazon EC2 instances. The application consists of a backend with multiple Amazon RDS databases. The application is causing a high number of reads on the databases. A solutions architect must reduce the number of database reads while ensuring high availability.

What should the solutions architect do to meet this requirement?

A. Add Amazon RDS read replicas.

B. Use Amazon ElastiCache for Redis.

C. Use Amazon Route 53 DNS caching

D. Use Amazon ElastiCache for Memcached.

Correct Answer: *A*

Question #194 *Topic 1*

A company needs to run a critical application on AWS. The company needs to use Amazon EC2 for the application's database. The database must be highly available and must fail over automatically if a disruptive event occurs.

Which solution will meet these requirements?

A. Launch two EC2 instances, each in a different Availability Zone in the same AWS Region. Install the database on both EC2 instances. Configure the EC2 instances as a cluster. Set up database replication.

B. Launch an EC2 instance in an Availability Zone. Install the database on the EC2 instance. Use an Amazon Machine Image (AMI) to back up the data. Use AWS CloudFormation to automate provisioning of the EC2 instance if a disruptive event occurs.

C. Launch two EC2 instances, each in a different AWS Region. Install the database on both EC2 instances. Set up database replication. Fail over the database to a second Region.

D. Launch an EC2 instance in an Availability Zone. Install the database on the EC2 instance. Use an Amazon Machine Image (AMI) to back up the data. Use EC2 automatic recovery to recover the instance if a disruptive event occurs.

Correct Answer: *C*

Question #195 *Topic 1*

A company's order system sends requests from clients to Amazon EC2 instances. The EC2 instances process the orders and then store the orders in a database on Amazon RDS. Users report that they must reprocess orders when the system fails. The company wants a resilient solution that can process orders automatically if a system outage occurs.

What should a solutions architect do to meet these requirements?

A. Move the EC2 instances into an Auto Scaling group. Create an Amazon EventBridge (Amazon CloudWatch Events) rule to target an Amazon Elastic Container Service (Amazon ECS) task.

B. Move the EC2 instances into an Auto Scaling group behind an Application Load Balancer (ALB). Update the order system to send messages to the ALB endpoint.

C. Move the EC2 instances into an Auto Scaling group. Configure the order system to send messages to an Amazon Simple Queue Service (Amazon SQS) queue. Configure the EC2 instances to consume messages from the queue.

D. Create an Amazon Simple Notification Service (Amazon SNS) topic. Create an AWS Lambda function, and subscribe the function to the SNS topic. Configure the order system to send messages to the SNS topic. Send a command to the EC2 instances to process the messages by using AWS Systems Manager Run Command.

Correct Answer: *D*

Question #196 *Topic 1*

A company runs an application on a large fleet of Amazon EC2 instances. The application reads and writes entries into an Amazon DynamoDB table. The size of the DynamoDB table continuously grows, but the application needs only data from the last 30 days. The company needs a solution that minimizes cost and development effort.

Which solution meets these requirements?

A. Use an AWS CloudFormation template to deploy the complete solution. Redeploy the CloudFormation stack every 30 days, and delete the original stack.

B. Use an EC2 instance that runs a monitoring application from AWS Marketplace. Configure the monitoring application to use Amazon DynamoDB Streams to store the timestamp when a new item is created in the table. Use a script that runs on the EC2 instance to delete items that have a timestamp that is older than 30 days.

C. Configure Amazon DynamoDB Streams to invoke an AWS Lambda function when a new item is created in the table. Configure the Lambda function to delete items in the table that are older than 30 days.

D. Extend the application to add an attribute that has a value of the current timestamp plus 30 days to each new item that is created in the table. Configure DynamoDB to use the attribute as the TTL attribute.

Correct Answer: *D*

Question #197 *Topic 1*

A company has a Microsoft .NET application that runs on an on-premises Windows Server. The application stores data by using an Oracle Database Standard Edition server. The company is planning a migration to AWS and wants to minimize development changes while moving the application. The AWS application environment should be highly available.

Which combination of actions should the company take to meet these requirements? (Choose two.)

A. Refactor the application as serverless with AWS Lambda functions running .NET Core.

B. Rehost the application in AWS Elastic Beanstalk with the .NET platform in a Multi-AZ deployment.

C. Replatform the application to run on Amazon EC2 with the Amazon Linux Amazon Machine Image (AMI).

D. Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Amazon DynamoDB in a Multi-AZ deployment.

E. Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Oracle on Amazon RDS in a Multi-AZ deployment.

Correct Answer: *BD*

Currently there are no comments in this discussion, be the first to comment!

Question #198 *Topic 1*

A company runs a containerized application on a Kubernetes cluster in an on-premises data center. The company is using a MongoDB database for data storage. The company wants to migrate some of these environments to AWS, but no code changes or deployment method changes are possible at this time. The company needs a solution that minimizes operational overhead.

Which solution meets these requirements?

A. Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 worker nodes for compute and MongoDB on EC2 for data storage.

B. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate for compute and Amazon DynamoDB for data storage

C. Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 worker nodes for compute and Amazon DynamoDB for data storage.

D. Use Amazon Elastic Kubernetes Service (Amazon EKS) with AWS Fargate for compute and Amazon DocumentDB (with MongoDB compatibility) for data storage.

Correct Answer: *D*

Question #199 *Topic 1*

A telemarketing company is designing its customer call center functionality on AWS. The company needs a solution that provides multiple speaker recognition and generates transcript files. The company wants to query the transcript files to analyze the business patterns. The transcript files must be stored for 7 years for auditing purposes.

Which solution will meet these requirements?

A. Use Amazon Rekognition for multiple speaker recognition. Store the transcript files in Amazon S3. Use machine learning models for transcript file analysis.

B. Use Amazon Transcribe for multiple speaker recognition. Use Amazon Athena for transcript file analysis.

C. Use Amazon Translate for multiple speaker recognition. Store the transcript files in Amazon Redshift. Use SQL queries for transcript file analysis.

D. Use Amazon Rekognition for multiple speaker recognition. Store the transcript files in Amazon S3. Use Amazon Textract for transcript file analysis.

Correct Answer: *C*

Question #200 *Topic 1*

A company hosts its application on AWS. The company uses Amazon Cognito to manage users. When users log in to the application, the application fetches required data from Amazon DynamoDB by using a REST API that is hosted in Amazon API Gateway. The company wants an AWS managed solution that will control access to the REST API to reduce development efforts.

Which solution will meet these requirements with the LEAST operational overhead?

A. Configure an AWS Lambda function to be an authorizer in API Gateway to validate which user made the request.

B. For each user, create and assign an API key that must be sent with each request. Validate the key by using an AWS Lambda function.

C. Send the user's email address in the header with every request. Invoke an AWS Lambda function to validate that the user with that email address has proper access.

D. Configure an Amazon Cognito user pool authorizer in API Gateway to allow Amazon Cognito to validate each request.

Correct Answer: *A*

Question #201 *Topic 1*

A company is developing a marketing communications service that targets mobile app users. The company needs to send confirmation messages with Short Message Service (SMS) to its users. The users must be able to reply to the SMS messages. The company must store the responses for a year for analysis.

What should a solutions architect do to meet these requirements?

A. Create an Amazon Connect contact flow to send the SMS messages. Use AWS Lambda to process the responses.

B. Build an Amazon Pinpoint journey. Configure Amazon Pinpoint to send events to an Amazon Kinesis data stream for analysis and archiving.

C. Use Amazon Simple Queue Service (Amazon SQS) to distribute the SMS messages. Use AWS Lambda to process the responses.

D. Create an Amazon Simple Notification Service (Amazon SNS) FIFO topic. Subscribe an Amazon Kinesis data stream to the SNS topic for analysis and archiving.

Correct Answer: *A*

Question #202 *Topic 1*

A company is planning to move its data to an Amazon S3 bucket. The data must be encrypted when it is stored in the S3 bucket. Additionally, the encryption key must be automatically rotated every year.

Which solution will meet these requirements with the LEAST operational overhead?

A. Move the data to the S3 bucket. Use server-side encryption with Amazon S3 managed encryption keys (SSE-S3). Use the built-in key rotation behavior of SSE-S3 encryption keys.

B. Create an AWS Key Management Service (AWS KMS) customer managed key. Enable automatic key rotation. Set the S3 bucket's default encryption behavior to use the customer managed KMS key. Move the data to the S3 bucket.

C. Create an AWS Key Management Service (AWS KMS) customer managed key. Set the S3 bucket's default encryption behavior to use the customer managed KMS key. Move the data to the S3 bucket. Manually rotate the KMS key every year.

D. Encrypt the data with customer key material before moving the data to the S3 bucket. Create an AWS Key Management Service (AWS KMS) key without key material. Import the customer key material into the KMS key. Enable automatic key rotation.

Correct Answer: *B*

Question #203 *Topic 1*

The customers of a finance company request appointments with financial advisors by sending text messages. A web application that runs on Amazon EC2 instances accepts the appointment requests. The text messages are published to an Amazon Simple Queue Service (Amazon SQS) queue through the web application. Another application that runs on EC2 instances then sends meeting invitations and meeting confirmation email messages to the customers. After successful scheduling, this application stores the meeting information in an Amazon DynamoDB database.

As the company expands, customers report that their meeting invitations are taking longer to arrive.

What should a solutions architect recommend to resolve this issue?

A. Add a DynamoDB Accelerator (DAX) cluster in front of the DynamoDB database.

B. Add an Amazon API Gateway API in front of the web application that accepts the appointment requests.

C. Add an Amazon CloudFront distribution. Set the origin as the web application that accepts the appointment requests.

D. Add an Auto Scaling group for the application that sends meeting invitations. Configure the Auto Scaling group to scale based on the depth of the SQS queue.

Correct Answer: *D*

Question #204 *Topic 1*

An online retail company has more than 50 million active customers and receives more than 25,000 orders each day. The company collects purchase data for customers and stores this data in Amazon S3. Additional customer data is stored in Amazon RDS.

The company wants to make all the data available to various teams so that the teams can perform analytics. The solution must provide the ability to manage fine-grained permissions for the data and must minimize operational overhead.

Which solution will meet these requirements?

A. Migrate the purchase data to write directly to Amazon RDS. Use RDS access controls to limit access.

B. Schedule an AWS Lambda function to periodically copy data from Amazon RDS to Amazon S3. Create an AWS Glue crawler. Use Amazon Athena to query the data. Use S3 policies to limit access.

C. Create a data lake by using AWS Lake Formation. Create an AWS Glue JDBC connection to Amazon RDS. Register the S3 bucket in Lake Formation. Use Lake Formation access controls to limit access.

D. Create an Amazon Redshift cluster. Schedule an AWS Lambda function to periodically copy data from Amazon S3 and Amazon RDS to Amazon Redshift. Use Amazon Redshift access controls to limit access.

Correct Answer: *D*

Question #205 *Topic 1*

A company hosts a marketing website in an on-premises data center. The website consists of static documents and runs on a single server. An administrator updates the website content infrequently and uses an SFTP client to upload new documents.

The company decides to host its website on AWS and to use Amazon CloudFront. The company's solutions architect creates a CloudFront distribution. The solutions architect must design the most cost-effective and resilient architecture for website hosting to serve as the CloudFront origin.

Which solution will meet these requirements?

A. Create a virtual server by using Amazon Lightsail. Configure the web server in the Lightsail instance. Upload website content by using an SFTP client.

B. Create an AWS Auto Scaling group for Amazon EC2 instances. Use an Application Load Balancer. Upload website content by using an SFTP client.

C. Create a private Amazon S3 bucket. Use an S3 bucket policy to allow access from a CloudFront origin access identity (OAI). Upload website content by using the AWS CLI.

D. Create a public Amazon S3 bucket. Configure AWS Transfer for SFTP. Configure the S3 bucket for website hosting. Upload website content by using the SFTP client.

Correct Answer: *C*

Question #206 *Topic 1*

A company wants to manage Amazon Machine Images (AMIs). The company currently copies AMIs to the same AWS Region where the AMIs were created. The company needs to design an application that captures AWS API calls and sends alerts whenever the Amazon EC2 CreateImage API operation is called within the company's account.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an AWS Lambda function to query AWS CloudTrail logs and to send an alert when a CreateImage API call is detected.

B. Configure AWS CloudTrail with an Amazon Simple Notification Service (Amazon SNS) notification that occurs when updated logs are sent to Amazon S3. Use Amazon Athena to create a new table and to query on CreateImage when an API call is detected.

C. Create an Amazon EventBridge (Amazon CloudWatch Events) rule for the CreateImage API call. Configure the target as an Amazon Simple Notification Service (Amazon SNS) topic to send an alert when a CreateImage API call is detected.

D. Configure an Amazon Simple Queue Service (Amazon SQS) FIFO queue as a target for AWS CloudTrail logs. Create an AWS Lambda function to send an alert to an Amazon Simple Notification Service (Amazon SNS) topic when a CreateImage API call is detected.

Correct Answer: *D*

Question #207 *Topic 1*

A company owns an asynchronous API that is used to ingest user requests and, based on the request type, dispatch requests to the appropriate microservice for processing. The company is using Amazon API Gateway to deploy the API front end, and an AWS Lambda function that invokes Amazon DynamoDB to store user requests before dispatching them to the processing microservices.

The company provisioned as much DynamoDB throughput as its budget allows, but the company is still experiencing availability issues and is losing user requests.

What should a solutions architect do to address this issue without impacting existing users?

A. Add throttling on the API Gateway with server-side throttling limits.

B. Use DynamoDB Accelerator (DAX) and Lambda to buffer writes to DynamoDB.

C. Create a secondary index in DynamoDB for the table with the user requests.

D. Use the Amazon Simple Queue Service (Amazon SQS) queue and Lambda to buffer writes to DynamoDB.

Correct Answer: *D*

Question #208 *Topic 1*

A company needs to move data from an Amazon EC2 instance to an Amazon S3 bucket. The company must ensure that no API calls and no data are routed through public internet routes. Only the EC2 instance can have access to upload data to the S3 bucket.

Which solution will meet these requirements?

A. Create an interface VPC endpoint for Amazon S3 in the subnet where the EC2 instance is located. Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.

B. Create a gateway VPC endpoint for Amazon S3 in the Availability Zone where the EC2 instance is located. Attach appropriate security groups to the endpoint. Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.

C. Run the nslookup tool from inside the EC2 instance to obtain the private IP address of the S3 bucket's service API endpoint. Create a route in the VPC route table to provide the EC2 instance with access to the S3 bucket. Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.

D. Use the AWS provided, publicly available ip-ranges.json file to obtain the private IP address of the S3 bucket's service API endpoint. Create a route in the VPC route table to provide the EC2 instance with access to the S3 bucket. Attach a resource policy to the S3 bucket to only allow the EC2 instance's IAM role for access.

Correct Answer: *B*

Question #209 *Topic 1*

A solutions architect is designing the architecture of a new application being deployed to the AWS Cloud. The application will run on Amazon EC2 On-Demand Instances and will automatically scale across multiple Availability Zones. The EC2 instances will scale up and down frequently throughout the day. An Application Load Balancer (ALB) will handle the load distribution. The architecture needs to support distributed session data management. The company is willing to make changes to code if needed.

What should the solutions architect do to ensure that the architecture supports distributed session data management?

A. Use Amazon ElastiCache to manage and store session data.

B. Use session affinity (sticky sessions) of the ALB to manage session data.

C. Use Session Manager from AWS Systems Manager to manage the session.

D. Use the GetSessionToken API operation in AWS Security Token Service (AWS STS) to manage the session.

Correct Answer: *A*